

O'REILLY®



Modern Linux Administration

HOW TO BECOME A CUTTING-EDGE LINUX ADMINISTRATOR

Sam R. Alapati

Modern Linux Administration

Sam R. Alapati

Beijing • Boston • Farnham • Sebastopol • Tokyo

O'REILLY®

Modern Linux Administration

by Sam R. Alapati

Copyright © 2016 Sam Alapati. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc. , 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://safaribooksonline.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com .

Editor: Brian Anderson

Production Editor: FILL IN PRODUCTION EDITOR

Copyeditor: FILL IN COPYEDITOR

Proofreader: FILL IN PROOFREADER

Indexer: FILL IN INDEXER

Interior Designer: David Futato

Cover Designer: Karen Montgomery

Illustrator: Rebecca Demarest

January -4712: First Edition

Revision History for the First Edition

2016-10-13: First Early Release

2016-11-30: Second Early Release

2017-01-12: Third Early Release

2017-02-07: Fourth Early Release

See <http://oreilly.com/catalog/errata.csp?isbn=9781491935910> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. Modern Linux Administration, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

While the publisher and the author(s) have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the author(s) disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-491-93591-0

[FILL IN]

Table of Contents

| | |
|---|-------------|
| Preface..... | xiii |
| 1. Modern Linux System Administration..... | 23 |
| Motivation for the new System Administration Strategies and Tools | 25 |
| Problems with Traditional Systems Administration | 25 |
| Monitoring | 26 |
| The Image Sprawl problem | 26 |
| Agile Development Methodologies and the System Administrator | 26 |
| Cloud environments | 27 |
| Impact of Big Data | 27 |
| Manual Operations without automation | 27 |
| Automated Infrastructure Management | 28 |
| Automating Redundant Configuration Work | 29 |
| Configuration Management | 30 |
| Infrastructure as Code | 31 |
| Modern Scripting Languages and Databases | 32 |
| Essential Programming Skills for the System Administrator | 32 |
| The Rise of NoSQL Databases | 33 |
| Caching | 33 |
| Content Delivery Networks | 34 |
| IT Orchestration | 34 |
| Provisioning with Vagrant | 35 |
| Vagrant and Configuration Management Tools | 35 |
| Provisioning - Automatic Server Deployment | 36 |
| Server (hardware) Virtualization | 36 |
| Containerization – the New Virtualization | 37 |
| Docker and Containerization | 38 |
| Docker Container Orchestration and Distributed Schedulers | 39 |

| | |
|---|-----------|
| Cluster Management and Cluster Operating Systems | 41 |
| Version Control Systems | 42 |
| Continuous Integration and Continuous Deployment | 43 |
| Benefits Offered by CI | 45 |
| Steps involved in CI | 45 |
| Continuous Integration and Continuous Deployment | 45 |
| Continuous Application deployment | 46 |
| Tools for Implementing Continuous Integration | 47 |
| Log Management, Monitoring and Metrics | 48 |
| Effective Metrics | 49 |
| Proactive Monitoring | 50 |
| Service Metrics | 50 |
| Synthetic Monitoring | 51 |
| Cloud Computing | 51 |
| Open Stack – the Open Source Cloud Platform | 52 |
| Software Defined Networking | 52 |
| Microservices, Service Registration and Service Discovery | 53 |
| Benefits of Microservices | 55 |
| Service Discovery | 55 |
| Service Registration | 56 |
| 2. Networking Essentials for a System Administrator..... | 23 |
| What the internet is | 58 |
| Packets | 58 |
| Packet Switches | 59 |
| Applications and APIs | 59 |
| Network Protocols | 59 |
| Network Messages and Message Formatting | 60 |
| Networking Essentials – Theory and Practice | 61 |
| Breaking up the real work into layers – the OSI Model | 61 |
| Protocol Layering | 62 |
| The Internet Protocol Stack | 64 |
| The Network Layer and TCP/IP Networking | 65 |
| The Forwarding Function | 66 |
| Routing Essentials and Routing Management | 75 |
| The Hypertext Transfer Protocol (HTTP) | 81 |
| Using the HTTP/2 Protocol for Enhanced Performance | 82 |
| Network Load Balancing | 83 |
| Benefits of Using a Network Load Balancer | 84 |
| Load Balancing with DNS | 84 |
| Enterprise Load Balancers | 85 |
| Software Based Load Balancing | 86 |

| | |
|---|-----------|
| Hardware Load Balancing | 88 |
| Using a Hosted Load Balancer Service | 89 |
| Modern Networking Technologies | 89 |
| Quality of Service (QoS) | 90 |
| Quality of Experience (QoE) | 90 |
| Routing and Network Congestion Control | 91 |
| Software-Defined Networking | 92 |
| Limitations of current networks | 92 |
| The three “Planes” in Networking | 93 |
| Defining Functions for the Network Control Plane | 95 |
| Network Functions Virtualization | 97 |
| The OpenFlow Protocol | 97 |
| 3. Scalability, Web Applications, Web Services, and Microservices..... | 23 |
| Scaling and Common Datacenter Infrastructures | 101 |
| The Front End Technologies | 102 |
| The Back End Technologies | 102 |
| Scalability of Applications | 103 |
| Content Delivery Networks | 105 |
| How large websites scale | 106 |
| Scaling Web Applications | 106 |
| Managing state at the front end | 107 |
| Other Types of State | 109 |
| Scaling the Web Services | 111 |
| Making Effective use of Third-party services | 113 |
| Working with Web Servers | 115 |
| Working with the Apache Web Server | 115 |
| The NGINX Web Server | 115 |
| Caching Proxies and Reverse Proxying | 117 |
| Handling Data Storage with Databases | 119 |
| Relational databases | 119 |
| Other Types of Databases | 119 |
| MongoDB as a Backend Database | 120 |
| Caching | 124 |
| HTTP-Based Caching (Browser caching) | 124 |
| Caching Objects | 128 |
| Asynchronous Processing, Messaging Applications and MOM | 131 |
| Messages and Message Queues | 132 |
| Components of a Messaging Architecture | 133 |
| Message Brokers and Message Oriented Middleware (MOM) | 134 |
| Messaging Protocols | 134 |
| Popular Message Brokers | 136 |

| | |
|--|-----------|
| The Model-View-Controller Architecture and Single Paged Applications | 137 |
| The Problem | 137 |
| MVC to the Rescue | 138 |
| Ruby on Rails | 139 |
| Full stack JavaScript Development with MEAN | 140 |
| Single Page Applications – the new Paradigm for Web Applications | 142 |
| Web Services | 145 |
| Web Service Basics | 146 |
| Simple Object Access Protocol (SOAP) | 146 |
| Two Types of Web Services | 148 |
| Service-Based Architectures and Microservices | 150 |
| Similarities between traditional SOAs and the Microservice approach | 151 |
| Differences between SOA and Microservices | 152 |
| Service Types | 153 |
| 4. Server Virtualization and Linux Containers..... | 23 |
| Linux Server Virtualization | 156 |
| The Architecture of Virtual Machines | 156 |
| The Virtual Machine Monitor (Hypervisor) | 157 |
| How VMs share Resources | 157 |
| Benefits offered by Virtual Machines | 159 |
| Drawbacks of Virtualization | 160 |
| Virtualization Types | 160 |
| Type of Hypervisors | 164 |
| Xen Virtualization | 165 |
| Kernel-Based Virtual Machines (KVM) | 166 |
| Considerations in Selecting the Physical Servers for virtualization | 169 |
| Migrating Virtual Machines | 170 |
| Application Deployment and Management with Linux Containers | 170 |
| Chroot and Containers | 172 |
| Applications and their Isolation | 173 |
| Virtualization and Containerization | 173 |
| Benefits offered by Linux Containers | 174 |
| Two Types of Uses for Linux Containers | 176 |
| The Building Blocks of Linux Containers | 177 |
| Namespaces and Process Isolation | 177 |
| Control Groups (cgroups) | 179 |
| SELinux and Container Security | 183 |
| Linux Containers versus Virtualization (KVM) | 185 |
| Linux Containers and KVM Virtualization – the Differences | 187 |
| Limitations of LXC | 187 |
| Container benefits | 187 |

| | |
|--|-----------|
| Linux Container Adoption Issues | 188 |
| Managing Linux Containers | 189 |
| 5. Working with Docker Containers..... | 23 |
| Docker Basics | 192 |
| When Docker Isn't Right for You | 193 |
| What Docker Consists of | 194 |
| The Docker Project | 194 |
| Docker Images and Docker Containers | 195 |
| What Linux Administrators should know in order to support Docker | 196 |
| Setting up the Docker Container Run-Time Environment | 198 |
| Getting Information about the Containers | 199 |
| Running Container Images | 200 |
| Managing Containers | 201 |
| Running Interactive Containers | 201 |
| Making your base Image Heftier | 202 |
| Committing a Container | 202 |
| Running Commands within a Container | 203 |
| Linking Containers | 203 |
| Running Services inside a Container | 203 |
| Running Privileged Containers | 205 |
| Building Docker Images | 205 |
| Building Images with a Dockerfile | 205 |
| Image Layers | 209 |
| Docker Image Repositories | 209 |
| Using a Private Docker Registry | 209 |
| New Operating Systems Optimized for Docker | 211 |
| Using CoreOS | 212 |
| Working with Atomic Host | 212 |
| The Docker Stack in Production | 213 |
| Provisioning Resources with Docker Machine | 214 |
| Docker Orchestration | 214 |
| Docker Orchestration and Clustering Tools | 215 |
| Distributed Schedulers for Docker Containers | 215 |
| Docker Containers, Service Discovery and Service Registration | 218 |
| Connecting Containers through Ambassadors | 219 |
| How Service Discovery Works | 220 |
| Zero-Configuration Networking | 220 |
| Service Discovery | 221 |
| Service Registration | 223 |

| | |
|---|---------------|
| 6. Automating Server Deployment and Managing Development Environments. | 23 |
| Linux Package Management | 226 |
| Using the rpm and dpkg commands | 226 |
| Package Management with YUM and APT | 226 |
| Fully Automatic Installation (FAI) | 227 |
| How FAI Works | 227 |
| How it Works | 228 |
| Setting up the Network Server | 228 |
| Setting up the PXE Server for a Network Installation | 228 |
| Using Kickstart | 229 |
| Automatically Spinning up Virtual Environments with Vagrant | 231 |
| Some Background | 233 |
| Vagrant and its Alternatives | 234 |
| Getting Started with Vagrant | 235 |
| Spinning up a New VM | 236 |
| The Vagrantfile | 238 |
| Vagrant Box | 238 |
| Provisioning Vagrant Boxes | 240 |
| Automated Provisioning | 242 |
| Creating Vagrant Base Boxes | 244 |
| Using Packer and Atlas for creating base boxes | 246 |
| Parallel Job-Execution and Server Orchestration Systems | 247 |
| Working with Remote Command Execution Tools | 247 |
| Server Provisioning with Razor | 252 |
| Server Provisioning with Cobbler | 254 |
| 7. Infrastructure as Code, Configuration Management and Orchestration Tools. | 23 |
| Infrastructure as Code | 258 |
| Benefits of Treating Infrastructure as Code | 259 |
| Cloud Environments and Infrastructure as Code | 260 |
| Tools for Treating Infrastructure as Code | 260 |
| Test-Driven Infrastructures | 261 |
| Configuration Management Automation | 263 |
| Need for Configuration Management Tools | 263 |
| Declarative Versus Imperative Programming | 264 |
| Configuration Management with Chef | 265 |
| Chef's Architecture and Key Components | 265 |
| Writing Chef Recipes | 267 |
| Using Test Kitchen to Manage Test Environments | 270 |
| Using Chef Client to Manage Nodes | 271 |
| Using Puppet for Configuration | 275 |
| Installing Puppet | 275 |

| | |
|---|-----------|
| Puppet Manifests | 275 |
| Creating Puppet Modules | 278 |
| Using a Puppet Server | 281 |
| Server Provisioning and Orchestration with Salt | 283 |
| Architecture of Salt – the Basics | 283 |
| Examples of Salt usage | 284 |
| Grains and Pillars | 285 |
| Basic Salt Commands | 286 |
| Execution Modules | 287 |
| Salt States and Configuration Management | 288 |
| Using Ansible for CM, Orchestration, and Provisioning | 290 |
| How Ansible is Different from the other CM Tools | 291 |
| The Guts of Ansible | 291 |
| How Playbooks Work | 293 |
| Modules | 293 |
| Inventories Describe your Infrastructure | 294 |
| 8. Version Control and Source Code Management..... | 23 |
| Version Control - An Introduction | 297 |
| Types of Version Control Systems | 297 |
| Revision Control and Infrastructure Management | 298 |
| Why Use Source Control? | 299 |
| Our First Foray into the World of Git | 299 |
| Introduction to Git | 301 |
| What Sets Git Apart from other CVS Systems? | 301 |
| Creating Local Repositories with Git | 302 |
| Creating a Repository | 302 |
| Checking the Repository Status | 302 |
| Checking the Logs and the Status of a Repository | 304 |
| Staging and Committing Files | 305 |
| Committing Multiple Files | 305 |
| Cloning Repositories | 306 |
| Remote Git | 307 |
| The Git Branching Model | 307 |
| Creating Branches | 308 |
| Merging Branches | 309 |
| Resolving Conflicts | 309 |
| Setting up Git on a Server – Hosting Git Repositories | 309 |
| Multiple Protocols | 309 |
| Setting up a Git Server | 311 |
| Enterprise Deployment of Git | 312 |
| Tools for Managing Git Repositories | 312 |

| | |
|---|-----------|
| Using GitHub | 312 |
| Setting up your Account | 313 |
| The GitHub Workflow | 313 |
| 9. Continuous Integration, Continuous Delivery, and Continuous Deployment. | 23 |
| Builds and Deployment | 316 |
| How Build Tools Help – Maven, Ant and Gradle | 318 |
| The Maven Build Tool | 318 |
| Using Ant to Perform Builds | 322 |
| Using Gradle as a Build Tool | 324 |
| An Overview of the Software Development Process | 324 |
| The Software Development Life Cycle | 325 |
| Following the Waterfall Model of Development | 325 |
| The Agile Software Development Process | 325 |
| The Scrum Framework | 326 |
| The Agile Software Development Process and Continuous Integration | 327 |
| Continuous Integration (CI) | 328 |
| The CI Process | 328 |
| Problems Detected by CI | 329 |
| Building Software at Every Change | 329 |
| Benefits of a CI System | 330 |
| Best Practices for CI | 330 |
| Requirements for a CI System | 332 |
| Features of a CI System | 332 |
| A Typical CI Scenario | 333 |
| Continuous Delivery | 335 |
| CD and CI | 335 |
| The Build (or Deployment) Pipeline | 336 |
| Continuous Delivery and DevOps | 337 |
| Continuous Integration with Jenkins | 337 |
| Architecture of Jenkins | 337 |
| Components of the Jenkins Framework | 338 |
| Jenkins and other Continuous Integration Tools | 338 |
| The Need for a Binary Repository Tool | 339 |
| A Typical Continuous Integration Strategy using Jenkins | 340 |
| Branching strategies | 340 |
| Creating the Continuous Integration Pipelines | 341 |
| Continuous Delivery with Jenkins | 341 |
| Continuous Deployment with Jenkins | 342 |
| Continuous Integration with TeamCity/GO | 343 |

| | |
|--|-----------|
| 10. Centralized Log Management and Analysis, and Handling Streaming Data..... | 23 |
| Introduction to Logging and Its Issues | 346 |
| What is a Log? | 346 |
| Varieties of Log Data | 346 |
| Issues in Performing Log Analysis | 347 |
| Centralizing Logs with Syslog | 349 |
| Using Rsyslog for Logging | 349 |
| SYSLOG-NG | 351 |
| Rotating Logs with Logrotate | 351 |
| Log Aggregation | 352 |
| Possible Strategies | 352 |
| Log Aggregation Workflow | 352 |
| Log Aggregation Alternatives | 353 |
| Effective Log Management and Log Analysis with ElasticSearch, Logstash, and Kibana | 353 |
| Searching Engines and Indexing - Elasticsearch | 354 |
| Basic Concepts of ElasticSearch | 355 |
| Understanding the ElasticSearch APIs | 358 |
| Elasticsearch Marvel and Shield | 360 |
| Logstash | 360 |
| Logstash Plugins | 361 |
| Logstash under the Hood | 362 |
| The Logstash Workflow | 363 |
| Configuring Logstash – Some Examples | 364 |
| Kibana | 366 |
| A Simple Example Showing How ELK Works | 367 |
| The Input Data | 367 |
| Configuring the Logstash Input | 367 |
| Using ElasticSearch in a Production Setting | 368 |
| Log Analysis with Splunk | 370 |
| Scaling and High Availability | 371 |
| Working with Data | 371 |
| Using the Splunk Processing Language | 372 |
| Advanced Analytics with Log Files | 373 |
| Splunk versus ELK | 374 |
| Saas solutions for Logging –Sumologic and Loggly | 374 |
| Handling Streaming Data | 375 |
| Apache Kafka | 377 |
| Benefits offered by Kafka | 378 |
| Kafka Architecture | 378 |
| How Kafka Works | 380 |

| | |
|---|-----------|
| 11. System Monitoring, Event Processing, and Tracking Application Performance. | 23 |
| Monitoring – an Overview | 385 |
| How Monitoring is generally done | 385 |
| Need for New Monitoring Tools | 386 |
| Do-it-yourself Versus using External Monitoring Services | 386 |
| Monitoring and the Mean Time to Recovery | 387 |
| How Metrics and Logs can help Proactive Behavior | 388 |
| Different Types of Metrics | 388 |
| Metric Tracking across Multiple Services | 389 |
| Devops and Monitoring | 389 |
| Time Series Databases | 389 |
| Using Cacti | 390 |
| Using Librato | 390 |
| Graphite | 391 |
| Components of Graphite | 392 |
| Statistics and Transformations | 393 |
| Event Processing Tools | 393 |
| Riemann | 394 |
| Big Panda | 396 |
| System Monitoring | 398 |
| Nagios | 398 |
| Ganglia | 402 |
| Using Munin | 403 |
| Using Collectd | 403 |
| Using Senu | 403 |
| Web and User Monitoring | 406 |
| Synthetic Monitoring | 407 |
| Real User Monitoring | 407 |
| Using Synthetic Monitoring and RUM Together | 408 |
| Pingdom | 408 |
| Application Performance Monitoring (APM) Suites | 409 |
| New Relic APM | 411 |
| AppDynamics | 411 |
| Comparison of New Relic and AppDynamics | 412 |
| On-Call Management Tools | 412 |

Who Should Read This Book

The quintessential reader for this book is someone who currently works as a Linux systems administrator, or wants to become one, having already acquired basic Linux admin skills. However, the books will be useful for all of the following:

- Developers who need to come to terms with systems concept such as scaling, as well as the fundamentals of important concepts that belong to the operations world – networking, cloud architectures, site reliability engineering, web performance, and so on.
- Enterprise architects who are either currently handling , or are in the process of creating new projects dealing with scaling of web services, Docker containerization, virtualization, big data, cloud architectures.
- Site reliability engineers (SREs), backend engineers and distributed application developers who are tasked with optimizing their applications as well as scaling their sites, in addition to managing and troubleshooting the new technologies increasingly found in modern systems operations.

In terms of Linux administration knowledge and background, I don't teach the basics of Linux system administration in this book. I expect the readers to know how to administer a basic Linux server and be able to perform tasks such as creating storage, managing users and permissions, understand basic Linux networking, managing files and directories, managing processes, troubleshooting server issues, taking backups, and restoring servers.

The overarching goal of this book is to introduce the reader to the myriad tools and technologies that a Linux administrator ought to know today to earn his or her keep. I do provide occasional examples, but this book is by no means a “how-to” reference for any of these technologies and software. As you can imagine, each of the technologies I discuss over the 16 chapters in this book requires one or more books dedicated

to that technology alone, for you to really learn that topic. There's no code or step-by-step instructions for the numerous newer Linux administration related technologies I discuss in this book, with a handful of exceptions. My goal is to show you what you need to know in order to understand, evaluate, and prepare to work with bleeding-edge Linux based technologies in both development and production environments.

Why I Wrote This Book

Let's say you want to learn all about the new containerization trend in application deployment and want to use Docker to make your applications portable. Just trying to come to grips with the wide range of technologies pertaining to Docker is going to make anybody's head spin – here's a (partial) list of technologies associated with just Docker containers:

- Docker project
- Docker Hub Registry
- Docker Images and Dockerfiles
- CoreOS and Atomic Host
- Cockpit
- Kubernetes
- Swarm
- Compose
- Machine
- Mesos
- Zookeeper
- Consul
- Eureka
- Smartstack
- OpenShift

And all this just to learn how to work with Docker!

No wonder a lot of people are baffled as to how to get a good handle on the new technologies, which are sometimes referred to as DevOps (however you may define it!), but really involves a new way of thinking and working with new cutting edge technologies. Many of these technologies were expressly designed to cope with the newer trends in application management such as the use of microservices, and newer ways of doing business such as cloud based environments, and new ways of data analysis such as the use of Big Data for example.

Over the past decade or so, there have been fundamental changes in how Linux system administrators have started approaching their work. Earlier, Linux admins were typically heavy on esoteric knowledge about the internals of the Linux server itself, such as rebuilding the kernel, for example. Other areas of expertise that marked one as a good Linux administrator were things such as proficiency in shell scripting, awk & sed, and Perl & Python.

Today, the emphasis has shifted quite a bit – you still need to know all that a Linux admin was expected to know years ago, but the focus today is more on your understanding of networking concepts such as DNS and routing, scaling of web applications and web services, web performance and monitoring, cloud based environments, big data and so on, all of which have become required skills for Linux administrators over the past decade.

In addition to all the new technologies and new architectures, Linux system administrators have to be proficient in new ways of doing business – such as using the new-fangled configuration management tools, centralized version control depositories, continuous development (CI) and continuous deployment (CD), just to mention a few technologies and strategies that are part of today's Linux environments

As a practicing administrator for many years, and someone who needs to understand which of the technologies out of the zillion new things out there really matter to me, it's struck me that there's a lack of a single book that serves as a guide for me to navigate this exciting but complex new world. If you were to walk into an interview to get hired as a Linux administrator today, how do you prepare for it? What are you really expected to know? How do all these new technologies related to each other? Where do I start? I had these types of concerns for a long time, and I believe that there are many people that understand that changes are afoot and don't want to be left behind, but don't know how and where to begin.

My main goal in this book is to explain what a Linux administrator (or a developer/architect who uses Linux systems) needs to understand about currently popular technologies. My fundamental thesis is that traditional systems administration as we know it won't cut it in today's technologically complex systems dominated by web applications, big data, and cloud-based systems. To this end, I explain the key technologies and trends that are in vogue today (and should hold steady for a few years at least), and the concepts that underlie those technologies.

There's a bewildering array of modern technologies and tools out there and you're expected to really know how and where to employ these tools. Often, professionals seeking to venture out into the modern systems administration areas aren't quite sure where exactly they ought to start, and how the various tools and technologies are related. This book seeks to provide sufficient background and motivation for all the key tools and concepts that are in use today in the systems administration area, so you can go forth and acquire those skill sets.

A Word on the New Technologies that are critical for Linux Administrators Today

In order to be able to write a book such as this, with its wide-ranged and ambitious scope, I've had to make several decisions in each chapter as to which technologies I should discuss in each of the areas I chose to cover in the book. So, how did I pick the topics that I wanted to focus on? I chose to reverse engineer the topic selection process, meaning that I looked at what organizations are looking for today in a Linux administrator when they seek to hire one. And the following is what I found.

Expertise in areas such as infrastructure automation and configuration management (Chef, Puppet, Ansible, SaltStack), version control (Git and Perforce), big data (Hadoop), cloud architectures (Amazon Web Services, OpenStack), monitoring and reporting (Nagios and Ganglia), new types of web servers (Nginx), load balancing (Keepalived and HAProxy), databases (MySQL, MongoDB, Cassandra), caching (Memcached and Redis), Virtualization (kvm), containers (Docker), server deployment (Cobbler, Foreman, Vagrant), source code management (Git/Perforce), version control management (Mercurial and Subversion), Continuous integration and delivery (Jenkins and Hudson), log management (Logstash/ElasticSearch/Kibana), metrics management (Graphite, Cacti and Splunk) .

Look up any job advertisement for a Linux administrator (or Devops administrator) today and you'll find all the technologies I listed among the required skillsets. Most of the jobs listed require you to have a sound background and experience with basic Linux system administration – that's always assumed - plus they need many of the technologies I listed here.

So, the topics I cover and the technologies I introduce and explain are based on what a Linux administrator is expected to know today to work as one. My goal is to explain the purpose and the role of each technology, and provide a conceptual explanation of each technology and enough background and motivation for you to get started on the path to mastering those technologies. This book can thus serve as your “road map” for traversing this exciting (but intimidating) new world full of new concepts and new software, which together have already transformed the traditional role of a system administrator.

Navigating This Book

This book is organized roughly as follows:

- Chapter 1 explains the key trends in modern systems administration, such as virtualization, containerization, version control systems, continuous deployment and delivery, big data, and many other newer areas that you ought to be familiar with, to succeed as a system administrator or architect today. I strive to drive

home the point that in order to survive and flourish as a system administrator in today's highly sophisticated Linux based application environments, you must embrace the new ways of doing business, which includes a huge number of new technologies, as well as new ways of thinking. No longer is the Linux system administrator an island until himself (or herself)! In this exciting new world, you'll be working very closely with developers and architects – so, you must know the language the other speaks, as well as accept that the other groups such as developers will be increasingly performing tasks that were once upon a long time used to belong to the exclusive province of the Linux administrators. Tell me, in the old days, did any developer carry a production pager? Many do so today.

- Chapter 2 provides a quick and through introduction to several key areas of networking, including the TCP/IP network protocol, DNS, DHCP, SSH and SSL, subnetting and routing, and load balancing. The chapter concludes with a review of newer networking concepts such as Software Defined Networking (SDN). Networking is much more important now than before, due to its critical role in cloud environments and containerization.
- Chapter 3 is a far ranging chapter dealing with the scaling of web applications and provides an introduction to web services, modern databases, and new types of web servers. You'll learn about web services and microservices and the differences between the two architectures. The chapter introduces you to concepts such as APIs, REST, SOAP and JSON, all of which play a critical role in modern web applications, which are a key part of any systems environment today. Service discovery and service registration are important topics in today's container heavy environments and you'll get an introduction to these topics here. The chapter also introduces you to modern web application servers such as Nginx, caching databases such as Redis and NoSQL databases (MongoDB).
- Chapter 4 discusses traditional virtualization and explains the different types of hypervisors. The chapter also introduces containers and explains the key ideas behind containerization, such as namespaces. SELinux and Cgroups (control groups), thus helping you get ready for the next chapter, which is all about Docker.
- Chapter 5 is one of the most important chapters in the book since it strives to provide you a through introduction to Docker containers. You'll learn about the role containerization plays in supporting application deployment and portability. You'll learn the basics of creating and managing Docker containers. The chapter explains the important and quite complex topic of Docker networking, both in the context of a single container as well as networking among a bunch of containers. The chapter discusses exciting technologies such as Kubernetes, which helps orchestrate groups of containers, as well as how to use Flannel to set up IP address within a Kubernetes cluster. I also show how to use Cockpit, a Web-based

container management tool, to manage containers running in multiple hosts in your own cloud. New slimmed down operating systems such as CoreOs and Red Hat Atomic Host are increasingly popular in containerized environments and therefore, I explain these types of “container operating systems” as well in this chapter.

- Chapter 6 shows how to automate server creation with the help of tools such as PXE servers, and automatic provisioning with Razor, Cobbler and Foreman. You’ll learn how Vagrant helps you easily automate the spinning up of development environments.
- Chapter 7 explains the principles behind modern configuration management tools, and shows how popular tools such as Puppet and Chef work. In addition, you’ll learn about two very popular orchestration frameworks – Ansible and Salt-stack.
- Chapter 8 discusses two main topics – revision control and source code management. You’ll learn about using Git and GitHub for revision control, as well as other revision control tools such as Mercurial, Subversion and Perforce.
- Chapter 9 is about two key modern application development concepts – continuous integration (CI) and continuous delivery (CD). The chapter explains how to employ tools such as Hudson, Jenkins, and Travis for CD and CI.
- Chapter 10 has two main parts. The first part is about centralized log management with the ELK (Elasticsearch, Logstash, and Kibana) stack. Performing trend analyses and gathering metrics with tools such as Graphite, Cacti, Splunk, and DataDog is the focus of the second part of the chapter.
- Chapter 11 shows how to use the popular OpenStack software to create an enterprise Infrastructure--as-a-Service. You’ll learn the architecture and concepts relating to the OpenStack cloud, and how it integrates with PaaS (Platform-as-a-Service) solutions such as Red Hat OpenShift and CloudFoundry.
- Chapter 12 is about using Nagios for monitoring and alerts and also explains the concepts and architecture that underlie Ganglia, an excellent way to gather system performance metrics. I also introduce two related tools – Sensu for monitoring and Zabbix for log management.
- Chapter 13 provides you a quick overview of Amazon Web Services (AWS) and the Google Cloud Platform, two very successful commercial cloud platforms.
- Chapter 14 consists of two main parts: the first part is about managing new types of databases such as MongoDB and Cassandra. The second part of the chapter explains the role of the Linux administrator in supporting big data environments powered by Hadoop. Hadoop is increasingly becoming popular and you need to know the concepts that underlie Hadoop, as well as the architecture of Hadoop 2,

the current version. The chapter shows how to install and configure Hadoop at a high level, as well how to use various tools to manage Hadoop storage (HDFS).

- Chapter 15 deals with security and compliance concerns in a modern systems environment. The chapter explains the unique security concerns of cloud environments, and how to secure big data such as Hadoop's data. You'll learn about topics such as identity and access management in AWS, virtual private networks, and security groups. The chapter closes by discussing Docker security, and how to make concessions to traditional security best practices in a containerized environment, and how to use super privileged containers.
- Chapter 16 is somewhat of a mixed bag! This final chapter is mostly about software reliability engineering (SRE) and it does by explaining various performance related topics such as enhancing Web Server performance, tuning databases and JVMs (Java Virtual Machines), and tuning the network. You'll learn about web site performance optimization using both RUM (real user monitoring) and through generating synthetic performance statistics.

If you're like us, you don't read books from front to back. If you're really like us, you usually don't read the Preface at all! Here are some basic guidelines as to how you may approach the book:

- Read Chapter 1 in order to understand the scope of the book and the lay of the land, so to speak. This chapter provides the motivation for the discussion of all the technologies and concepts in the remaining chapters.
- Quickly glance through Chapter 2, if you think you need a refresher course in essential networking concepts for a Linux administrator. If your networking chops are good, skip most of Chapter 2, except the very last part, which deals with modern networking concepts such as software defined networks (SDN).
- You can read the rest of the chapters in any order you like, depending on your interest and needs – there are really no linkages among the chapters of the book!
- Remember that conceptual overview of the various tools and software and explanation of the technical architectures are the real focus of the book – if you need to drill deep into the installation and configuration of the various tools, you'll need to read the documentation for that tool (or a book on that topic).

I hope you enjoy each of the chapters as much as I've enjoyed writing the chapters!

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Indicates new terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, databases, data types, environment variables, statements, and keywords.

Constant width bold

Shows commands or other text that should be typed literally by the user.

Constant width italic

Shows text that should be replaced with user-supplied values or by values determined by context.



This icon signifies a tip, suggestion, or general note.



This icon indicates a warning or caution.

Using Code Examples

PROD: Please reach out to author to find out if they will be uploading code examples to oreilly.com or their own site (e.g., GitHub). If there is no code download, delete this whole section. If there is, when you email digidist with the link, let them know what you filled in for title_title (should be as close to book title as possible, i.e., learning_python_2e). This info will determine where digidist loads the files.


Supplemental material (code examples, exercises, etc.) is available for download at https://github.com/oreillymedia/title_title.

This book is here to help you get your job done. In general, if example code is offered with this book, you may use it in your programs and documentation. You do not need to contact us for permission unless you're reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from this book does not require permission. Selling or distributing a CD-ROM of examples from O'Reilly books does require permission. Answering a question by citing this book and quoting example code does not require permission. Incorporating a significant amount of example code from this book into your product's documentation does require permission.

We appreciate, but do not require, attribution. An attribution usually includes the title, author, publisher, and ISBN. For example: “*Book Title* by Some Author (O’Reilly). Copyright 2012 Some Copyright Holder, 978-0-596-xxxx-x.”

If you feel your use of code examples falls outside fair use or the permission given above, feel free to contact us at permissions@oreilly.com.

Safari® Books Online

 **Safari®** Safari Books Online is an on-demand digital library that delivers expert content in both book and video form from the world’s leading authors in technology and business.

Technology professionals, software developers, web designers, and business and creative professionals use Safari Books Online as their primary resource for research, problem solving, learning, and certification training.

Safari Books Online offers a range of product mixes and pricing programs for organizations, government agencies, and individuals. Subscribers have access to thousands of books, training videos, and prepublication manuscripts in one fully searchable database from publishers like O’Reilly Media, Prentice Hall Professional, Addison-Wesley Professional, Microsoft Press, Sams, Que, Peachpit Press, Focal Press, Cisco Press, John Wiley & Sons, Syngress, Morgan Kaufmann, IBM Redbooks, Packt, Adobe Press, FT Press, Apress, Manning, New Riders, McGraw-Hill, Jones & Bartlett, Course Technology, and dozens more. For more information about Safari Books Online, please visit us online.

How to Contact Us

Please address comments and questions concerning this book to the publisher:

O’Reilly Media, Inc. 1005 Gravenstein Highway North Sebastopol, CA 95472
800-998-9938 (in the United States or Canada) 707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at <http://www.oreilly.com/catalog/<catalogpage>>.

Don’t forget to update the link above.

To comment or ask technical questions about this book, send email to bookquestions@oreilly.com.

For more information about our books, courses, conferences, and news, see our website at <http://www.oreilly.com>.

Find us on Facebook: [*http://facebook.com/oreilly*](http://facebook.com/oreilly)

Follow us on Twitter: [*http://twitter.com/oreillymedia*](http://twitter.com/oreillymedia)

Watch us on YouTube: [*http://www.youtube.com/oreillymedia*](http://www.youtube.com/oreillymedia)

Acknowledgments

Fill in...

Modern Linux System Administration

Linux (and other) system administration has changed tremendously since the advent of internet based web applications and the proliferation of Big Data based systems, and the rush to cloud based systems. A quick perusal of job postings will reveal that organizations are looking for administrators who can handle the seismic changes in IT systems over the past decade. To be successful in this milieu, you need to understand how to work with the new computing paradigms such as cloud based systems, continuous integration and delivery, microservices, modern web application architectures, software based networks, big data, virtualization and containerization.

Old line systems administration isn't obsolete by any means, but as organizations keep moving to the public cloud, there's less need for traditional system administration skills. Today, cloud administrators will take care of the mundane system administration tasks – one need not spend as much time tweaking the Linux kernel as the old Linux ninjas used to, maybe you'll never have to muck with the kernel in most cases, since that's all done by the teams that manage the server farms in the cloud. “The times are a changing”, and it's important to move along with the changing times and learn what makes the modern system administrator tick. For example, if you don't know the Ruby programming language, you're at a huge disadvantage today, since many important tools that are in vogue today (and will be tomorrow as well!) are scripted with Ruby.

As our goal is to understand the main concepts and tools involved in modern system administration, this book doesn't presume to explain the fundamentals of traditional system administration. For new system administrators or users, this book shows you what you need to know after you learn the basics of system administration. For the experienced users, this book shows what you need to know to stay relevant in today's world of system administration.

The speed of innovation and the heightened competition due to the ease of starting up new web based businesses has been the primary catalyst behind many of the changes in system administration. In order to survive and flourish in today's environment, organizations need to make changes incredibly fast. The fast pace means that new software and enhancements to existing software don't have the luxury of time as in the past.

The giants of Web based businesses today, such as Netflix, Facebook and the rest all can make very swift changes to their software as part of their routine operations – each of these companies introduces hundreds and even thousands of changes to their software every single day.

The proliferation of cloud based computing means that most system administrators may never even get to set foot in a data center! Traditional system administration involved ordering hardware, and “racking and stacking” that hardware in a data center. Now, you increase your capacity by issuing an API call or by merely clicking a button on the cloud provider's web page.

Devops is a term that has come increasingly to the forefront in the past few years. Devops is meant to be a way for operational and development teams to work together to speed up the creation and deployment of software. While I don't explicitly address this book to devops professionals, many, if not all of the topics that I discuss in this book are highly relevant to devops professionals. Sysadmins have always worked with devops, through legend has it that the relationship was somewhat acrimonious, with sysadmins being accused of undue zealotry in maintaining their fiefdoms. Today, there's absolutely no chance of such a vast schism between the two groups: either the two groups swim together, or they both sink!

If you're indeed a devops person who's interested in the development side, I'd like to clarify that while I do discuss in depth several new concepts and tools geared towards development, the book is squarely aimed at systems administrators, or those who want to learn about the systems admin side of the equation.

The main thrust of the book is to discuss and explain the key principles and the processes underlying the way modern businesses are architecting and supporting robust highly scalable infrastructures. Sometimes I do show how to install and get going with a tool, but the focus is really on the role the tools play, and how to integrate them into your work as an effective Linux systems administrator.

The best way to benefit from the book is to absorb the main conceptual principles underlying modern systems administration – any tools I discuss in the book are there mostly to illustrate the concepts. Progress is very rapid in this area and new techniques and new tools are introduced all the time. Tools which are popular today for performing certain tasks may be easily supplanted by better tools with short notice. Thus, focusing on the conceptual side of things will help you in the long run, by

showing you how to solve major problems confronted by organizations in developing software and managing web sites, scaling and performance, etc.

Motivation for the new System Administration Strategies and Tools

Several important goals are at the root of most of the newer developments in the software building process and therefore, in systems administration, such as:

- Speed to market (reduce the software cycle time)
- Better software (increase the software quality)
- Cost aware architectures (lower the costs of deploying and maintaining infrastructure)

Modern system administration concepts and strategies that I enumerate throughout this book, such as rapid infrastructure deployment, continuous integration of applications, automated testing and push-button deployments are all a means to achieving these key goals.

Traditional IT practices are conservative and risk averse, often shying away from any changes that might increase the volatility of their operations. Today, Companies like Etsy and Facebook deploy to production numerous times every single day. Netflix has its Chaos Monkey tool that lets it randomly terminate production instances or introduce latency into the application. Having realized that frequent changes bring more benefits than pain, these and other organizations have endeavored in recent years to make their systems antifragile.

Strong testing and QA pipelines and continuous deployment strategies are what enable modern companies to reduce the risk of individual changes. Amazon for example can automatically rollback the software changes in case a deployment doesn't quite pan out. Obviously the company's doing something right, since less than 0.001% of Amazon's deployment result in an outage, although on average it deploys a new change almost every 10 seconds.

Problems with Traditional Systems Administration

Traditional systems administration concepts date back to several decades, predating the advent of major technological innovations such as the internet, cloud computing, newer networking models, and many others. While the guts of administering systems remains effectively the same as always, the job requirements and what management expects from system administrators have slowly but irrevocably changed over the past few years.

Let's review the changes in some areas of traditional systems administration to learn why one ought to change their basic approach to systems administration in many ways compared to the traditional way they administered systems, and how modern tools and techniques are transforming the very nature of the system administrator's role in a modern IT environment.

Monitoring

Modern systems management requires monitoring, just as the traditional systems administration did. However, traditional monitoring was mostly limited to tracking the uptime of servers and whether key services were running on it. Whereas in the past, one focused more on system metrics, today, application or service metrics have come to play an equal or even larger role in ascertaining the health and well-being of systems – keeping the end user happy is today's corporate mantra in all areas of business, including the area of systems administration.

Faster and Frequent Deployments

In today's fast paced web environments, deployments aren't something that are massive and infrequent. It's actually the opposite - deployments are small and quite frequent. Following is a sample of the production environment for Amazon (circa 2014): 11.6 seconds: Mean time between deployments (weekday) 1,079 : Max # of deployments in a single hour 10,000 : Mean # of hosts simultaneously receiving a deployment 30,000 : Max # of hosts simultaneously receiving a deployment

The Image Sprawl problem

While using "golden images" is definitely a superior approach to traditional deployments which are always done from scratch, they do tend to exacerbate the problem of image sprawl. Image sprawl is where multiple images are in deployment, usually in different versions. Images become unwieldy and management becomes chaotic. As the number of images grows, you'll find yourself performing regular manual changes, which tend to lead to deviation from the gold standard. A gold standard in this context refers to a known set of good configuration. Configuration management, the primary objective underlying gold standard usage, becomes hard over time.

Agile Development Methodologies and the System Administrator

Agile Operations is the counterpart to agile development practices, which involve strategies like Kanban and scrum, along with frequent, small code rollouts. The high frequency of code changes means that the operations teams can't be in relative isolation from the development teams, as in the days past. The rigid barriers between the two teams have been gradually coming down due to the high degree of cooperation

and interaction between development and operations that the agile development methodologies require.

Cloud environments

Systems administration practices that work well for a company's data center aren't really usable in toto when you move to a cloud environment. Storage and networking both are fundamentally different in the cloud, especially an external one. System administrators are expected to understand how to work with both public clouds such as AWS, Azure and Rackspace, as well as know how to set up their own private cloud based environments using something like OpenCloud.

Impact of Big Data

Traditional warehouses can't scale beyond a certain point, regardless of how much hardware and processing capacity you throw at them. The advent of the Web and the consequent deluge of data required a different paradigm, and distributed processing turned out to be the best approach to solving problems posed by big data. Hadoop is here to stay as a platform for storing and analyzing big data. Administrators must know how to architect and support Hadoop and other big data environments.

Manual Operations without automation

Traditional systems are to the most extent still run with a heavy dose of manual operations. While script based systems administration has been around for many years, most administrators still perform their duties as they did 30 or even 40 years ago – by hand, one operation after the other. Consequently, change management is slow and there are plenty of opportunities to make mistakes.

New trends in system administration and application development include the following:

- Infrastructure automation (infrastructure as code)
- Automated configuration management
- Virtualization and containerization
- Microservices
- Increasing use of NoSQL and caching databases
- Cloud environments– both external and internal
- Big data and distributed architectures
- Continuous deployment and continuous integration

In the following sections, I briefly define and explain the concepts behind the trends in modern system administration. In the following chapters of the book, I discuss most of these concepts and the associated tools in detail.

Automated Infrastructure Management

Configuring the environment in which applications run is just as important as configuring the application itself. If any required messaging systems aren't configured correctly for example, an application won't work correctly. Configuring the operating system, the networks, database, and web application servers is critically important for an application to function optimally.

The most common way of configuring systems is to do so as you go. That is, you first install the software and manually edit the configuration files and settings until the darn software works correctly. No record is made of the prior state of the configurations as you iterate through successive configuration states. This means that you can't revert easily to the last "good" configuration if any changes go bad.

A fully automated process offers the following benefits:

- It keeps the cost (in terms of effort and delays) of creating new environments very low.
- Automating the infrastructure creation (or rebuilding) process means then when someone playing a critical role in a team leaves, nothing really stops working or stymies your attempts to fix it.
- Automating things imposes an upper bound on the time it takes to get back to a fully functioning state of affairs.
- A fully automated system also helps you easily create test environments on the fly – plus, these development environments will be exact replicas of the latest incarnation of the production environment.
- An upgrade to your current system won't automatically lead to an upheaval. You can upgrade to new versions of systems with no or minimal downtime

You can use a tool such as Chef, Puppet (or Ansible) to automate the configuration of the operating system and other components. As explained earlier, all you need to do is to specify which users should have access to what, and which software ought to be installed. Simply store these definitions in a VCS, from where agents will regularly pull the updated configuration and perform the required infrastructure changes. You gain by not having to manually do anything, plus, since everything is flowing through the VCS, the changes are already well documented, providing an effective audit trail for the system changes.

Automating Redundant Configuration Work

Setting up a new infrastructure or using an existing but unwieldy infrastructure setup isn't a trivial task for new system administrators. While the laying out of the infrastructure itself is pretty straightforward, it usually involves steps that are inherently prone to simple errors. And once you set up an infrastructure with built-in errors, a lot of times you're forced to live with those errors as long as that infrastructure is in place.

Redundant work and duplication of tasks occupies the scarce time of administrators. Manual installation and configuration of infrastructure components such as servers and databases isn't really practical in large scale environments that require you to setup hundreds and even thousands of servers and databases.

Configuration management software grew out of the need to eliminate redundant work and duplicated efforts. The configuration tools help automate infrastructure work. Instead of manually installing and configuring applications and servers, you can simply describe what you want to do in a text-based format. For example, to install an Apache Web Server, you use a configuration file with the declarative statement:

```
All web servers must have Apache installed.
```

Yes, as simple as this statement is, that's all you'd have to specify to ensure that all web servers in a specific environment have the Apache web server installed on them.

Automating infrastructure and application deployment requires more than one simple tool. There's some overlap among the different types of automation tools, and I therefore briefly define the various types of tools here:

- Configuration management tools let you specify the state description for servers and ensure that the servers are configured according to your definition, with the right packages, and all the configuration files correctly created.
- Deployment tools: deployment tools generate binaries for the software that an organization creates, and copies the tested artifacts to the target servers and starts up the requested services to support the applications.
- Orchestration tools: orchestration of deployment usually involves deploying to remote servers where you need to deploy the infrastructure components in a specific order.
- Provisioning tools: provisioning is the setting up of new servers, such as spinning up new virtual machine instances in the Amazon AWS cloud.

Some tools are purely meant for one of the four purposes I listed here, such as deployment for example, where a tool such as Jenkins or Hudson performs purely integration and deployment related functions. Most tools perform more than one

function. A tool such as Ansible for example, can perform all four of these things – configuration management, deployment, orchestration, and provisioning very well.

Configuration Management

Configuration management is how you store, retrieve, identify, and modify all artifacts and relationships pertaining to a project. Modern configuration management practices are an evolution of the traditional system administration strategies to manage complex infrastructures.

Until recently, scripting was the main means of automating administrative tasks, including configuring systems. As infrastructure architectures become ever more complex, the number and complexity of the scripts used to manage these environments also grew complex, leading to more ways for scripted procedures to fail. In 1993, the first modern configuration management (CM) system, CFEngine, was started to provide a way to manage UNIX workstations. In 2005 Puppet was introduced and soon become the leader in the market until Chef was introduced in 2009. Most CM systems share the same basic features:

- Automation of the infrastructure (infrastructure as code)
- Ruby (or a similar scripting language) as the configuration language
- Extensibility, customizability, and the capability to integrate with various other tools
- Use of modular, reusable components
- Use of thick clients and thin servers – the configuration tools perform most of the configuration work on the node which is being configured, rather than on the server that hosts the tools
- Use of declarative statements to describe the desired state of the systems being configured

Good configuration management means that you:

- Can reproduce your environments (OS versions, patch levels, network configuration, software and the deployed applications) as well as the configuration of the environment
- Easily make incremental changes to any individual environment component and deploy those changes to your environment
- Identify any changes made to the environment and be able to track the time when the changes were made, as well as the identity of the persons who made the changes.

- Easily make necessary changes and also obtain the required information, so as to decrease the software cycle time



when you handle server configuration like software, naturally you can take advantage of a source code management system such as Git and Subversion to track all your infrastructure configuration changes

Popular configuration management tools include the following:

- Chef
- Puppet
- Ansible
- Capistrano
- SaltStack

Infrastructure as Code

The term infrastructure as code is synonymous with configuration management. Tools such as Chef transform infrastructure into code. Other CM tools such as Chef, Puppet, Ansible and Saltstack also use the infrastructure as code approach to configure the infrastructure. All these tools are automation platforms that let you configure and manage an infrastructure, both on-premises as well as in the cloud.

Infrastructure as code is really a simple concept that lets your infrastructure reap the benefits of automation by making an infrastructure versionable, repeatable, and easily testable. CM tools let you fully automate your infrastructure developments as well as automatically scale infrastructure, besides handling infrastructure repairs automatically (self-healing capabilities).

Applications typically lock down their configuration and don't allow untested adhoc changes over time. Why should the administrator do any different? You must learn to treat environmental changes as sacrosanct, and work through a structured and formal build, deploy, and test process the same way as developers treat their application code.

While you can't always build expensive test systems that duplicate fancy production systems, you do need to deploy and configure these test environments the same way as you do the production systems.

Let me illustrate the dramatic difference between the old approach and the configuration tools methodology. Following is a Chef recipe that shows how to change permissions on a directory:

```
directory '/var/www/repo' do
  mode '0755'
  owner 'www'
  group 'www'
end
```

If you were to do the same thing using a scripting language, your script would need to incorporate logic for several things, such as checking for the existence of the directory, and confirming that the owner and group information is correct, etc. With Chef, you don't need to add this logic since it knows what to do in each situation.

Site reliability engineering is a term that is increasingly becoming popular in the system administration world. In conventional shops, if a disaster occurs, you need to rebuild everything from scratch. In fact, almost every company maintains an up-to-date disaster recovery plan and even performs dry runs on a regular basis to test their disaster recovery mechanisms. CM helps you quickly restore services following a disaster, since it automates all deployments. Same goes for scheduled upgrades as well, as CM tools can build out the upgraded systems automatically. Bottom line: shorter outages and shorter downtimes for the end users.

Modern Scripting Languages and Databases

Often system administrators in the current milieu wonder what programming languages they ought to learn. A good strategy would be for administrators to become adept at a couple, or even one scripting language, say Python or Ruby. In addition, they should learn new languages and frameworks because in today's world you'll be dealing with numerous open source tools. In order to work efficiently with these tools and adapt them to your environment and even make enhancements, you need to know how these tools are constructed.

There's really no one scripting language that can serve as a do-it-all language. Some languages are better for handling text and data while others maybe be more efficacious at working with cloud vendor APIs with their specialized libraries.

Essential Programming Skills for the System Administrator

Traditionally system administrators used a heavy dose of shell scripting with Bash, Awk and Sed, to perform routine Linux administration operations such as searching for files, removing old log files and managing users, etc. While all the traditional scripting skills continue to be useful, in the modern Linux administration world, a key reason for using a programming language is to create tools. Also, most of the exciting open source tools such as Chef and Puppet are written in languages such as

Ruby, so it's imperative that you learn the modern languages such as Ruby and GO (Golang), if you want to be an effective modern system administrator.



New technologies such as Linux Containers, Docker, Packer and etcd all use GO for their internal tooling.

The Rise of NoSQL Databases

Traditional infrastructures and applications mostly relied (and a lot of them still do) on relational databases such as Oracle, MySQL, PostgreSQL and Microsoft SQL Server for running both online transaction processing systems as well as for data warehouses and data marts.. In large environments, there are usually dedicated database administrators (DBAs) hired expressly to manage these databases (and these databases do require a lot of care and feeding!), so sysadmins usually got into the picture during installation time and when there was some kind of intractable server, storage, or network related performance issue that was beyond the typical DBA's skill set.

Two of the most common relational databases that are used today on the internet are MySQL and PostgreSQL. In addition to these traditional databases, modern system administrators must also be comfortable working with NoSQL databases. NoSQL databases have become common over the past decade, owing to the need for handling document storage and fast clustered reading

Today, it's very common for organizations, especially small and medium sized companies, to expect their system administrators to help manage the NoSQL databases, instead of looking up to a dedicated DBA to manage the databases. For one thing, the new genre of NoSQL databases requires far less expert database skills (such as SQL, relational modeling etc.). Furthermore, companies use a bunch of different databases for specialized purposes and the size of the databases in most organizations doesn't call for a dedicated DBA for each of these databases. So, managing these databases, or at the least, worrying about their uptime and performance is falling more and more in the lap of the system administrator.

Chapter 3 explains the key concepts behind the NoSQL databases such as MongoDB.

Caching

In addition to NoSQL databases, the use of caching has come to occupy a central place. Modern application stacks contain several points where you can use caching to vertically scale the applications.

Setting up remote caching clusters using open source caching solutions such as Memcached or Redis is pretty straightforward but scaling and managing failure recovery scenarios isn't. You can let Amazon Elastic Cache to set up a hosted cached cluster with Memcached or Redis. Microsoft Azure handles your Redis replication and automatic failover setup in a few easy clicks. Chapter 3 explains caching and the use of external distributed caching servers such as Memcached and Redis. See

Content Delivery Networks

Content Delivery Networks (CDNs) are a network of distributed servers that deliver web content to users based on the geographical location of the user and the origin of the web pages. CDNs are efficient in speeding up the delivery of content originated by websites with a heavy traffic and a wide reach. CDNs also protect you against large increases in the web traffic. CDNs use content delivery servers that cache the page content of popular web pages. When users access these pages, the CDN redirects the user to a server that's closest to that user.

Chapter 3 explains the concepts underlying a CDN and how it works.

IT Orchestration

IT Orchestration goes beyond configuration, and enable you to reap the benefits of configuration management at scale. Orchestration generally includes things such as the following:

- Configuration
- Zero downtime rolling updates
- Hot fixes Let's say you're configuring multiple instances of a web server with a CM tool, each of them with a slightly different configuration. You may also want all the servers to be aware of each other following a boot up. Orchestration does this for you. An orchestration framework lets you simultaneously execute configuration management across multiple systems.

You can use a CM tool for orchestration, but a more specialized tool such as Ansible is excellent for IT orchestration. If you were to use just a CM tool, you'd need to add on other tools such as Fabric and Capistrano for orchestration purposes. With a dedicated orchestration tool such as Ansible, you won't need the other tools. Ansible also helps you manage cloud instances without modification. Note that while Ansible is really an orchestration tool, since it permits the description of the desired state of configuration, you can also consider it a CM tool.

Often companies use Ansible along with Puppet or Chef as the CM tool. In this case, Ansible orchestrates the execution of configuration management across multiple sys-

tems. Some teams prefer to use just Ansible for both configuration and orchestration of an infrastructure.

Provisioning with Vagrant

It's quite hard to develop and maintain web applications. The technologies, code and the configuration keep changing all the time, making it hard to keep your configuration consistent on all servers (prod, staging, testing, and dev). Often you learn that a configuration for a web server or a message queue is wrong, only after a painful production screw up. Virtualized development environments are a great solution for keeping things straight. You can set up separate virtualized environments for each of your projects. This way, each project has its own customized web and database servers, and you can set up all the dependencies it needs without jeopardizing other projects.

A virtual environment also makes it possible to develop and test applications on a production-like virtual environment. However, virtual environments aren't a piece of cake to set up and require system administrator skillsets. Vagrant to the rescue!

Vagrant is a popular open source infrastructure provisioning tool. Vagrant lets you easily create virtualization environments and configure them using simple plain text files to store the configuration. Vagrant makes your virtual environments portable and shareable with all team members. With Vagrant, you can easily define and create virtual machines that can run on your own system.

You can generate virtual machines with Oracle's VirtualBox or VMware as the providers. Once you've

```
$ vagrant up
```

Since Vagrant integrates with various hypervisor and cloud providers, you can use it to provision both an on premise virtual infrastructure as well as a full blown cloud infrastructure.

Vagrant and Configuration Management Tools

Earlier, you saw how you can use CM tools such as Chef and Puppet to provision and configure infrastructure components such as web servers and databases. You can use Vagrant with a configuration management tool such as Chef to bring up a complete environment, including a virtual infrastructure. Once you provision the cloud or virtual infrastructure through Vagrant, you can use Chef to deploy and configure the necessary application servers and databases on the virtual (or cloud) servers.

A common use case for Vagrant is the fast creation of a disposable infrastructure and environment for both developers and testers.

Provisioning - Automatic Server Deployment

With the proliferation of server farms and other large scale server deployment patterns, you can rule out old-fashioned manual deployments. PXE tools such as Red Hat Satellite for example, help perform faster deployments, but they still aren't quite adequate, since you need to perform more work after deploying the operating system to make it fully functional.

You need some type of a deployment tool (different from a pure configuration tool such as Puppet or Chef), to help you with the deployment. For virtual servers, VMware offers CloudForms, which helps with server deployment, configuration management (with the help of Puppet and Chef) and server lifecycle management. There are several excellent open source server deployment tools such as the following:

- Cobbler
- Razor
- Crowbar
- Foreman

In Chapter 6, I show how tools such as Razor and Cobbler help with the automatic deployment of servers.

Server (hardware) Virtualization

In the old days, all applications ran directly on an operating system, with each server running a single OS. Application developers and vendors had to create applications separately for each OS platform, with the attendant increase in effort and cost. Hardware virtualization provides a solution to this problem by letting a single server run multiple operating systems or multiple instances of the same OS. This of course lets a single server support multiple virtual machines, each of which appears as a specific operating system, and even as a particular hardware platform.

Most application libraries and server delivery platforms usually were installed at the system level and thus disparate applications would often find conflicts that were resolvable only by using dedicated systems for each application. Virtual machines solve this problem by hosting dedicated operating systems as virtual environments on top of another OS.

Hardware virtualization separates hardware from a single operating system, by allowing multiple OS instances to run simultaneously on the same server. Hardware virtualization simulates physical systems, and is commonly used to increase system density and hike the system utilization factor. Multiple virtual machines can share the

resources of a single physical server, thus making fuller use of the resources you've paid for.

Containerization – the New Virtualization

Virtual machine technology has been around for a good while and folks understand how it works: virtual machines contain a complete OS and run on top of the host OS. You can run multiple virtual machines, each with a different OS, on the same host. Containers share some features with virtualization but aren't the same as virtual machines. Under containerization, both the host and the container share the same kernel, meaning containers are based on the same OS as their host. However, since containers share the kernel with the host, they require fewer system resources.

Container virtualization is newer than virtualization and uses software called a virtualization container that runs on top of the host OS to provide an execution environment for applications. Containers take quite a different approach from that of regular virtualization, which mostly use hypervisors.

The goal of container virtualization isn't to emulate a physical server with the help of the hypervisor. Containers are based on operating system virtualization – all containers share the same OS kernel and their isolation is implemented within that one kernel. All containerized applications share a common OS kernel and this reduces the resource usage since you don't have to run a separate OS for each application that runs inside a container on a host. Processes running inside a container have a small hook into the underlying OS kernel.

Virtualization systems such as those supported by VMware let you run a complete OS kernel and OS on top of a virtualization layer known as a hypervisor. Traditional virtualization provides strong isolation among the virtual machines running on a server, with each hosted kernel in its own memory space and with separate entry points into the host's hardware.

Since containers execute in the same kernel as the host OS and share most of the host OS, their footprint is much smaller than that of hypervisors and guest operating systems under traditional virtualization. Thus, you can run a lot more containers on an OS when compared to the number of hypervisors and guest operating systems on the same OS.

Containers are being adopted at a fast clip since they are seen as a good solution for the problems involved in using normal operating systems, without the inefficiencies introduced by virtualization. New lightweight operating systems such as CoreOS have been designed from the ground up to support the running of containers.



Containers don't have the main advantage provided by hardware virtualization such as the virtualization provided by VMware, which can support disparate operating systems. This means, for example, that you can't run a Windows application inside a Linux container. Containers today are really limited to the Linux operating system only.

Docker as a solution for Image sprawl

Docker is primarily a solution for managing image sprawl. Docker is lightweight and its images are layered, allowing you to easily iterate on the images. While Docker can fix many of the problems inherent in the golden image strategy, it doesn't supplant configuration management tools such as Chef and Puppet. After all, you can use the CM tools to install and configure Docker itself, as well as the Docker containers. Containers require orchestration and deployment support, which the CM tools excel in.

Docker and Containerization

Docker makes it easy to overcome the headaches involved in managing containers. Docker is an application that offers a standard format in which to create and share containers. Docker didn't materialize out of thin air: it extends the contributions made by LXC, cgroups and namespaces to simplify the deployment and use of containers.

As mentioned earlier, it was Google that started developing CGroups for Linux and also the use of containers for provisioning infrastructure. In 2008 the LXC project was started, combining the technology behind CGroups, kernel namespaces and chroot, as a big step forward in the development of containers. However, using LXC to run containers meant a lot of expert knowledge and tedious manual configuration.

It was left to Docker, however, to complete the development of containers to the point where companies started adopting them as part of their normal environment, starting around 2012. Docker brought containers from the shadows of IT to the forefront.

Docker did two main things: it extended the LXC technology and it also wrapped it in user friendly ways. This is how Docker became a practical solution for creating and distributing containers. Docker makes it easy to run containers by providing a user friendly interface. In addition, Docker helps you avoid the reinventing of the wheel, by providing numerous prebuilt public container images you can download and get started.

Docker Container Orchestration and Distributed Schedulers

Managing processes on a single server is easy with the help of the Linux kernel and the init system. Managing and deploying a large number of containers isn't a trivial concern. When you deploy containers on multiple hosts, you not only need to worry about the deployment of the containers, but you also need to concern yourself with the complexities of inter-container communications and the management of the container state (running, stopped, failed, etc.). Figuring out where to start up failed servers or applications and determining the right number of containers to run, all are complex issues.

Message Oriented Middleware (Message Buses)

The prevalence of service oriented architectures means that administrators build systems that connect multiple services, in a loosely coupled format. These systems are distributed by definition, and hence need a way for the individual components to communicate efficiently at high speeds. Messaging buses (also called message brokers or MOM (Message Oriented Middleware), are communication platforms that enable applications and services to communicate in distributed architectures. Streams of messages pass through the bus and are acted upon by the services (worker nodes).

Popular message buses include the following:

- RabbitMQ
- ZeroMQ
- Celery

Containers decouple processes from the servers on which the containers run. While this offers you great freedom in assigning processes to the servers in your network, it also brings in more complexity to your operations, since you now need to:

- Find where a certain process is running right now
- Establish network connections among the containers and assign storage to these containers
- Identify process failures or resource exhaustion on the containers
- Determine where the newly started container processes should run

Container orchestration is the attempt to make container scheduling and management more manageable. Distributed schedulers are how you manage the complexity involved in running Docker at scale. You simply define a set of policies as to how the applications should run and let the scheduler figure out where and how many instances of the app should be run. If a server or app fails, the scheduler takes care of

restarting them. All this means that your network becomes a single host, due to the automatic starting and restarting of the apps and the servers by the distributed scheduler.

The bottom line is to run the application somewhere without concerning yourself with the details of how to get it to run somewhere. Zero downtime deployments are possible by launching new application versions along the current version, and by gradually directing work to the new application.

There are several container orchestration and distributed scheduling tools available, as the following sections explain.

FLEET

Fleet works with the system daemon on the servers in order to perform as a distributed init system. Thus, it's best used in operating systems with systemd, such as CoreOS and others. In addition, you need etcd for coordination services.

KUBERNATES

Kubernetes, an open source tool initiated by Google, is fast becoming a leading container orchestrator tool. Kubernetes helps manage containers across a set of servers. Unlike Fleet, Kubernetes demands far fewer requirements to be satisfied by the OS and therefore you can use it across more types of operating systems than Fleet. Kubernetes contains features that help with deploying of applications, scheduling, updating, scaling, and maintenance. You can define the desired state of applications and use Kubernetes's powerful "auto features" such as auto-placement, auto-restart, and auto-replication to maintain the desired state.

Apache Mesos

Apache Mesos is considered by some as the gold standard for clustered containers. Mesos is much more powerful than Kubernetes but requires you to make more decisions to implement it as well! Mesos has been around even before Docker became popular. Mesos is a framework abstraction and lets you run different frameworks such as Hadoop as well as Docker applications (there are projects in place to let even Kubernetes run as a Mesos framework) on top of the same cluster of servers. Mesosphere's Marathon framework and the Apache Aurora project are the two frequently used Mesos frameworks that support Docker well.

Apache Mesos is a mature technology used by well-known internet companies to support their large scale deployments. Kubernetes and Apache Mesos provide somewhat similar features as what the public cloud vendors themselves offer with their proprietary technology.

Swarm

Docker itself provides a native clustering tool named Swarm, which lets you deploy containers across a large pool. Swarm presents a collection of Docker hosts as a single resource. While Swarm is really lightweight and has fewer capabilities than either Kubernetes or Mesos, it's adequate for many purposes. You can use a single Docker Swarm container to create and coordinate container deployment across a large Docker cluster.

Cluster Management and Cluster Operating Systems

Increasingly, distributed services are using clusters for both redundancy as well as the scaling benefits offered by the cluster based architectures. Clusters provide numerous benefits but also bring their own unique problems. Chief among these is the efficient allocation and scheduling of resources to the various services running in the cluster. Cluster operating systems are the answer to this problem, with Apache Mesos being the most well-known of these types of systems. Here is a short list of the most important cluster operating systems.

- Mesos is a kernel/operating system for distributed clusters. It supports both .war files and Docker containers.
- Marathon is a scheduler that helps schedule and run jobs in a Mesos cluster. It also creates a private Platform –as-a-Service on top of the Mesos framework.
- Fleet is for Docker containers
- YARN (Yet Another Resource Negotiator) is the processing component of Apache Hadoop and together with the storage component of Hadoop - HDFS (Hadoop Distributed File System), forms the foundation of Hadoop 2.

Each framework in a cluster has different computing requirements. Companies must therefore run these different frameworks together and share the data and resources among them. Any cluster resource manager in a cluster should be able to support the goals of isolation, scalability, robustness and extensibility.

Mesos, which became a top-level Apache project in 2013, is becoming increasingly popular as a cluster manager that helps improve resource allocation by enabling the dynamic sharing of a cluster's resources among different frameworks. Twitter and Airbnb are among the production users of Mesos.

Mesos does the same for the data center as what the normal operating system kernel does for a single server. It provides a unified view and easy access to of all the cluster's resources. You can use Mesos as the center piece for your data center applications, and make use of its scalable two-phase scheduler, which avoids the typical problems you experience with a monolithic scheduler. Chapter 10 discusses Mesos in detail.

Distributed jobs run across multiple servers. System administrators, for course, use tools to monitor a set of servers and use configuration tools to perform the server configuration updates etc. You can also use CI (Continuous Integration) tools such as Jenkins to manage some general infrastructure management tasks. But the best way to schedule both ad-hoc and scheduled jobs across multiple servers is to use a remote command execution tool. All the following can help you execute commands across multiple hosts:

- Fabric
- Capistrano
- Mcollective

MCollective for example, is a great tool for parallel job execution that you can use to orchestrate changes across a set of servers in near real-time.

There are simple parallel execution tools based on the Parallel Distributed Shell (pdsh), an open source parallel remote command execution utility, or you can even script one yourself. However, these tools are primitive because they loop through the system in order, leading to the time drift issue. They also can't deal with deviations in responses and make it hard to track fatal error messages. Finally you can't integrate these tools with your CM tools. A tool such as Mcollective overcomes all these drawbacks of the plain vanilla parallel execution tools and allows you execute commands parallelly on thousands of servers belonging to different platforms. Chapter 6 shows how MCollective works.

Version Control Systems

Version control is the recording of changes to files so you can recall specific versions of the files at a later time. Although version control systems are mostly used by developers (and increasingly by system administrators), source code isn't the only thing you can version – the strategy of keeping multiple versions of the same document in case you'll need them later on is applicable to any type of document.

Version control systems offer numerous benefits such as the following:

- Take a project (or some of its files) back to a previous state
- Find out which user made the modifications that are giving your team a headache
- Easily recover from mistakes and the loss of files NOTE: Key practices such as continuous integration and automated deployments depend critically on the usage of a central distributed version control repository. Really, with benefits

such as these and very little overhead, what's there to argue about version control systems?

Application development teams have been using version control systems for a number of years now, to track and maintain application code. However, it's important to version your infrastructure, configurations and databases, along with your application code. You must script all your source artifacts.

Just as the application developers do with their application code, you must have a single source of truth for your systems, including servers, databases and web servers. This results in letting you quickly resolve problems and easily recreate known good states of your infrastructure components. You don't need to waste precious time figuring out which versions of your application code should be paired with which environment.

The most commonly used version control tools today are the following:

- Git
- Perforce
- SVN (subversion)

Of the three tools, Git has become enormously popular, for several reasons.

When you store your configuration definitions in a version control system, you can also set up things so that any configuration changes you commit automatically trigger actions to test and validate changes. This is how continuous integration is triggered for infrastructure changes, so the changes are automatically tested and validated. You therefore integrate a CI or CD tool such as Jenkins, or TeamCity with the CVS. Since configuration definitions are inherently modular, you can run implementation actions on only those parts of the infrastructure code that you've changed.

Version control systems (VCSs) are nothing new – most of us are quite familiar with these systems, variously referred to as source control systems, revision control systems or source code management systems. However, early version control systems such as CVS (Concurrent Version System) have been supplanted over the past several years by far more powerful open source tools such as Git, Mercurial and Subversion.

Chapter 7 discusses version control systems, including Git and GitHub.

Continuous Integration and Continuous Deployment

One of the ways in which system administrators can help developers be more productive is by helping developers spend most of their time writing code and creating or enhancing software, instead of running around fixing problems caused by bad

builds and merges. Broadly speaking, administrators can help developers in the following two ways:

- Provide performance metrics and test results to developers so they can evaluate their work
- Help build environments for developers to work in, ideally with the same configuration as the production systems



The terms rollout, releasing, and deploying are often used synonymously to refer to the process of making changes available to internal or external end users.

In traditional application deployment, there are numerous steps between generating code and the product into the hands of the users (internal or external). These steps include:

- Writing new source code or modifying existing code
- Committing the source code
- Building binaries from the modified code
- Performing quality assurance (QA) testing
- Staging the application
- Deploying the application to production

Not only are there several clearly demarcated steps along this process, with different owners and stakeholders, there's also quite a bit of manual work involved in those processes. This has the rather unsettling effect that while you're fixing some bugs in your code, the tedious manual processes you employ for testing the changes and deploying it may introduce their own bugs into the software! Automation of course, is the way around these potential errors and it also speeds up the entire pipeline. Continuous integration is the broad name given to this automating of the software build and deployment process.



the primary purpose behind a CI or CD system is to catch bugs when they're young, that is, fairly early in the development and testing process.

Benefits Offered by CI

Continuous integration offers several benefits, as summarized here:

- Enhanced reliability of the artifacts that you deploy
- Promotion of automation and reduction in the number of manual processes
- Early identification and remediation of software bugs
- More frequent software updates

Steps involved in CI

Continuous integration involves a sequence of steps that must be performed in order, to get to the final stage of deploying the software. In a nutshell, CI involves the following steps or procedures:

- Developers check in code to a common VCS on a regular basis, usually at least once a day. Everything that an application needs to run – source code, database migration scripts, etc. are checked in.
- CI tools such as Jenkins and TeamCity run automated builds whenever there are changes in the VCS.
- The automated builds include the testing of the code checked in by the developers, such as unit tests, code coverage and functional tests.
- A build is classified as a fully integrated build after it passes all tests.
- The CI server provides the tested artifacts (same as binaries or executables) for download. You can alternatively use an external artifact repository such as Artifactory, Nexus or Nuget.
- The CI server provides visibility through a dashboard into the broken builds and test failures, and automates communications to the team members for actions to fix the problems.
- One of the key requirements of CI is to automate deployments to test environments, where you can run automated functional tests by deploying the previously generated tested artifacts.

Continuous Integration and Continuous Deployment

There's often some confusion between the two similar sounding terms Continuous Integration (CI) and Continuous Delivery (CD). Here's how the two terms differ:

- Continuous integration, also referred sometimes as Continuous Staging, involves the continuous building and the acceptance testing of the new builds. The new

builds aren't automatically deployed into production. They're introduced into production manually, after approval of the acceptance testing process.

- Continuous delivery is where new builds are automatically pushed into production after acceptance testing.

In general, organizations use CI when they first start out along the road of continuous testing and deployment. They need to develop an extremely high degree of confidence in their build strategies and pipelines before moving to a CD approach. In any case, since both CI and CD really use the same set of processes and tools, you often are OK regardless of what you call the deployment approach (CI or CD).

Developers unit test the applications when they build or modify applications. Regardless of whether they use Test Driven Development (TDD) strategies, unit testing help verify small chunks of code at a time with tools such as Junit and RSpec, which speed up the writing of unit tests. A CI tool such as Jenkins, Hudson or Travis Continuous Integration (Travis CI) lets you create efficient build pipelines that consist of the build steps for the individual unit tests. In order to be promoted to the staging phase, the source code, after it's built into a binary package, must pass all the unit tests you define in the test pipeline. The compiled code without errors is called a build artifact and is stored in a directory that you specify.

Continuous Application deployment

Application deployment tools let you automate releases. These tools are at the center of continuous delivery, along with the continuous integration frameworks such as Jenkins. Capistrano is a deployment library that's highly popular as an application deployment tool. You can also consider other application deployment automation tools such as:

- Ansible
- Fabric
- Jenkins

Automating deployment speeds up the flow of software deployment. Many organizations have automated their software deployments in a way that enables them to introduce software changes several times every day. Automated deployment reduces the software cycle times by removing the human error component from deployments. The end result is fast and frequent, high quality deployments. Anyone with the appropriate permissions can deploy software as well as its environment by simply clicking a button.

It's quite common for project members to delay acceptance testing until the end of the development process. Developers may be making frequent changes and even

checking them in, and running automated unit tests. However, in most cases there's no end to end testing of the application until the very end of the application. Even the unit tests are of dubious value at times, since they aren't tested in a production-like environment. Instead, the project managers often schedule an elaborate and time consuming integration testing phase when all the development concludes. Developers merge their branches and try to get the application working so the testers can put the app through the paces of acceptance testing.

What if you can spend just a short few minutes after adding new changes to see if the entire app works? Continuous integration (CI) makes this possible. Every time one of the developers commits a change, the complete application is built and a suite of automated tests is run against the updated complete application. If the change broke the app, the development team needs to fix it right away. The end result is that at any given time, the entire application will function as it's designed.

So, how is the “continuous” part of CI defined? Continuous in this context simply means that every time your team changes something and commits the change to your version control system. A continuous integration tool such as Jenkins integrates with a VCS such as Git, which lets you automatically submit code to Jenkins for compiling and testing executions soon as the developers commit new code to the VCS. Later, the developers and others can check the stages of the automated test and build cycles using the job history provided by Jenkins as well as other console output.

Continuous Integration results in faster delivery of software as your app is in a known functioning state at all times. When a committed change messes the app up, you'll know immediately and can fix it immediately as well, without having to wait for the lengthy integration phase of testing at the very end of development. It's always cheaper in terms of effort and time to catch and fix bugs at an early stage.

Tools for Implementing Continuous Integration

Although CI is really more a set of practices than a specific tool, there are several excellent open source tools such as the following:

- Hudson: has a large set of plugins that allow you to integrate it with most tools in the build and deployment environment.
- CruiseControl: this is a simple, easy to use tool.
- Jenkins: the most popular CI tool out there today with a huge community of users and numerous plugins for almost anything you may want to do with respect to CI/CD.

In addition, you can check out the following commercial CI servers, some of which have free versions for small teams.

- TeamCity (JetBrains): contains numerous out-of-the box features to let you get started easily with CI
- Go (ThoughtWorks): Is an open source tool that descends from one of the earlier CI server named CruiseControl. Delivery Pipelines are the strong feature of Go, and the tools excels at visualization and configuring the pipelines.
- Bamboo (Atlassian)

In chapter 8, which deals with continuous integration, I explain the concepts behind the popular CI tools Jenkins and Hudson.



Checking into the trunk or mainline means that you aren't checking into branches. CI is by definition impossible if you aren't checking into the trunk, since any code you check into a branch isn't integrated with the rest of the existing application code or the changes being made by the other developers.

Applications use both unit testing and acceptance testing on a regular basis to test software changes. Unit tests are for testing the behavior of small portions of the application and don't work their way through the entire system such as the database or the network. Acceptance tests, on the other hand, are serious business: they test the application to verify its full functionality as well as its availability and capacity requirements.

How did companies perform integration before they started adopting the new CI tools? Well, most teams used the nightly build mechanism to compile and integrate the code base every day. Over time, automated testing was added to the nightly build processes. Further along the road came the process of rolling builds, where the builds were run continuously instead of scheduling a nightly batch process. Continuous builds are just what they sound like: as each build completes, the next build starts, using the application version stored in the VCS.

Log Management, Monitoring and Metrics

Traditional log management has involved poring over boring server logs with the intention of divining the root cause of system failures or performance issues. However, the advent of the web means that you now have other types of logs to deal with as well. Companies have lots of information that's trapped by the web logs of users, but most users aren't really ready to mine these humongous troves of data.

Traditional systems administrators might think it's odd to consider logs as the fount of business intelligence, but that's exactly what they've become in today's web based world. Logs today don't mean just old fashioned server and platform logs – they mean so much more. Log management in the contemporary sense mostly refers to the

management of logs of user actions on a company's websites, including clickstreams etc. These logs yield significant insights into user behavior, on which so much of a business's success depends. Thus, it's the business managers and product owners that are really the consumers of the products of the log management tools, not the system administrators.



The term logs is defined quite broadly in the context of log management, and not limited to server and web server logs. From the viewpoint of Logstash, any data with a time stamp is a log.

If you want to assess the stability and performance of your system, you need the system logs. A distributed system can produce humongous amounts of logs. You definitely need to use tools for collecting the logs as well as for visualizing the data.

There are great log management tools out there today that you can use to manage and profit from logs. Although originally tools such as Logstash were meant to aggregate, index and search server log files, they are increasingly being used for more powerful purposes. Logstash combined with Elasticsearch and Kibana (ELK) is a very popular open source solution for log analysis. ELK is increasingly being used as a real time analytics platform. It implements the Collect, Ship+Transform, Store and Display patterns. Here's what the three tools in the ELK stack do:

- ElasticSearch: a real-time distributed search and analytics engine optimized for indexing and based on Lucene, an open source search engine known for its powerful search capabilities
- Logstash: a log collection, analyzing and storing framework, used for shipping and cleaning logs
- Kibana: a dashboard for viewing and analyzing log data

Effective Metrics

Graphite is a popular open source tool for storing metrics. Graphite helps you build dashboards lightening quick for just about any metric that you can think of. Graphite is highly flexible, scalable, and is extremely easy to use.

Graphite lets you perform various types of actions with metrics, by providing a time-stamp and value for the metrics. You can graph the metric and set alarms for it. Graphite offers an API that's quite easy to use. Developers can add their self-service metrics to the dashboards.

Proactive Monitoring

You need to have two entirely different types of monitoring for your systems. The first type of monitoring is more or less an extension of traditional server monitoring. Here, tools such as Nagios and New Relic Server provide you visibility into key areas bearing on system performance, such as capacity, CPU and memory usage, so you can fix problems when they rear their ugly head.

The second and more critical type of monitoring, especially for those organizations who are running complex web applications, are the application performance monitoring tools. Tools such as New Relic APM enable code-level identification and remediation of performance issues.

The goal in both types of monitoring is to let everybody view the available application and server performance data – so they can contribute to better decisions to improve/fix things.

Tools for Monitoring

There are several popular monitoring tools, including Nagios, Ganglia, statsd and Graphite. Let me briefly describe the Nagios and Ganglia tools here – I describe these and tools in detail in Chapter 7.

Nagios is a system and networking monitoring application that helps you monitor web applications, as well as resource utilization such as disk, memory and CPU usage. Nagios is very popular as an open-source monitoring tool. Nagios doesn't include any monitoring scripts when you install it. It's simply a plug-in scheduler and execution program. You turn Nagios into a true monitoring system by using various plug-ins that check and report on various resources. Nagios uses exit codes from the various plug-ins to determine subsequent actions. When a plug-in script exits with a non-zero exit code, Nagios generates and sends alert notifications to you.

Ganglia is a distributed monitoring system with powerful graphical capabilities to display the data it collects. While Nagios shows you if an application or server is failing preconfigured health checks, Ganglia shows you the current and historical picture as to the resource usage and performance of applications and hosts. As with Nagios, Ganglia is highly configurable and lets you add custom metrics to the Ganglia graphs. Its ability to visualize resource usage trends helps with your capacity planning exercises. Nagios and Ganglia work very well together. If Nagios emits alerts about a host or application, Ganglia helps you immediately find out if CPU or resource utilization, and system load have anything to do with the alerts sent by Nagios.

Service Metrics

System administrators regularly track system metrics such as OS and web server performance statistics. However, service metrics are very important too. Among other

things, service metrics reveal how your customers are using your services, and which areas of a service can benefit from enhancements. Apache Hadoop has numerous built-in counters that help you understand how efficient MapReduce code is. Similarly, Codehale's Metric x Library provides counters, timers or gauges to for JVM. It also lets you send the metrics to Graphite or another aggregating and reporting system.

Synthetic Monitoring

While system metrics have their uses, it's hard to figure out how a service or an application will react when it's confronted by an unusual set of circumstances. Organizations sometimes insert fake events into their job queues, called synthetic transactions, to see how the system behaves. Synthetic monitoring often yields more meaningful metrics than low-level traditional metrics. Often, synthetic monitoring is used to measure availability of the applications and real-user monitoring is used to monitor performance. Chapter 16 delves into application performance monitoring and Real User Monitoring (RUM) and synthetic application monitoring.

Cloud Computing

In the past few years, especially over the past 5 years, there has been an explosive increase in the outsourcing of hardware and platform services, through the adoption of cloud computing. The reasons for the headlong rush of companies to cloud based computing aren't hard to figure out. When you move to the cloud, you don't need to spend as much to get a new infrastructure in place and hire and train the teams to support it. Applications such as ERM (Enterprise Risk Management) that are notoriously difficult to implement successfully, can be had at a moment's notice by subscribing to a cloud based application. Enhanced performance, scalability and reliability and speed of implementation are some of the main drivers of the move to cloud-based computing.



Platform-as-a-Service means that vendors can support the entire stack required for running mission critical applications, such as the web servers, databases, load balancers etc. You can monitor and manage the cloud based infrastructure through web based interfaces

Cloud based providers such as Amazon's AWS and the Rackspace Cloud have wowed both system administrators and developers with their ability to provide huge amounts of on-demand computing power at the mere push of a button! Using tools such as knife, you can spin up and spin down server capacity as your demand fluctuates. You can use the configuration tools not only with the public cloud providers but also with private cloud platforms such as VMware's vSphere.

However, for many organizations, especially small to middle sized ones, it makes a lot of sense to simply outsource the entire task of running the infrastructure and building a new data center to a service such as Amazon Web Services (AWS). AWS (as well as Microsoft's Azure and Rackspace Cloud) is a highly optimized and efficient cloud infrastructure and if the cost analysis works for you, it may be the best way to move to a cloud-based environment. After all, if you want to go the route of something such as AWS, all you need is a valid credit card and you could be setting up a cloud environment and running a web application or a big data query in a matter of a few hours!

Let me make sure and point out that cloud computing isn't an unmixed blessing for everybody. Cloud computing comes with its own shortcomings, chief of which are concerns such as the inability of service providers to offer true SLAs (service level agreements), data integration between private and public cloud based data, the inability to audit the provider's systems or applications, the lack of customization, and inherent security issues due to the multi tenancy model of public clouds.

Amazon Web Services (AWS) is a truly amazing set of services that enable you to get on the cloud with minimal effort and pain. AWS includes several exciting components such as auto scaling and elastic load balancing. I discuss AWS in detail in Chapter 10.

Even when you aren't using a cloud computing vendor such as AWS, Microsoft Azure or Google App engine, you're likely to be virtualizing your on premise infrastructure through private clouds using something like OpenStack. OpenStack is hot right now, so let me introduce you to it in the following section.

Open Stack – the Open Source Cloud Platform

OpenStack is a popular open source platform that enables you to build an IaaS cloud (Infrastructure as a Service). OpenStack is designed to meet the needs of both private and public cloud providers and can support small to very large sized environments. OpenStack is free Linux based software that provides an orchestration layer for building a cloud data center. It provides a provisioning portal that allows you to provision and manage servers, storage and networks.

The goal of OpenStack is to provide an open standard, easy to use, highly reliable cloud computing platform for organizations so they can build either a public or a private cloud.

Software Defined Networking

Traditional network administration involves administrators manually configuring and maintaining complex physical network hardware and connectivity. Software defined networking (SDN) is a fast evolving area in networking. SDN aims to apply

the virtualization concepts currently used at the storage and processing areas to networking. By combining virtualization and distributed architectures, SDN aims to remove the need to manage physical network devices on an individual basis. SDN allows administrators to manage network services in an abstract fashion and facilitates the automating of network tasks.

SDN is becoming increasingly popular, and modern cloud computing platforms such as Open Cloud rely on it. Cloud computing regards SDN as an essential and key foundation. OpenFlow and Frenetic/Pyretic are some of the new tools associated with SDN. Chapter 2 discusses SDN and Network Flow Virtualization (NFV) in detail.

Microservices, Service Registration and Service Discovery

Microservices are a fast growing architecture for distributed applications. Microservice architectures involve the breaking up of an application into tiny service apps, which each of the apps performing a specialized task. You build up applications by using microservices as your building blocks.

You can deploy microservices independent of each other since the services are loosely coupled. Each microservice performs a single task that represents a small part of the business capability. Since you can scale the smaller blocks per your requirements, fast-growing sites find it extremely convenient to adopt a microservice based architecture. Microservices help deliver software faster and also adapt to changes using newer technologies.

Right now, many organizations use Node.js, a cross-platform runtime environment for developing server-side web applications, to create the tiny web services that comprise the heart of a microservice based architecture. Following are some of the common tools and frameworks used as part of creating and managing microservices.

- Node.js
- Zookeeper (or etcd or consul)
- Doozer
- Serf
- Skydns/skydock



Some folks use the two-pizza rule to as a guideline for determining if a service really qualifies as a microservice. The Two-Pizza rule states that if you can't feed the team that's building the microservice with two pizzas, the microservice is too big!

Microservices are just what they sound like – they’re small services that are geared towards a narrow functionality. Instead of writing one long monolithic piece of code, in a microservice based approach, you try to create independent services that work together to achieve specific objectives.

If you’re finding that it’s taking a very long time and a large amount of effort to maintain code for an app, the app may very well be ready for breaking up into a set of smaller services. As to how small a microservice ought to be, there’s no hard and fast rule. Some define a microservice in terms of the time it takes to write the application – so, for example, one may say that any service that you can write within two weeks qualifies as a microservice.



A PaaS such as Cloud Foundry is ideal for supporting microservice architectures, since it includes various types of databases, message brokers, data caches and other services that you need to manage.

The key thing is that the services are independent of each other and are completely autonomous, and therefore, changes in the individual services won’t affect the entire application. Each of the services collaborates with other services through an application programming interface (API). Microservices in general have the following features:

- Loosely coupled architecture: microservices are deployable on their own without a need to coordinate the deployment of a service with other microservices
- Bounded Context: any one microservice is unaware of the underlying implementation of the other microservices
- Language Neutrality: all microservices don’t have to be written in the same programming language. You write each microservice in the language that’s best for it. Language neutral APIs such as REST are used for communications among the microservices
- Scalability: you can scale up an application that’s bottlenecked without having to scale the entire application

One may be wondering as to the difference between microservices and a service-oriented architecture (SOA). Both SOA and microservices are service architectures that deal with distributed set of services communicating over the network. The focus of SOA is mostly on reusability and discovery, whereas the microservices focus is on replacing a monolithic application with an agile, and more effective incremental functionally approach. I think microservices are really the next stage up for the principles behind SOA. SOA hasn’t always worked well in practice due to various reasons.

Microservices are a more practical and realistic approach to achieving the same goals as SOA does.

Benefits of Microservices

Following is a brief list of the key benefits offered by microservice based architectures.

- **Heterogeneous Technologies:** you don't have to settle for a common and mediocre technology for an entire system. Microservices let you choose the best of breed approach to technologies for various functionalities provided by an application.
- **Speed:** it's much faster to rewrite or modify a tiny part of the application rather than make sure that the entire application is modified, tested and approved for production deployment.
- **Resilient systems:** you can easily isolate the impact of the failure of small service components.
- **Scaling:** you can scale only some services that actually need scaling and leave the other parts of the system alone.
- **Easy Application Modifications and Upgrades:** with microservices, you can easily replace and remove older services since rewriting new ones is almost trivial.

Both service registration and service discovery play a key role in the management of microservices, so let me briefly explain these two concepts in the following sections.

Service Discovery

Communication among the microservice applications is a major issue when you use the microservice architecture. Server Registration and Service Discovery are the solutions for this issue. As the number of microservices proliferates, both you and the users need to know where to find the services. Service discovery is the way to keep track of microservices so you can monitor them and know where to find them.

There are multiple approaches to discovering services, the simplest strategy being the use of the Domain Name Service (DNS) to identify services. You can simply associate a service with the IP address of the host that runs the service. Of course this means that you'd need to update the DNS entries when you deploy services. You could also use a different DNS server for different environments. However, in a highly dynamic microservice environment, hosts keep changing often, so you'll be stuck with having to frequently update the DNS entries.

The difficulties in handling dynamic service discovery through DNS have led to other approaches, which involve the concept of service registration.

Service Registration

Service registration is the identification and tracking of services by having services register themselves with a central registry that offers a look up service to help you find the services. There are several options for implementing service registration, as I explain in the following sections.

Zookeeper

Zookeeper is a well-known coordination service for distributed environments, and happens to be a critical component of highly available Hadoop driven big data systems. Zookeeper offers a hierarchical namespace where you can store your service information. You could set up monitors and watches so you and the clients are alerted when there are any service related changes.

EtcD

EtcD is a distributed key-value store that provides shared configuration and service discovery for clusters. EtcD is configured to run on each node in a cluster and handles the master election during the loss of the current master. In a CoreOS cluster, for example, application containers running in a cluster read and write data into etcD. The data include things such as the database connection details and cache settings.

Consul

Consul is more sophisticated than Zookeeper when it comes to service discovery capabilities. It also exposes an HTTP interface for discovering and registering services and also provides a DNS server. Consul also has excellent configuration management capabilities.

Chapter 5 discusses service discovery and service registration in more detail, in the context of using Docker containers.

Networking Essentials for a System Administrator

Networking is complex – and my goal is to simplify it a bit so you can be an effective system administrator. Networking has always been a key component of a Linux system administrator’s skill set. However, now more than ever, it has become really one of the, if not the most important skill sets you need to master. The reason is that today, with heavy internet driven workloads and the proliferation of cloud based systems, networking has become the fulcrum of everything. When dealing with a cloud environment such as Amazon’s AWS, you’ll find that most of the complex work involves the setting up of the cloud DNS, internet gateways, virtual private networks and so on.

As a system administrator, one of your key tasks is to ensure that your system is able to accept and send data, both from within your own corporate network as well as from anywhere on the internet. It’s therefore essential that you understand the key components of a network and grasp how data is moved through the network. You must also understand the important protocols that are in use today to move data across networks.

The TCP/IP (Transmission Control Protocol/Internet Protocol) protocols are the two key protocols that stand at the center of modern networking. This chapter builds on basics such as packets and frames and leads up to a thorough explanation of the TCP/IP protocol. In addition, I also talk about several key networking concepts such as the domain naming service (DNS), dynamic host protocol (DHCP), and similar useful networking concepts.

What the internet is

Before we delve into the intricacies of modern networking, it probably is a good idea to discuss the Internet and find out how it works. This gives us the opportunity to define various key concepts of a network, such as the following:

- Packets
- Headers
- Routers
- Protocols

The internet is the largest network in the world, and connects billions of computing devices across the world. These devices are called hosts or end systems. The systems are called hosts since they host applications such as web server or an E-mail server, and they're referred to as end system since they sit at the edge of the Internet.

End systems access the Internet via Internet Service Providers (ISPs). These end systems are connected together via a network of communication links and packets switches. When a host needs to send data to another host, it segments the long messages (such as email messages, a JPEG image or a MP3 audio file) into smaller chunks of data called packets, and adds header bytes to the segments. The resulting packages of data (packets) are then sent through the network to the destination host, where they are reassembled so as to get the original data.

Packets

Network packets are the smallest chunk of data transmitted by a network – each of the packets contains the data we are transmitting across the network as well as some control data that helps the networking software and protocols determine where the packet should go.

A frame is synonymous with a packet in common usage, but it's actually different from a packet – a frame is the space in which the packet goes. Each frame has a preamble and post-amble that tells hardware where the frame begins and ends. The packet is the data contained in a frame.

Each type of network protocol uses a header to describe the data it's transmitting from one host to another. Packets are much small in size compared to the actual data – for example in a packet that's 1500 bytes in size, the headers for TCP, IP and Ethernet together take up only 54 bytes, leaving 1446 bytes for the payload.

When a network card on a host gets a packet, it verifies that it's indeed supposed to accept it by looking up the destination address in the packet header. Next, it generates an interrupt to the operating system. The operating system will then request the

device driver or the NIC to process the packet. The device driver examines the packet to find the protocol it's using, and places the packet where the protocol's software or the stack can find it. Most likely, the packet will be placed in a queue and the stack grabs the packets from the queue in the order they arrived.

As the packet passes through the network layers, the driver strips appropriate headers at each of the layers. For example, IP strips the IP headers and TCP strips the TCP headers. Finally, the packet will be left with just the data (payload) that it needs to deliver to an application.

Packet Switches

A packet switch grabs the packets arriving on its incoming communication links and forwards the packets to its outgoing communication links. The two most commonly used packet switches are routers and link-layer switches. Routers are more common in the network core and the link-layer switches are usually employed in access networks. An access network is one that physically connects an end system to the very first router (also called an edge router) on a path to other end system. Access networks could be a home network, a mobile network, an enterprise network, or a regional, national, or global ISP. Thus, a company's devices are linked by an Ethernet switch which then uses the company's routers to access the company's ISP and through it, to the Internet.

The path travelled by a packet through all the communication links and the packet switches between the source and destination end systems is known a route through the networks.

Applications and APIs

While you can look at the Internet as a set of hardware and software components that use a set of protocols, routers and switches to move around vast amount of packets across the world, what we're really interested in is how the Internet helps applications we all use. Thus, it's probably a good idea to view the Internet as a linked infrastructure that provides various services needed by applications that run on the end systems. Here, by applications, I'm referring to things such as E-Mail, instant messaging, Web surfing, social networks, video streaming, Voice-over-IP (VoIP), distributed gaming, and so on.

Network Protocols

Network protocols are at the heart of the Internet, so let's delve into network protocols a bit. Network protocols are similar in a way to the normal protocols we all follow in a civilized world. Protocols in the networking context are sets of rules and procedures that allow different devices and systems to communicate with each other. In the OSI Reference Model, a protocol is a set of rules that governs communication

among entities at the same layer. TCP/IP is often referred to a protocol – it’s actually a protocol suite, since it’s really a set of protocols.

A network protocol helps hardware or software components of a device such as a computer exchange message and take actions. More formally, we can say the following: a protocol defines the format and order of messages exchanged between two systems, as well as the actions taken during the transmission and the receipt (or non-receipt) of the messages.

A hardware-implemented protocol for example, controls the flow of the actual bits on the links (“writes”) between two network interface cards. Similarly, a software-implemented protocol in a router specifies a packet’s route from the source to the destination system. When a router forwards a packet that comes into its communication link, it forwards that packet to one of its outgoing communication links. How does the router know to which link it should forward the packet? The exact method depends on the type of computer network you’re dealing with.

In the Internet, every end system has an IP address that identifies it. The end system sending a packet to a destination end system includes the destination’s IP address in the packet’s header portion. Each router has a forwarding table that contains mappings of destination addresses to that router’s outbound communication links. When packets arrive, the router examines their addresses and searches its forwarding table, using the packet’s destination addresses to find the correct outbound communication link. The router then forwards the packet to this link.

The Internet uses a set of special routing protocols that are used to automatically set the forwarding tables. The routing protocols use algorithms such as those that employ the shortest path between the router and a destination – they then compute the shortest path and configure the forward tables in that router.



You can use the traceroute program (www.traceroute.org) to trace a route for accessing any host on the internet.

Network Messages and Message Formatting

So, networks represent a set of connected servers that exchange information. How does this information actually flow through the network? Networks send and receive data in the form of structures called messages. When information needs to be sent through a network, it’s sent by breaking the data into these small chunks (messages).

There are several types of network messages, some of which are the following:

- **Packets:** Messages sent by protocols operating at the network layer of the OSI model are called packets, or IP packets, to be more accurate.
- **Datagrams:** this is more or less synonymous with a packet and refers to the network layer as well. Often this term is used for messages sent at higher OSI levels.
- **Frames:** these are messages that are sent at the lower levels of the OSI model, most commonly the data layer. Frames take higher-level packets or datagrams and add the header information needed at the lower levels.

Each of the networking protocols uses messages and each protocol uses a different formatting method to specify the structure of the message it carries. For example a message connecting a web server to a browser will be different from a message that connects Ethernet cards.

While the formatting of the message depends on the specific protocol sending the messages, all messages contain headers and a footer, along with a data element which represents the actual data being transmitted in a message. The OSI model refers to the message handled by a protocol as its PDU (Protocol Data Unit) and the data part as its SDU (Service Data Unit). A lower layer protocol's SDU then is the PDU of a higher-layer protocol.

With this quick introduction to the Internet, let's move on to learn the important elements of modern day networking.

Networking Essentials – Theory and Practice

A network is a group of connected servers that use specialized hardware and software to exchange information. The size of a network can range from really tiny affairs such as a home network, to massive, complex infrastructures, such as the networks managed by Google, for example. The biggest network of course, is something you're all familiar with – the internet itself.

Networks can be classified into classes such as Local Area Networks (LANs), Wide Area Networks (WANs), and Metropolitan Area Networks (MANs). A subnetwork or subnet is a portion of a network - a network that's part of a large network. A segment is a small section of a network and usually smaller than a subnetwork. An internet-work (or Internet) is a large network formed by connecting many smaller networks.

Breaking up the real work into layers – the OSI Model

Networking involves several technologies, which interact in a complex and sophisticated manner to perform their job, which is to move data along the network. To make sense of networking, it's essential that you start at a conceptual level, by going to the fundamental networking model, known as the OSI Reference Model (OSI).

The OSI model provides a basic framework that explains how data moves through a network, by organizing the components of the network into several interlinked networking layers. The OSI model is a great learning tool for understanding networks – and you’ll frequently find references to the various layers of the OSI when dealing with various aspects of networking.

Many of you may have come across the OSI model as part of a Computer Science college course, or during a network training class, and therefore think that it’s something theoretical, with little or no practical use. Networking models such as the OSI seem, let’s face it, boring! What possible benefit could there be to learning about a theoretical model such as this? After all, the OSI model was conceived a very long time ago, when a LAN (local area network) was what a network mostly meant. The model doesn’t adapt very well to even WAN technologies, let alone the internet.



The network layers at various network layers together are called the protocol stack.

Protocol Layering

Networks and network protocols, including the hardware and software that supports the protocols, are organized into multiple layers so as to provide an efficient structure to the network design. Each of the many network protocols belongs to a specific network layer. Each layer offers a specific set of services to the layer above it and consumes services provided by the layer beneath it, and therefore we call this a service model of layering. As an example, a layer named *n* will consume a service such as an unreliable message delivery provided by a layer below it, named *n-1*. In its turn, layer *n* will provide the service of a reliable message delivery, made possible by adding its functionality that detects and retransmits lost messages.

As mentioned earlier, protocol layering makes it simpler to update network components in a specific layer. However, they do carry the drawback that several of the network layers duplicate the functionality of a lower layer.

A networking model such as OSI is a framework that shows the multitude of tasks involved in making a network function, by splitting the work into a series of connected layers. Each of the layers has hardware and software running on it, whose job it is to interact with the hardware and software on other layers.

The OSI model, as mentioned earlier, consists of seven distinct layers, with the lowest layer being where the low level software and hardware are implemented. The layers divide up the work involved in implementing a network. The highest layer, Layer 7, is the application layer and this layer deals with user applications and operating system

software. As you proceed from the bottom most layer (Layer 7) to the upper most layer (Layer 1, application), you move up along increasing levels of abstraction. The lower layers deal with both hardware and software, and by the time you get to the top layers, they're dealing exclusively with software. The lower levels do all the work so the upper level layers, especially Layer 7, can serve application users.

It's easier to understand the layers of the model by grouping the seven layers into just two simple parts:

- Lower layers (Layers 1, 2, 3, and 4): these are the physical, data link, network and transport layers and are responsible for formatting, packaging, addressing, routing, and the transmitting data over the network. These layers aren't concerned with the meaning of the data – they just move it.



There's no hard and fast rule of how you classify the seven OSI layers into higher and lower layers. For example, I chose to place the transport layer into the set of lower layers, since it primarily provides services to the higher layers by moving data. However, this layer is often implemented as software and so some folks may want to group it with the higher layer.

- Upper Layers (Layers 5, 6, and 7): these are the session, presentation and application layers – they concern themselves with software application related issues. The upper layers interact with the users and implement the applications that run over the network. The upper layers aren't concerned with the low level details of how data moves over the network, instead depending on the lower layers to do that for them.

The TCP/IP Protocol Suite

TCP/IP (transaction communications protocol/internet protocol) is a set of networking communication protocols, and is the king of modern networking. It's also the underlying language of the internet, and it's thus important that you understand this protocol suite in some detail. In the TCP/IP protocol suite, the TCP protocol provides network layer functionality and the IP protocol the transport layer functionality. TCP/IP is ubiquitous – if you're using the HTTP protocol, which all of us do when we use the internet, you're using TCP/IP, since the HTTP protocol is part of the TCP/IP suite of protocols. Among the many reasons for its widespread use is the fact that the TCP/IP protocols are highly scalable, and are universally accepted and used across the internet.



There are two IP protocols – IPv4 and IPv6. IPv6 was created because IPv4 was going to run out of IP addresses. Although IPv4 did run out of IP addresses, it's still going to be around. IPv6 isn't really used for most end customers of Internet Service Providers (ISPs).

When you send out a packet it first goes to your default gateway. Routers keep forwarding the packet until it reaching its destination. IP provides simplicity and efficiency to message transmission but it doesn't offer delivery guarantees. Here's where TCP comes in – while IP delivers and routes packets efficiently, TCP guarantees the delivery and the order of packet transmission.

TCP protocols are stateful, meaning that they manage the state of the connection. UDP is a stateless protocol and this doesn't involve any handshaking. UDP is a faster protocol as a result, since there's no mechanism in the form of state to control how data is transmitted, track the data transmission, and retransmit the data if there are any errors during the initial transmission. TCP effortlessly handles all of these aspects of message transmission, as a result of setting up a stateful session following the initial handshake between clients and servers.

In Chapter 3, you'll learn about the wonderful new HTTP/2 protocol that'll eventually replace the current HTTP/1.1 protocol – one of its best features is that it uses a single TCP connection to perform multiple transactions.

The Internet Protocol Stack

The Internet Protocol stack, also called the TCP/IP network protocol stack, uses just four network layers to represent the work done by the top 6 layers of the OCI Model. The seventh layer in the OCI model, the Physical Layer, isn't part of TCP/IP, since the interface between the TCP/IP protocol stack and the underlying network hardware occurs at the Data Link Layer (Layer 2 in the OCI model).

Following is a simple description of the four TCP/IP layers and how they correspond to the OCI layers.

- **Network Interface Layer:** this is the lowest layer in the TCP/IP model and where the TCP/IP protocols from its upper layer interface with the local network. On most TCP/IP networks there's no Network Interface Layer since it's not necessary when you run TCP/IP over Ethernet.
- **Internet (also called network) Layer:** corresponds to the network layer in the OCI model and is responsible for data packaging, delivery, and routing. It's in this layer that the well-known IP protocol, the guts of TCP/IP, works. Routing protocols such as RIP (Routing Information Protocol) and BGP (Border Gateway Protocol)

that determine the routes to be followed by datagrams between source and destination end systems operate in the network layer as well.



The internet and the transport layer are the guts of the TCP/IP network layer based architecture – they contain almost all the key protocols that implement TCP/IP networking.

- **Host-to-Host Transport Layer:** Usually called just the Transport Layer, this layer provides the end-to-end communications over the internet. It allows logical connections to be made between network devices and ensures that transport-layer packets (called segments) are sent and received, and resent when necessary. Key transport protocols are TCP and UDP. NOTE: Real time applications such as video conferences and Internet based phone services are slowed down by TCP's congestion control mechanism and hence often run their applications over UDP rather than TCP.
- **Application Layer:** The application layer is where the application-layer protocols reside. This layer encompasses layers 5-7 of the OSI model and includes well known protocols such as HTTP, FTP, SMTP, DHCP and DNS. Over the past several years, the rise of the internet and the widespread use of the TCP/IP protocol has led to an obscuring of the original seven layers of the OSI model, although the latter is still relevant. TCP/IP protocols have basically supplanted the OSI model.

The Network Layer and TCP/IP Networking

Of all the TCP/IP (or Internet) network layers, the network layer is probably the most complex, so let's spend some time learning how the network layer functions and the services that this layer provides.

The job of the network layer is seemingly simple: it moves packets from the source host to the destination host. The two key functions provided by the network layer for our purposes are forwarding and routing. While forwarding deals with the transferring of a packet from an incoming communication link to an outgoing communications link within the same router, routing involves all routers within a network.

It's the interactions among the routers through various routing protocols that determine the end-to-end path that a network packet follows as it vends its way from a source host to the destination host.

The Internet's network layer is a minimalist network-layer service layer: it doesn't provide a guaranteed delivery service which is one of the many service types a network layer can provide, instead, it provides a best-effort service, where the packets aren't guaranteed to be received in the same order in which they're sent, and the timing between packets isn't guaranteed to be preserved – even the delivery of the transmitted packets isn't guaranteed!



Routing and forwarding are often used as synonyms, but there are crucial differences between these two network layer functions, as explained in this section.

The Internet's network layer consists of three major protocols (or sets of protocols):

- The IP Protocol: You've learned about the IP protocol earlier in this chapter. The IP protocol determines the datagram (a network-layer packet) format, addressing and packet handling conventions.
- Routing protocols: these determine the best path for network packets to traverse. Routing protocols determine the values in a router's forwarding table.
- The Internet Control Message Protocol (ICMP): this is the network layer's error reporting and router signaling protocol. This protocol reports datagram errors and responds to requests for some types of network related information.

The Forwarding Function

When a network packet arrives at a router, its input communication link receives the packet and the router must then move the packet to the correct output link – this is called forwarding the packet. Note that the terms forwarding and switching are often used as synonyms. All routers use forwarding tables.

A router forwards a network packet by using the header value (such as the destination address of the packet) of an arriving packet to look up the corresponding values stored in the forwarding table. The corresponding values in the forwarding table for a specific header show the router's outgoing link interface to which the router must forward the packet.

Routing algorithms determine the values that should be added to the forwarding tables. Routers use the routing protocol messages that they receive to configure their forwarding tables.

A router in this case is a type of a packet switch. A packet switch transfers packets from input link interfaces to output link interfaces. Since routers are network layer (layer 3) devices, they base their forwarding decisions on the values in the fields per-

taining to the network layer. Routers are store-and-forward packet switches that forward network packets using the network layer addresses. Switches, on the other hand, while they're also store-and-forward packet switches, are very different from routers since they forward packets using MAC addresses.

A router has four main components: input ports, output ports, a switching fabric that connects the input/output ports and actually switches (forwards) packets from input ports to output ports and finally, a routing processor. The input/output ports and the switching fabric are usually implemented in hardware. It's these three components that perform the forwarding function, sometimes together referred to as the forwarding data plane. This layer needs to perform very fast, and the speed (in nanoseconds) at which it should function is too fast for a software implementation.

Forwarding and Addressing Using the IP Protocol

Forwarding and addressing in the Internet are determined by the IP protocol in the network layer. Today, there are two IP protocols in use, the most common being IPV4 (IP protocol version 4). IPV6 (IP protocol version 6), a solution introduced to overcome the limited number of IP addresses made available by the IPV4 protocol, will be the IP protocol to which everybody will be moving to over time. However, this isn't going to happen anytime soon, and IPV4 is going to stick around for a while.

An IPV4 address is simple, and is also called a dotted quad number, such as 192.68.1.10, for example. It's called dotted quad because it's made up of four decimal numbers ranging from 0-255, with each number separated by a period.

In what follows, I explain IPv4 and the Internet's address assignment strategy, called Classless Interdomain Routing (CIDR).

CIDR – The IPv4 Addressing Strategy

This section deals with the details of IPv4 network addressing. Typically hosts have a single link into the network and IP uses this link to send its datagrams. The boundary between the physical network link and the host is called an interface. The boundary between a router and its input/output links is also called an interface. Since IP requires each host and router interface to have its own IP address, an IP address then is really associated with an interface than with the actual host server or the router.

An IP address is four bytes (32 bits) long, and has a hierarchical structure. There are a total of about 4 billion possible IP addresses. An IP address looks like the following: 121.6.104.88, where each period separates the four bytes, which are expressed in the dotted- decimal notation from 0-255. Each byte in the address is written in the decimal form and is separated from the other bytes in the IP address by a period (dot). Let's say you've an IP address such as 192.36.224.12. In this address, 192 is the decimal equivalent of the first 8 bits of the address, and so on. An IP address is hierarchical since as you go from the left to the right portions of the address, you get more and

more specific information about the hosts located in the Internet. That is, you'll get a more specific sense of the network where the host with this IP address lives.

Hosts and Networks

Each component of a network is assigned a network address. Usually the network address is 32 bits long. The network address is always split between the network component and the hosts that belong to that network. In the Internet, every host and every router must have a unique IP address. The first several bits of the network address of a host or interface are taken up by the network component (actually, the subnet, as you'll see shortly) and the rest are used to enumerate the hosts in a network.



In the Internet lingo, a subnet is also called an IP network, or just a network.

As I mentioned earlier, the first several bits in the network address are assigned to the network. An octet is 8 bits which means the number before the dot in a dotted decimal IP address notation, as in 192.168.1.42. In this IP address, the first octet is 192 and there are three other octets following this.

Networks are organized by size into various classes. The first octet of the IP address determines the class to which the network address belongs to. Here are the four network classes and the octet ranges they're assigned in a network address:

| Class | Octet Range |
|-------|-------------|
| A | 0-126 |
| B | 128-192.167 |
| C | 192.169-223 |

Following is what you need to remember regarding the address ranges:

- The 127.0.0.1 address is special: it's a loopback address that every host (that uses IP) uses to refer to itself.
- The following are all private IP address blocks that you can't allocate to anyone on the Internet, and are meant to be used for internal networks only.
- All IP addresses in the 10.0.0 network
- The networks ranging from 172.16 to 172.31
- The network 192.168

Subnetting and Netmasks

For performance and management reasons, it's common to create subnetworks under an organization's main network. Subnetting refers to the creation of the subnetwork. If your corporate network is 10.0.0.0, you can create smaller class C networks under it such as 10.1.1.0, 10.1.2.0 and so on.

Each of the subnets will have a 24-bit network component and an 8-bit host component in its 32 bit network address. The first 8 bits are designated for the organization's main network and the last 8 bits are for designating the hosts in the network (the host component of an IP address is set to all zeroes. This makes it easy for you to tell which part of the network address corresponds to the network itself and which addresses correspond to hosts). In this example, the 16 bits allocated to the subnet allow you to create potentially 65,534 sub networks (calculated by raising 2 to the power of 8).

To really grasp network addressing, you must understand binary numbering, which I explain in the following section.

Binary Numbers

In normal counting, you use decimal numbers, where each digit ranges from 0-9. In a number such as 1234, for example, the ones place has 4 as its value. The next number, 3, is in the tens place, the number 2 is a higher power of 10 (100) and the 1 is a yet higher power of 100 (1000). In a binary number, on the other hand, the values can be either 0 or 1. Thus, 1111 is a binary number, which has 1's in all places. You start at the right with a value of 1, and each digit to its left side is a higher power of two. Since binary numbers have just two possible values for each digit, (0 or 1), they're also referred as base 2 numbers.

Here's how binary numbers correspond to decimal numbers – make sure you calculate the value of the binary numbers by computing and adding the values of the digits from right to left.

| | | | | | | | | | | | |
|-------------|---|---|----|----|-----|-----|-----|-----|------|------|------|
| ** Decimal: | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | |
| ** Binary: | 0 | 1 | 10 | 11 | 100 | 101 | 110 | 111 | 1000 | 1001 | 1010 |

To calculate the value of a binary number, you add the values for each binary digit place where you see a 1 – a 0 means there's no corresponding binary value. The binary number 111 evaluates to 7 since it has a value for the 4, 2 and 1 places. Similarly the binary number 1101 evaluates to the decimal number 13 (8+4+0+1).

The Netmask

The netmask serves to let the IP stack know which portion of a network address belongs to the network and which part to the hosts. Depending on this, the stack

learns whether a destination address is local (on the LAN), or it needs to be sent to a router for forwarding outside the LAN.

According to the definition of a netmask, the bits in the network address that are zero belong to the host. Let's say you're looking at the network address 192.168.1.42 with the netmask of 255.255.255.0. For easier understanding, let's show the binary representation of the IP addresses, which are in their dotted decimal form:

| Decimal | Binary | | | |
|---------------|----------|----------|----------|----------|
| 192.168.1.42 | 11000000 | 00101000 | 10000001 | 00101010 |
| 255.255.255.0 | 11111111 | 11111111 | 11111111 | 00000000 |

Given the IP address (192.168.1.42) and its netmask (255.255.255.0), how do you tell which part of the IP address is the network and which part, the host? Remember that the bits that are zero in the netmask belong to the host. Thus, the three octets (255.255.255) are allocated to the network and the last octet (0) represents the host.

Given an IP address and its netmask, you use the bitwise AND operation to figure out the network address. Bitwise AND operations are simple: if both bits in the corresponding place are 1, the result of the AND is a 1. If either bit or both bits evaluate to 0, the result of the AND operations is a zero. If you then perform the bitwise AND operations on the IP address 192.168.1.42 and its netmask of 255.255.255.0 using their binary representation, here's what you get:

| | | | |
|----------|----------|----------|----------|
| 11000000 | 10101000 | 00000001 | 00101010 |
| 11111111 | 11111111 | 11111111 | 00000000 |
| ----- | | | |
| 11000000 | 10101000 | 00000001 | 00000000 |

The resulting bit pattern after performing the bitwise AND operation (11000000 10101000 00000001 00000000) is the same as 192.168.1.0 in the dotted decimal format. It's painful to write out the complete netmask each time you want to refer to it. Luckily, there's a shortcut – you simply write out the network address followed by a slash (/) and the number of bits in the netmask. Thus, the network address 192.168.1.0 with a netmask of 255.255.255.0 will translate to 192.168.1.0/24. This is so since there are 24 bits on the network address portion of the netmask (remember that all bits that are non-zeros belong to the network and the zeros go to the host).

Let's say you're dealing with the Internet and the IP protocol. In this case, if you've three host interfaces and a router interface connecting to them, the four interfaces form a subnet, often simply called a network in the Internet terminology. The address assigned to this subnet, say, is 224.1.1.0/24. You can have additional subnets such as 224.1.2.0/24 and 224.1.3.0/24 as well.

The /24 notation in the addressing scheme is called a subnet mask, which means that the leftmost 24 bits (out of the total of 32 bits) of the network address define the subnet address. In this case, the subnet 224.1.1.0/24 consists of the three host interfaces with the addresses 224.1.1.1, 224.1.1.2, 224.1.1.3 and a router interface with the

address 224.1.1.4. Additional hosts you attach to the subnet 224.1.1.0/24 would all be required to have addresses of the form 224.1.1.xxx.

The Internet address assignment strategy is called Classless Interdomain Routing (CIDR), and is a way to generalize subnet addressing that I described earlier. Under CIDR, a 32-bit IP address has the dotted-decimal form a.b.c.d/x and is divided into two parts, with the x showing the number of bits in the network portion (first part) of the IP address. The x number of leading bits are also called the network prefix (or just prefix) of the IP address.

Routers outside your company's network will consider only the x leading prefix bits (remember that there are a total of 32 bits in an IP address). So, when a router in the Internet forwards a datagram to a host inside your network, it's concerned only with the leading x bits of the IP address, which means that the routers need to maintain a much shorter forwarding table than otherwise. For your company, the router has to maintain a single entry of the form a.b.c.d/x – that's it! Let's say you're dealing with the subnet 224.1.1.0/24. Since x is 24 here, there are 8 bits left in the IP address. These 8 bits are considered when forwarding packets to routers inside your company.

How You Get Your IP Addresses

When your organization asks for IP addresses, it's assigned a block of contiguous IP addresses, which means a set of address with a common prefix and all your company's IP addresses will thus share a common prefix. All addresses are managed by the Internet Corporation for Assigned Names and Numbers (ICANN).

The non-profit ICANN not only allocates the IP addresses but also manages the DNS root servers and assigns domain names. To get a block of IP addresses for one of your organization's subnets, your network admin needs to contact your Internet Service Provider (ISP). Let's say the ISP was assigned the address block 200.23.16.0/20. The ISP will break this block into chunks and assign each chunk of addresses to a different organization. Let's say our ISP wants to break its address block into 8 blocks for assigning to various organizations. This is how the breakup would look like:

| | | | | | |
|-------------|----------------|----------|----------|----------|----------|
| ISP's block | 200.23.16.0/20 | 11001000 | 00010111 | 00010000 | 00000000 |
| Company 1 | 200.23.16.0/23 | 11001000 | 00010111 | 00010000 | 00000000 |
| Company 2 | 200.23.18.0/23 | 11001000 | 00010111 | 00010000 | 00000000 |
| ... | | | | | |
| Company 8 | 200.23.30.0/20 | 11001000 | 00010111 | 00010000 | 00000000 |

Note that I've underlined the subnet portion of each organization's subnet.

Configuring IP Addresses with the Dynamic Host Configuration Protocol (DHCP)

Once your organization is assigned a block of IP addresses, it must assign IP addresses out of that pool of addresses to all hosts and router interfaces in your network. If you've just a handful of servers for which you need to configure IP addresses,

it's easy to do so manually. The sys admin manually configures the IP addresses with a network management tool. However, if you need to configure IP addresses for a large number of servers, you're better off using the Dynamic Host Configuration Protocol (DHCP), which allows hosts to be automatically assigned an IP address).



DNS and DHCP are related in many ways. DHCP is a way to provide a network device with an IP address, and whereas DNS maps IP addresses to hostnames. You can integrate DHCP and DNS, whereby when DHCP hands a network device an IP address, DNS automatically updated to reflect the device name.

There's a client and a server DHCP that you need to configure for using DHCP to handle IP address assignment. You can configure DHCP so it assigns either a temporary IP address which is different each time a host connects to the network, or have it assign the same IP address each time. When you start the DHCP client, it requests the network for an IP address and the DHCP server responds to that request by issuing an IP address to help complete the client's network configuration. NOTE: Since DHCP automates the network related aspects of connecting the hosts, it's called a plug-and-play protocol.

The response that a DHCP server issues consists of the IP address that it assigns to a host, and may also include the address of the local DNS server, the address of the first-hop router (known as the default gateway), and a network mask. Client use this information to configure their local settings. The IP addresses granted by the DHCP server have a lease attached to them. Unless the clients renew the lease on the address, the address goes back into the server's address pool to satisfy requests from other clients.



Internet IP addresses are network-layer addresses. Link layer addresses, are also called physical addresses, or MAC addresses. The Address Resolution Protocol (ARP) translates between these two types of addresses, IP addresses are in the dotted decimal notation, and MAC addresses are shown in the hexadecimal notation

Network Address Translation (NAT)

Network Address Translation (NAT) is a way to map your entire network to a single IP address. You use NAT when the number of IP addresses assigned to you is fewer than the total number of hosts that you want to provide Internet access for. Let's say you've exhausted the entire block Of IP addresses assigned to you by your ISP. You can't get a larger block of IP addresses allocated, because your ISP has no contiguous portions of your current address range available. NAT comes to your rescue in situations like this.

Using NAT means that you take advantage of the reserved IP address blocks and set up your internal network to use one or more of these network blocks. The reserved network blocks are:

| | |
|----------------|---------------------------------|
| 10.0.0.0/8 | (10.0.0.0 - 10.255.255.255) |
| 172.16.0.0/12 | (172.16.0.0 - 172.31.255.255) |
| 192.168.0.0/16 | (192.168.0.0 - 192.168.255.255) |



Address spaces such as 10.0.0./8 (and the other two listed here) are portions of the IP address space reserved for private networks, or a realm with private addresses, such as home network, for example. A realm with private addresses on a network whose addresses have relevance only to devices living on that network

A system requires a minimum of two network interfaces to perform NAT translation. One of these interfaces will be to the Internet and the other to your internal network. NAT translates all requests from the internal network so they appear to be coming from your NAT system.

How NAT Works

NAT uses IP address and port forwarding to perform its job. When a client in your network contacts a host on the Internet, it is required to send IP packets for that host. These IP packets contain the address information required to enable the packet to reach their destination hosts. NAT uses the following bits of information from this information set:

- The IP address of the source (ex. 192.184.1.36)
- The TCP (o UDP) port of the source (2448)

When the packets from your internal host pass through the NAT gateway, NAT modifies the packets so they appear as if they were sent by the NAT gateway, and not the clients. In our case, NAT will make the following changes:

- Replaces the IP address of the source with the external IP address of the gateway (ex. 198.52.100.1)
- Replaces the source port with a random unused port on the gateway (ex.52244)

The NAT gateway records the address/port changes in a “state” table so it can reverse the changes it made for packets being returned to the internal hosts from the Internet. To the Internet, the NAT-enabled router looks as a single device with a single IP address. All Internet traffic leaving your organization’s router to the Internet will share the same source IP address of 192.52.100.1 and all traffic entering that router

from the Internet will have the same destination address of 192.52.100.1. So, NAT in effect hides your internal network from the rest of the world!

It's important to remember that neither the internal hosts, nor the Internet host are aware that NAT is performing this "behind the scenes" address translation – the internal host thinks the NAT system is just an Internet gateway and the Internet host thinks the packets come directly from the NAT server. So, the Internet host sends its replies to the packets to the NAT gateway's external IP (198.52.100.1) at the translation port 52244.

The NAT gateway receives a large number of these replies destined for various internal hosts. How then does it know to which host it ought to send a specific reply packet? That's where the "state" table comes in. The state table is formally called a NAT translation table, and includes the IP addresses and port numbers of the internal hosts as entries in the table. NAT looks up the state table to see which established connection matches the reply packets. When it finds a unique IP/port combination match, telling it that the packets belong to a connection initiated by the internal host 192.184.1.36, it reverses the changes it made originally for the outgoing network packets, and will forward the reply packets coming from the Internet host to the correct host.

The New IP Protocol – Ipv6

About 20 years ago, efforts began to develop a new IP protocol to replace the IPv4 protocol, since the 32-bit IP address space (maximum of 4 billion IP addresses) was starting to get used up fast. Why is this a big deal? If there aren't any IPv4 addresses to allocate, new networks can't join the internet! The new IP protocol IPv6 that'll eventually replace IPv4 is designed with a much larger IP address space, as well as several other operational improvements compared to IPv4.

Besides containing a larger IP address space, the IPv6 datagram has several changes, including the following:

- **Expanded addressing capability:** the size of the IP address is increased from 32 to 128 bits to ensure that we never run out of IP addresses. Earlier, you had only unicast and multicast addresses. IPv6 adds the anycast address, which lets a datagram be delivered to any member of a group of hosts.
- **Streamline Headers:** Several IPv4 header fields have been removed or made optional, making the streamlined shorter header field (40 bytes) allow much faster processing of an IP datagram.
- **Handling fragmentation and reassembly:** IPv6 removes the time consuming operations of datagram fragmentation and reassembly from the routers and places it on the end systems (source and destination), thus speeding up IP for-

warding inside the network. Other costly operations in IPv4 such as performing header checksums have been eliminated as well.

While new IPv6 capable systems are backward compatible and can send and receive IPv4 datagrams, currently deployed IPv4 systems can't handle IPv6 datagrams. The big question right now is how to migrate all the current IPv4 based systems to IPv6. Ipv6 adoption is growing fast, but it's still going to take a while to switch completely over to IPv6 from IPv4. Changing a network layer protocol such as IP is much harder than changing protocols at the application layer such as HTTP and P2P file sharing, for example.

The IPsec Protocol

IPv4, as well as its replacement, IPv6, weren't designed with security in mind. IPsec is a popular secure network layer protocol that's backward compatible with both IPv4 and IPv6. The IPsec protocol is commonly deployed in Virtual Private Networks (VPNs).

You can use IPsec to perform secure private communications in the unsecure public Internet. Many organizations use IPsec to create virtual private networks (VPNs) that run over the basically insecure public Internet.

Routing Essentials and Routing Management

It's quite easy for two hosts within the same LAN to talk to each other – each host needs to send an ARP (address resolution protocol) message and get the MAC address of the other host and that's it. If the second host is part of a different LAN, then you need the two LANs to talk to each other, and you use a router to make this possible.

The router is a network device that sits between two networks and redirects the network packets to the correct destination. If a host is connected to a network with multiple subnets, it needs a router or a gateway to communicate with the hosts on the network. While hosts know the destination of a packet, they don't know the best path to the destination. Either the router knows where the packet ought to go, or it knows of another router than can make that determination.

The router's job is to know the topology of networks plugged into it. When you want to communicate with a server from a different LAN, you don't talk to the other LAN directly. You do specify the destination IP of the host you want to send data to, but also specify the destination MAC address for the router. So the router will get your packets first and check the destination IP and send it to the other LAN, since it knows that the destination IP isn't part of your LAN. And it does the opposite for packets coming from distant LANs to your LAN.

Earlier, I explained how the network layer consists of two major functions: forwarding and routing. When a router receives a packet, it indexes a forwarding table and finds the link interface to which it must send the packet. Routers are where packet forwarding decisions are made in a network. Routing algorithms are used by the router to configure the forwarding tables. Routing is the determining if the best paths (route) from a sender to the receiver, through a network consisting of other routers.

Routing Algorithms

Most hosts are connected to a single router, which is known as the default router (or first-hop router) of a host. The default router of the host that sends a packet is called the source router, and the default router on the receiving host is called the destination router. The purpose of routing algorithms is to find the most efficient path from a source router to the destination router. It's routing algorithms that are at the heart of routing in modern networks. While you might think that the algorithms with lowest cost path is the best, it isn't always so, as you'll learn soon.

Routing algorithms use graphs to formulate a routing problem. The goal of a routing algorithm is to find the least-cost path between a source and a destination. If all links in a graph have the same financial costs and link speeds, then the least-cost path is the same as the shortest path.

There are numerous routing algorithms out there, and you can classify the algorithms according to several criteria:

- Global versus decentralized routing algorithms: Global routing algorithms compute least-cost paths between routers using complete knowledge about the network, such as information pertaining to all the node connectivity between the nodes and all the link costs. Since each node maintains a vector of estimates of the costs of traversing to all other nodes, the decentralized routing algorithm is called a distance-vector (DV) routing algorithm. The link-state and distance-vector algorithms are the two major routing algorithms used in today's Internet.



In the LS algorithm, each node in a network broadcasts to all the other nodes, letting them only the cost of only those links that are directly connected to it. The DV algorithm, on the other hand, means that each node speaks to just its directly connected neighbors, to whom it provides least-cost estimates to traverse from itself to all the nodes it knows about in the network.

- Static versus dynamic algorithms: a static routing algorithm means the routes change mostly because you manually edit a router's forwarding table. A dynamic

routing algorithm changes routing paths periodically, or due to a change in the link costs or a change in the network topology.



Currently used Internet routing algorithms such as RIP, OSPF, and BGP are load-insensitive,

- Load-sensitive versus load-insensitive routing algorithms: In a load-sensitive routing algorithm, the current level of congestion in a link is taken into account - congested links may have a high cost and the algorithm will try to find routes around these congested links. A load-insensitive algorithm doesn't with the current (or past) level of congestion in a link as part of its computation of the link's cost.

Hierarchical Routing

Routers today all use either the LS or the DV algorithm to find the best path for routing the datagrams. With millions of routers in the Internet, the overhead of computing, storing and transmitting routing information among these routers quickly becomes very expensive. There's a different issue as well with maintaining a bunch of routers, all independent from each other - organizations want to administer their routers and their network as they wish, without exposing their network architecture to the entire world. You thus need way to organize routers into a hierarchy to lower the cost of managing the routing information, while allowing companies to maintain autonomous groups of routers (belonging to just their own network).

An administrative system (AS) is a way to organize routers so a group of routers is under the admin-



An intra-AS routing protocol determines routing within an AS and an inter-AS routing protocol determines routing behavior between two ASs.

All routers within an AS run an identical routing algorithm such as the LS algorithm for example, and have complete information about all the other routers in that AS. The algorithm running in an AS is called an intra-autonomous system routing protocol. The routers that belong to a specific AS use the AS's intra-AS routing protocol to forward packets along an optimal path to other routers within the same AS. Note that different ASs may run different intra-AS routing protocols (LS or DV). The concept of an administrative system (AS) of groups of related routers under the same admin-

istrative control and using the same routing protocol among themselves is a crucial feature of how routing works in the Internet today.

One or more routers in each AS act as a gateway router to connect an AS to another AS. If the AS has but a single link that connects it to other ASs, the single gateway router receives the packets from each router and forwards the packet to other AS. However, if there are multiple links and hence multiple gateway routers, each router needs to configure its forwarding table to manage the external AS destinations. The inter-AS routing protocol helps get information from neighboring ASs that they're "reachable", and send this information to all routers internal to an AS. An inter-AS routing protocol is for communications between two ASs, and therefore both ASs must run the same protocol – most commonly, the BGP4 protocol is used for this purpose.

Now that you know what an autonomous system of routers is, IT'S time to learn about the key Routing protocols of which there are three today: RIP, OSPF and BGP.

Common Routing Protocols

As I'd explained earlier, routing protocols determine the path taken by IP datagrams between a source and a destination host. Let me quickly summarize the three important routing protocols that are used in the Internet – in the following discussion, the RIP and the OSPF routing protocols are intra-AS routing protocols, also called interior gateway protocols:

- **Routing Information Protocol (RIP):** This is the one of the oldest routing protocols, and is used extensively in the Internet, typically in lower layer ISPs and enterprise networks. RIP is a distance-vector protocol that's quite similar to the DIV protocol described earlier. A hop in the RIP terminology refers to the number of subnets traversed along the shortest-path. RIP is limited to the use of autonomous systems that are shorter than 15 hops in diameter. Under RIP each router maintains an RIP table, also called a routing table. The routing table contains the router's estimates of its distance from various other routers, as well as the router's forwarding table. In Linux systems, RIP is implemented as an application layer process, and the process named `routed` ("route dee") executes RIP and maintains the routing information and works with other routed process running on nearby routers.
- **Open Shortest Path First (OSPF):** OSPF was designed to replace RIP and contains several advanced features such as enhanced security through authentication, and support for a hierarchical autonomous system within a single routing domain. OSPF is a link-state protocol that uses a least-cost path algorithm and also the flooding of link-state information,
- **Border Gateway Protocol (BGP):** Unlike RIP and OSPF, BGP is an inter-AS routing protocol. It's also known as BGP4. BGP provides an AS a way to obtain

reachability information regarding neighboring ASs and determine the best route to subnets based on reachability and AS policies. In this protocol, routers exchange information over semi-permanent TCP connections.

Static Routing

How does a router know about the networks that are plugged into it? A routing table contains the network information. You can manually edit the routing table, and the routing information you configure here will stay fixed until someone edits the table manually – this routing table is consequently called a static routing table.

A routing table consists of three things:

- Network addresses
- Netmask
- Destination interface

When a network packet gets to the router with a routing table such as this, it'll apply each of the netmasks in the routing table to each destination IP address. If the network address it thus computes matches any of the network addresses in the routing table, the router will forward the packet to that network interface.

For example, let's say the router gets a packet with the destination address of 192.168.2.233. The router finds that the first entry in the routing table doesn't match this destination address. However, the 2nd entry does, as the network address that you get when you apply the netmask 255.255.255.0 to the destination address 192.168.2.233 is 192.168.2.0. The router forwards the packet to Interface 2. If a packet's destination address doesn't match any of our three routes, the packet will automatically match the default address and will be sent to the destination interface for that address - usually this is a gateway to the internet.

A Linux host is aware of multiple standard routes, such as the following:

- The loopback route, which points to the loopback device
- A route to the local area network (LAN), which is used to direct network packets directly to the hosts in the same LAN
- The default route, which is used for packets that are to be sent to destinations outside the LAN.

You use the traditional route command to manage routes. The route command modifies the routing table. Therefore, it's a good idea to first copy the current routing table before issuing the route command! The following command shows how to set the default route for a host with a default gateway at 192.168.1.2:

```
# route add default gw 192.168.1.2 eth1
```

This command adds a default route and sets the gateway to 192.168.1.2. It also sets the interface to this connection through eth1 (a network adapter). Once you issue this route command, all incoming traffic (other than any that were stipulated through other routes) will be sent via your web server, which is 192.168.1.2. Any routes that you add won't survive a system restart, so if you want to make the new route permanent, add the route command with the word up preceding it, to the /etc/network/interfaces file, as shown here:

```
# up route add default gw 192.168.1.2 eth0
```

In order to see if your changes stuck, you don't need to restart the server – just restart the network, as shown here:

```
# /etc/init.d/networking restart
```

On the other hand, if you changed your mind about the new route, you can remove it with the route del command as shown here:

```
# route del default gw 192.168.2 eth0
```

Using the newer ip command instead, you'd add a default route to a host as follows:

```
-# ip route add default via 192.168.1.1
```

You can view the routing table using one of the following commands:

- netstat
- route
- ip route

Dynamic Routing

Static routing is OK when dealing with small networks. As the network size grows and you need to deal with a large number of subnets, the overhead involved in manually managing the routing tables becomes prohibitive. Dynamic routing is the solution for managing routing in large networks. Under dynamic routing, each router needs to know only about the networks close to it. It then transmits information about these nearby networks to all the other routers connected to it.

Dynamic routing can be based on one of the three routing protocols you learned about earlier in this chapter: the Routing Information Protocol (RIP), the Open Shortest Path First (OSPF) protocol, and the Border Gateway Protocol (BGP). RIP is simple and is fine for small networks due to its simplicity. This protocol chooses the best route based on the fewest number of routers it needs to traverse – also called the number of hops.

The problem with RIP is that it doesn't take into account network congestion on the routes. OSPF, on the other hand, bases its decision not on the number of hops but how quickly a router can communicate with another router. In doing this, it automatically takes into account any congestion along the route. OSPF is more complex and expensive than RIP and BGP, and is suitable for larger networks.

Network Performance – Bandwidth and Latency

We usually use terms such as “network speed” and “how fast” the network is, when talking about network efficiency. The speed or “fastness” of a network can be really better understood through examining two related but different concepts – bandwidth and latency. Bandwidth is the amount of data a connection can move within a specific amount of time, and is measured in bits per second. Latency is the time it takes for a request to receive the response it's seeking. Latency and bandwidth are related, but they're quite different things – one refers to the capacity of the network and the other, to its speed – obviously, the interaction between the two determines the ‘speed’ or “fastness” of the network.

The Hypertext Transfer Protocol (HTTP)

Of all the protocols that are supported by the application layer of the network, HTTP is probably the most important, in light of its importance in our Internet dominated world. Well, I do understand that everyone knows what HTTP is! However, I start from the basics so I can explain the critical role played by it in how web applications work. This background will be very helpful during our discussion of load balancing web applications in the next chapter, as well as in the discussion of performance tuning of web applications in Chapter 16.

The Web is the most famous and most popular Internet application. The HyperText Transfer Protocol (HTTP) is the Web's application-layer protocol, and is the foundation of the Web. HTTP is implemented in a pair of client and server programs that communicate through HTTP messages. The HTTP protocol defines the message structure and how exactly the client and the server exchange the messages.

HTTP is a request-response protocol that enables clients and services to communicate with each other. The client can be a web browser and the server can be an application that runs on the server hosting a website. When a client submits a HTTP request to a server, the server returns a response that contains the request status, and most likely, the requested data as well.

Using the HTTP/2 Protocol for Enhanced Performance

HTTP/2, the evolving alternative to the current HTTP protocol (HTTP/1.1) lets you create simpler, faster, and more robust web applications. If you want to provide the same features through HTTP/1.1, you'll need to depend on multiple workarounds, all of which leads to a much more complex and less robust application.

HTTP/2 offers several new optimizations without altering the application semantics of HTTP. Familiar concepts such as HTTP methods, URIs, HTTP headers and status codes remain the same. You can use HTTP/2 instead of HTTP/1.1 in your current applications without having to modify the applications. The most important changes introduced in HTTP/2 are in the framing (within the network transport layer - please see Chapter 2 for a detailed explanation) layer. HTTP/2 provides its optimizations through changing the way data is framed (formatted) and transported between the requester and the responder.

To really cut through the thicket, what HTTP/2 does is to break up the traditional HTTP protocol communications into smaller binary encoded frames. These frames are then mapped to messages that belong to specific streams, which can be multiplexed within a single TCP connection by interleaving the independent frames and reassembling them at the other end. This ability to breakdown/interleave/reassemble messages is the biggest optimization offered by HTTP/2.

While a developer can still deliver a working application with traditional HTTP, the newer protocol HTTP/2 offers vastly superior performance. HTTP/2 reduces latency by the following techniques:

- Using compression to reduce the protocol overhead: In HTTP/1.1, the HTTP headers that contain metadata about the resources being transferred are sent as plain text and add significant overhead (about 500-800 KB on average). HTTP/2 reduces the overhead and improves performance by compressing the request/response metadata using the efficient JHPACK compression format.
- Enabling full request and response multiplexing: In HTTP/1.1, if clients want to improve performance by making parallel requests, they must use multiple TCP connections. HTTP/2 removes this limitation by allowing full request and response multiplexing. All HTTP/2 connections are persistent and just a single connection is required per origin. Overhead is reduced significantly compared to HTTP/1.1, because an overwhelming majority of active connections carry just a single HTTP transaction.
- Facilitating response prioritization: HTTP/2 breaks up the HTTP messages into individual frames and allows each stream to have a different weight. Using these weights, clients can construct and communicate a “prioritization tree” that shows its preferences for receiving the response to its web requests. Servers use the pri-

oritization information to prioritize the processing of streams through appropriate resource allocation.

- Facilitating server push: HTTP/2 lets a server push additional resources to clients on top of the response to the original request. Server initiated push workflows enhance performance since the server uses its knowledge of the resources required by the client and sends them out to the client instead of waiting for the client to request those resources.

Switching to HTTP2

Before the whole world can switch to HTTP/2 and partake of its superior performance capabilities, millions of servers need to be updated to use binary framing. All the clients that access these web servers must also update their browsers and networking libraries. While all modern browsers have enabled HTTP/2 support with minimal intervention for most users, server upgrades to HTTP/2 aren't really trivial in terms of the time and effort necessary to convert them over to HTTP/2.

As of October 2015, although only 5% of all SSL certificates in Netcraft's SSL survey were on web servers that supported SPDY or HTTP/2, 29% of the most SSL sites in within the top 1000 most popular sites do support SPDY/HTML/2. Also, 8% of the top million sites do support the new protocol. This makes sense since the busiest web sites gain the most by optimizing connections with HTTP/2. Widespread use of HTTP/2 is still several years away since browser vendors currently support HTTP/2 only over encrypted TLS (Transport Layer Security) connections. This means that a large proportion of non-HTTPS sites will continue using HTTP 1.1 well into (many years) the future.

Network Load Balancing

Load balancing is the distribution of a system's load over more than one system, so it can be handled concurrently and faster than if just one system handled all the work. Database servers, web servers and others use load balancing architectures all the time, and so do networks.

A load balancer is either software or a piece of hardware that can distribute traffic arriving at an IP address over multiple servers. Load balancers strive to spread the load evenly among the servers and allow you to dynamically add and remove the servers. The servers stay hidden behind the load balancer and since clients can see just the load balancer, you can add web servers anytime, without disrupting your services.

Benefits of Using a Network Load Balancer

A network load balancer offers the following benefits.

- **Security:** Routing all traffic between the users and the web servers through a load balancer hides the data center from the users.
- **Transparent server maintenance:** You can take servers out the load balancer pool without affecting clients (after making sure the active connections are completed). This helps you perform rolling upgrades of your infrastructure without incurring expensive downtime.
- **Easily add servers to the pool:** when you add new servers to the pool, the load balancer immediately starts sending it connections without any delay, as is the case in DNS load balancing. When we discuss auto-scaling in the context of Amazon or Open Stack in Chapter 9, you'll see how a load balancer can let you scale automatically without any adverse impact on the users.
- **Easily manage server failures:** you can easily yank web servers with performance issues out of the load balancer pool, to keep connections from being sent out to that server.
- **Reduce web server resource usage:** if your compliance requirements allow it, you can use Secure Sockets Layer (SSL) offloading (also called SSL termination), to lower the web server resource usage. By doing this, you let the load balancer handle all SSL encryption/decryption. This also means that you can avoid running the SSL based web servers, and have all the web servers processing just HTTP connections and not HTTPS connections, which arrive over SSL.

Load Balancing with DNS

Domain Name Service (DNS) based load balancing (also called the poor person's load balancing) is very easy to set up – all you need to do is provide multiple IP addresses for a domain. That is, you need to add multiple A records for the same domain. When a client attempts a connection to the domain, DNS will hand the client a different IP address sequentially, in a round robin fashion.

Let me illustrate how to set up DNS load balancing with a simple example that uses three IP addresses.

| | | | |
|-------------|----|---|-----------------|
| example.com | IN | A | 126.126.126.130 |
| example.com | IN | A | 126.126.126.131 |
| example.com | IN | A | 126.126.126.132 |

Or you can satisfy the domain name just once:

| | | | |
|-------------|----|---|-----------------|
| example.com | IN | A | 126.126.126.130 |
| | IN | A | 126.126.126.131 |
| | IN | A | 126.126.126.132 |

You can test your DNS load balancing configuration by issuing the `nslookup` command:

```
# nslookup example.com 127.0.0.1
Server:          127.0.0.1
Address:         127.0.0.1#53
Name:            example.comAddress:          126.126.126.130
Name:            example.com
Address:         126.126.126.131
Name:            example.com
Address:         126.126.126.132
```

DNS load balancing is easy to configure, simple to understand, and very easy to debug as well. However, there are some serious drawbacks as well to DNS based load balancing. One or more of the hosts could end up with a lopsided distribution of the load. The load balancer also keeps sending connection requests to a server that's no longer up.

DNS load balancing was more common a while ago, when load balancers were much more expensive than now. While DNS load balancing is good for a quick spin to learn about load balancing, you need a much more mature and robust load balancing mechanism in a real world production system with heavy internet traffic.

Enterprise Load Balancers

When a client resolves a domain name to an IP address through DNS, the client needs to connect to that IP address to request a web page (or a web service endpoint). To enhance the scalability of your applications, it's a good idea to let a load balancer act as the entry point to your data center, by letting users connect to the load balancer directly instead of to your data center. This will help you scale up easily and be in a position to make infrastructure changes that are transparent to your customers.

DNS load balancing, while simple and cheap (especially in the past when commercial load balancers were much more expensive), isn't really load balancing in the real sense, although it does distribute traffic among a set of servers. DNS has several issues, as summarized here:

- DNS isn't transparent to users. If you remove a server, it may lead to problems because the users may have cached its IP address
- Even when you add a new server, users may not use it since they may have cached your DNS records and so, will keep connecting to the old server(s)
- DNS makes it difficult to recover from failures

For all these reasons, DNS isn't a really viable proposition for production environments.

There are many types of load balancers from which you can choose, with the following three types being the most popular options:

- Software based load balancing
- Hardware based load balancing
- Using a Hosted Server for load balancing

In the following sections, I describe the three load balancing options.



Chapter 9 discusses AWS and ELB in detail.

Software Based Load Balancing

If you don't want to use a load balancer service such as ELB, you can use one of the many open-source load balancers that are available. Open source load balancers are software based. Two of the most popular software-based load balancing options today are HAProxy, a TCP/HTTP load balancer, and the NGINX web server. In the following sections, I explain these two options.

Using HAProxy for Load Balancing

Since HAProxy is session-aware, you can use it with web applications that use sessions, such as forums, for example. High Availability Proxy (HAProxy) is a pure load balancer, and you configure it in two ways:

- You can configure HAProxy as a layer 4 (see chapter 2 for an explanation of network layers) load balancer. In this case, HAProxy uses just the TCP/IP headers to distribute traffic across the servers in the load balancer pool. This means that HAProxy can distribute traffic for all protocols, and not just for HTTP (or HTTPS), which enables it to distribute traffic for services such as message queues and databases as well, in addition to web servers.
- When configured as a layer 7 load balancer, HAProxy contains enhanced load balancing capabilities such as SSL termination for example, but consumes more resources to inspect HTTP related information. It's suitable for web sites crawling under high loads while requiring the persistence of Layer7 processing.

HAProxy offers three types of services:

- Load balancing
- High availability

- Proxying for TCP and HTTP-based applications

The high availability part of HAProxy is possible because HAProxy, besides performing load balancing by distributing requests among multiple web servers, can also check the health of those servers. If a server goes down, HAProxy will automatically redirect traffic to the other web servers. On top of this, you can use it to help the servers monitor one another, and automatically switch a slave node to the master role if the master node fails.

Using Nginx for Load Balancing

NGINX, the new hotshot web server, also performs load balancing functions. NGINX is more than a pure load balancer – it also contains reverse HTTP proxy capabilities. This means it can cache HTTP responses from the web servers. NGINX is a great candidate when you need a reverse proxy as well as a load balancer

HAProxy versus Nginx

HAProxy offers some benefits when compared to NGINX, as a load balancer:

- HAProxy as its name indicates, provides high availability support and is more resilient, and is easier to recover from failures.
- HAProxy is in many ways simpler to configure
- HAProxy performs better than NGINX, especially when you configure it as a layer 4 load balancer.

NGINX's reverse HTTP proxy capabilities enable it to offer the following benefits:

- Caching capabilities that reduce resource usage of your web services layer
- Easy scaling of the web services layer by adding more servers to the NGINX pool

Scaling the Load Balancer

For many small to medium sized applications, you can run a single HAProxy or NGINX service to handle the workload. For high availability, you may want to configure a hot standby as well on the side. While both NGINX and HAProxy can handle thousands of requests per second and a large number of concurrent connections (tens of thousands), ultimately they both have a hard capacity limit.

When you do reach the capacity limit of your software based load balancer, it's time to scale out the load balancers horizontally, by creating multiple load balancers. When you deploy multiple load balancers, each of them will receive some of the arriving traffic, using a round-robin DNS to apportion the traffic among the load balancers.

While there are several drawbacks to using round-robin DNS in assigning traffic to web servers, it's a harmless way to distribute work among the load balancers themselves, since they are more stable entities in the sense that you don't often replace them, or move them around.

Note that you can use both HAProxy and NGINX even in a hosted environment, if for any reason you want to use your own load balancing setup and not the provider's load balancing service.

Hardware Load Balancing

Hardware load balancers are meant for use within your own data center, and are useful when hosting heavily used websites. Hardware load balancers are the real deal: unlike software load balancers, which you can only scale horizontally (by adding more of them), a hardware load balancer is easier to scale vertically. Typically, a hardware load balancer can handle hundreds of thousands, or even millions of simultaneous connections. These load balancers provide high throughput and a very low latency, all of which dramatically increase the power of your network. Big-IP from F5 and Netscaler from Citrix are the leading hardware based load balancers.

A data center network connects its internal hosts with each other, and also connects the data center to the Internet. The hosts do all the heavy lifting, by serving web pages, storing email and performing computations, and so on. The hosts are called blades and they resemble pizza boxes stacked in a tray, and are usually generic commodity servers. Hosts are stacked in racks, with each rack containing anywhere from 20-40 blades. On top of the rack is the main switch called the Top of Rack (ToR) switch that connects the hosts with each other as well as with other switches in that data center. Each host has its own data center specific internal IP address. Hosts usually use a 1 Gbps Ethernet connection to the ToR switch, although 10 Gbps connections have become common of late.

Each application running in a data center is associated with a public IP address to which external clients send their requests and receive their responses from. When you use load balancing, the external requests are directed first to the load balancer. It's the load balancer's job to balance the workload by distributing requests among the hosts.

A large data center may use multiple load balancers, each dedicated to a set of cloud based applications. A load balancer such as this is often called a "layer-4 switch" since it makes its load balancing decisions based on the destination port number and destination IP address in the packets – both of these entities belong to the network layer (layer 4). In addition to balancing the workload, a load balancer also provides you security benefits, since it provides a NAT-like function by translating the public external IP addresses into the internal IP address of the hosts.

Hardware load balancers are expensive and since they're pretty sophisticated pieces of hardware, require appropriately trained personnel to manage them. However, for non-hosted, heavily trafficked web sites, hardware load balancing is the best approach.

Using a Hosted Load Balancer Service

Managing load balancers gets harder as the size of your infrastructure grows. As with other infrastructure services, it's a good idea to consider using a hosted load balancer service such as the Elastic Load Balancer (ELB) service offered by Amazon AWS. You can benefit from a third-party service such as ELB regardless of whether you host your own applications or host them on Amazon AWS (or Azure).

If you're using a hosted service such as Amazon Web Service (AWS), it's better to use the host's load balancer such as AWS's Elastic Load Balancing (ELB) service than to set up your own. Everything is done for you by AWS and all you have to do is to configure ELB to work with the set of instances that you choose from AWS's dashboard. A load balancing service such as ELB has several major advantages:

- It's simple to set up, as everything is done for you by the service provider (you can do the minimal configuration from your side via a console such as the AWS console)
- Ability to automate load balancer configuration changes
- It's easy to scale up using auto-scaling, with the click of a button, so to speak
- Connection draining features that let you take web servers out the pool without disconnections – when you take a server out, it waits for existing users to disconnect
- A load balancing service such as ELB allows SSL termination
- No upfront infrastructure expenses, and you only pay for what you use
- Easy failure management, as failed web servers are easily replaced

The load balancers I've been describing thus far are public facing load balancers, but Amazon and other cloud providers also let you configure internal load balancers as well. Internal load balancers sit between your front-end web servers and your internal services, and let you reap all the benefits of load balancing inside your stack.

Modern Networking Technologies

The widespread use of cloud provider supported data centers as well as large enterprise data centers driven by big data has led a large number of interconnected servers, with an overwhelming portion of the network traffic occurring within the data center network itself. Current computing trends such as the increasing popularity of big

data, cloud computing, and mobile network traffic have contributed to a tremendous increase in the demand for network resources. In the following section, I briefly discuss two important concepts that let you quantify a network's performance: quality of service (QoS) and quality of experience (QoE).

Quality of Service (QoS)

QoS is a set of measurable performance characteristics of a network service, which a service provider can guarantee via a service level agreement (SLA). QoS commonly includes the following performance properties:

- **Throughput:** this is measured in bytes per second or bits per second for a given connection or traffic flow
- **Delay:** also called latency, this represents the average or maximum delay
- **Packet Jitter:** the maximum allowable network jitter (jitter is the variation in the delay of received packets due to network congestion, improper queuing, or configuration errors)
- **Error rate:** usually the maximum error rate in terms of bits delivered in error
- **Packet loss:** defined by the fraction of packets not dropped
- **Availability:** expressed as a percentage of time
- **Priority:** a network can also assign priorities for various network traffic flows to determine how the network handles the flows

When you move to a cloud environment, putting a QoS agreement in place means that the business applications get some type of minimal performance guarantees. QoS is a major determinant of resource allocation in a cloud environment.

Quality of Experience (QoE)

While QoS is a measurable set of performance properties, QoE is more intuitive and subjective, being the impression of quality felt and reported by users. It's how end users perceive the network performance and the quality of service, regardless of the actual QoS metrics. By capturing the end user's perception of network quality, QoE provides a different, and in many ways, a more usable measure of quality. However, while it's fairly easy to measure QoS, QoE is not so amenable to quantification and it's hard to come up with accurate measures for this component.

Both QoS and QoE depend to a great extent on network routing and congestion control, which I discuss in the following section.

Cloud Networking

Obviously one uses the Internet as part of a cloud setup with an external provider, but that's usually only a small part of the networking infrastructure needed to support cloud based operations. Often you also need high performance, highly reliable networking to be established between the cloud provider and the subscriber. Here, at least some traffic between the organization and the provider bypasses the internet, using dedicated private network facilities managed by the cloud service provider. In general, cloud networking includes the network infrastructure required to access a cloud, including connections over the internet. The linking of company data centers to the cloud also includes the setting up of firewalls and other network security mechanism to enforce security policies.

Routing and Network Congestion Control

Routing and network congestion are two key aspects of a network that have a major bearing on how efficiently a network transmits network packets. Efficient transmission of network traffic has a direct bearing on both QoS and QoE.

The internet's main job is to move packets from one place to another, using alternate paths or routes. The routing function is essential since there's often more than one path for a packet to travel to its destination. Routers usually employ an algorithm based on performance, such as least cost routing, which minimizes the number of hops through the route. Cost could be based on throughput or time delay. The router's job is to accept network packets and forward them to other routers or to their final destination, using what are called forwarding tables.

Routers base their routing decisions on various routing protocols in order to forward packets through an interconnected set of networks. An autonomous system is a set of routers and networks managed by an organization, with the routers exchanging information through a common networking protocol. Within an autonomous system, a shared routing protocol called an interior router protocol (also referred to as an interior gateway protocol) passes the routing information between the various routers. OSPF is one such protocol.

A protocol for passing routing information between routers belonging to different autonomous systems is called an exterior router protocol (or an exterior gateway protocol), with BGP (Border Gateway Protocol) being an example of this type of protocol.

A router's control function includes activities such as executing routing protocols and maintaining routing tables and handling congestion control policies. Network congestion occurs when internet traffic exceeds the capacity or if the network doesn't manage the traffic in an efficient fashion. When a router is overwhelmed by incoming network packets, it can either discard the newly arrived packets (not good!), or manage the traffic using flow control, which means neighboring routers will be forced to share the excess demand. As the network load increases, the queue lengths at various nodes get longer and nodes start dropping packets. This of course leads to the sources retransmitting these packets, and in the extreme case, the capacity of the network in effect eventually diminishes to almost nothing.

Two new network innovations that many service and application providers are adopting these days to combat network congestion are software-defined networking (SDN) and network functions virtualization (NFV).

Software-Defined Networking

Thus far in this chapter, I've been busy describing protocols, network layers, packets and frames, but the buzz these days in networking is all about something called software-defined networking (SDN). Although there's still more hype regarding SDN than actual deployments, it's clear that traditional data centers are slowly metamorphosing into software based data centers. Software-defined networking is a way of organizing the functionality of computer networks. With SDN, you can virtualize your network, which enables you to exert greater control over it and also provides support in traffic engineering.

SDN is the outcome of a search initiated around the year 2004 for a better and modern network management paradigm, and all indications are that it's slowly and quietly replacing traditional networking as it provides a way to meet the increasing demands made on networks by new trends such as cloud, big data, mobility and social networking.

Limitations of current networks

Networking models today are the product of architectural and protocol decisions made around 40 years ago. Networks were expected to be static entities and their topologies weren't expected to change significantly over time. Until the advent of virtualization, each of the applications was associated with a specific physical server which had a fixed location in the network. Virtualization has changed all of this. Applications are now usually distributed across multiple virtual machines, all of which can be moved around to optimize or balance the server workload.

Traditional networking uses the concept of subnetted network segments and routing, along with addressing schemes and namespaces. The migrating of VMs across a network causes the physical end points of a network flow to change and creates chal-

allenges for traditional networking. In addition to the problems introduced by virtualization, networks are used today to deliver various types of services such as voice, video and data. Traditional legacy networks are static in nature, and can't dynamically keep up with the fast changing network traffic and application demands.

Networks today are built by connecting large numbers of complex network routers which accept data packets from applications and forward them to the next router on the path to the packet's final destination. The router's operating system controls the packet forwarding function and is designed to work with the vendor's hardware platform. Specialized routing protocols use the operating system and the proprietary packet-forwarding hardware to send the packets along to their destinations.

There are inherent problems with the traditional networking model. In order to change the network behavior, you need to configure each of the routers and issue commands in the proprietary language of the router vendor. In this closed environment, it's difficult for routers to interact easily with the rest of the network components.

The three "Planes" in Networking

Computer networks use three distinct planes to perform their tasks – the data plane, the control plane and the management plane. Let me explain these three planes briefly:

- **Data Plane:** processes the packets received by a network by using the packet header information to determine whether the packet should be forwarded or dropped. If it decides to forward the packets, it must determine the host and port to which it should send the packet.
- **Control plane:** this plane computes the forwarding state used by the data plane to forward network packets, by calculating the state using distributed or centralized algorithms.
- **Management plane:** this plane coordinates the interactions between the data and the control planes.

While the data plane uses abstractions such as TCP/UDP and IP, thus creating a reliable transport layer out of a set of unreliable transport layers, the control plane has no such abstractions. There's no unifying abstraction underlying the control plane, with a bunch of control plane mechanisms created to achieve different goals. For example, a routing control plane may be used to provide access control lists or firewalls, by influencing the routing of packets through controlling the calculation of the forwarding state. Each of the control planes solves a specific narrow problem, but they all can't be used together.

The limited functionality of the control planes means there's no one solution that contains all the control plane functionality. Why do we need to worry about the fact that the different "small bore" control planes can't talk to each other? Well, this type of architecture has an impact on the bottom line.

In a legacy network, when you have a large amount of data that needs to be moved, by default the fastest network link is employed, although it's the most expensive. Since the applications can't really communicate with the network due to the way the control plane is designed in traditional networks, it never finds out about the availability of slower but cheaper alternate network links.

In a traditional network, it's not just the routers that perform the data, control and management functions in an integrated fashion – all the other network devices, such as a network bridge, packet switch etc. perform these functions in the same fashion as a router.

SDN constitutes the set of instructions that govern the control plane in a modern network. SDN's approach is quite different from that of today's distributed network routing protocols. SDN essentially involves the simple computation of a function. It computes functions on an abstract view of the physical network layer, ignoring the actual physical infrastructure of the network. This allows the network engineers to manage network traffic by ignoring the physical network design. The networking operating system takes the SDN generated function and distributes it to all the network switches.

Server virtualization is a big reason for the advent of SDN. While virtualization allows you to easily migrate servers across virtual machines for load balancing or for high availability purposes, it doesn't play too well with traditional networking. For example, the virtual LAN (VLAN) that's used by a virtual machine must be assigned to the same switch port as the host server. Since you may often move virtual machines it means you need to reconfigure the VLAN every time you move the VMs around. Since traditional switches perform both the control and data plane functions, it's hard to modify network resources and profiles with these switches. The need to provide a rapid network response to the use of virtual servers has been a big motivating force in the move to SDN.

In traditional networks, the control plane needs to compute the forwarding state of network packets, consistent with the low level hardware and software and the entire network topology. In addition it must ensure that it inserts the forwarding state into all the physical forwarding boxes in a network. A new protocol means that the control plane has to redo all of the work, which isn't very efficient.

Under SDN, the traditional three plane functionality of the router is broken up. Router hardware has less work to do in this architecture: all it needs to do is provide

the data plane functionality – that’s it! A software application running on a separate platform and connected to the router provides the control and management planes.

Performing all the routing functions within the router itself means that all routers must implement the same routing and control protocols. In SDN, a central controller will perform the tasks such as routing, naming, declaration of policies and security.

The SDN control plane consists of one or more SNC controllers, which define data flows that occur

Defining Functions for the Network Control Plane

When designing the network control plane, you need three types of functions:

- A forwarding function
- A network state function
- A configuration function

I describe the three functions briefly in the following sections.

The Forwarding Function

The packet-forwarding function needs to be implemented differently from how the network switch is implemented. This function abstracts the details of low level hardware and software involved in the creation of the network switch. The switch should be able to use any underlying low level mechanism and the software that runs in the switch shouldn’t impact the forwarding function.

The OpenFlow interface, which is a set of application programming interfaces (APIs), is one way to link the control plane software and the network switch.

The Networking Function

The goal of the network state function is the presenting of a graph-like global network view, with necessary network information such as network delay and the recent loss rate. This global network view functionality is provided by a network operating system (NOS), and the controlling software uses the network graph to make decisions. The global network view is continuously updated with the changes occurring with the network switches. The NOS has the capability to update the switches to control the packet forwarding.

In a traditional network design, network switches are responsible for routing packets of user information among themselves. The switches use the information to update their view of the global network, and modify how they forward the network packets.

Legacy network control planes implement peer-based distributed routing algorithms, such as the shortest path first (SPF) algorithm. The problem is that the algorithms are

complex, and the distributed decisions by the switches are based on only a partial knowledge of the network's global state.

In an SDN based network, there's a general purpose software algorithm that runs on the NOS servers and determines the network topology by polling the network switches. Based on the responses from the switches, the software creates a global network view. The global network view acts as the basis for various types of control programs such as routing or traffic control programs. Under SDN, instead of distributed decision making based on imperfect network information, a centralized control program based on the global network view is used to determine how switches ought to forward the packets.

With the centralized decision making through the control program, redesigning the network to route packets differently is easy – all you need to do is modify the control program. You don't have to create a brand new distributed routing algorithm, as is the case in a traditional network.

The Configuration Function

The control program doesn't concern itself about how the actual routing behavior is implemented. That is, the control program doesn't have to configure the routing tables on all network nodes to implement a routing algorithm. A virtualization layer is added to the SDN model to translate the control program's routing instructions into configuration commands for the switches in the physical network. This virtualization layer updates the global network view to reflect the routing decisions made by the control program. The layer hands the updated global network view to the NOS, which actually configures all the switches in the underlying physical network. It maps the packet-routing control commands and maps them to the physical switches. The OpenFlow interface facilitates the interaction between the NOS and the switches.

Since routers are left with just the data plane functionality, they don't need to be complex things any longer. You can now make do with basic packet-forwarding hardware based routers to communicate with the NOS through the OpenFlow (or a similar) protocol. Also, since under SDN the networking intelligence has been separated and moved to the SDN layers, switches can be mere commodity hardware as well.

Probably the most important benefit of SDN is the fact that by enabling network virtualization, SDN lets you migrate easily between your current network and the cloud. SDN makes it easy to export your virtual topology to the cloud and ignore the physical design of the network. Network teams can create a network policy statement using their network topology and replicate this policy statement in the cloud. Under SDN, the control plane is not on hardware but exists in software. You can easily perform large scale simulation of the control plane and test new network designs before moving into production.

In traditional networks, the control plane is part of a proprietary switch or a router box, whereas in SDN it's a program, with software determining how to forward network packets. Instead of designing a network, you'll be programming networks based on the abstracted view of the physical network. The software you use to program the network will be independent of the hardware used by the network. Troubleshooting SDN based networks is also simpler since you can simulate the programs once you identify the source of a network related problem such as incorrectly forwarded network packets, for example.

Network Functions Virtualization

In today's computing world, storage and servers have already been virtualized. Network functions virtualization (NFV) allows the virtualization of the physical networks. Once it virtualizes the network, NFV enables software applications that use the network to reconfigure the network as they see fit, thus letting the network provide the best service possible to the applications.

NFV breaks away the set of essential network functions such as routing, firewalling, NAT and intrusion detection from vendor controlled proprietary platforms and implements these functions in software, using virtualization. If NFV sounds very similar to SDN, it's no mystery, as they share several features and objectives. Both SDN and NFV believe in:

- Using software instead of hardware to provide network functions
- Using commodity hardware instead of proprietary network platforms
- Using standardized APIs

The key thing to remember is that both SDN and NFV decouple or break up components of the traditional network architecture. While SDN decouples the data and control planes of network traffic control, NFV removes the network functions from proprietary hardware through virtualization and implements the network functions as software. You can use either SDN or NFV by itself, but of course, employing both together will get you the most benefits. If you use SDN by itself, you'll be implementing the control plane functionality on a dedicated SDN platform. If you use NFV as well, you can implement the network data plane functionality on a VM and the control plane functionality on either a dedicated SDN platform or an SDN VM.

The OpenFlow Protocol

The OpenFlow protocol is a new network protocol that has been created to enable software-defined networks. The OpenFlow protocol provides structure for the messaging between the control and data planes and provides external applications access to the forwarding plane of a network switch or router.

Traditional networks tend to stay static over time as a result of their complexity. In order to move a device to a different location on the network, the network administrators must modify multiple switches, routers, firewalls and updated ACLs, VLANs and other protocol based mechanisms that work on the device and the link levels. The OpenFlow protocol was designed to solve the problems posed by legacy network protocols. In SDN architectures OpenFlow is the interface between the control and forwarding layers and lets you directly access the network devices such as switches and routers in the forward plane. Standard network protocols can't offer this type of functionality, and thus OpenFlow is a necessary ingredient to move the control part out of networking switches and into centralized control software.

With this review of networking behind us, let's turn to how modern web applications are architected, and how you achieve scalability in a world of modern web applications, web services and microservices.

Scalability, Web Applications, Web Services, and Microservices

Supporting web sites and web applications is often a key function of Linux system administrators. An overwhelming majority of the world's web sites and web applications run on Linux. In the old days, when it came to supporting web based applications, all an administrator had to know was how to set up a web server such as the Apache HTTP server, and a few things about DNS and how to connect the web front end to the backend databases.

Over the past several years, there have been truly revolutionary changes on many fronts, changes which require you to be conversant with a lot more technologies that have come to play a critical role in driving web applications. The rise of web based applications and the consequent concurrency issues it gave rise to due to massive user bases have called for innovations in virtually all areas of the traditional web application architecture. In addition to newer application architectures, there are also vastly more moving pieces in a typical web application today than in the applications from the previous generations. This chapter has two major goals:

- Explain the concept of scalability and how you can enhance it using modern approaches, in all areas of a web application, such as the front end and back end web servers, databases, caching servers, etc.
- Introduce you to several modern innovations such as new application architectures, NoSQL databases, modern caching concepts, asynchronous processing models, high availability, and newer, more efficient web servers

Scalability is the ability of applications to handle a large number of users, data, transactions or requests without a slowdown in response time. Modern application architectures utilize several strategies to achieve scaling for each of the components of the

application stack. Throughout this chapter, our focus is squarely on scaling using modern techniques.

Web applications are often deluged with large numbers of simultaneous users, making concurrency a tricky issue. Over the past 10 years or so, the need for higher scalability has driven several significant developments in web architectures, with the adoption of newer concepts such as web services and microservices. This chapter provides an introduction to these architectures.

Traditionally, web applications have used the open source database MySQL and the open source Apache web server to power them. While the established user base of MySQL and the Apache Web server is still very high, the pursuit of scalability and concurrency has led to newer database and web servers that have been explicitly designed to provide scalability.

Modern web servers designed with high concurrency in mind, such as the NGINX web server and modern databases that are based on a non-SQL approach to data processing such as MongoDB and Cassandra have come to become major players in today's application architectures. This chapter explains the way these modern tools work, and how they provide the benefits that make them significant components in the modern many application architectures.

This chapter discusses several concepts that are used by developers, but administrators need to know about these crucial concepts so they can better serve their clients, both within and outside their organization. For example, learning about popular programming concepts such as Model-View-Controller (MVC) architecture and the reasons for the popularity of the MEAN stack for building web applications helps understand what the developers in your organization are doing, and why. Similarly, understanding the basics of web performance helps you provide more efficient services for your external customers (Chapter 16 discusses web performance optimization in detail)

Two of the most popular web application frameworks today are what are generally referred to as the MEAN Stack and Ruby on rails. MEAN stands for a framework that uses the following components together to drive web applications: MongoDB, Express, Angular JS and Node.js. Ruby on Rails is a very popular framework for developing web applications, and relies on the principle of “convention over configuration”.

This chapter explains the Model-View-Controller (MVC) architecture that underlies modern web applications. Following this, it explains both the MEAN application stack as well as Ruby on Rails, which is a highly popular modern web application framework. You'll also learn about the new paradigm in web applications, the Single Page Application (SPA).

Microservices are increasingly becoming the norm, more or less supplanting the previously popular concept of web services. You'll learn about both web services and microservices in this chapter, and also find out the key differences between these two approaches.

Scaling and Common Datacenter Infrastructures

Before a client request even arrives at a company's data center, it usually traverses through third party entities such as a GeoDNS or a CDN service. Once it arrives at the data center, web traffic is handled by multiple technologies, some of which fall under the front end type and others the back end architecture components. Let's list the typical components of front end and back end technologies.

While managing a data center with the multiple layers shown here does add to the complexity, it helps you scale your environment. For example, load balancers reduce the load on the web application servers in the front. Similarly, backend web services need to deal with a lower amount of traffic since the application level caches and the message queues handle some the load.

While I list a fairly large number of components for a data center, by no means do you need all of these in your environment. If your applications are simple, they may not need all the complex layers, such as frontend cache server, message servers, search servers and backend cache servers, for example. The objective here in listing these components is to make you aware of the potential complexity of a data center architecture.

Scaling web applications (and your entire computing environment for that matter) means knowing how to scale all the components of a typical architecture, including both the front line and backend technological components of a data center.

Front End and Back End Web Technologies

You probably are familiar with the commonly used terms “front end” and “back end” when dealing with web applications. The front end in this context refers to the application code, HTTP and web servers, load balancers, reverse proxies and Content Delivery Networks (CDNs). The backend layers include web services, object caches, message queues and databases. I discuss how you can enhance scalability and performance by using several strategies at the front end as well as the back end of your applications. Here's a listing of the main technologies I explain in this chapter:

- Load Balancers
- Web servers (web application servers)
- Caching servers (from caching servers)
- Reverse Proxy Servers

- Content Delivery Networks (CDNs)
- Web services and Microservices
- Object caches
- Messaging (queuing) services

The Front End Technologies

When dealing with web applications, front-end technologies are responsible for rendering the user interface and handling the user connections. Typically, the front-end technologies are the following:

- Reverse proxies
- Content Delivery Networks (CDNs)
- Hardware (or software) load balancers
- Domain Naming Service (DNS)
- Front cache servers (optional)
- Front-end web application servers

The web application layer is the presentation layer and its job is to simply handle user

The Back End Technologies

Back end technologies include the following:

- Web services: Web services constitute the heart of the application and contain the business logic. These services use either REST or SOAP (I explain both of these in this chapter) to communicate with the front-end applications. If you keep the web services stateless, scaling them involves adding more servers to run the web services.
- Cache servers: these are used both by the front-end application services as well as the backend web services to reduce latency by storing partially computed results.
- Message queues: both front end applications servers and the web service machines can send messages to the queues. The message queues help postpone some processing to later stages and also delegate work to worker machines which process the message queues. The queue worker machines perform work such as taking care of asynchronous notifications and order fulfillment which typically require significant time to complete.
- Databases: a data center can contain both traditional relational databases such as MySQL and several big data stores such as Hadoop, in addition to NoSQL database such as MongoDB and Cassandra. Regardless of the type of database, scala-

bility is a running theme common to all databases, since data sizes are usually large, regardless of the type of data that's stored and processed.

- Other servers: These can include batch processing servers, search servers, file servers and lock servers.
- Third-party services: these include services from both CDN providers and the providers of various business services such PayPal and Salesforce, for example.

Scalability of Applications

We all know in general what scalability means – it's the ability to handle an increasing number of users and their requests without deterioration in the performance of the applications. Scalability contains the following two major dimensions:

- Ability to handle higher concurrency: as applications become more popular, the number of users simultaneously accessing the web sites starts rising, bringing to the fore a special set of problems, such as an increasing number of open connections and active threads, not to speak of higher I/O for the storage subsystem as well as the network, and a higher load on CPU. Latency and response time become crucial measures of the ability of a system to handle high concurrency usage patterns.
- Processing More Data: As a website becomes more popular, more users will be making requests and your system will be handling increasingly larger amounts of data. Processing data involves disk I/O as well as network I/O, and therefore the system has to keep up with the higher data loads without slowing down the speed of processing.

Scalability and performance are intertwined, and present two sides of the same coin. A system that can't process things fast enough can't scale very well. There are two major approaches to scaling applications – vertical scaling and horizontal scaling, as I explain in the following sections.

Scaling Vertically

Vertical scaling is usually the first response to meeting higher scalability requirements. Under vertical scaling, you seek to upgrade the processing power of your systems without reconfiguring your applications or the components of your application architecture. When it comes to upgrading the processing power of servers, the most common approaches are to increase the RAM allocated to the servers, or sometimes, even switch to different servers with more and/or faster CPUs. The following are all ways in which you can scale up.

- Increase RAM to reduce the disk I/O – computations in memory are much faster as compared to computations on disk, so for most database dependent applications, allocating more RAM leads to dramatic improvements in performance.
- Get faster network interfaces: you can often gain throughput by upgrading your network adapters.
- Move to more powerful servers: you can move your databases and web servers to servers with more CPUs and virtual cores. The CPUs themselves can also be faster than the CPUs you're currently using.
- Improve the disk I/O: One of the first things you can do here is to see if you can invest in speedier hard disk drives (higher RPMs). Web applications in particular predominantly use random reads and writes, which are much faster when you use solid-state drives (SSDs). If it makes sense from the cost point of view (while they were extremely expensive to begin with, SSD costs have been continuously dropping), SSDs could dramatically speed up disk I/O.

While vertical scaling can and does enhance scalability in most cases, and offer the benefit that it doesn't require you to make any changes to your application architectures, it has inherent limitations over the long run. Scaling up is usually extremely expensive and not always cost effective. You'll find that you get a smaller payoff by spending extra money on buying these muscular infrastructure components (ultra-fast disk drives, faster CPUs, and more powerful RAID controllers, for example) with higher specs.

Even assuming that you can afford the increasing cost (the cost increases nonlinearly) of the heftier hardware and other components to support vertical scaling, you'll eventually hit a performance wall due to the sheer fact that there are hard limits on how much memory a system can handle, and how fast disk drives can spin.



Content Delivery Networks (CDNs) also help significantly improve scalability by taking some load off your web servers for providing static web content. I discuss CDNs in the following sections.

Scaling Horizontally

While vertical scaling takes the approach of bulking up the existing infrastructure so it can process more things faster, horizontal scaling enhances your infrastructure's processing capacity by adding more servers to your infrastructure. You can add more database servers or more web servers, for example, to support a growing volume of business.

The principle of horizontal scaling underlies the massive infrastructures built up by behemoths such as Amazon and Google, which run data centers with hundreds of

thousands of commodity servers, as well as the increasing popularity of big data processing models such as Hadoop and Mesos.

Content Delivery Networks

One of the best ways to offload some of an application's web traffic is to employ a third-party content delivery network (CDN) service, such as Akamai, CloudFare, Rackspace (Rackspace Cloud Files Service), and Amazon's Elastic Compute Cloud (EC2). Currently Akamai is the worldwide market leader. Content delivery networks are hosted services using a distributed network of caching servers that work somewhat similar to the way caching proxies work, and are typically used for distributing an application's static files such as images, CSS, JavaScript, PDF documents and videos.

Users that need to download the static content connect directly to the CDN provider instead of your web server. The CDN will serve the requested static content to the users, and if it doesn't have the requested content, it gets that content from your web servers and caches it, and so it can send it along to the users directly for all subsequent requests.

CDNs are highly cost effective for most organizations. They do things at a fraction of what you'd have to spend if you were to set up complex and hard to maintain network infrastructures across the globe to cache the content in your own private infrastructure.

While CDNs are primarily meant for providing static content, they aren't limited to it. CDNs can also serve dynamic content for your websites, offloading more work from your own data center to the CDN provider to process content. As Chapter 9 explains, you can configure Amazon's CloudFront service to provide both static and dynamic content. Note that in addition to simply serving static web content, CDNs can also work as a HTTP proxy, as I explain later.

If your needs aren't big, you can set up a simple Apache server to deliver your static files so you can take needless load off your web application servers. You can use the single server CDN setup until your needs require you to scale up, and that's when you can splurge for a third-party CDN provider. In addition to scalability, third-party CDNs also offer faster performance since a globally distributed CDN can serve its content from the nearest server.

Since CDNs use caching to do their work, they are similar to proxy servers for caching purposes. In fact, CDNs use HTTP headers to cache content, in the same way as proxy servers. Large web applications benefit immensely from a CDN since CDNs not only reduce your network bandwidth usage, but also decrease the network latency by serving content to users from a the nearest cache server from their distributed network of cache servers.

As mentioned earlier, you can configure a CDN to cache both static and dynamic content. For example, Amazon CloudFront can serve both static and dynamic content for your applications. Typically, however, web applications use CDNs to cache static files such as images, CSS, PDF documents and JavaScript.

How large websites scale

When scaling for large audiences, organizations use multiple data centers. Spreading the data centers across the world enhances user experience and also provides higher availability. Two of the most common strategies used to service a global audience are the use of a GeoDNS and the use of edge-cache servers.

GEODNS

Normal DNS servers resolve domain names to IP addresses. A GeoDNS takes this concept a bit further: they serve IP addresses (of the web servers or most commonly, the IP addresses of the load balancers) based on the location of the client, allowing clients to connect to web servers running closer to where they're located. Obviously, minimizing network latency is the goal in using a GeoDNS strategy.

Edge-cache Servers

Large companies can also reduce network latency by hosting several edge-cache servers across the world. Edge-cache servers are HTTP cache servers that are situated closer to the clients, enabling them to (partially) cache their HTTP content. The edge-cache servers handle all requests from a client browser and can serve entire web pages if they have it in the cache, or they can use cache fragments of HTTP responses to assemble complete web pages by sending background requests to the company's web servers.

Both GeoDNS and edge-cache servers can be used together with CDNs. For example a company's Asian customers will resolve the company's domain name to an IP address of an Asian edge server. The customers will be served results cached by an edge-cache server or from one of the company's own web application servers. The company's CDN provider loads static files such as JavaScript files from the nearest data center of the CDN provider in Asia. The bottom line is to keep latency low, and cut the costs in achieving that goal.

Scaling Web Applications

Scalability of both the front end and the backend of web applications is crucially affected by the concept of state. To understand how this is so, you need to know the difference between stateful and stateless services:

- A stateful service holds data relating to users and their sessions. Data in this context could relate to user session data, local files or locks.
- A stateless service (or server) doesn't hold any data. Data in this case captures the state and stateless services let external services such as a database handle the maintenance of a client's state. Since they don't have any data, all the service instances are identical.



Stateless services don't retain state or any knowledge relating to users between HTTP requests made by the users. Since a stateless service stores the state in an external shared storage (such as a database), all stateless servers are virtually the same - there's no difference among them.

Having to maintain any type of state has a negative impact on your system's ability to scale. For example, keeping web servers (both front-end and backend web service servers) stateless lets you scale the servers simply by adding more servers. State has relevance both in the front-end layer as well as on the backend, as the following sections explain.

Managing state at the front end

For front end servers such as web servers, there are two main types of state that need to be managed – HTTP sessions and file storage. Let's review how web applications handle these two types of state.

HTTP Sessions and State

While the HTTP protocol itself is stateless, web applications create sessions on top of HTTP to keep track of user activity. There are several ways to track session state, as explained here:

- Using session cookies: Web applications usually implement their sessions with cookies, and the web server adds new session cookie headers to each response to a client's HTTP request, enabling it to track requests that belong to specific users.
- Using an external data store: Using cookies is a simple and speedy approach when data is small, but if the session's scope gets large, cookies can slow down the web requests. A better strategy may be to store the session data in an external data store such as Memcached, Redis or Cassandra. If the web application is based on a Java type JVM-based language, you can store the session data in an object clustering technology such as Terracotta instead. The web server remains stateless as far as the HTTP session is concerned, by not storing any session data

between web requests. Instead, the web application grabs the session identifier from the web requests and loads the pertinent session data from the data store.

- Implementing “sticky sessions” through a load balancer: In this strategy, the load balancer takes over the responsibility for routing session cookies to the appropriate servers, by inspecting the request headers. The load balancer uses its own cookies to track the assignments of the user sessions to various web servers. Note that under this strategy the web servers do store the local state by storing session data.

File Storage and State

In addition to the state of the applications, file storage is another important type of web application state handled by web servers. If you’re wondering how file storage plays a role in managing state, it’s useful to remind yourself that typically, users can both upload content to the web servers and also download various files generated by the web applications.

Instead of reinventing the wheel, you can set up your own file storage and delivery system using open-source data stores and related tools to scale your file storage system. This lets you take advantage of scalability features such as partitioning and replication that are built into these data stores. For example, GridFS, which is an extension that’s part of MongoDB, can split large files and store them inside MongoDB as MongoDB documents. Similarly, Netflix’s Astyanax Chunked Object Store takes advantage of another open source data store, Cassandra, by using its partitioning, failover and redundancy features, and by building file storage features on top of the data model.

Auto Scaling

Auto scaling is the automatic adding and removing of processing power, such as adding or removing of web servers based on changing workloads. While scalability in this chapter mostly deals with scaling out to efficiently address increasing workloads, auto scaling is the automation of an infrastructure’s capacity by automatically adding to or removing servers from the infrastructure based on the workloads. Auto scaling is really more relevant in a hosted environment such as Amazon AWS, where, by using Amazon’s auto scaling feature one can considerably reduce their hosting costs. For a client that’s running a web application by hosting it on AWS, handling a holiday rush is a breeze with auto scaling – and once the holiday rush business is taken care of, you won’t need the beefed up infrastructure, so the system automatically scales itself down. Just as you can use a service such as Amazon’s Elastic Load Balancer to auto scale your front end web servers, you can also auto scale the web servers hosting your web services.

While you may be able to manage your proprietary file storage system by leveraging the open source data stores, a far better strategy would be to simply delegate the whole area of distributed file storage and delivery to inexpensive third-party services such as Amazon's Simple Storage Service (S3). You can use these types of storage services even if you aren't using AWS to host your web applications.

Regardless of whether you use a service such as Amazon's S3 or handle the file storage yourself, it's a good idea to employ a content delivery network (CDN) to send files to the users. The CDN can use Amazon S3 or public web servers to download and cache public content. If the files that users need to download are private files, you can store them in Amazon S3 again, but in private instead of public buckets, or use private file servers if you're handling the data storage yourself. The web application servers will then fetch the private files from S3 or your own private file servers, enabling users to download them.



Web servers can be front-end web servers or servers that host the web services in the backend.

Other Types of State

In addition to session data and file storage, other types of state such as local server caches, resource locks and application in-memory state negatively impact state. Caching in particular is tricky since it does enhance performance and reduce the workload of the database and web services, as you'll see later in this chapter.

Resource locks synchronize access to shared resources but it's very difficult to implement those locks correctly in practice, and you should instead try to move the state out of your application servers. You may want to consider a distributed locking service such as ZooKeeper or Consul. As this chapter explains later on, ZooKeeper not only is used for distributed locking (for Java based applications), but also for leader election and managing runtime cluster membership information. For scripting languages such as Ruby, you can use a simpler lock implementation based on operations within a NoSQL database such as Memcached or Redis.

Regardless of whether you use ZooKeeper or Memcached or a different implementation of a locking service, the principle here is that you remove the state related functionality from the web server and move it to a separate, independent service. This makes it easy to scale out your web applications.

Scaling the DNS

1. As Chapter 8 (Cloud and Amazon AWS) shows, Amazon's Route 53 services is integrated extremely well with Amazon AWS and lets you easily configure DNS through your AWS web interface. Route 53 uses latency based routing to route clients to the data center nearest to them. In some ways this is similar to GeoDNS which I explained earlier in this chapter, and is probably even better in some ways than GeoDNS, since it bases its routing decisions on latency instead of location.

If you aren't hosting your infrastructure on Amazon, not to worry, as there are several DNS services out there such as easydns.com and dyn.com that provide services similar to those offered by Amazon's Route 53 service.



A CDN is commonly used for serving static files such as images, CSS and JavaScript files, but you can use a CDN to proxy all web requests arriving at the web servers, if you so wish.

Scaling the Load Balancers

Amazon's Elastic Load balancing (ELB) is part of Amazon's auto scaling feature that automatically adds servers to meet rising workloads and also replaces the web servers automatically when the servers fail. Since ELB scales transparently and is highly available, you don't need to worry about managing it.

In addition to Amazon's AWS, other cloud providers such as Microsoft's Azure offers the Azure Load balancer and Rackspace offers the Cloud Load Balancers and OpenStack comes with LbaaS. All of these cloud providers offer similar load balancing services as ELB. You can scale the load balancing easily with any one of these services, regardless of whether you host your infrastructure on those providers or not.

Scaling the Web Servers

The programming language in which you implement your front end applications, as well as the type of web servers to use, are both crucial decisions you need to make when it comes to architecting the front-end web servers, which provide the presentation logic as well as serve as an aggregation layer for the web service results. Although using the same technology stack across all of a web application's multiple layers offers several benefits, in practice, it's common for different layers to use different architectures to solve their unique problems.

In terms of the actual web servers, you have a choice between traditional web services such as Apache and Tomcat (or JBoss) and the newer web servers such as NGINX, which I discuss later in this chapter. In terms of increasing horizontal scalability, the

real concern here isn't the language framework or the type of web server. It really is a question of whether you've ensured that the front-end web servers are truly stateless.

Using Caching to scale the front-end

Caching helps you scale by reducing the workload of web services and databases in the backend. There are multiple ways in which you can manage caching at the front end of your web applications:

- Using a Proxy server: As you'll learn shortly, you can use a Content Delivery Network (CDN) to cache web content. However, CDNs aren't ideal for caching entire web pages. You can use a reverse proxy server such as NGINX (or Varnish) to better control both the type and the duration of the cached web content.
- Caching in the web browser (HTTP caching): Today's web browsers can store large amounts of data, ranging into multiple megabytes in size. Modern web applications such as Single Page Applications (SPAs) can access the local browser storage from the application's (JavaScript) code, thus reducing the workload of the web servers. You normally implement HTTP caching in the front-end layer and in the backend web services layer in a similar fashion.
- Using an object cache: As I explained earlier in this chapter, an object cache allows you to store fragments of a web server's responses. Both Redis and Memcached are very powerful shared distributed object caches that can help you scale applications to millions of users through clusters of caching servers.

Later sections in this chapter will explain the concepts of a proxy server, browser caching, and object caching in more detail.

Scaling the Web Services

You can scale your web services using multiple strategies such as making the servers that are hosting the web services stateless, caching the service's HTTP responses and through functional partitioning. The following sections explain each of these strategies.

Making the Servers Stateless

As with the rest of your architecture, keeping state out of the servers hosting your web services plays a significant role in scaling your web services. You can keep state off the servers by letting it store just auto expiring caches and by storing all other persistent data representing the application state in external data stores, caches and message queues. Keeping the servers stateless will let you employ load balancers to distribute client work over all the servers hosting web services, and allows you to take servers out of the load balancing pool when they crash. You can add and remove servers, as well as perform upgrades transparently without affecting your users. Scal-

ing is as simple as just adding more web service hosts to the load balancer pool to handle more connections.

Using HTTP Protocol Caching

A REST web service's GET responses are cacheable, but in order to make the most of this, you must ensure that all the GET method handles are read-only, since it means that issuing GET requests would leave the state of the web service unaltered. If all GET handles are strictly read-only in nature, you can add caching proxies in front of the web services to handle most of the incoming web traffic.

In order to implement HTTP caching at the web service layer to promote scaling, you add a layer of reverse proxies between the web service clients such as the front-end servers and the web and your web services. All requests will now pass through a reverse proxy. This layer of reverse proxy HTTP cache will distribute requests among multiple web service servers.

Patterns such as using local object caches would make it difficult to cache the GET responses, since the web servers can end up with different versions of the data. You can make the resources public to ensure that you only have a single cached object per URL and thus enhance the efficiency of HTTP caching, and make it truly lower the workload of the web services. Making the GET handles public lets you reuse the objects easily.

Using Microservices to Enhance Scaling

You can break up a large web service into multiple independent microservices, with each component addressing a specific functionality. For example, you can create microservices such as a `ProductCatalogService` and a `RecommendationService` and other similar services, each addressing a separate subsystem, in order to break up the work that was being handled by a single web service.



Web services help web applications scale since the web services can use different technologies and you can scale each of the web services independently, by adding more web servers to host a specific web service, for example. APIs also help you scale because you can break up the web services layer into much smaller microservices.

Creating multiple, independent microservice based subsystems will distribute the total work among more databases and other services, which enhances scaling. Since different microservices can have different usage patterns, with some microservices needing to perform more read requests, and others with mostly write requests, for example, you can scale each of the services independently. For example, some services may require different types of databases or a different way of caching the content – you can employ the “best of breed” strategy for each of the microservices.

Making Effective use of Third-party services

Organizations with small to medium size operations can significantly enhance their competitive situation by using services offered by third-party entities instead of building everything from scratch, besides having to devote considerable expense and effort to maintain the services. Three of the most useful such services are:

- Content Delivery Networks (CDNs)
- Site analytics
- Client logging

Let me briefly describe how outsourcing each of these services helps you develop more effective web services and SPAs.

Site Analytics

As with traditional websites, web services and SPAs can benefit from using analytic services such as Google Analytics and New Relic to help them understand the usage patterns of their web sites and unearth performance bottlenecks. SPAs can be enhanced in two different ways to take advantage of a service such as Google Analytics:

- Use Google Events to track the hashtag changes: Google Events uses parameters passed along at the end of requests to process the information about that event. Developers can organize and track various events. For example, if the following call shows up in the reports for an SPA it tells you that a chat event has occurred and that the user sent a message on the game page.

```
_trackEvent( 'chat', 'message', 'game' );
```

Let's see how Google Analytics can help you optimize your website performance. You can use a Google Analytics report to correlate the traffic statistics for your top 10 web pages with the average time it takes to render the web pages. This helps you create value based rankings of the web pages such as the following:

| Value Rank | Page | Average Seconds (Response Time) | Requests per Hour |
|------------|----------|---------------------------------|-------------------|
| 1 | / | 0.33 | 1000000 |
| 2 | /mypage1 | 0.99 | 200000 |
| 3 | /mypage2 | 0.60 | 95000 |
| ... | | | |

Ranking the web pages as shown here lets you quickly estimate the potential bang for the buck, so to speak. It lets you see, for example that improving their home page load times by just 5 ms offers vastly greater improvement in overall performance as compared to increasing the performance of the web pages mypage1 or mypage2 by 10 ms or 20 ms. Web optimization, as with other optimizations, requires the calculation of

the costs and benefits of the optimization efforts, and data such as the simple set of metrics shown here are the best way to prioritize your efforts.

Server-Side Google Analytics: You can also track the server side with something like Node.js, to see the types of data being requested from the server. This will help during the troubleshooting of slow server requests. However, tracking the server side isn't as useful as tracking the client-side actions.

CLIENT LOGGING

In single page applications (SPAs), which are increasingly becoming the norm for web applications, unlike in traditional websites, when a client hits an error on the server, the error isn't recorded in a log file. While you can write code to track all errors, it's better to use a third-party service to collect the client errors. You of course, can add your own enhancements or tracking features, even with a third-party service in place.

There are many useful third-party services that collect and aggregate errors and metrics generated by new applications. All client logging services work the same underneath: they catch errors with an event handler (`window.onerror`) and trap the errors by surrounding the code with a try/catch block.

The company NewRelic happens to be the most popular of these services, and is the current standard for monitoring web application performance. A service such as NewRelic provides considerable insight into the performance of the server as well the clients. NewRelic provides not only error logging, but also performance metrics for each step of the request and response cycle, so you can understand if the application response is slow because the database processing is slow, or if the browser isn't rendering the CSS styles fast enough.

Besides NewRelic, you can also check out the error logging services such as Errorception (for JavaScript error logging), Bugsense (good for mobile apps) and Airbrake (specializes in Ruby on Rail applications).

CONTENT DELIVERY NETWORKS (CDNS)

You've already learned about CDNs earlier in this chapter. CDNs provide scalability for static content. In a SPA context, where Node.js is commonly used as the web server, you're at a disadvantage when serving static content, especially large static content files that contain images, CSS and JavaScript since they can't take advantage of the asynchronous nature of Node.js. Apache Web server does a far better job with the delivery of such files due to its pre-fork capability.

In our next section, let's review application scalability, which is the heart of most if not all modern enhancements in web application technologies.

Working with Web Servers

Organizations of every size use web servers, some to serve simple web sites and others to support heavy duty work by playing the role of backend web servers. There are several web servers out there, all open source products, of which the most popular ones today are the Apache Web Server and NGINX, with the latter still with substantially smaller usage numbers than the former, but whose usage is growing at a much faster rate. Let's take a quick look at these two popular web servers and the uses for which they're best suited.

Working with the Apache Web Server

The Apache Web Server (Apache from here on) is the most well-known of all web servers. It's important to understand that although it's quite easy to get started with Apache out of the box, it's not optimized for using CPU and RAM efficiently. Apache is quite bulky, with several modules that provide various chunks of functionality.

While Apache does an admirable job in serving dynamic sites (Linux + Apache + MySQL + PHP constitute the well-known application framework known as LAMP), it's really not ideal for sites that just want to serve a lot of static files at high speed. You can remove most of the modules and streamline Apache for these types of sites, but why do so when you can use a web server that's been explicitly designed to do that type of work (serve static files at high speeds)? In this type of scenario, NGINX might be the better choice.

The NGINX Web Server

While the Apache web server is still the most common web server used in the world, there are questions as to its scalability and performance when dealing with modern website architectures. The NGINX (pronounced "Engine-X") web server, an open source web server introduced in 2004, is fast becoming the go to web server for providing high performance and high concurrency, while using a low amount of memory. Right now, NGINX is the number two web server on the internet, behind the Apache web server, and moving up very fast.

Note that NGINX isn't just a high performance HTTP server – it's also a reverse proxy, as well as an IMAP/POP3 proxy server.

Benefits Offered by Nginx

The NGINX web server or hardware reverse proxy both offer great functionality and performance. NGINX offers the following properties which are highly useful:

- Nginx can work as a load balancer and contains advanced features such as Web-Sockets and throttling.

- You can also configure NGINX to override HTTP headers, which lets you apply HTTP caching to applications that don't implement caching headers correctly (or even use it to override configured caching policies).
- Since Nginx is also FastCGI server, you can run web applications on the same web server stack as the reverse proxies.
- Nginx is highly efficient since it uses asynchronous processing always: this allows the server to proxy tens of thousands of simultaneous web connections with an extremely minimal connection overhead.

Nginx is getting more popular day by day when compared to other web servers, and its' share of the internet uses is growing at a fast clip. Netcraft reported that in September 2015, while Apache and Microsoft both lost market share, NIGIX grew its share of total web sites on the Internet to about 15%. This is a truly phenomenal number for a web server as new as NGINX.

Why High Concurrency is Important

Web sites aren't any longer the simple affairs they used to be a few years ago – today's web applications need to be up 24/7, and in addition to e-commerce, provide entertainment (movies, gaming etc) and various types of live information to users.

Concurrency has always been a challenge when architecting these complex web applications. Increasing social web and mobile usage means that there are a large number of simultaneously connected users using a web application. On top of this, modern web browsers enable you to open multiple simultaneous connections to the same website to enhance the page loading speed.

Clients often stay connected to web sites to reduce the hassles of having to open new connections. The web server must therefore account for enough memory to support a large number of live connections. While faster disks, more powerful CPUs, and more memory certainly do help in providing higher concurrency levels, the choking point for high demand web applications is going to be the web server, which should be able to scale nonlinearly as the number of simultaneous connections and requests per second creeps higher and higher.

The Apache web server, under its traditional architecture, doesn't allow a website to scale nonlinearly. Each web page request leads to Apache spawning a new process (or thread). Thus, to scale, Apache spawns a copy of itself for every new connection, making it a platform not easily amenable to scaling, due to the high memory and CPU demanded by its architecture.

In 2003, Daniel Kegel came up with the C10K manifesto. The C10K manifesto was a call to architect new types of web servers that can handle ten thousand simultaneous connections. The NGINX server is a direct outcome of the attempts to answer the call of the C10K manifesto, and certainly meets the challenge posed by Kegel.

The NGINX server was architected with scalability as its main focus. NGINX is event based, so it doesn't spawn new processes/threads for each new web page request. As the memory and CPU usage doesn't increase linearly with workloads, the web server can easily support tens of thousands of concurrent connections.

I might add here that Apache has also enhanced its capabilities with the release of the Apache 2.4x branch. New multiprocessing core modules and new proxy modules designed to enhance scaling, performance and resource utilization have been added to compete better with event-driven web servers such as NGINX.

Apache and NGINX have their relative strengths, so it's not really one or the other in most cases. Since NGINX is great at serving static content, you can deploy it in the front end to serve pre-compressed static data such as HTML files, image files and CSS files, etc. You may even take advantage of its load balancing capabilities with its easy to set up reverse proxy module. The reverse proxy module lets NGINX redirect requests for specific MIME types of a web site to a different server. At the same time, you can deploy one or more Apache web servers for delivering dynamic content such as PHP pages. The reverse proxy module will redirect the request for dynamic data to the Apache server.

Caching Proxies and Reverse Proxying

Web proxy caching is the storing of copies of frequently used web objects such as documents and images, and serve this information to the user. The goal is to speed up service to the clients, while saving on the Internet bandwidth usage. There are two types of proxies – caching proxies and reverse proxies, as explained in the following sections.

Caching Proxies

A caching proxy is a server that caches HTTP content. Local caching proxy servers sit in the network, in between your web servers and the clients and all HTTP traffic is routed through the caching servers so they can cache as many web requests as possible. Caching proxies use what's called a read-through cache to store web responses which can be reused by sending them to multiple servers instead of generating new web traffic for each user. More than corporate networks, it's Internet Service Providers (ISPs) that install caching proxies to reduce the web traffic generated by users.

Local proxy servers are much less popular these days since network bandwidth has become way cheaper over time. Also, most websites use the Secure Sockets Layer (SSL) exclusively to serve content. Since SSL communications are encrypted, caching proxies can't intercept those requests, since they can't encrypt and decrypt messages without the required security certificates.

Reverse Proxies

A reverse proxy doesn't reverse anything – it does something as a caching proxy! These types of proxy servers are called reverse proxy servers because of their location – you place them inside your data centers to cache the response of your web servers and thus reduce their workload. Reverse proxies help you scale in many ways:

- You can employ a bank of reverse proxy servers since each proxy is independent of the others, thus promoting your application's scalability
- They can reduce the requests reaching the web servers, if you are able to cache full pages
- A reverse proxy helps you control both the request types to be cached and the duration for which the requests are cached.
- By placing a bank of reverse proxy servers between the front end web servers and the web servers hosting the web services, you can make the web service layer much faster by caching the web service responses.

As explained elsewhere in this Chapter (in the section on the NGINX web server), NGINX has several features which make it an excellent proxy server. If you've chosen to use a hardware load balancer in your data center, you can make it perform double duty by letting it act as a reverse proxy server as well. If not, an open-source reverse proxy technology solution such as NGINX, Varnish or Apache Traffic Server will work very well. You can scale any of these by simply adding more instances.

High availability and Keepalived

High availability means that an application restarts or reroutes work to another system automatically when the first server encounters a failure. In order to setup a highly available system, the system must be able to redirect the workload, and there also must be a mechanism to monitor failures and automatically transition the system to a healthy server when service interruptions are sensed.

Keepalived is a Linux –based routing software that lets you achieve high availability. Keepalived assigns multiple nodes to a virtual IP and monitors those nodes, and fails over when a node goes down. Keepalived uses the VRRP (Virtual Router Redundancy Protocol) protocol, for supporting high availability.

VRRP is a network protocol that automatically assigns available IP routers to participating hosts. This lets you increase the availability of routing paths through automatic default gateway selections on an IP subnetwork. VRRP allows Keepalived to perform router failover and provide resilient infrastructures. The VRRP protocol ensures that one of the nodes in the web server infrastructure is the master, and assigns it a higher priority than the other nodes. The others are deemed backup nodes, and listen for multicast packets from the master.

If the backup nodes don't receive any VRRP advertisements for a specific interval, one of the backup nodes takes over as the master and assigns the configured IP to itself. If there are multiple backup nodes with the same priority, the node with the highest IP wins the election. You can set up a highly available web service with Keepalived. Basically what you're doing is configuring a floating IP address that's automatically moved between web servers. If the primary web server goes down, the floating IP is automatically assigned to the second server, allowing the service to resume. Keepalived is often used together with HAProxy to provide redundancy, since HAProxy doesn't come with its own redundancy capabilities.

Handling Data Storage with Databases

Databases are ubiquitous in any IT environment and are used for storing various types of information. There are several types of databases, based on the type of information they store as well as the types of data retrieval they allow, as explained in the following sections. Most organizations today use multiple types of databases, with a mix of relational and NoSQL databases.

Relational databases

A relational database management system (RDBMS) is the most well-known type of database, and still remains the predominant type of database in use, even with the advent of several newer types of databases. Most applications such as shopping, sales, and human resources are handled best by a relational database.

Relational databases store information in the form of relationships among various entities (such as employees and managers), and are extremely efficient for querying and storing transactional data. Oracle, IBM's DB2 and Microsoft SQL Server are the leading commercial relational databases and MySQL, PostgreSQL, MariaDB (very similar to the MySQL database) are the leading open source relational databases.

Commercial relational databases such as Oracle require full-fledged administrators owing to their complexity. On the other hand, it's not uncommon for organizations to look to the Linux system administrator to help out with the administration of open source databases such as MySQL and PostGres (and especially the newer NoSQL databases such as MongoDB, CouchDB and Cassandra), and perform tasks such as installing, backing up and recovering the database.

Other Types of Databases

In today's world, while relational databases are still important, there are several new types of databases that are increasingly becoming important, as explained in the following sections.

NoSQL Databases

NoSQL databases are useful for handling unstructured data and are especially good at handling large quantities of such data. In this chapter, I discuss MongoDB and Cassandra, two very popular NoSQL databases that you ought to be familiar with. Later chapters discuss both of these databases in more detail. In addition, there are several other popular NoSQL databases such as CouchDB.

Caching Databases

Memcached and Redis are two well-known caching databases. I explain the role of both of these data stores in this chapter, and discuss them in more detail in Chapter 9.

Cloud Databases

Cloud databases are hosted by cloud providers such as Amazon and Google for their cloud customers. Amazon offers Amazon DynamoDB, a powerful NoSQL database, and Google, the Google Cloud Storage and Google Cloud Datastore. Chapter 9 explains Amazon's DynamoDB and RDS Datastore. Cloud databases take the complexities of database management out of an organization's hands, but they've their own headaches, such as security issues, for example, since the data is hosted by the cloud provider.

In addition to the types of databases listed here, there are many other less well known database types, such as object databases (for example, db4o), and NewSQL databases such as VoltDB and Clustrix, as well.

It's important to understand that most organization today use several types of databases. This happens to be true not only for a large environment, where you expect many databases of many different types, but even for smaller organizations, especially those that deal with web applications.

MongoDB as a Backend Database

MongoDB is a highly scalable open source NoSQL database that provides high performance for document-oriented storage. The following sections explain the salient features of MongoDB.

A Document Database

Document-oriented storage means that instead of using traditional rows and columns like a relational database, MongoDB stores data in the JSON document format. Instead of tables, you have collections, and documents play the role of rows.

MongoDB uses collections to store all documents. Although collections are called the counterparts of relational database tables, they are defined entirely differently since

there's no concept of a schema in MongoDB. A collection is simply a group of documents that share common indexes.

Dynamic Schemas

Relational databases are all about schemas – the data you store in the database must conform to an existing schema, which describes the type of data you can store, such as characters and integers, and their length. MongoDB doesn't use schemas – you can store any type of JSON document within any collection. This means that a collection can have disparate types of documents in it, each with a different structure. You can also change the document structure by updating it – no problems there.

Automatic Scaling

MongoDB lets you scale horizontally with commodity servers, instead of having to go the more expensive route of vertical scaling, as is the case with a relational database. Although you can create clusters of relational databases, that's mostly for high availability, and not for enhancing performance. Scaling with multiple servers in a MongoDB database, on the other hand, is expressly designed for enhanced performance. There are two aspects to the automatic (or horizontal) scaling:

- Automatic sharding: helps distribute data across a cluster of servers
- Replica sets: provide support for low-latency high throughput deployments

High Performance

MongoDB supports high performance data storage, by supporting embedded data models that reduce I/O activity. It also supports indexing to support faster queries – you can include keys from embedded documents and even arrays.

High Availability

MongoDB uses a replication facility called replica sets, to provide both automatic fail-over and data redundancy. A replica set is defined as a set of MongoDB servers that store replicas of the same data set, enhancing both redundancy and data availability.

For Single Page Applications (SPAs), which I discuss elsewhere in this chapter, MongoDB's document-

MongoDB and the CAP Theorem

As with other document databases and other databases that fall under the broad umbrella of NoSQL Databases, MongoDB doesn't offer support for the well-known ACID properties provided by traditional relational databases. ACID properties are the hallmark of all modern transactional databases (such as MySQL and Oracle), and refer to the following set of principles expected of a database:

- **Atomicity:** This is the all-or-none principle which requires that if even one element of a transaction fails, the entire transaction fails. The transaction succeeds only if all of its tasks are performed successfully.
- **Consistency:** This property ensures that transactions are always fully completed, by requiring that the database must be in a consistent state both at the beginning and at the end of a transaction.
- **Isolation:** Each transaction is considered independent of the other transactions. No transaction has access to any other transaction that's in an unfinished state.
- **Durability:** Once a transaction completes, it's "permanent". That is, the transaction is recorded to persistent storage and will survive a system breakdown such as a power or disk failure.

While the ACID requirements served traditional relational databases just fine for many years, the advent and eventual popularity of non-relational data such as unstructured data, non-relational data, and the proliferation of distributed computing systems led to new views about the required transaction properties that modern databases must satisfy. The Consistency, Availability and Partition Tolerance (CAP) theorem sought to refine the requirements to be met for implementing applications in modern distributed computing systems. The CAP theorem actually stands for the following three principles:

- **Consistency:** this is the same as the ACID consistency property. Satisfying this requirement means that all the nodes in a cluster see all the data they're supposed to be able to view.
- **Availability:** the system must be available when requested.
- **Partition Tolerance:** failure of a single node in a distributed system mustn't lead to the failure of the entire system. The system must be available even if there's some data loss or a partial system failure.

The problem is, at any given time, a distributed system can usually support only two out of the three requirements listed here. This means that tradeoffs are almost always inevitable when using distributed data stores such as the popular NoSQL databases.

For database reliability purposes, meeting the Availability and Partition Tolerance requirements is absolutely essential, of course. That means that the consistency requirement is often at risk. However, leading NoSQL databases such as Cassandra and Amazon's DynamoDB deal with the loss of consistency just fine. How so? This is possible due the adoption by these databases of something called the BASE system, which is a modified set of ACID requirements to fit modern NoSQL and related non-relational databases. Here's what BASE stands for:

- **Basically Available:** the system guarantees the availability of data in the sense that it'll respond to any request. However, the response could be a “failure” to obtain the request data set, or the data set returned may be in an inconsistent or changing state.
- **Soft:** the state of the system is always “soft, in the sense that the “eventual consistency” (the final requirement) may be causing changes in the system state at any given time
- **Eventually Consistent:** The system will eventually become consistent once it stops receiving new data inputs. As long as the system is receiving inputs, it doesn't check for the consistency of each transaction before it moves to the next transaction.

Amazon's DynamoDB, for example, which lies behind Amazon's shopping cart technology, stresses high availability, meaning it can afford to go easy on the consistency angle. These types of databases eschew the complex queries necessary to support consistency in the traditional sense, settling instead for the eventual consistency goal. Eventual consistency in this context means that in a distributed system, not all nodes see the same version of the data – at any given time the state may diverge between nodes – that is, it's possible for some nodes to serve stale data. However, given sufficient time, the state will come to be the same across the system.

MongoDB, a popular NoSQL database, on the other hand, favors consistency and partition tolerance over high availability.

The main point to take away from this discussion of CAP and BASE is that while NoSQL databases have their advantages, particularly in the way they support horizontal scaling and the efficient processing of non-relational data, they do come with unique drawbacks and involve crucial sacrifices in terms of simultaneous support for traditional principles such as data consistency and availability.

Consistency, while it's a laudable objective that a database can satisfy, has a negative impact on cost effective horizontal scaling. If the database needs to check the consistency of every transaction continuously, a database with billions of transactions will incur a significant cost to perform all the checks.

It's the principle of eventual consistency that has allowed Google, Twitter and Amazon, among others, to continuously interact with millions of their global customers, keeping their systems available and supporting partition tolerance. Without the principle of eventual consistency, there wouldn't be all these systems today that deal successfully with the exponential rise of data volumes due to cloud computing, social networking and related modern trends.

Caching

A cache is simply a set of data that you store for future use. Caching is a key technology that web applications use to increase both performance and scalability. Caching can play a significant role in scaling your applications, since you need fewer computing resources to serve an ever growing customer base. Caching is a common technique that's used at almost every level of the application stack, including operating systems, databases, HTTP browser caches, HTTP proxies and reverse proxies, as well as application object caches.

Not everything can be cached – and should be cached! Good candidates for caching are objects that won't become invalidated with the passage of a short period of time. If users are frequently updating the data, it's not feasible to cache that data as the data becomes quickly invalidated.

Regardless of the specific caching strategy you employ, application code that seeks to reuse cached objects for multiple requests or several users is always going to make the responses faster, besides saving resources such as server CPU/RAM/, and network bandwidth. Even a design that only caches just page fragments instead of entire web pages will provide terrific performance benefits.

For web applications, you need to be concerned about two main types of caches – HTTP-based caches and custom object caches.

HTTP-Based Caching (Browser caching)

In a web application stack, the HTTP-based cache is the most common and predominant component. In order to understand HTTP based caching, it's important to first understand something about the all-important HTTP caching headers. HTTP caching is common, and often a web request makes use of several HTTP caches linked to each other.

HTTP caches are read-through caches. In a read-through caching architecture, clients connect to the cache first to request resources such as web pages or CSS files. The cache will return a resource to the client directly if it's in the cache, and if it's not there, contacts the originating server and fetches the data for the client. Read-through caches are transparent to the client, who's unaware as to where the resource is being sent from, and they can't distinguish between a cached object and an object they get by connecting directly to the service.

Letting the browser cache the content dramatically increases performance in many cases. You are storing the files on the client machine itself in this case. Sometimes the client server may even store previously rendered web pages, completely bypassing the network connection to the server (this is the reason why when you press the “back” arrow, the previous page appears so quickly on your screen!) and HTTP caches help

lower the load on you data center and move it to external servers that are located near the users. Not only that, but caching also results in faster responses to user requests.



HTTP headers not only help cache web pages and static resources, but also the responses of web services. You can insert an HTTP cache between the web services and client. In fact, REST based web services excel at caching web service responses.

How HTTP Caching is Used

At one extreme, you can cache responses indefinitely. Static content such as images and CSS usually is cached “forever”. You can cache static content such as this forever and use new URLs for newer versions of the static files, thus ensuring that users are always using the correct (that is, compatible) HTML, CSS and JavaScript files. On the other extreme, you can ensure that a HTTP response isn’t ever cached, by using a HTTP header such as Cache-Control: no-cache.

Types of HTTP Caching

You can configure two types of HTTP caches, as explained here:

- **Browser cache:** All web browsers use a browser cache based on disk storage as well as memory, to avoid resending HTTP requests for resources that have already been cached. This caching helps load web pages faster and keeps resource usage low.
- **Caching proxies:** Caching proxies (also called web proxies) are servers that sit between the internet and the users (that is, the web browsers) and use a read-through cache to reduce the web traffic generated by the network’s clientele.

A Web Cache is also called a proxy server, and is an entity that satisfies HTTP requests on behalf of an origin Web server. Rather than each browser connecting directly to the internet, they relay their requests through the proxy. The proxy server is the only one that initiates internet connections. The clients connect to the proxy server, which’ll forward the requests to the target web server only if it doesn’t find the content in its cache.

Both ISPs (Internet Service Providers) and corporate networks (local proxy servers) can use a caching proxy, with ISPs seeking to maximize the caching of web requests to lower the web traffic generated by users.

Local proxy servers have one big advantage over the ISP proxy servers - when you use SSL, the external caching servers can’t intercept the requests since they lack the required security certificates to read the messages.

Good proxy servers can dramatically reduce the network bandwidth usage in an organization, as well as improve performance – so you get to run things faster, but at a lower cost! Web caches can significantly cut down on the Web traffic, keeping you from having to upgrade your bandwidth, and improving application performance. Usually, an ISP purchases and installs the Web Cache. A company for example, may install a Web cache and configure all of its user browsers to point to the Web cache, so as to redirect all the user HTTP requests from the origin servers to the Web Cache.



Content Delivery Networks (CDN) are functionally equivalent to the caching proxies, since they also rely on HTTP headers for caching – the big difference is that the CDN service provider manages all the caching.

- **Reverse proxies:** These are servers that work quite similar to caching proxies, with the difference being that the caching servers are located within the data center itself to reduce the load on the web servers. The reverse proxy server intercepts web traffic from clients connecting to a site and forwards the requests to the web server if they can't serve it directly from their own cache. Reverse proxies are great for overriding HTTP headers and manage request caching easily. They also help you scale REST based web services by sitting between the front-end and the web-service servers, lowering the amount of requests that your web service web servers need to serve.

Benefits of a Web Cache

Web caches help you in two ways; they reduce the response time for client requests and also reduce the amount of traffic flowing through your company's access link to the Internet. Let's use some hypothetical numbers to breakdown the total response time of a web request to understand how caching can help you big time. Total response time is the time it takes to satisfy a browser's request for an object, and is composed of the following three types of delays:

- LAN delay
- Access delay
- Internet delay

Let's see how a Webcache can make a difference in keep Web response times very low. Let's assume the following:

- Your company's network is a high speed LAN and a router in the Internet connects to a router in the Internet via a 15 Mbps link.

- The average size of an object is 1 Mbits.
- On average, about 15 requests are made by your company's user browses to the origin Web servers (no Webcache at this point)
- The HTTP messages are on the average quite small, and therefore, there's negligible traffic between your router and the Internet router.
- The Internet side router takes on average 2 seconds to forward an HTTP request and receive the response – this is called the Internet delay.

You can measure the LAN delay by looking at the traffic density on the LAN:

$$(15 \text{ requests/second}) * (1 \text{ Mbits/request}) / (100 \text{ Mbps}) = 0.15$$

A low traffic intensity such as 0.15 means that you can ignore the LAN delay, which will be in tens of milliseconds duration at most.

The access delay however is significant, as the traffic density on the access link which is from your own router to the Internet router, is:

$$(15 \text{ requests/second}) * (1 \text{ Mbits/request}) / (15 \text{ Mbps}) = 1$$

An access delay of 1 is bad – it means that the access delays are large and grow without bound. You're going to end up with an average response time for HTTP requests that'll be several minutes long, which is something you just can't have.

One way to bring down the total response time is to upgrade your access link from 15 Mbps to 100 Mbps (quite expensive), which will lower the traffic intensity to 0.15. Thus means that the access delay will be insignificant, just as the LAN delay is, meaning that the response time will be comprised of just the Internet delay, which is 2 seconds.

The other alternative doesn't involve an expensive upgrade of your access links – you just install a Web cache in your network! You can use open source software and commodity software to do this.

Let's say the hit ratio of the cache is 0.4. That means 40% of all web requests are satisfied directly from the cache and only 60% of the requests are sent to the origin servers. Since the access servers now handle only 60% of the traffic as compared to its earlier traffic volume, the traffic intensity on the access link goes down to 0.6 from its previous value of 1.0. On a 15 Mbps link, a traffic intensity that's less than 0.8 means the access delay is negligible (in tens of milliseconds).

The average response time now is:

$$0.4 * (0.01 \text{ seconds}) + 0.6 * (2.00 \text{ seconds}) = 1.2 \text{ seconds (approx.)}$$

As you can see, a Web Cache leads to dramatically low response times, even when compared to a solution that requires a faster (and much more expensive) access link.

There really isn't much for you to do if you farm out HTTP caching to a CDN. However, since reverse proxy servers are run from your own data center, you need to manage them yourself. A hardware load balancer, as mentioned earlier, can also provide reverse proxying services. You can also use an open source solution such as NGINX or Varnish. NGINX for example is famous for handling tens of thousands of requests every second per each instance. The essential keys to effective management of proxying then are the type and size of the cached responses and the length of time for which you want to cache them.



A Content Distribution Network (CDN) uses distributed web caches throughout the Internet, helping localize the web traffic. You can use a shared CDN such as Akamai and Limelight to speed up the response times of web requests.

Caching Objects

Caching objects is another key way to boost web application performance. Object caching falls under the category of application caching, where the application developers assume the responsibility for the content to be cached, as well as the time for which the content must be cached.

Caching stores such as Memcached let you set expiry times for web documents and also let you update Memcached when new web pages are added to your site. For example, if you store a web page in Memcached and that content remains valid for one minute, all your users will be sent data directly from Memcached and your web servers won't have to handle that load. The web servers simply generate a new page every minute and users get that new page when they refresh their browser. Since only the application developer knows the application well enough to determine the content and duration of the cached objects, object caching isn't something that's as simple as browser caching.



A cache is a server or a service geared towards reducing both the latency and the resources involved in generating responses by serving previously generated and stored content. Caching is a critical technique for scaling an infrastructure.

Caching application objects is done in a different way than HTTP content caching. Applications exposed over HTTP can be cached in a variety of ways. As with HTTP caches, there are several types of object caches: client-side caches, local caches and distributed object caches. The first two are simple affairs, as described here:

- **Client-side Caches:** Today's web browsers, which all use JavaScript, can store application data such as user preferences, for example, on a user's laptop or on mobile devices. Caching data on the user's devices of course makes web applications run faster and also lowers the load on your servers. JavaScript running in the browser can retrieve the cached objects from the cache. Note that the users can always wipe the cache clean. Single-page applications (SPAs) that we learned about earlier in this chapter, benefit from the client-side cache since they run a lot of code within the browser, and also rely on asynchronous web requests (AJAX). This is especially true for SPAs explicitly designed for mobile devices.



Caching objects in local memory lets applications access resources extremely fast, since there's no network latency.

- **Local Caches:** A local cache is located on the web servers. There are a few different ways to cache objects locally, but all of them employ the same strategy: store the objects on the same server where the application code is running. One way is for the front end and back end applications to use local caches by caching application objects in a pool. The application incurs virtually no cost when accessing cached objects since they're stored right in the application's process memory. Some languages such as PHP (but not Java, which is multithreaded) can also use shared memory segments that allow multiple processes to access the stored objects in the cache. Alternately, you can deploy a local caching server with each web server.

Memcached and Redis are two very popular caching servers and you'll learn more about them in the next section. When you're just starting out, or if you have to manage a small web application, you can run both the webserver and the caching server on the same machine to cut costs and make things run faster since the network roundtrips are very short.

Local Caching

Local caching, as you guessed, doesn't involve an external caching database. Local caching, regardless of how you want to implement it in practice, is very easy to set up and doesn't involve too much complexity. They're also very low latency solutions and don't involve any issues with locking etc. however, there's no synchronization between the local application caches when you employ a bunch of them.

Lack of synchronization among the caches on different servers means that you run the risk of inconsistent data because of different cached objects representing different

values for the same information. The application servers don't coordinate the storing of the cached objects, leading to redundancy and duplication of the objects stored in the cache. Fortunately, there's a caching solution that addresses all of these issues with local object caches – a distributed object cache, which I discuss next.

Distributed Object Caching – Memcached and Redis

A distributed object cache works exactly the same as a local object cache – both use a simple key-value store where clients can store objects for a set duration of time. The big difference is that unlike a local cache, a distributed cache is remotely hosted. Redis and Memcached are popular open source solutions that serve a wide variety of uses. Memcached is a fast key/value store that's very simple and lacks too many bells and whistles, making it very easy to implement. It's used in several of the busiest web sites in the world today.

Distributed caching servers are easy to work with through any programming language. Here's an example showing a caching interface in PHP:

```
$m = new Memcached();  
$m->addServer('10.0.0.1', 11211);  
$m->set('UserCount', 123, 600);
```

This code does the following:

- Sets Memcached as the caching server
- Sets the IP for the cache server
- Sets the caching data: the name of the object to be cached, its key value and the TTL (duration for which it'll stay in the cache) for that object

A cache server such as Memcached or Redis is really a database, and thus offers most of the capabilities of a key-value store, such as replication, query optimization and efficient use of memory.

Caching servers aren't an alternative to the regular databases – all you're doing is when you generate new database content from a relational database for example, you store the content in Memcached or Redis for future use. For subsequent requests for that data, you first check the cache and if the data is there, you send the content back without having to call the database.

While caching is easy to setup and manage, there are situations where it isn't ideal. Transactions constantly modify, or add and delete the data in the database tables. If the user must absolutely, positively get the very latest version of the data, caching becomes problematic, as it tends to lag behind the current state of the database.



If you're using Amazon AWS to host your applications you can use Amazon Elastic Cache, which is Amazon's hosted cache cluster that uses Memcached or Redis (you get to pick), or set up your own caching service on your EC2 instances.

You've learned the basics of caching in this section, and how it helps you scale as well as make your applications run faster. While there are several types of caching, not all caches are equal, however! The earlier in the request process for a web page or data an application satisfies the request by retrieving it from a cache, the more beneficial is the cache.

Caching servers are typically deployed on dedicated machines, and you usually deploy them in clusters, with multiple caching servers. You can implement replication or data partitioning to scale, once you exhaust the limits of adding more memory to the caching servers. For example, if you're using Memcached as your object cache, you can partition the data among multiple Memcached servers in order to scale up.

When formulating a caching strategy it's important to remember that you can't (and don't want to) cache everything. You must choose which web pages or services you need to cache. To do that, use some type of metrics. Refer to the metrics example I provided earlier in this chapter (page 11), which showed how third-party services such as Google Analytics can help you maximize the potential gains from caching.

Asynchronous Processing, Messaging Applications and MOM

In traditional software execution, synchronous processing is the standard way to perform operations. A caller such as a function, a thread, or a process sending a request to another process, or an application making requests to a remote server won't proceed to the next step until it gets the response for its request. In other words, the caller waits for a response or responses before continuing further execution. Under an asynchronous processing model, however, a caller doesn't wait for the responses from the services it contacts. It sends requests and continues processing without ever being blocked.

Let me take a simple example here to demonstrate the difference between the synchronous and asynchronous approaches. Let's say an application's code sends out an e-mail to a user. Under the synchronous processing model, the code needs to wait for the e-mail service to do what it takes to send that email out: resolve the IP addresses, establish the network connection, and finally send the email to an SMTP server. The email service needs to also encode and transfer the message and its attachments, all of which takes some time (a few seconds). During this time, the execution of the appli-

cation code pauses, a pause that's often referred to as blocking, since the code is waiting on an external operation to complete.



Platforms such as Node.js come with built in asynchronous processing capabilities, making messaging and message brokers less useful than for other platforms such as Java and PHP, which aren't asynchronous.

Obviously, the synchronous processing model isn't conducive to building responsive applications: not only is this a slow method but it's also highly resource intensive, since blocked threads continue to consume resources even when they're just waiting. Total execution time is the sum of the time taken to perform each of the operations in a service, and because your app is doing things serially rather than parallelly, blocking operations slow down your application response times. The benefits that accrue from asynchronous processing are really due to the fact that your applications and services perform work parallelly instead of sequentially, thus bringing into play vastly higher resources to process the workloads. Execution times are rapid and user interest is unlikely to flag under such architectures.

Responsive applications are really what it's all about, as asynchronous processing helps build these types of applications since it doesn't involve blocking operations.

Messages and Message Queues

Essentially, a message is a piece of code written in XML or JSON that contains instructions for performing the asynchronous operation. A message queue (managed by a message broker) allows you to reap the benefits of synchronous processing. Asynchronous processing can be beneficial in a situation where you want to keep clients from waiting for time consuming tasks to complete – ideally, the client should be able to continue their execution with no blocking.

Here are some typical use cases for using asynchronous processing and message queues:

- Any operation that requires a heavy amount of resources such as the generation of heavy reports can be sent to a message queue.
- Anytime you need to perform operations on a remote server that takes some time to complete, an asynchronous model of processing can speed up things.
- Any critical operations such as placing orders or processing payments can't be expected to wait for less critical operations to complete. You can divert the less important parts of the operation to a message queue so they can be processed asynchronously by a separate message consumer.

Components of a Messaging Architecture

Messaging revolves around the following three crucial constituent operations:

- **Producing:** this simply means the process of sending messages. A producer is a program that creates and sends messages to a message queue. Message producers are often also called message publishers.
- **Queuing:** a message queue is a location to store messages, and is part of the message broker. Multiple producers can send messages to a queue and multiple consumers can retrieve messages from a queue.
- **Consuming:** means the same as receiving messages. A message consumer is a program that waits to receive messages sent by a producer. It's the message consumer that actually performs the asynchronous operation for the message producer. It's common for producers, consumers, and queues to reside on different servers, and even use different technologies.

Message queues are the heart of asynchronous processing, so let's learn more about them in the next section.

Message Queues

The message queue stores and distributes asynchronous requests. As the producers create new messages and send them to the queue, the queue arranges those messages in a sequence and sends them along to the consumers. Consumers then will act on the messages, by performing the asynchronous actions.

Message Queues and Asynchronous Processing

Message queues provide the essence of asynchronous processing – non-blocking operations - and therefore, non-blocking I/O. The message producers and message consumers work independently, neither of them blocking the other, nor with either being aware of the other. One of them dedicates itself to creating messages and the other to the processing of those message requests.

Message queues not only enable asynchronous processing, but are useful in evening out spurts in message traffic and in isolating failures. During times of heavy traffic, the message queues continue to accept the high traffic and keep queuing the messages. The front end application may produce a large number of messages which the back end component will consume, but the user isn't affected since the front end servers aren't waiting for the operations to complete.

A key advantage of message queues is that they promote scalability – you can employ banks of message producers and consumers on dedicated servers and just keep adding additional servers to handle a growing workload.

How Message Queues are Implemented

At its simplest, a message queue is just a thread running within the main application process. You can also implement it by storing the messages in the file system or in a relational database such as MySQL.

For heavy duty message processing as well as for tackling on additional features such as high availability on to the message queue, you need a full-fledged specialized messaging application called a message broker or message oriented middleware, which is the topic of our next section.

Message Brokers and Message Oriented Middleware (MOM)

A message broker is a dedicated application that provides not only message queuing, but also the routing and delivery of the messages.

Message brokers relieve you from having to custom write code for providing critical messaging functions – you simply configure the broker to get the functionality you desire.

Message broker software is also often called message-oriented middleware (MOM) or an enterprise service bus (ESB). However, a Message Broker is really a higher level concept that's built on top of MOM (or an ESB), with the MOM providing the underlying services such as message persistence and guaranteed delivery.



A message broker adds others things to a MOM, such as rules based business process integration, content based routing, data transformation engine, etc.

What Message Brokers do

A message broker accepts and forwards messages. Just as a post office accepts your mail and delivers it to whomever you address the mail to, a message broker such as RabbitMQ accepts, stores, and forwards data, with the data in this context being messages.

The message broker performs a critical function in asynchronous processing – it separates the consumers from the producers. One of the great things that message brokers do very well is the fast queuing and dequeuing of large volumes of messages, since they're optimized for high throughput.

Messaging Protocols

You can specify how clients connect to the broker and how messages are transmitted, by selecting a correct messaging protocol. Messaging protocols control the transmis-

sion of messages from producers to consumers. I briefly describe the most commonly employed protocols here:

- **Streaming Text–Oriented Messaging Protocol (STMP):** This is a rudimentary messaging protocol that has low overhead, but has no advanced features such as those offered by other commonly used protocols.
- **Advanced Message Queuing Protocol (AMQP):** AMQP is a comprehensive messaging protocol that has been well accepted as an industry standard for messaging. The best thing about AMQP is that it's a standardized protocol, and hence is easy to integrate into your software. AMQP offers features such as guaranteed delivery of messages, transactions and many other advanced features.
- **Java Messaging Service (JMS):** JMS is a powerful messaging standard but it's confined to Java technologies such as Java or Scala.

Of the three messaging protocols listed here, AMQP is the most commonly used protocol among enterprises.

Pull versus Push Messaging

Developers can configure message consumers to work in two different ways: pull or push:

- The pull model, also known as a cron-like model, lets the message consumer pull messages off the message queue. For example, in an Email service, the consumer picks up the messages from the queue and uses SMTP to forward the emails to the mail servers. Applications like PHP and Ruby use a pull model whereby the consumers connect to the queue once in a while and consume the available messages (or a set number of messages). Once they consume the messages, the consumers detach themselves from the queue.
- In the push model, the message consumer is always connected to the message broker, and maintains an idle wait by blocking on the socket read operation. The message broker will push new messages through the permanent connection, as fast as the consumer can handle them. Applications such as Java and Node.js which use persistent application containers to maintain constant connections to the message brokers.

Subscription Methods

Message consumers can be configured to use various subscription methods to pull messages off the message queues maintained by the message broker. Here are the two common subscription methods:

- **Publish/subscribe (pub/sub):** In the pub/sub subscription method, the broker publishes messages not to a message queue, but to a topic. A topic in this context is something like a channel. Consumers connect to the broker and let it know the topic or topics they want to subscribe to. The broker transmits all messages published to those topics to the consumer that requested them, and the consumers receive the messages in their private queue. Thus, in this mode, there's a dedicated message queue for each consumer, containing just those messages that were published on a specific topic or topics.
- **Direct worker queue:** Producers send all messages to a single queue, with each message routed to a specific consumer. So, multiple consumers share the single worker queue. Tasks that take considerable time such as sending emails and uploading content to external services to multiple consumers benefit from this method.

Popular Message Brokers

There are several powerful, robust and scalable open-source message brokers and two of the most popular of these are RabbitMQ and ActiveMQ. If you are hosting your operations in the Amazon cloud, you can also use Amazon's Simple Queue Service (SQS), which I discuss in Chapter 9.

Your choice of a message broker solution should be guided by your use cases and on key criteria such as the volume of message, message size, rate of message consumption, concurrent producers/consumers. Two important factors are whether you want to ensure that messages are safeguarded from loss by storing them, and whether you need message acknowledgment.

RabbitMQ is an open source message broker written in the Erlang programming language, and is currently the leading open source messaging broker.

One of the best things that sophisticated message brokers such as RabbitMQ and ActiveMQ offer is their ability to create custom routes. Routes determine which messages are sent to a queue. Earlier, you learned about the publish/subscribe and direct worker queue subscription methods. RabbitMQ lets you create flexible routing rules by matching text patterns. For example, you can create queues that capture all the error messages from a system and a consumer that will then consume these messages and send notifications to appropriate teams. Alternately, a consumer could write messages from a queue to a file.

Both RabbitMQ and ActiveMQ contain roughly the same features, and offer similar performance benefits.

Powerful message brokers such as ActiveMQ and RabbitMQ come with out of the box capabilities for setting up custom routing schemes and other things. Instead of

rewriting or modifying the producer/customer code, all you have to do is just configure the message broker to take advantage of these features.

As chapter 9 explains, Amazon's SQS (Amazon Simple Queue Service), while not perfect, offers benefits not available with either RabbitMQ or ActiveMQ.

The Model-View-Controller Architecture and Single Paged Applications

Most web applications retrieve data from a data store of some type (relational database, NoSQL database or a flat file) and send it to the users. The application also modifies stored data or adds new data or removes data according to what the user does in the user interface.

Since on the face of it everything revolves around the interactions between the user interface and the data store, an obvious design strategy is to link these two together directly to minimize programming and improve performance. Not so fast! While this seems to be “logical” approach, the fact that the user interfaces change frequently and given the fact that applications often incorporate business logic that goes well beyond the mere transmission of raw data from the data stores, means that one ought to look at a much more sophisticated application design.

The Problem

By combining the presentation logic with the business logic, you lose one of the biggest advantages of web based applications, which let you change the UI anytime you want, without worrying about having to redistribute the applications. It's well known that in most web applications, user interface logic changes much more often than business logic. If you combine the presentation code and business logic together, each time you make a slight change in the UI, you'll need to test all your business logic! It's very easy to introduce errors into such a web application design.

Following are the drawbacks of directly tying together the presentation and the business logic portions of a web application.

- Testing user interfaces is tedious and takes a much longer time, as compared to the testing of the business logic. Therefore, the less non-UI code you link to the UI code, the better you can test the application.
- User interface activity may include simple presentation of data by retrieving it from the data store and displaying it to the user in an attractive fashion. When the user modifies data, however, the control is passed to the business logic portion of the application, which contacts the data source and modifies the data.

- Whereas business logic is completely independent of devices, user interface code is extremely device-dependent. Instead of requiring major changes in the UI to make the application portable across different devices and all the testing and related efforts that this involves, you can simply separate the UI and business logic, both to speed up the migration to different devices, as well as to reduce the errors inherent in such a process.
- It takes different skill sets to develop great HTML pages as compared to coming up with ingenious business logic, thus necessitating the separation of the development effort for the two areas.
- A single page request tends to combine the processing of the action associated with the link chosen by the user and the rendering of that page.

MVC to the Rescue

The Model-View-Controller architecture (formally introduced in 1988, although the idea existed in a simpler form since the 1970's) separates the modeling of the application domain, the presentation of the data, and the actions based on user input into three separate classes:

- **Model:** the model manages the behavior and data of the application and responds to requests for information and to instructions to change the state of the model (modifying the data)
- **View:** manages the display of information
- **Controller:** interprets the user inputs and informs the model and/or the view to modify data when necessary.

In a modern web application, the view is the browser and the controller is the set of server side components that handles the HTTP requests from the users. In this architecture, the view and the controller depend on the model, but the model is independent of the two. This separation lets you build and test the model independent of how you present the data.

All modern web applications such as those based on Node.js and the Ruby on Rails framework follow the MVC design pattern to separate user interface logic from business logic. The MVC design pattern avoids all the problems with traditional application design listed in the previous section. It also lets the model be tested independently of the presentation, by separating the model from the presentation logic.

MVC is now the well accepted architecture for web applications. Most of the popular web application frameworks follow the MVC pattern. There's some difference in how different web application frameworks interpret the MVC pattern, with the difference

being how they apportion the MVC responsibilities between the client and server. In recent years, there have been tremendous improvements in the capabilities of clients and newer web app frameworks such as AngularJS, EmberJS and Backbone let MVC components to partially execute on the client itself.

Ruby on Rails

Ruby on Rails (Rails from here on) is an extremely popular web application framework, created by David Heinemeier Hansson around 2004 as part of a project for his web application software development company named 37signals (now known as basecamp). As its name indicates, Rails is a web development framework written in the Ruby programming language. Rails is currently the #1 tool for building dynamic web applications.

The enormous popularity of Rails owes quite a bit to the use by Rails of the Ruby language, which acts as a kind of domain-specific language for developing web applications. This is what makes it so easy to perform typically complex web application programming tasks such as generating HTML and routing URLs so effortlessly, and in a highly compact and efficient manner.

The MVC Pattern and Rails

The standard Rails application structure contains an application directory called `app/` with three subdirectories: `models`, `views`, and `controllers`. This isn't a coincidence – Rails follows the Model-View-Controller (MVC) architectural pattern. As you can recall, MVC enforces separation between the business (domain) logic and the input and presentation logic associated with the user interface. In the case of web applications, the business logic is encapsulated by data models for entities such as users and products, and the user interface is nothing but the web pages in a web browser.

When a web browser interacts with a Rails application, it sends a request that's received by a web server and is passed along to a rails controller. The controller controls what happens next. Sometimes the controller may immediately render a view, which is a template that gets converted to HTML and sent to the browser as a response.

In dynamic web applications, which most web apps are today, the controller will interact with a model. A model is a Ruby object that represents a site entity such as a user. The model is responsible for contacting the backend database and retrieving the results requested through the browser. The controller invokes the model, gets the result through the model and renders the view and returns the complete web page to the browser as HTML.

Controllers, Actions, and Routes

The controllers include what are called actions inside them. The actions are actually functions that perform specific tasks such as retrieving results from a database or printing a message on the screen.

Rails uses a router that sits in front of the controller and determines where to send requests that are coming from web browsers. There's a root route that specifies the page served on the root URL. The root URL is of the format `/` ("slash"). You specify various routes in the Rails routes file (`config/routes.rb`). Each route consists of the controller name and the associated action. Each route is an instruction to Rails as to which action to perform, and thus which web page to create and send back to the browser.

RAILS AND REST

Earlier in this chapter, you learned about Representational State Transfer (REST), which is an architecture for developing web applications (as well as distributed networked systems). Rails uses the REST architecture, which means that application components such as users and microposts are modeled as resources which can be read or modified. More precisely, these resources can be created, read, updated and deleted – these operations correspond to the well-known create, select, update and delete operations of relational databases – and to the POST, GET, PATCH (early Rails versions used PUT for the updates) and DELETE requests of the HTTP protocol.

Full stack JavaScript Development with MEAN

Traditionally web programming required felicity with several programming languages. For client side programming, one was expected to know HTML (markup), CSS (styling) and JavaScript (functionality). On the server side, developers needed to know a language such as Java, PHP, Perl, plus, of course, SQL, which is a full-fledged language in its own right. In addition, when dealing with web applications, one had to know data formats well too, such as XML and JSON. The range of programming languages and the complexity involved in each of those languages (and data formats) has led to specialization among developers, into front-end and back-end development teams.

JavaScript has simplified things considerably by using a single language throughout the development stack, leading to the birth of the moniker “full stack development”.

The MEAN web application stack is a framework that uses MongoDB, Express, Angular JS and Node.js for building modern web applications (SPAs). In the following sections, I briefly explain the various components that make up the MEAN stack.

MONGODB

MongoDB is highly reliable, extremely scalable and performs very well, and although it's a NoSQL database, contains some key features that are normally found only in a relational database. The big thing about MongoDB in the context of building SPAs is that it lets you use JavaScript and JSON through entire application that your developers are building.

MongoDB's command line interface uses JavaScript for querying data, meaning you can use the same expression to manipulate data as in a browser environment. Furthermore, MongoDB uses JSON as its storage format, and all the data management tools are designed with JSON in mind.

Express

Express serves as the web framework in a MEAN application. Express, built with the Ruby language (using the Sinatra framework), describes itself as a minimalist framework for Node.js, and provides a thin layer of fundamental web application features.

AngularJS

AngularJS is an extremely popular frontend framework for creating SPAs, and uses the MVC approach to web applications. HTML wasn't really designed for declaring the dynamic views that web applications require, although it's still great for declaring static documents. AngularJS helps extend the HTML vocabulary for web applications.

Node.js

Node.js is a server side platform for building scalable web applications. Node.js leverages Google Chrome's V8 JavaScript engine and uses an event driven non-blocking I/O model that helps you build real-time web applications.

Node.js isn't something that's totally alien – it's simply a headless JavaScript runtime and uses the same JavaScript engine (V8) found inside Google Chrome, with the big difference that Node.js lets you run JavaScript from the command line instead of from within a web browser.

Node.js is a popular framework that helps create scalable web applications, and it's really Node.js that has helped make JavaScript a leading alternative for server-side programming. Node.js enables JavaScript API to move beyond the browser environment by letting you use JavaScript to perform server side operations such as accessing the file system and opening network sockets.

Node has been around only for about 6 years now (starting in 2009) and has been embraced with great fervor by developers who have been able to increase throughput by using the key features of Node.

Node is especially suitable when a web application keeps connections open with a large number of users when there's no active communication between users and the server for long stretches of time. Or, there may be an exchange of just a tiny bit of data between the client and the server over a period of time. Node excels in these types of environments by letting a single host running a Node.js server support a far greater number of concurrent connections than alternative technologies.

Conciseness is a hallmark of Node.js, as can be seen from the following code, which implements a web server with very little work on your side:

```
var http = require('http');
http.createServer(function (req, res) {
  res.writeHead(200, {'Content-Type': 'text/plain'});
  res.end('Hello World\n');
}).listen(1337, '127.0.0.1');
console.log('Server running at http://127.0.0.1:1337/');
```

This code does the following:

- Starts a server that listens on port 1337
- When the server receives a connection, it sends the message “Hello World”.
- The code prints a message to the console to tell you it's running on port 1337.

The bottom line with the MEAN (or any similar framework) is that a developer is able to create a production grade web application with just HTML5, CSS3 and JavaScript. You can consider the MEAN stack as a modern day alternative to the traditional well known LAMP stack for building web applications, which consists of Linux, Apache web server, MySQL and PHP.

Single Page Applications – the new Paradigm for Web Applications

Up until just a few years ago, all web applications were what are called multipage web applications. Recently, single-page-applications (SPAs) have become very popular. In addition, modern web applications sometimes come as a mix of both the traditional multi-page and the modern single-page application paradigm. In a world where the users are very sophisticated and expect easy communication and immediate responsiveness from the websites they use, SPAs are the new standard for web applications, replacing the old clunky websites that re-render entire pages after each user click.

Let me briefly explain how the traditional approach differs from the modern way of building web applications.

The Traditional Approach

Originally when the web was created and people started building web based applications, you used just HTML, a web server, and a language such as PHP to code the

application to build the application. While you can still build web applications just fine with the traditional approach, they aren't very scalable web sites. Single-page applications have basically supplanted the traditional web application model for all heavily used websites.

The Rise of the Single-Page Application

An SPA is an application delivered to the web browser that doesn't reload the page during its use. The applications function on just a single page in the browser and they give the appearance and feeling of a desktop application, since they're smoother than traditional applications.

Like any other web application, a SPA helps users complete a specific task, say, the reading of a document. Instead of reloading the web page, it lets you read a file from the same web page where you are right now. Think of the SPA as a fat client loaded from a web server. Several developments in the recent years have helped to make single-page applications popular, the most important being the adoption of Java Script, AJAX and HTML5.

Benefits of a SPA

SPAs offer several benefits to your users as compared to traditional websites. For one, they deliver the best parts of both a desktop application and a website, since an SPA can render like a desktop application. The SPA needs to draw only those parts of the interface that change, whereas a traditional website redraws the complete page on every single user action, which means there's a pause and a "flash" during the retrieval of the new page from the server by the browser, and the subsequent redrawing of the new page.

If the web page is large or if there's a slow connection or a busy server, the user wait might be long, whereas the SPA renders rapidly and provides immediate feedback to the user.

Like a desktop application, an SPA can keep the user posted of its state by dynamically rendering progress bars or busy indicators, whereas traditional websites have no such mechanism to tell the user when the next page will be arriving.

Users can access SPAs from any web connection, with just a browser, such as smart phones, tablets, etc. They also are cross-platform like websites, whereas desktop applications aren't always so.

SPAs are highly updateable, just as a website is, since all the users need to do to access the new version of an SPA is to just refresh their browser. Some SPAs are updated multiple times in a single day, whereas updating desktop applications is no trivial affair – besides taking too long to deploy the new versions, often there could be long

intervals, sometimes as long as several years(!), between the old and the new versions of a desktop application.

JavaScript

While older technologies such as Java applets and Flash were also SPA platforms, it's JavaScript SPAs that have made SPAs popular. Up until a few years ago, Flash and Java were the most widely adopted SPA client platforms because their speed and consistency was far better than that of JavaScript and browser rendering. Even in the early days (about 10 or so years ago), JavaScript offered the following advantages compared to Flash and Java:

- A no-plugin architecture that reduces development and maintenance efforts
- The no-plugin architecture also means JavaScript needs fewer resources to run
- A single client language (JavaScript) is used for everything, instead of a bunch of languages
- More fluid and interactive web pages

However, you couldn't rely on JavaScript for consistently providing crucial capabilities on most browsers. Today, most of the early weaknesses of JavaScript have been either removed or diminished in importance, thus bringing the advantages of JavaScript and browser rendering to the fore.

JavaScript has been around for a while, being the standard for client-side scripting in its previous incarnation. On the client side, JavaScript is the only language supported by all the popular browsers. In its early days, JavaScript provided simple page interactions by performing tasks such as changing an image's attributes on mouse overs and supporting collapsible menus, thus providing functionality missing in HTML. Over time, JavaScript has become more general purpose, moving beyond the client-side usage patterns of its early days.

AJAX and SPAs

Asynchronous JavaScript and XML (AJAX) is the technology that underlies a SPA application. AJAX is a non-blocking way for clients to communicate with a web server without reloading the page. Static content is where the web pages don't change at all, and are the mainstay of simple web sites with a few web pages.

Web applications use dynamic content, where the web pages are generated on the fly in response to search requests or button clicks by the users. AJAX lets the web browser poll the server for new data. In an application that has a contact form that uses AJAX to submit the form, the page won't reload when I submit the form. Instead it merely shows me a confirmation text on the page indicating that my response was

sent. When you don't use AJAX during the form submission, the entire page will reload or may bring up a completely different page with the confirmation message.

As the use of high speed internet grew several years ago, Ajax applications became more popular. The applications made background requests instead of fully reloading the web page, in order to make the page more responsive to requests. As Ajax requests grew in importance, applications consequently started used fewer page loads, culminating in the birth of the SPA, which uses just a single page load and uses Ajax calls to request all further data.

Traditional applications use a full response (POST) to the web server each time you send a form such as what I described earlier. An SPA makes asynchronous calls to update bits and pieces of the content. It's AJAX that enables the process of making these partial requests/responses to the server. In this context, it's important to note that rather than polling the server for new data, newer technologies like WebSockets make the browser maintain an open connection with the server so that the server can send data on demand.



Almost all SPAs use AJAX to load data and for other purposes, but AJAX has other uses as well. In fact, the majority of AJAX uses are non-SPA related.

You can look at SPA as a paradigm for building web applications, while AJAX represents the actual set of techniques that allows JavaScript based applications to support SPAs.

Web Services

A web service is a software component stored on a server, which can be accessed by an application or software component from a different server over a network. Web services promote the reusability of software components in internet based applications. Web services communicate through technologies such as XML, JSON and HTTP.

Two key Java APIs facilitate web services:

- JAX-WS is based on the Simple Object Access protocol (SOAP) that allows web services and clients to communicate even if they are written in different languages.
- JAX-RS: uses the Representational State Transfer (REST) network architecture based on the traditional web request and response mechanisms such as GET and POST, which I explained earlier in this chapter.

Web services, since they are platform and language independent, allow organizations to work together whether their hardware, software and communication technologies are compatible or not.

Amazon, eBay, PayPal and Google and others make their server-side applications available to their partners through web services. Using Web services, businesses can spend less time developing everything from scratch, and focus on more innovative services to provide enhanced shopping experiences for their customers.

An online music site for example, may have links to the websites of companies that sell concert tickets. In this case, the online music store is said to be consuming the concert ticket web service on its site. By consuming the concert ticket service, the music store provides additional services to its customers, besides benefiting financially from the web services. When an application consumes a web service, it means that it invokes the web services running on servers running elsewhere on the internet.



In the Java programming language, a web service is a Java class that allows its methods to be called by applications running on other servers through common data formats and protocols, such as XML, JSON and HTTP.

Web Service Basics

The server on which a web service lives is called the web service host. In the Java programming language, a web service is implemented as a class that lives in the host.

Publishing a web service is making the web service available to receive client requests. Consuming a web service is the using of a web service from a client application. Client applications send requests to the web service host and receive responses from the server. Now you can see how an application can retrieve data through a web service, without having direct access to the data. Same is the case where an application without massive processing power can utilize another server's resources to perform computations.

Web services can use one of two protocols to do their work - SOAP or REST, as explained in the following sections.

Simple Object Access Protocol (SOAP)

Simple Object Protocol (SOAP), the original technology used to communicate with web services, is a platform independent protocol that uses XML to interact with web services. The SOAP protocol describes how to mark up requests and responses so they can be sent via protocols such as HTTP. The SOAP message is XML markup that tells the web service how to process the message. Since SOAP messages are written in

XML, they're platform independent. Since firewalls allow HTTP traffic, SOAP based services can easily send and receive SOAP messages over HTTP connections.

When an application invokes a SOAP based web service's method, the request and additional information is packaged in a SOAP message within a SOAP envelope and sent to the web service host. The web service host processes the message contents by calling the method the client wishes to execute with the arguments specified by the client, and sends back the results to the client as another SOAP Message, to be parsed by the client in order to receive the results.

Representational State Transfer (REST)

Representational State Transfer (REST) is an alternative to SOAP and provides a different architectural style for implementing web services, with the web services referred to as RESTful web services. RESTful web services adhere to web standards and use traditional request and response mechanisms such as the GET and POST request methods.

In a RESTful web service, each operation is identified by a distinct URL, which lets the server know which operation to perform when it receives a request. Typically, RESTful web services return the data in the XML or JSON format, but they can also return it via HTML or plain text.

RESTful web services can be used directly from a web server or embedded in programs. When used in browser based applications, the browser can locally cache the results of REST operations when the web server is invoked via a GET request. Later requests for the operations are faster because they can be loaded from the browser's cache. Amazon's web services (aws.amazon.com) are a good example of RESTful web services.

Understanding APIs

APIs help expose data from legacy systems and also build an application without starting from scratch. Using APIs, a business can easily share information with its customers and suppliers, and customers can perform tasks such as checking the availability of a product from anywhere, and hailing a cab ride from a smartphone.

APIs are behind much of today's online work. Around 2001 companies started sharing their Web-based APIs with external users. Today, APIs are the operating system of the Web as well as mobile activity. By helping mash together data and services, APIs connect systems previously isolated from each other, and help create new applications.

APIs reduce the need to build apps from scratch – you can simply acquire APIs from providers to provide features such as payment processing and authentication. A company can bring in data from other entities to add services such as geographical maps

and credit card processing to their applications without having to write code for all the functionality. Thus, APIs reduce the cost of starting up a business and hence also reduce the barriers to entering a market. Many web sites (as well as web services) offer APIs that allow one to explicitly request data in a structured formats, thus saving you the trouble of having to scrape those sites for that data.

In addition to powering applications that customers use, APIs enhance communications between servers, thus helping automate processes and predicting problems, thus making businesses run more efficiently.

API usage has increased steadily over the past several years, driven by the increasing importance of web applications and the deluge of mobile applications. Early on, all web applications were simple monolithic designs, and used basic HTML, JavaScript and communicated over HTTP. In the past 10 years, APIs have proliferated, as alternative ways to interact with web applications. APIs help interactions with web applications and the need for organizations to integrate their systems on the web has led to the widespread use of APIs. It's mobile applications, however, that opened the floodgates for APIs. APIs help mobile applications easily access the data and functions of web applications without developers having to rewrite everything new for the mobile platform.

Two Types of Web Services

There are two main architectural styles to building web services – function centric services and resource centric services, as I explain here.

Function Centric Web Services

Function centric web services have been around for quite a while, since the early 1980s. Function centric web services can be thought of as services that would come into play where the application's code calls any function. The web service will transparently send the code and the data required for executing the function to a remote server and retrieves the results from that server, without the application being aware of the fact that the function was executed on the remote server.

Function centric technologies include technologies such as Common Object Request Broker Architecture (CORBA), Distributed Component Object Model (DCOM) and the Simple Object Access Protocol (SOAP), with SOAP being the most commonly employed technology.

SOAP uses XML to encode messages and uses HTTP to transport the requests and responses to and from the clients and servers. A web service provider uses two types of XML resources – Web Service Definition language (WSDL) files and XML Schema definition (XSD) files to describe the available methods and the definition of the data structures to be interchanged between the provider and the web application. These

two XML resources constitute a web services contract and (assuming you're using Java in this case) service developers use this information from the provider to create a Java client library and end up with a set of Java classes to implement the web service.

The SOAP server manages the SOAP libraries that perform tasks such as authentication, and error handling. Client developers on the other hand use the native Java client library generated based on the web service contract and compile and deploy the client application with the web service client code incorporated into that client application.

SOAP is very complex to implement and isn't conducive for many beginning efforts in web application development, or even for scaling web services. While the XML documents that are used for SOAP requests contain the request parameters and method names, the URLs don't contain the information required for making remote procedural calls and hence you can't cache the responses based on the URLs. Furthermore, the various web service specifications of SOAP (called the ws-* specifications) such as those relating to transactions, for example, make web service protocols stateful, thus preventing you from making the web servers stateless, which ought to be the key goal in scaling web applications.

Developing SOAP web services with dynamic languages such as Ruby and Python ran into many integration issues. Web technologies required a viable integratable option to SOAP, and this led to the creation of JavaScript Object Notation (JSON) based REST web services. Before we delve into the resource based web services, it probably is a good idea to learn a bit about JSON documents.

Resource Based Web Services

REST based web services have supplanted SOAP based function-centric web services as the leading web service architecture a few years ago, and are the current architectural standard for web services. Instead of using functions as their basis, resource centric web services revolve around the concept of a resource, and only a limited set of operations are allowed to be performed on the resource objects.

REST services use URLs to identify resources, and hence can use the HTTP methods I described earlier in this chapter, such as GET, PUT, POST and DELETE. In a REST based web service, the GET method fetches information about the resource and the POST method updates a resource or adds an entry, while DELETE will remove an object.

Compared to a typical SOAP web service architecture, with its many standards and ws-* specifications, all of which make it hard to work with and also to integrate, REST based web services are easy to set up since you deal with just four basic HTTP request methods. The REST framework needs to support only this small amount of functionality, meaning that the web stack need not be complex at all. You also don't need to

manage API contract artifacts such as the WSDL and XSD files required by SOAP based web services.

In addition to their inherent simplicity, REST web services are stateless and their GET method operations can be cached by HTTP caches between the client and the web services. This allows you to offload heavy web traffic over to reverse proxy servers to lighten the load on your web services as well as the databases.

However, it's not all roses with REST web service architectures – they do have some important drawbacks. Since they are so simple to implement, REST services require clients to authenticate using a security mechanism such as OAuth2 (a popular authorization framework that lets applications obtain limited access to user accounts on an HTTP service such as Facebook, for example), and use HTTP's transport layer security (TLS) to encrypt messages.

REST is definitely an easier architecture for a new firm to get started with compared to SOAP based services and is much easier to integrate with other web technologies, in any web stack you are likely to employ. If your needs don't require a sophisticated architecture with numerous features, using REST based web services is an easy decision to make, due to its enormous simplicity.

Service-Based Architectures and Microservices

Starting around the year 2005, service-oriented architecture (SOA) became the most popular architecture to promote the reuse of business functionality, and to enable business groups to communicate and collaborate in a better fashion.

SOA is an architecture that uses loosely coupled and highly autonomous services that each focuses on solving a specific business need. Loosely coupled in this context means that the different components are independent from each other and know only a minimal amount of other components. On the practical side, this means fewer dependencies between the components, which means changes in one component are less likely to adversely affect other components.

Decoupling also lets the people who work with the components specialize in those systems and also that you can scale each component independently. The high autonomy means individual web services act like an application themselves. For example the following would all be web services that work together:

- Product Catalog Service
- Recommendation Service
- Payment Processing Service

Web services use functional partitioning to divide a complex system into independent loosely coupled applications, with each service handling a small portion of the total functionality.

In an SOA architecture, the goal is to create generic services that are highly autonomous and decoupled from other services, and can be strung together to build complex applications. Service orchestration and service policies are used to build the complex applications.

In the microservices architecture, instead of building a single monolithic architecture based application, developers build a suite of components, which work together over the network. Each of the components in a microservice architecture can be written in the language that's best suited for a task, and the components can be deployed independently of the others. In addition, when you need to scale a component, you can do so independently for just that component, without worrying about the rest of the components.

Horizontally scalable applications benefit significantly from a microservices architecture. Take the case of a financial trading firm that deals with options and futures contracts. Its application may have components such as the following:

- A user interface for the traders
- Code for setting trades that interacts with the stock exchanges and the order management system
- A proprietary pricing system

When the company modifies its pricing algorithms, only the pricing system is touched and the user interface and the back end interfaces are left alone. This obviously leads to faster changes and a more agile business.

SOA architectures typically use the Simple Object Access Protocol (SOAP), which is a set of technologies that help define, discover, and use web services. Prior to this, humongous monolithic applications were the only way to architect applications, and SOA seemed like a great alternative to this unwieldy way of architecting applications. SOA implementation however turned out to be a nightmare for many folks, since it's a complex architecture and has resulted in numerous failed projects, thus tarnishing its image.

Similarities between traditional SOAs and the Microservice approach

Both microservices and SOA are service based, meaning that services are at the center of everything. Microservices architecture is quite similar to SOA in its basic premise of breaking up inefficient monolithic applications into modular components.

Microservices are connected with each other through a thin layer of simple APIs and the well-known standards of HTTP. Services are used to implement business functionality in both approaches. They both are also distributed architectures, where the services are accessed remotely through remote access protocols such as Representational State Transfer (REST), Simple Object Access Protocol (SOAP), Java message Service (MS), Advanced Message Queuing Protocol (AMQP), Remote Method Invocation (RMI) and other similar protocols.

Distributed applications, while complex to implement, offer key benefits when compared to monolithic applications, such as increased scalability and a more focused development of the applications. The use of self-contained applications is common within a distributed architecture, helping enhance the reliability and speed of these applications, while simultaneously making it easier to maintain them.

Modularity is at the center of distributed service based architectures. Modularity means that the application is broken up into small self-contained services that you design, develop, and deploy separately, with no dependence on other application components. Whereas traditional monolithic applications usually require massive unwieldy rewrites or refactoring to incorporate changes dictated by the business, modular architectures let you rewrite the small self-contained services from scratch, thus keeping the application “fresh” in terms of its design and functionality.

While some may think that microservices are just SOA in a new garb (“Microservices are SOA done right”), there are key differences between the two architectures. Since microservices is what everybody seems to be doing these days, I’ll focus on those services in this chapter and won’t go into the details of traditional SOA architectures.

There are two major types of services that microservice architectures focus on: functional services that support specific business functions such as sales and marketing, and infrastructure services that support nonfunctional tasks such as auditing, authorization and logging. You can view the functional services as external facing services and the infrastructure services as private, internal shared services. Microservices adopt a share-as-little-as-possible approach, meaning that the services are sealed units with little or no dependency on other components.

Differences between SOA and Microservices

While both SOA and microservices are based in the idea of modular services that break up a monolithic application, there are some significant differences between the two approaches, as explained in the following sections.

Service Types

Microservices generally have fewer service types, with functional and infrastructure services being the two major service types. The typical microservices architecture looks like the following:

```
client requests => API layer => functional service => infrastructure service
```

SOA architectures are different from the microservice service types. There are more service types in SOA, such as enterprise services and application services. The typical SOA architecture resembles the following:

```
client requests => messaging middleware => enterprise services =>
application services => infrastructure services
```

In this architecture:

- Business services are abstract high level services that capture the core enterprise level business functions. These functions are usually represented through XML, Web Service Definition Language (WSDL), or Business Process Execution language (BPEL). Typically these services don't involve implementation.
- The enterprise services are the actual steps that provide the implementation for the functionality offered by the business services.
- The messaging middleware helps the business services and the corresponding enterprise services to communicate.
- Application services are even more fine-grained than enterprise services and provide narrow business functionality. These services can be invoked directly or through enterprise services.

In general, SOA means that services tend to cover a large amount of business functionality, such as a claims processing service for example. Microservices typically are very fine-grained with services that are much more narrowly focused with typical service names such as `UpdateCustomerAddress` and `GetCreditRating`, for example.

Need for Coordination

Since microservices involve fewer service types, usually the same application development team owns both the functional and the infrastructure services. However, in SOA, the larger number of service types means that there are different service owners for each of the service types. For example, business services are owned by the business users and application services by app development teams, and infrastructure services by the infrastructure service teams. Therefore, there's a need for coordination among the multiple groups to satisfy the business requests. Microservices, on the other hand, typically don't require this coordination since only one group is involved.

Time to Market

The smaller number of service types and the fact that there's no need for coordination among multiple teams means that microservices can be developed, tested, deployed and maintained with less effort and cost, and are also much faster to get to market.

Sharing the Components

SOA in general shares components among the services, with its share-as-much-as-possible architecture, whereas microservices typically consist of standalone services that are independent of each other. Microservices change and evolve independent of other services in the enterprise. While sharing repeated service functionality among multiple services does reduce duplication of business functionality, changes in those shared services over time could lead to problems, since the change may not impact all the services that share the changed service in a uniform fashion.

Server Virtualization and Linux Containers

Server virtualization, which involves running multiple virtual servers on a single physical machine, is ubiquitous and is one of the key foundational layers of modern cloud computing architectures. Virtualization lets a physical server's resources such as CPU, RAM and storage to be shared among several virtual servers. Virtualization helps substantially lower your costs of supporting a complex computing environment, besides speeding up deployments, as virtual machines can be spun up in a fraction of the time it takes to order, receive and configure physical servers.

In this chapter, I discuss the basic architecture of server virtualization first, and follow it up by explaining the concept of a hypervisor, which is the key piece of software that serves as a resource allocation and hardware abstraction layer between the physical server and the virtual servers you create on the physical server. I also explain the different types of virtualization such as full and paravirtualization.

This chapter isn't limited to traditional hardware virtualization. Container virtualization is relatively new and is quite different from hardware virtualization. Unlike hardware virtualization, container virtualization doesn't mimic a physical server with its own OS and resources. This type of virtualization is all about enabling applications to execute in a common OS kernel. There's no need for a separate OS for each application, and therefore the containers are lightweight, and thus impose a lower overhead compared to hardware virtualization.

This chapter introduces the Linux Containers technology. Linux containers keep applications together with their runtime components by combining application isolation and image-based deployment strategies. By packaging the applications with their libraries and dependencies such as the required binaries, containers make the applications autonomous. This frees up the applications from their dependence on various components of the underlying operating system.

The fact that containers don't include an OS kernel means that they're faster and much more agile than VMs (virtual machines). The big difference is that all containers on a host must use the same OS kernel. The chapter delves into the Linux technology that makes possible containers – namely, namespaces, Linux Control Groups (Cgroups) and SELinux. Chapter 5 continues the discussion of containers, and is dedicated to container virtualization including Docker, currently the most popular way to containerize applications.

Linux Server Virtualization

Linux server virtualization is the running of one or more virtual machines on a physical server that's running the Linux operating system. Normally a server runs a single operating system (OS) at a time. As a result, application vendors had to rewrite portions of their applications so they'd work on various types of operating systems. Obviously, this is a costly process in terms of time and effort.

Hardware virtualization, which lets a single server run multiple operating systems, became a great solution for this problem. Servers running virtualization software are able to host applications that run on different operating systems, using a single hardware platform as the foundation. The host operating system supports multiple virtual machines, each of which could belong to the same, or a different OS.

Virtualization didn't happen overnight. IBM mainframe systems from about 45 years ago started allowing applications to use a portion of a system's resources. Virtualization became mainstream technology in the early 2000's when technology made it possible to offer virtualization on x86 servers. The awareness that the server utilization rate was extremely low, as well as the rising cost of maintaining data centers with their high power costs, has made virtualization wide spread. A majority of the servers running across the world today are virtual – virtual servers way outnumber physical servers.

The Architecture of Virtual Machines

As you can guess, unlike a physical machine, a virtual machine (VM) doesn't really exist – it's a software artifact that imitates or mimics a physical server. That doesn't mean that a VM is something that's only in our minds – it actually consists of a set of files.

There's a main VM configuration file that specifies how much memory and storage is allocated to the VM. The configuration file also names the virtual NICs assigned to the VM, as well as the I/O it's allowed to access. One of these files is the VM configuration file, which specifies how many CPUs, how much RAM and which I/O devices the VM can access. The configuration files show the VM storage as a set of virtual disks, which are actually files in the underlying physical file system.

When an administrator needs to duplicate a physical server, a lot of work is required to acquire the new server, install the OS and application files on it, and copy the data over. Since a VM is just a set of files, you can get one ready in literally minutes after making just a handful of changes in the VM configuration file. Alternatively, you can provision new VMs through VM templates. A template contains default settings for hardware and software. Provisioning tools can simply use a VM template and customize it when they deploy new servers.

Virtualization in the early x86 was based purely software-based virtualization. Although Pope and Goldberg in their seminal paper “Formal Requirements for Virtualizable Third Generation Architectures” specified the three key properties for a virtual machine monitor (efficiency, resource control and equivalence), it wasn’t until the mid-2000’s that the x86 architecture started satisfying these three requirements. Hardware-assisted virtualization is the way that these ideal requirements started being realized.

Software-based virtualization has inherent limitations. The x86 architecture employs the concept of privilege levels (also called privilege reigns) for processing machine instructions

The Virtual Machine Monitor (Hypervisor)

The key software that makes virtualization possible is the virtual machine monitor (VMM), actually known by its other name, hypervisor. A hypervisor is the software that does the heavy lifting in virtualized systems – it coordinates the low-level interaction between virtual machines and the underlying host physical server hardware. The hypervisor sits between the VMs and the physical server and allows the VMs to partake of the physical server’s resources such as disk drives, RAM and CPU.



Virtualization lets a powerful physical server appear as several smaller servers, thus saving you space, power and other infrastructure expenses. A big advantage of virtualizing an environment is resource sharing among the servers, meaning that when one of the virtual servers is idle or almost so, other servers running on the same physical server can use the idle resources granted to the first server and speed up their own processing.

How VMs share Resources

Virtualization lets the resources of the host server such as CPU, RAM, physical storage and network bandwidth be shared among the virtual servers running on top of a physical server. Often, even on a non-virtualized server shortage of any one of these resources can slow applications down. How, then, can multiple servers share a single set of resources without bringing the system to a halt? Let’s learn how virtualization

typically handles the sharing of these resources among the VMs, to avoid bottlenecks and other performance issues.

In the context of resource allocation to the virtual machines, it's important to understand the key concept of overcommitting. Overcommitting is the allocation of more virtualized CPUs or memory than there's available on the physical server, and refers to the fact that you assume that none of the virtual servers will use their resources to the full extent on a continuous basis. This allows you to allocate the physical server's resources in a way that the sum of the allocated resources often exceeds the physical resource limit of the server. Using virtual resources in this fashion allows you to increase guest density on a physical server.

Disk Storage

Storage is not as much virtualized as the other server resources are. You simply allocate a chunk of the host storage space to each of the VMs, and this space is exclusively reserved for those VMs. Multiple VMs writing to the same storage disk might cause bottlenecks, but you can avoid them through the use of high performance disks, RAID arrays configured for speed, and network storage systems, all of which increase the throughput of data.

When it comes to storage, virtualization often uses the concept of thin provisioning, which lets you allocate storage in a flexible manner so as to optimize the storage available to each guest VM. Thin provisioning makes it appear that there's more physical storage on the guest than what's really available. Thin provisioning is different from overprovisioning, and applies only to storage and not to CPU or RAM.

CPU

CPU sharing is done on the basis of time slicing, wherein all the processing requests are sliced up and shared among the virtual servers. In effect, this is the same as running multiple processes on a non-virtualized server. CPU is probably the hardest resource to share, since CPU requests need to be satisfied in a timely fashion. You may at times see a small waiting time for CPU, but this is to be expected, and is no big deal. However, excessive waiting times can create havoc with the performance of applications.

Virtualized CPUs (vCPUs) can be overcommitted. You need to be careful with overcommitting vCPUs, as loads at or close to 100% CPU usage may lead to requests being dropped, or slow response times. You're likely to see performance deterioration when running more vCPUs on a VM than are present on the physical server. Virtual CPUs are best overcommitted when each guest VM has a small number of vCPUs when compared to the total CPUs of the underlying host. A hypervisor such as KVM can easily handle switches between the VMs, when you assign vCPUs at a ratio of five CPUs (on 5 VMs) per on physical CPU on the host server.

Network Bandwidth

Network bandwidth can be overprovisioned since it's unlikely that all VMs will be fully utilizing their network bandwidth at all times.

Memory

It's possible to overcommit memory as well, since it's not common to see the RAM being fully used by all the VMs running on a server at any given time. Some hypervisors can perform a “memory reclamation”, whereby they reclaim RAM from the VMs to balance the load among the VMs.

It's important to remember that applications that use 100% of the allocated memory or CPU on a VM can become unstable in an overcommitted virtual environment. In a production environment, it's critical to test extensively before overcommitting memory or CPU resources, as the overcommit ratios depend on the nature of the workloads.

Benefits offered by Virtual Machines

Virtualization offers several benefits to an IT department, such as the following:

Lower Costs

Virtualization lowers the cost of hardware purchases and maintenance, power and cooling, data center space and involves far lesser administrative and management effort.

Server Consolidation

Server consolidation is probably the most common, and one of the biggest motivating factors behind the drive to virtualize systems. Consolidation means that you reduce the footprint of your physical servers, saving not only capital outlays but also operating costs in terms of lower energy consumption in the data center. For example, you need fewer floor switches and networking capacity with virtualization when compared to physical servers.

Isolation

Since the guest operating systems are fully isolated from the underlying host, even if the VM is corrupted, the host is still in an operating state.

Easy Migration

You can move a running virtual machine from one physical server to another, without disconnecting either the client or the applications. You can move a running VM to a different physical server without impacting the users, using tools such as

vMotion (VMware) and Live Migration (RedHat Linux), both of which enhance the uptime of the virtualized systems.

Dynamic Load Balancing

You can move VMs from one physical server to another for load balancing purposes, so you can load balance your applications across the infrastructure.

Higher Availability

You can quickly restart a failed VM on a different physical server. Since virtual guests aren't very dependent on the hardware, and the host provides snapshot features. You can easily restore a known running system in the case of a disaster.

To summarize, virtualization offers several compelling benefits, which led to its widespread usage in today's IT environments. Reduced capital outlays for purchase and support since you need to purchase fewer physical servers, faster provisioning, the ease of supporting legacy applications side by side with current applications, and the fact that virtualization gets you attuned to the way things are done in modern cloud based environments, have all been factors for its widespread use.

Drawbacks of Virtualization

Virtualization isn't a costless solution – you do need to keep in mind the following drawbacks of virtualization:

- There's often a performance overhead for the abstraction layer of virtualization
- Overprovisioning is always a potential problem in a virtualized environment and this could lead to performance degradation, especially during peak usage.
- Rewriting existing applications for a virtual environment may impose a stiff upfront cost
- Losing a single hypervisor could mean losing all the VMs based on that hypervisor
- Administrators need specialized training and expertise to successfully manage the virtualized environments.

Virtualization Types

In addition to sharing the CPU and RAM of the parent server, VM guests share the I/O as well. The classification of hypervisors into different types is based on two basic criteria: the amount of hardware that's virtualized and the extent of the modifications required of the guest system. Modern virtualization is hardware based and doesn't use traditional software I/O virtualization (emulation) techniques. Software virtualization

uses slow techniques such as binary translation to run unmodified operating systems. By virtualizing at the hardware level, virtualization seeks to deliver native performance levels. The following sections explain the two popular I/O virtualization techniques – paravirtualization and full virtualization.

Paravirtualization

Paravirtualization, as the name itself indicates, isn't really "complete virtualization" since the guest OS needs to be modified.

The paravirtualization method presents a software interface to the VM that's similar to that of the host hardware. That is, instead of emulating the hardware environment, it acts as a thin layer to enable the guest system to share the system resources.

Under paravirtualization, the kernel of the guest OS running on the host server is modified, so it can recognize the virtualization software layer (hypervisor). Privileged operations are replaced by calls to the hypervisor, in order to reduce the time the guest OS will spend performing operations that are more difficult to run in the virtual environment than in the non-virtualized environment. Costly operations are performed on the native host system instead of on the guest's virtualized system. The hypervisor performs tasks on behalf of the guest OS and provides interfaces for critical kernel operations such as interrupt handling and memory management. Both Xen and VMWare are popular examples of paravirtualization.



A big difference between fully virtualized and paravirtualized architectures is that you can run different operating systems between guest and host systems under full virtualization, but not under paravirtualization.

Since paravirtualization modifies the OS, it's also called OS-assisted virtualization, with the guest OS being aware that it's being virtualized. Paravirtualization offers the following benefits:

- Under paravirtualization, the hypervisor and the virtual guests communicate directly, with the lower overhead due to direct access to the underlying hardware translating to higher performance. VMs that are "aware" that they're virtualized offer higher performance.
- Since paravirtualization doesn't include any device driver at all, it uses the device drivers in one of the guest operating systems, called the privileged guest. You therefore aren't limited to the device drivers contained in the virtualization software.

Paravirtualization, however, requires you to modify either the guest OS or the use of paravirtualized drivers. It therefore imposes the following limitations:

- You're limited to open source operating systems and proprietary operating systems where the owners have consented to make the required code modifications to work with a specific hypervisor. Paravirtualization isn't very portable since it doesn't support unmodified operating systems such as Microsoft Windows.
- Support and maintainability issues in production environments due to the OS kernel modifications needed for paravirtualization.

Paravirtualization can cover the whole kernel or just the drivers that virtualize the I/O devices. Xen, an open source virtualization project, is a good example of a paravirtualized environment. Xen virtualizes the CPU and memory by modifying the Linux kernel and it virtualizes the I/O with custom guest OS device drivers.



In addition to full and paravirtualization, there's also something called software virtualization, which uses emulation techniques to run unmodified virtual operating systems. Linux distributions such as RedHat Linux don't support software virtualization.

Full Virtualization

Full virtualization is a technique where the guest operating system is presented a simulated hardware interface by a hardware emulator. In full virtualization, the virtualization software, usually referred to as a hypervisor (guest OS drivers) emulates all hardware devices on the virtual system. The hypervisor creates an emulated hardware device and presents it to the guest operating system. This emulated hardware environment is also called a Virtual Machine Monitor or VMM, as explained earlier.

Guests use the features of the underlying host physical system to create a new virtual system called a virtual machine. All components of that the virtual machine presents to the operating system are virtualized. The hypervisor simulates specific hardware. For example when QEMU simulates an x86 machine, it provides a virtual Realtek 8139C+PCI as the network adapter. This means that the guest OS is unaware that it's running on virtual, and not on real hardware.

he VM allows the guest OS to run without any modifications and the OS behaves as if it has exclusive access to the underlying host system. Since the physical devices on the host server may be different from the emulated drivers, the hypervisor needs to process the I/O before it goes to the physical device, thus forcing the I/O operations to move through two software layers. This means not only slower I/O performance but also higher CPU usage.



In paravirtualization, the virtualization software layer abstracts only a portion of the host system's resources, and in full virtualization, it abstracts all of the host system resources.

Since the guest OS is a full emulation of the host hardware, this virtualization technique is called full virtualization.

You can run multiple unmodified guest operating systems independently on the same box with full virtualization. It's the hypervisor that helps run the guest operating systems without any modification, by coordinating the CPU of the virtual machine and the host machine's system resources.

The hypervisor offers CPU emulation to modify privileged and protected CPU operations performed by the guest OS. The hypervisor intercepts the system calls made by the guest operating systems to the emulated host hardware and maps them to the actual underlying hardware. You can have guest systems belonging to various operating systems such as Linux and Windows running on the same host server. Once again, the guest operating systems are completely unaware of the fact that they're virtualized and thus don't require any modifications.



Full virtualization requires complete emulation, which means more resources for processing from the hypervisor.

QEMU (which underlies KVM, to be discussed later in this chapter), VMWare ESXi and VirtualBox are popular fully virtualized hypervisors. Full virtualization offers many benefits, as summarized here:

- The hypervisor offers a standardized environment for hardware for the guest OS. Since the guest OS and the hypervisor are a consistent package together, you can migrate this package across different types of physical servers.
- The guest OS doesn't require any modification.
- It simplifies migration and portability for the virtual machines
- Applications run in truly isolated guest operating systems
- The method supports multiple operating systems which may be different in terms of their patch level or even completely different from each other, such as the Windows and Linux operating systems

The biggest drawback of full virtualization is that since the hypervisor needs to process data, some of the processing power of the host server is commandeered by the hypervisor and this degrades performance somewhat.

Type of Hypervisors

The hypervisor, by presenting virtualized hardware interfaces to all the VM guests, controls the platform resources. There are two types of Hypervisors, based on where exactly the hypervisor sits relative to the operating system and the host, named Type 1 and Type 2 hypervisors.

Type 1 Hypervisors

A Type 1 hypervisor (also called a native or bare metal hypervisor) is software that runs directly on the bare metal of the physical server, just as the host OS does. Once you install and configure the hypervisor, you can start creating guest machines on the host server.

Architecturally, the Type 1 hypervisor sits directly on the host hardware and is responsible for allocating memory to the virtual machines, as well as providing an interface for administration and for monitoring tools. VMWare ESX Server, Microsoft Hyper-V and several variations of the open source KVM hypervisor are examples of a Type 1 hypervisor.

Due to its direct access to the host server, the Type 1 hypervisor doesn't require separate CPU cycles or memory for the VMs and thus delivers greater performance.

It's important to understand that most implementations of a bare metal hypervisor require virtualization support at the hardware level through hardware assisted virtualization techniques (explained later in this chapter), and VMWare and KVM are two such hypervisors.

Type 2 Hypervisors

A Type 2 hypervisor (also called a hosted hypervisor) is deployed by loading it on top of the underlying OS running on the physical server, such as Linux or Windows. The virtualization layer runs like a hosted application directly in top of the host OS. The hypervisor provides each of the virtual machines running on the host system with resources such as a virtual BIOS, virtual devices and virtual memory. The guest operating systems depend on the host OS for accessing the host's resources.

A Type 2 hypervisor is useful in situations where you don't want to dedicate an entire server for virtualization. For example, you may want to run a Linux OS on your Windows laptop – both VMWare Workstation and Oracle VM Virtual Box are examples of Type 2 hypervisors.

Traditionally, a Type-1 hypervisor is defined as a “small operating system”. Since a Type 1 hypervisor directly controls the resources of the underlying host, its performance is generally better than that of a Type 2 hypervisor, which depends on the OS to handle all interactions with the hardware. Since Type 2 hypervisors need to perform extra processing (‘instruction translation’), they can potentially adversely affect the host server and the applications as well.

You can pack more VMs with a Type 1 hypervisor because this type of hypervisor doesn’t compete with the host OS for resources.

Kernel Level Virtualization (Hardware Assisted Virtualization)

Under kernel level virtualization, the host OS contains extensions within its kernel to manage virtual machines. The virtualization layer is embedded in the operating system kernel itself. Since the hypervisor is embedded in the Linux kernel, it has a very small footprint and disk and network performance is higher in this mode. The popular open source Kernel Virtual Machine (KVM) virtualization model uses kernel level virtualization (hardware-assisted virtualization method).

Bare Metal versus Hosted Hypervisors

Virtualization solutions that use a Type-2 hypervisor such as VirtualBox are great for enabling single users or small organizations to run multiple VMs on a single physical server. VirtualBox and similar solutions run as client applications and not directly on the host server hardware. Enterprise computing requires high performance virtualization strategies that are closer to the host’s physical hardware. Bare metal virtualization involves much less overhead and also exploits the built in hardware support for virtualization better than Type-2 hypervisors.

Most Linux systems support two types of open-source bare-metal virtualization technologies: Xen and Kernel Virtual Machine (KVM). Both Xen and KVM support full virtualization, and Xen also supports the paravirtualization mode. Let’s start with a review of the older Xen technology and then move on to KVM virtualization, which is the de facto standard for virtualization in most Linux distributions today.

Xen Virtualization

Xen was created in 2003 and acquired later on by Citrix, which announced in 2013 that the Xen Project would be a collaborative project between itself, Xen’s main contributor, and the Linux foundation. Xen is very popular in the public cloud environment with companies such as Amazon Web Services and Rackspace Cloud using it for their customers.

Xen is capable of running multiple types of guest operating systems. When you boot the Xen hypervisor on the host physical hardware, it automatically starts a primary virtual machine called Domain 0 (or dom0), or the management domain. Domain 0

manages the systems and by performing tasks such as creating additional virtual machines, and managing the virtual devices for the virtual machines, as well as tasks such as suspending, resuming, and migrating virtual machines, the primary VM will provide the virtual management capabilities for all other VMs, called the Xen guests. You administer Xen through the `xm` command-line suite.

The Xen daemon, named `xend`, runs in the `dom0` VM and is the central controller of virtual resources across all VMs running on the Xen hypervisor. You can manage the VMs using an open source virtual machine manager such as `OpenXenManager`, or a commercial manager such as `Citrix XenCenter`.

Xen Architecture

Xen is a Type 1 hypervisor and so it runs directly on the host hardware. Xen inserts a virtualization layer between the hardware and the virtual machines, by creating pools of system resources, and the VMs treat the virtualized resources as if they were physical resources.

Xen uses paravirtualization, means the guest OS must be modified to support the Xen environment. The modification of the guest OS lets Xen use the guest OS as the “most privileged software”. Paravirtualization also enables Xen to use more efficient interfaces such as virtual block devices to emulate hardware devices.

Xen’s Benefits and Drawbacks

Xen offers highly optimized performance due to its combination of paravirtualization and hardware assisted virtualization. However, it has a fairly large footprint and integrating it isn’t easy and could overwhelm the Linux kernel over time. It also relies on third-party products for device drivers as well as for backup and recovery and for fault tolerance. High I/O usually slows down Xen based systems.

While Xen offers a higher performance than KVM, it’s the ease of use of KVM virtualization which has led to it’s becoming the leading virtualization solution in Linux environments.

KVM supports native virtualization on processors that contain extensions for hardware virtualization. KVM supports several types of processors and guest operating systems, such as Linux (many distributions), Windows, and Solaris. There’s also a modified version of QEMU that uses KVM to run Mac OS X virtual machines.

Kernel-Based Virtual Machines (KVM)

Linux KVM (Kernel-based Virtual Machine) is the most popular open-source virtualization technology today. Over the past few years, KVM has overtaken Xen as the default open source technology for creating virtual machines on most Linux distributions.

Although KVM has been part of the Linux kernel since the 2.6.20 release (2007), until release 3.0, you had to apply several patches to integrate KVM support into the Linux kernel. Post 3.0 Linux kernels automatically enable KVM's integration into the kernel, allowing it to take advantage of improvements in the Linux kernel versions. Being a part of the Linux kernel is a big deal, since it means frequent updates and a lower Total Cost of Operation (TCO). In addition, KVM is highly secure since it's integrated with SELinux in both RedHat Linux and CentOS.

KVM differs from Xen in that it uses the Linux kernel as its hypervisor. Although a Type-1 hypervisor is supposed to be similar to a small OS, the fact that you can configure a custom lightweight Linux kernel and the availability of large amounts of RAM on today's powerful 64-bit servers means that the size of the Linux kernel isn't a hindrance.

Just as Xen has its `xm` toolset, KVM has an administrative infrastructure that it has inherited from QEMU (short for Quick Emulator) a Linux emulation and virtualization package which achieves superior performance by using dynamic translation. By executing the guest code directly on the host CPU, QEMU achieves performance close to the native OS.

QEMU supports virtualization while executing under the Xen hypervisor or by using the Linux KVM kernel module. Red Hat has developed the `libvirt` virtualization API to help simplify the administration of various virtualization technologies such as KVM, Xen, LXC containers, VirtualBox and Microsoft Hyper-V. As an administrator, it's great to learn `libvirt` because you can manage multiple virtualization technologies by learning a single set of commands (command line and graphical) based on the `libvirt` API.

In order to support KVM virtualization, you need to install various packages, with the required package list depending on your Linux distribution.

Using Storage Pools

When you create one or two KVM based VMs, you can use disk images that you can create on the local disk storage of the host. Each VM will in essence be a disk image stored in a local file. However, for creating enterprise wide virtualization environments, this manual process of creating VMs is quite tedious and hard to manage. The `libvirt` package lets you create storage pools to serve as an abstraction for the actual VM images and file systems.



The `libvirt` package provides standard, technology independent administrative commands to manage virtualization environments.

A storage pool is a specific amount of storage set aside by the administrator for use by the guest VMs.

Storage pools are divided into storage volumes which are then assigned to guest VMs as block devices.

A storage pool can be a local directory, physical disk, logical volume, or a network file system (NFS) or block-level networked storage managed by libvirt. Using libvirt, you manage the storage pool and create and store VM images in the pool. Note that in order to perform a live migration of a VM to a different server, you should locate the VM disk image in an NFS, block-level networked storage, or in HBA (SCSI Host Bus Adapter) storage that can be accessed from multiple hosts.

Creating the Virtual Machines

The libvirt package contains the virsh command suite that provides the commands to create and manage the virtualization objects that libvirt uses, such as the domains (VMs), storage pool, networks, devices, etc. Following is an example that shows how to create an NFS based (netfs) storage pool:

```
virsh pool-create-as NFS-POOL netfs \  
--source-host 192.168.6.248 \  
--source-path /DATA/POOL \  
--target /var/lib/libvirt/images/MY-NFS-POOL
```

In this command, MY-NFS-POOL is the name of the new storage pool and the local mount point that'll be used to access this NFS based storage pool is /var/lib/libvirt/images/MY-NFS-POOL. Once you create the storage pool as shown here, you can create VMs in that pool with the virt-install command, as shown here:

```
virt-install \  
--name RHEL-6.3-LAMP \  
--os-type=linux \  
--os-variant=rhel6 \  
--cdrom /mnt/ISO/rhel63-server-x86_64.iso \  
--disk pool=My-NFS-POOL,format=raw,size=100 \  
--ram 4096 \  
--vcpus=2 \  
--network bridge=br0 \  
--hvm \  
--virt-type=kvm \  

```

Here's a summary of the key options specified with the virt-install command:

- --os-type and --os-variant: indicate that this VM will be optimized for the Linux RedHat Enterprise Linux 6,3 release,
- --cdrom: specifies the ISP image (virtual CDROM device that will be used to perform this installation)

- `--disk`: specifies that the VM will be created with 100GB of storage from the storage pool named NFS-01.
- `--ram` and `--vcpus`: specify the RAM and virtual CPUs for the VM
- `--hvm`: indicates that this is a fully virtualized system (default)
- `--virt-type`: specifies `kvm` as the hypervisor (default)



RedHat Enterprise Virtualization (RHEV) is based on KVM.

Considerations in Selecting the Physical Servers for virtualization

The choice of the physical servers for a virtualized environment is critical, and you've several choices, as explained in the following sections.

Build Your Own versus Purchasing

You can create your own servers by purchasing and putting together all the individual components such as the disk drives, RAM and CPU. You should expect to spend less to build your own systems, so that's good, but the drawback is the time it takes to get your systems together. In addition, you're responsible for maintaining these systems with partial or no service contracts to support you, with all the attendant headaches.

If you're considering putting together your own systems, it may be a good idea to check out the Open Compute project (<http://opencompute.org>), which creates low cost server hardware specifications and mechanical drawings (designs). The goal of Open Compute is to design servers that efficient, inexpensive and easy to service. Consequently these specs contain far fewer parts than traditional servers. You can try and purchase hardware that meets these specifications to ensure you're getting good hardware when you're trying to keep expenses low.

Purchasing complete systems from a well-known vendor is the easiest and most reliable way to go, since you don't need to worry about the quality of the hardware and software, in addition to getting first class support. It also gets maintenance off your hands. However, as you know, you're going to pay for all the bells and whistles.

Rack and Blade Servers

Blade servers are commonly used for virtualization since they allow for a larger number of virtual machines per chassis. Rack servers don't let you create as many VMs per chassis.

The choice of the type of server depends on factors such as their ease of maintainability, power consumption, remote console access, server form factor, and so on.

Migrating Virtual Machines

Migrating virtual machines means the moving of a guest virtual machine from one server to another. You can migrate a VM in two ways: live and offline, as explained here:

- **Live Migration:** this process moves an active VM from one physical server to another. In Red Hat Enterprise Linux, this process moves the VM's memory and its disk volumes as well, using what's called live block migration.
- **Offline Migration:** During an offline migration, you shut down the guest VM and move the image of the VM's memory to the new host. You can then resume the VM on the destination host and the memory previously used by the VM on the original host is released back to the host.

You can migrate VMs for the following purposes:

- **Load Balancing:** You can migrate one or more VMs from a host to relieve its load
- **Upgrades to the Host:** When you're upgrading the OS on a host, you can avoid downtime for your applications by migrating the VMs on that host to other hosts.
- **Geographical Reasons:** Sometimes you may want to migrate a VM to a different host in a different geographical location, to lower the latency of the applications hosted by the VM.

Application Deployment and Management with Linux Containers

A Linux container (LXC) is a set of processes that are isolated from other processes running on a server. While virtualization and its hypervisors logically abstract the hardware, containers provide isolation, letting multiple applications share the same OS instance.

You can use a container to encapsulate different types of application dependencies. For example, if your application requires a particular version of a database or scripting language, the containers can encapsulate those versions. This means that multiple versions of the database or scripting language can run in the same environment, without requiring a completely different software stack for each application, each with its own OS. You don't pay for all of this with a performance hit, as containerized

applications deliver roughly the same performance as applications that you deploy on bare metal.

Linux Containers have increasingly become an alternative to using traditional virtualization, so much so that containerization is often referred to as the “new virtualization”.

On the face of it both virtualization and containerization seem to perform the same function by letting you run multiple virtual operating systems on top of a single OS kernel. However, unlike in traditional virtualization, a container doesn’t run multiple operating systems. Rather, it “contains” the multiple guest operating systems in their own userspace, while running a single OS kernel.

At a simple level, containers involve less overhead since there’s no need to emulate the hardware. The big drawback is that you can’t run multiple types of operating systems in the same hardware. You can run 10 Linux instances in a server with container based virtualization, but you can’t run both Linux and Microsoft Server guests side by side.



This chapter introduces you to Linux container technology and the principles that underlie that technology, and also compares traditional virtualization with containerization. Chapter 5 is dedicated to Docker containers and container orchestration technologies such as Kubernetes.

Linux Containers (LXC) allow the running of multiple isolated server installs called containers on a single host. LXC doesn’t use offer a virtual machine – instead it offers a virtual environment with its own process and network space.

Linux containers have analogies in other well known ‘Nix operating systems:

- FreeBSD: Jails
- SmartOS: Zones
- Oracle Solaris: Zones

Linux containers (through Docker) are radically changing the way applications are built, deployed and instantiated. By making it easy to package the applications along with all of their dependencies, containers accelerate application delivery. You can run the same containerized applications in all your environments – dev, test, and production. Furthermore, your platform can be anything: a physical server, a virtual server, or the public cloud.

Containers are designed to provide fast and efficient virtualization. Containerization provides different views of the system to different processes, by compartmentalizing

the system. This compartmentalization ensures guaranteed access to resources such as CPU and IO, while maintaining security.

Since each container shares the same hardware as well as the Linux kernel with the host system, containerization isn't the same as full virtualization. Although the containers running on a host share the same host hardware and kernel, they can run different Linux distributions. For example, a container can run CentOS while the host runs on Ubuntu.

NOTE Linux Containers (LXC) constitute a container management system that became part of the Linux Kernel 2.6.24 in August 2008. As with Docker (see Chapter 5), Linux Containers make use of several Linux kernel modules such as cgroups, SELinux, and AppArmor.

Linux Containers combine an application and all of its dependencies into a package which you can make a versioned artifact. Containers provide application isolation while offering the flexibility of image-based deployment methods. Containers help isolate applications to avoid conflicts between their runtime dependencies and their configurations, and allow you to run different versions of the same application on the same host. This type of deployment provides a way to roll back to an older version of an application if a newer version doesn't quite pan out.

Linux containers have their roots in the release of the chroot tool in 1982, which is a filesystem specific container type virtualization tool. Let's quickly review chroot briefly to see how it compares to modern containerization.

Chroot and Containers

Linux containerization is often seen as an advancement of the chroot technique, with dimensions other than just the file system. Whereas chroot offers isolation just at the file system level, LXC offer full isolation between the host and a container and between a container and other containers.

The Linux `chroot()` command (pronounced “cha-root”) lets a process (and its child processes) redefine the root directory from their perspective. For example, If you chroot the directory `/www`, and when you issue the command `cd`, instead of taking you to the normal root directory (`/`), it leaves you at `/www`. Although `/www` isn't really the root directory, the program believes that it's so. In essence, chroot restricts the environment and that's the reason the environment is also referred to as a jail or chroot jail.

Since a process has a restricted view of the system, it can't access files outside of its directory, as well as libraries and files from other directories. An application must therefore have all the files that it needs right in the chroot environment. The key prin-

ciple here is that the environment should be self-contained within a single directory, with a faux root directory structure.

Linux containers are similar to chroot, but offer more isolation. Linux containers use additional concepts beyond chroot, such as control groups. Whereas chroot is limited to the file subsystem, control groups enable you to define groups encompassing several processes (such as “sshd” for example) and control resource usage for those groups for various subsystems such as the file system, memory, CPU and network resources and block devices.

Applications and their Isolation

Isolating applications is a key reason for using container technologies. In this context, an application is a unit of software that provides a specific set of services. While users are concerned just with the functionality of applications, administrators need to worry about the external dependencies that all applications must satisfy. These external dependencies include system libraries, third-party packages and databases.

Each of the “dependencies” has its own configuration requirements and running multiple versions of an application on a host is difficult due to potential conflicts among these requirements. For example, a version of an application may require a different set of system libraries than another version of the same application. While you can somehow manage to run multiple versions simultaneously through elaborate work-arounds, the easiest solution to managing the dependencies is to isolate the applications.

Virtualization and Containerization

Both containerization and virtualization help address the issues involved in efficient application delivery, where applications in general are much more complex yet must be developed with lower expense and delivered faster, so they can quickly respond to changing business requirements.

At one level, you can view both containers and traditional virtualization as allowing you to do the same thing: let you run multiple applications on the same physical servers. How then, are containers a potentially better approach? Virtualization is great for abstracting from the underlying hardware, which helps lower your costs through consolidating servers, and make it easy to automate the provision of a complete stack that includes the OS, the application code and all of its dependencies.

However, great as the benefits of virtualization are, virtual machines have several limitations:

- By replacing physical servers with virtual servers, you do reduce the physical server units – yet server sprawl doesn't go away – you're simply replacing one type of sprawl with another!
- Virtual technology isn't suitable for microservices that can use hundreds of thousands of processes, since each OS process requires a separate VM.
- Virtual machines can't be instantiated very quickly – they take several minutes to spin up, with means inferior user experience. Containers on the other hand can be spun up blazingly fast- within a few short seconds!
- Lifecycle management of VMs isn't a trivial affair – every VM has a minimum of two operating systems that need patching and upgrading – the hypervisor and the guest OS inside the VM. If you have a virtualized application with 20 VMs, you need to worry about patching 21 operating systems (20 guests systems + 1 hypervisor).

While traditional virtualization does offer complete isolation, once you secure containers with Linux namespaces, CGroups and SELinux, you can get virtually (no pun) the same amount of isolation. Linux containers offer a far more efficient way to build, deploy, and execute applications in today's modern application architectures that use microservices and other new application paradigms. Linux containers offer an application isolation mechanism for lightweight multitenancy and offer simplified application delivery.

Containers, as I've mentioned earlier, have been around for over a decade, and things like Solaris Zones and FreeBSD Jails have been with us for even longer. The new thing about current containerization is that it's being used to encapsulate an application's components, including the application dependencies and required services. This encapsulation makes the applications portable. Docker has contributed substantially to the growth of containerization, by offering easy to use management tools as well as a great repository of container images. RedHat and others have also offered smaller footprint operating systems and frameworks for management, as well as containerization orchestration tools such as Kubernetes (please see Chapter 5).

Benefits offered by Linux Containers

Linux containers enhance the efficiency of application building, shipping, deploying, and execution. Here's a summary of the benefits offered by containers.

Easier and Faster Provisioning

While VMs take several minutes to boot up, you can boot up a containerized application in mere seconds, due to the lack of the overhead imposed by a hypervisor and a guest OS. If you need to scale up the environment using a public cloud service, the ability to boot up fast is highly beneficial.

Lower Costs

Due to their minimal footprint, many more containers fit on a physical server than virtual machines.

Better Resource Utilization

You can monitor containers easily since they all run on a single OS instance. When idle, the containers don't use any server resources such as memory and CPU unlike a virtual machine, which grabs those resources when you start it up. You can also easily remove unused container instances and prevent a virtual machine like sprawl.

Easier Lifecycle Management

Even if you have a large number of applications running on a containerized server, you need to patch and upgrade just a single operating system, regardless of how many containers run on it, unlike in the case of virtual machines. Since there are fewer operating systems to take care of, you're more likely to upgrade than to apply incremental patches.

Greater Application Mobility

Containers make it easy to move application workload between private and public clouds. Virtual machines are usually much larger than containers, with sizes ranging often in the Gigabytes. Containers are invariably small (a few MBs) and so it's easier to transport and instantiate them.

Quicker Response to Changing Workloads

Containers speed up application development due the testing cycles being shorter, owing to the containers including all the application dependencies. You can build an app once and deploy it anywhere.

Easier Administration and Better Visibility

You have far fewer operating systems to manage since multiple containers share the same OS kernel. You also have better visibility into the workload of a container from the host environment, unlike with VMs, where you can't peek inside the VM.

Containers aren't a mere incremental enhancement of traditional virtualization. They offer many ways to speed up application development and deployment, especially in the areas of microservices, which I discussed in Chapter 3. Since microservices can startup and shutdown far quicker than traditional applications, containers are ideal for them. You can also scale resources such as CPU and memory independently for microservices with a container based approach.

Isolation of Services

Let's say your organization has a sensitive application that makes use of SSL to encrypt data flowing through the public internet. If you're using a virtualized setup, the application image includes SSL and therefore, you'll need to modify the application image whenever there are SSL security flaws. Obviously, your application is down during this time period, which may be long, since you'll need to perform regression testing after making changes to SSL.

If you were using a container based architecture on the other hand, you can separate the SSL portion of the application and place it in its own container. The application code isn't intertwined with SSL in this architecture. Since you don't need to modify the application code, there's no need for any regression testing of the application following changes in SSL. A huge difference!

Two Types of Uses for Linux Containers

There are two different ways in which you can employ Linux containers. You can use containers for sandboxing applications, or you can utilize image-based containers to take advantage of the whole range of features offered by containerization. I explain the two approaches in the following sections.

Host Containers

Under the host containers use case, you use containers as lightweight application sandboxes. All your applications running in various containers are based on the same OS as the host, since all containers run in the same user space as the host system. For example, you can carve a RHEL 7 host into multiple secure and identical containers, with each container running a RHEL 7 userspace. Maintenance is easy since updates need to be applied just to the host system. The disadvantage to this type of containerization is that it's limited to just one type of OS, in this example a RHEL runtime.

Image-Based Containers

Image based containers include not just the application, but also application's runtime stack. Thus, the container runs an application that has nothing to do with the host OS. Both the container and application run times are packaged together and deployed as an image. The containers can be non-identical under image based containerization. This means you can run multiple instances of the same application on a server, each running on a different OS platform. This is especially useful when you need to run together application versions based on different OS versions such as RHEL 6.6 and RHEL 7.1, for example. Docker, which we discuss extensively in Chapter 5, is based on image based containers. Docker builds on LXC and it includes the userspace runtime of applications.



You can deploy containers both on bare metal, and on virtualized servers.

The Building Blocks of Linux Containers

Linux containers have become increasingly important as application packaging and delivery technology. Containers provide application isolation along with the flexibility of image-based deployment. Linux containers depend on several key components offered by the Linux kernel, such as the following:

- Cgroups (Control groups) – allow you to group processes for optimizing system resource usage by allocating resources among user-defined groups of tasks
- Namespaces – isolates processes by abstracting specific global system resources and making them appear as a distinct instance to all the processes within a namespace
- SELinux – securely separates containers by applying SELinux policy and labels.

Namespaces, cgroups and SELinux are all part of the Linux kernel, and they provide the support for containers, which run the applications. While there's a bunch of other technologies used by containers, namespaces, cgroups and SELinux account for most of the benefits you see with containers.

In the following sections, let's briefly review the key building blocks of Linux containers:

- Process isolation – namespaces provide this
- Resource management – Cgroups provide resource management capabilities
- Security – SELinux takes care of security

Namespaces and Process Isolation

The ability to create multiple namespaces enables process isolation. It's namespaces that make it possible for Linux containers to provide isolation between applications, with each namespace providing a boundary around applications. Each of these applications is a self-contained entity with its own file system, hostname and even a network stack. It's when it's running within a namespace that an application is considered to be running within a container.

A namespace makes a global system resource appear as a dedicated resource to processes running within that namespace. This helps different processes see different

views of the system, something which is similar to the concept of zones in the Solaris operating system. This separation of the resource instances lets multiple containers simultaneously use the same resource without conflicts. Namespaces offer lightweight process virtualization, although without a hypervisor layer as in the case of OS virtualization architectures such as KVM.



Mount namespaces together with chroots help you create isolated Linux installations for running non-conflicting applications.

In order to isolate multiple processes from each other, you need to isolate them at every place they may bump into each other. For example, the file system and network are two obvious points of conflict between two applications. Application containers use several types of namespaces as described here, each of which helps isolate a specific type of resource, such as the file system or the network.

Namespaces isolate processes and let you create a new environment with a subset of the resources. Once you set up a namespace, it's transparent to the processes. In most Linux distributions, the following namespaces are supported, with RHEL 7 and other distributions adding the user namespace as well.

- mnt
- net
- uts
- pid
- ipc

Let's briefly discuss the namespaces listed here in the following sections.

Mount Namespaces

Normally, file system mount points are global, meaning that all processes see the same set of mount points. Mount namespaces isolate the set of filesystem mount points viewed by various processes. Processes running within different mount namespaces, however, can each have different views of a file system hierarchy. Thus, a container can have a different /tmp directory from that of another container. The fact that each application sees a different file system means that dependent objects can be installed without conflicts among the applications.

UTS Namespaces

UTS namespaces let multiple containers have separate hostnames and domain names, thus providing isolation of these two system identifiers. UTS namespaces are useful when you combine them with network namespaces.

PID Namespaces

PID namespaces allow processes running in various containers to use the same PID, so each container can have its own init process, which is PID1. While you can view all processes running inside the containers from the host operating system, you can only see a container's own set of processes from that container. All processes, however, are visible within the “root” PID namespace.

Network Namespaces

Network namespaces allow containers to isolate the network stack, which includes things such as the network controllers, firewall, iptable rules, routing tables, etc. Each container can use separate virtual or real devices and have its own IP address. Network namespaces remove port conflicts among applications, since each application uses its own network stack, with a dedicated network address and TCP port.

IPC Namespaces

IPC (inter process communication) namespaces allow interprocess communication (IPC) resource isolation, which allows containers to create shared memory segments and semaphores with identical names, although they can't influence those resources that belong to other containers. Inter process communication environment includes things such as message queues, semaphores and shared memory.

Control Groups (cgroups)

Control groups (cgroups for short) let you allocate resources such as CPU time, block IO, RAM and network bandwidth among groups of tasks that you can define, thus providing you fine-grained control over system resources. Using cgroups, the administrator can hierarchically group and label processes and assign specific amounts of resources to these processes, thus making for an efficient allocation of resources.

The Linux `nice` command lets you set the “niceness” of a process, which influences the scheduling of that process. The nice values can range from -20 (most favorable scheduling) to a value of 19 (least favorable to the process). A process with a high niceness value is accorded lower priority and less CPU time, thus freeing up resources in favor of processes with a lower niceness value. Note that niceness doesn't really translate to priority – the scheduler is free to ignore the nice level you set. In traditional systems, all processes receive the same amount of system resources, and so an application with a larger number of processes can grab more system resources com-

pared to applications with fewer running processes. The relative importance of the application should ideally be the criterion on which resources ought to be allocated, but it isn't so, since resources are allocated at the process level.

Control groups let you move resource allocation from the process level to the application level. Control groups do this by first grouping and labeling processes into hierarchies, and setting resource limits for them. These cgroup hierarchies are then bound with the systemd unit tree, letting you manage system resources with the `systemctl` commands (or by editing the system unit files).

A Cgroup is a kernel provided filesystem, and is usually mounted at `/cgroup`, and contains directories similar to `/proc` and `/sys` that represent the running environment and kernel configuration options. In the following sections, I explain how cgroups are implemented in RedHat Enterprise Linux (and Fedora).

Cgroup Hierarchies

You organize cgroups in a tree-based hierarchy. Each process or task that runs on a server is in one and only one of the cgroups in a hierarchy. In a cgroup a number of tasks (same as processes) are associated with a set of subsystems. The subsystems act as parameters that can be assigned and define the “resource controllers” for memory, disk I/O, etc.

In RHEL 7 (and CentOS), the `systemd` process, which is the parent of all processes, provides three unit types for controlling resource usage – services, scopes and slices. `Systemd` automatically creates a hierarchy of slice, scope and service units that provide the structure for the cgroup tree. All three of these unit types can be created by the system administrator or by programs. `Systemd` also automatically mounts the hierarchies for important kernel resource controllers such as devices (allows or denies access to devices for tasks in a group), or memory (sets limits on memory usage by a cgroup's tasks). You can also create custom slices of your own with the `systemctl` command.

Here's a brief description of the three unit types provided by `systemd`:

- **Service:** services let `systemd` start and stop a process or a set of processes as a single unit. Services are named as `name.service`.
- **Scope:** processes such as user sessions, containers and VMs are called scopes and represent groups of externally created processes. Scopes are named as `name.scope`. For example, Apache processes and MySQL processes can belong to the same service but to different scopes – the first to the `apache` scope and the second to the `Mysql` scope.
- **Slice:** a slice is group of hierarchically organized scopes and services. Slices don't contain any processes – it's the scopes and services that do. Since a slice is hierarchical in nature, the name of a slice unit corresponds to its path in the hierar-

chy. If the slice name is `parent-name.slice`, the slice named `parent-name.slice` is a subslice of the parent slice.

The kernel creates the following four slices by default to run the system:

- `-.slice` – root slice
- `system.slice` – default location for system services (systemd automatically assigns services to this slice)
- `user.slice` – default location for user sessions
- `machine.slice` – default location for VMs and Linux containers

Slices are assigned to scopes. Users are assigned implicit subslices and you can define new slices and assign services and scopes to those slices. You can create permanent services and slice units with unit files. You can also create transient service and slice units at runtime through issuing API calls to PID 1. Transient services and slice units don't survive a reboot, and are released after they finish.

You can create two types of cgroups: transient and persistent. You can create transient cgroups for a service with the `systemd-run` command, and set limits on resources that the service can use. You can also assign a persistent cgroup to a service by editing its unit configuration file. The following example shows the syntax for creating a transient cgroup with `systemd-run`:

```
systemd-run --unit=name --scope --slice=slice_name command
```

Once you create the cgroup, you can start a new service with the `systemd-run` command, as shown here:

```
# systemd-run --unit=toptest --slice=test top -b
```

This command runs the `top` utility in a service unit, within a new slice named `test`.

You can override resources by configuring the unit file or at the command line as shown here:

```
# systemctl set-property httpd.service CPUShares=524 MemoryLimit=500M
```

Unlike in virtualized systems, containers don't have a hypervisor to manage resource allocation, and each container appears as a regular Linux process from the point of view of the OS. Using cgroups helps allocate resources efficiently since you're using groups instead of processes. A CPU scheduler, for example, finds it easy to allocate resources among groups rather than among a large number of processes.

Systemd stores the configuration for each persistent unit in the `/usr/lib/systemd/system` directory. To change the configuration of a service unit you must modify the configuration file either manually by editing the file, or with the `systemctl set-property` command.

Viewing the Hierarchy of Control Groups

You can view the hierarchy of the control groups with the `systemd-cgls` command in RHEL 7, as shown in the following output from the command, which also shows you the actual processes running in the cgroups.

```
├─1 /usr/lib/systemd/systemd --switched-root --system --deserialize 20
├─user.slice
│   └─user-1000.slice
│       └─session-1.scope
│           ├─11459 gdm-session-worker [pam/gdm-password]
│           └─11471 gnome-session --session gnome-classic
│           ...
└─system.slice
    ├─systemd-journald.service
    │   └─422 /usr/lib/systemd/systemd-journald
    ├─bluetooth.service
    │   └─5328 /usr/lib/systemd/systemd-locale
    ├─sshd.service
    │   └─1191 /usr/sbin/sshd -D
    ...
```

As the output reveals, slices don't contain processes – it's the scopes and services that contain the processes. The `.slice` is implicit and identifies with the hierarchy's root.

Cgroup Subsystems (Resource Controllers)

In older versions of Linux, administrators used the `libcgroup` package and the `cgconfig` command to build custom cgroup hierarchies. In this section, I show how RHEL 7 moves resource management from the process to the application level, by binding the cgroup hierarchies with the `systemd` unit tree. This lets you manage the resources either through `systemctl` commands or by editing the `systemd` unit files. You can still use the `libcgroup` package in release 7, but it's there only to assure backward compatibility.

A cgroup subsystem is also called a resource controller and stands for specific resources such as CPU or memory. The kernel (`systemd`) automatically mounts a set of resource controllers, and you can get the list from the `/proc/cgroups` file (or the `listsubsys` command). Here are the key controllers in a RHEL 7 system:

- `cpu`: provides cgroup tasks access to the CPU
- `cpuset`: assigns individual CPUs to tasks in a cgroup
- `freezer`: suspends or resumes all tasks in a cgroup when they reach a defined checkpoint
- `memory`: limits memory usage by a cgroup's tasks

Systemd provides a set of parameters with which you can tune the resource controllers.

Optimizing Resource Usage with CGroups

Now that you have a basic idea about cgroups and resource controllers, let's see how you can use cgroups to optimize resource usage.

Let's say you have two MySQL database servers, each running within its own KVM guest. Let's also assume that one of these is a high priority database and the other, a low priority database. When you run the two database servers together, by default the I/O throughput is the same for both.

Since one of the database services is a high priority database, you can prioritize the I/O throughput by assigning the high priority database service to a cgroup with large number of reserved I/O operations. At the same time, assign the low priority database server to a cgroup with a lower number of reserved I/O operations.

In order to prioritize the I/O throughput, you must first turn on resource accounting for both database servers:

```
# systemctl set-property db1.service BlockIOAccounting=true
# systemctl set-property db2.service BlockIOAccounting=true
```

Next, you can set the priority by setting a ratio of 5:1 between the high and low priority database services, as shown here:

```
# systemctl set-property db1.service BlockIOWeight=500
# systemctl set-property db2.service BlockIOWeight=100
```

I employed the resource controller BlockIOWeight in this example to prioritize the I/O throughput between the two database services, but you could also have configured block device I/O throttling by setting the blkio controller to achieve the same result.

SELinux and Container Security

The two major components of the container architecture that you've seen thus far – namespaces and cgroups aren't designed for providing security. Namespaces are good at making sure that the /dev directory in each container is isolated from changes in the host. However, a bad process from a container can still potentially hurt the host system. Similarly, cgroups help with the avoiding of denial of service attacks since they limit the resources any single container can use. However, SELinux, the third major component of modern container architectures is designed expressly to provide security, not only for containers, but for also normal Linux environments.

RedHat has been a significant contributor to SELinux over the years, along with the Secure Computing Corporation. The National Security Agency (NSA) developed

SELinux to provide the Mandatory Access Control (MAC) framework often required by the military and similar agencies. Under SELinux, processes and files are assigned a type and access to them is controlled through fine-grained access control policies. This limits potential damage from well-known security vulnerabilities such as buffer overflow attacks. SELinux significantly enhances the security of virtualized guests, in addition to the hosts themselves.

How SELinux Works

SELinux implements the following mechanisms in the Linux kernel:

- Mandatory Access Control (MAC)
- Multi-level Security (MLS)
- Multi-category security (MCS)

The sVirt package enhances SELinux and uses Libvirt to provide a MAC system for containers (and also for virtual machines). SELinux can then securely separate containers by keeping a container's root processes from affecting processes running outside the container.

Enabling SELinux

SELinux can be disabled or made to run in a permissive mode. In both of these modes, SELinux won't provide sufficient secure separation among containers. Following is a brief review of SELinux modes and how you can enable it.

SELinux operates in two modes (three if you want to add the default “disabled” mode) – enforcing and permissive. While SELinux is enabled in both the Enforcing and the Permissive modes, SELinux security policies are enforced only in the Enforcing mode. In the Permissive mode, the security policies are read but not applied. You can check the current SELinux mode with the `getenforce` command:

```
# getenforce
Enabled
#
```

There are several ways to set the SELinux mode, but the easiest way is to use the `setenforce` command. Here are ways you can execute this command:

```
# setenforce 1
# setenforce Enforcing
# setenforce 0
# setenforce Permissive
```

The first two options set the mode to Enforcing and the last two to Permissive.

sVirt

RHEL 7 provides Secure Virtualization (sVirt), which integrates SELinux and virtualization, by applying MAC (Mandatory Access Controls) when using hypervisors and VMs. sVirt works very well with KVM, since both are part of the Linux kernel.

By implementing the MAC architecture in the Linux kernel, SELinux limits access to all resources on the host. In order to determine which users can accept a resource, each resource is configured with a SELinux context such as the following:

```
system_u:object_r:httpd_sys_content_t:s0
```

In this example, here's what the various entities in the SELinux context stand for:

- `system_u`: is user
- `object_r`: role
- `httpd_sys_content_t`: type
- `s0`: level

The goal of sVirt is to protect the host server from a malicious instance, as well to protect the virtual instances themselves from a bad instance. The way sVirt does this by configuring each VM created under KVM virtualization to run with a different SELinux label. By doing this, it creates a virtual fence around each of the VMs, through the use of unique category sets for each VM.

It's common for Linux distributions to ship with Booleans that let you enable or disable an entire set of allowances with a single command, such as the following:

- `virt_use_nfs`: controls the ability of the instances to use NFS mounted file systems.
- `virt_use_usb`: controls the ability of the instances to use USB devices.
- `virt_use_xserver`: controls the ability of the instances to interact with the X Windows system

Linux Containers versus Virtualization (KVM)

Virtualization (such as KVM virtualization) and containers may seem similar, but there are significant differences between the two technologies. The most basic of the differences is that virtualization requires dedicated Linux kernels to operate, whereas Linux containers share the same host system kernel.

Your choice between virtualization and containerization depends on your specific needs, based on the features and benefits offered by the two approaches. Following is a quick summary of the benefits/drawbacks for the two technologies.

Benefits and Drawbacks of Virtualization

Assuming you're using KVM virtualization, you can run operating systems of different types, including both Linux and Windows, should you need it. Since you run separate kernel instances, you can assure separation among applications, which ensures that issues with one kernel don't affect the other kernels running on the same host. In addition security is enhanced due to the separation of the kernels. On top of this, you can run multiple versions of an application on the host and the VM, besides being able to perform virtual migrations as I explained earlier.

On the minus side, you must remember that VMs need more resources and you can run fewer VM on a host compared to the number of containers you can run.

Benefits and Drawbacks of Containers

Containers help you isolate applications, but maintaining the applications is a lot easier than maintaining them on virtual machines. For example, when you upgrade an application on the host, all containers that run instances of that application will benefit from that change. You can run a very large number of containers on a host machine, due to their light footprint. Theoretically speaking, you can run up to 6000 containers on a host, whereas you can only run a few VMs on a host.

Containers offer the following additional benefits:

- **Flexibility:** since an application's runtime requirements are included with the application in the container, you can run containers in multiple environments.
- **Security:** since containers typically have their own network interfaces and file system that are isolated from other containers, you can isolate and secure applications running in a container from the rest of the processes running on the host server
- **Performance:** typically containers run much faster than applications that carry the heavy overhead of a dedicated VM
- **Sizing:** since they don't contain an entire operating system unlike VMs, containers of course, are very compact, which makes it quite easy to share them.
- **Resource allocation:** LXC helps you easily manage resource allocations in real time.
- **Versatility:** You can run different Linux distributions on the same host kernel using different containers for each Linux distribution.

Linux Containers and KVM Virtualization – the Differences

The key difference between KVM virtual machines and Linux containers is that KVM VMs require a separate kernel of their own while containers share the same kernel from the OS.

You can host more containers than VMs for a given hardware, since containers have a light footprint and VMs are resource hungry.

KVM Virtualization lets you:

- Boot different operating systems, including non-Linux systems.
- Separate kernels mean that terminating a kernel doesn't disable the whole system.
- Run multiple versions of an application on the same host since the guest VM is isolated from changes in the host.
- Perform live migrations of the VMs

Linux Containers: **Are designed to support the isolation of applications.** Since system wide changes are visible inside all containers, any change such as an application upgrade will automatically apply to all containers that run instances of the application. ** The lightweight nature of containers means that you can have a very large number of them on a host, with the maximum number running into 6000 containers or more on some systems.

Limitations of LXC

Unlike a fully virtualized system, LXC won't let you run other operating systems. However, it's possible for you to install both a virtualized (full or para) system on the same kernel as the LXC host system and run both the virtualized guests and LXC guests simultaneously. Virtualized management APIs such as libvirt and ganeti are helpful if you wish to implement such as hybrid system.

Container benefits

As organizations move beyond monolithic applications to microservices, new application workloads involve a connected matrix put together to server specific business needs, but easily rearrangeable into a different format. Containers are a key part of this new application architecture. For developers who create applications, containers offer these benefits:

- Better quality releases
- Easier and faster scalability of the applications
- Isolation for applications

- Shorter development and test cycles and fewer deployment errors

From the point of the IT operations teams, containers provide:

- Better quality releases
- Efficient replacement of full virtualization
- Easier management of applications When a container is instantiated, the processes execute within a new userspace created when you mount the container image. The kernel ensures that the processes in the container are limited to executing system calls only from their own namespaces such as the mount namespace and the PID namespace. The namespaces are containerized in this case.

You can take the same containerized image and run it on a laptop or servers in your datacenter, or on virtual machines in the cloud.

A virtual machine packages virtual hardware, a kernel and a user space. A container packages just the userspace – there's no kernel or virtual hardware in a container.

Linux Container Adoption Issues

Linux containers are starting to be used widely, with some large cloud service providers already using them at scale. Following are some concerns that have led to a slower than expected adoption of Linux containers.

- Security: security issues are a concern with enterprise adoption, due to the fact that kernel exploits at the host OS level will mean that all containers living on that host are at risk. Vendors are therefore fine tuning security techniques such as mandatory access control (MAC) to tighten things up security wise. SELinux already offers Mandatory access controls, and a different project named libseccomp lets you eliminate syscalls, which prevent hacked containers from compromising the kernel supporting them.
- Management and orchestration: vendors are working on creating frameworks for managing container images and orchestrating their lifecycle. New tools are being created for supporting containers. Docker is a great framework that makes it very easy to create, manage and delete containers. Docker can also limit the resource containers can consume, as well as provide metrics on how containers are using the resources. New tools for building and testing containers are in the wings as well. Linux container adoption will accelerate by agreeing on a standard for inter-container communications, using solutions such as virtual switching ng, hardware enhanced switching and routing and so on.

Managing Linux Containers

While SELinux, Cgroups and namespaces make containerization possible, a key missing piece of Linux containers is the ability to manage them – Docker solves this problem very nicely, as explained in the next chapter, which is all about Docker containers.

LXC is a userspace interface for providing containment features for the Linux kernel. It enables you to easily create and manage both system and application containers. It also helps you easily automate container deployment. LXC seeks to create a Linux environment that's close to the standard Linux installations, but without using a separate kernel. LXC containers are usually regarded as a midway solution between a chroot and a full-fledged virtual machine.

LXC is free software, with most of the code releases under the terms of a GNU license. While LXC, is quite useful, it does have some requirements such as the following:

- Editing of configuration files for controlling resources
- It (LXC) maybe implemented differently between distributions, or even among different releases of the same distribution

Docker offers a much more efficient and powerful way to create and manage Linux containers. Docker is an application that enables almost effortless management of Linux containers through a standard format. Docker isn't a totally new technology – it builds on the concepts you've seen earlier in this chapter, such as namespaces, cgroups and LXC to go beyond what's possible with userspace tools such as LXC. Besides helping you efficiently manage containers, Docker, since it is a type of image based containerization, makes containers portable, thus making it easy to share them across hosts.

With this useful background of Linux containers, let's turn to a discussion of Docker containers in the next chapter.

Working with Docker Containers

Docker is the most well-known container platform that's increasingly being deployed for developing and running portable distributed applications. A Linux system administrator (or a developer) can use Docker to build, ship and run applications on various platforms, including the cloud, both on physical as well as virtual machines.

Chapter 4 explained the foundations of generic containers, which are Cgroups, SELinux and namespaces. Where Docker comes in is that it lets you easily package, run and maintain containers with handy tools, both from the command line and through HTTP APIs. This is what makes it possible for you to easily package applications and their runtime environments into self-contained images. A simple Dockerfile is what you use to package the application and its runtime environment together.

Few technologies are affecting and disrupting established ways of running applications as Docker, which provides a sophisticated way to package and run your applications. As a Linux administrator, you can get Docker up and running as a service in a few short minutes. Not only that, it's really easy to manage and monitor the containers that you run in a development or production environment.



At the center of modern containerization is the Open Container Initiative launched in 2016. The initiative is managed by the Linux Foundation, to create standards around container formats and their runtime environment.

A key reason for the wide spread success of containerization and Docker, especially, is that containerization lets you efficiently deploy applications in a cloud based environment. Docker has made it extremely easy to containerize applications. However, it's also important to understand that a lot of things are still evolving when it comes to

containerization – only about 40% of Docker’s customers are actually running the containers in production.



Containers make it possible to “overwrite” a container that’s part of a running application, which means less downtime when introducing changes.

Single containers are pretty easy to run, but managing a set of related containers is a complex affair. Fortunately, there are several tools such as Kubernetes, Swarm and others that allow you to do precisely this, and I show how you can take advantage of these container orchestration technologies.

Containers isolate applications from other apps running on a host, but don’t offer the same extent of isolation as VMs – since VMs are ahead of the curve in times of security, a lot of people run containers inside VMs, thereby getting the best of both the worlds, so to speak.

You can run Docker in a number of Linux systems, such as Ubuntu, Fedora, RHEL, and CentOS. Containerization means doing things in a small way and hence you really don’t need a traditional bloated Linux system to run Docker. In addition to stripped down versions of the standard Linux distributions, you can run Docker on specially designed lightweight Linux systems optimized for container deployments, such as CoreOS and Atomic Host.

Containers include both the application code and its configuration and associated dependencies. Since the Linux operating system that supports these containers doesn’t have to support all the app’s dependencies any longer, stripped down container-oriented operating systems such as CoreOs Red Hat’s Atomic Host are increasingly becoming popular as vehicles for running containers in production. I explain the CoreOS and Atomic Host operating systems later in this chapter.

Docker Basics

Docker containers package software in a complete filesystem that includes everything an application needs to run. This means the app will always run the same way no matter the environment. Docker containers offer the following key benefits:

- **Lightweight** – the images are built from layered file systems, so multiple images share common files, thus reducing the disk usage. They also share the same OS kernel, thus using RAM more efficiently.
- **Open Standards**: The open standard model means containers run on all Linux distributions.

- **Security:** Containers isolate applications from each other as well as from the infrastructure itself.

The main use of Docker is to create and run applications designed and configured to run within a container. Here, it's good to pause and understand what a containerized application is. Traditional applications run directly on a host operating system. They use the Linux server's file system, process table, network interfaces and ports and so on. The problem with this set up is that you can't run multiple versions of a software package on the same server, as this would cause conflicts. It's also hard to move an application since you'd have to move the code as well as its dependencies, which isn't easy to do.

When you create an application container, you want to make it easy to share the containers with others and distribute it. Docker is a container runtime engine that makes it easy to package applications and push them to a remote registry, from where other users can pull those applications.

Often people are a bit confused as to how Docker containers compare to traditional virtual machines. It's simple: a VM emulates a foreign environment while a container is a lightweight OS designed to make self-contained applications portable. Due to their extreme lightweight nature, applications running in containers incur very little overhead.

A key issue in developing and deploying applications is the packaging together of all the required dependencies for a software application, such as its code, the runtime libraries, etc. With Docker, you simply package all the required code and dependencies into a single Docker image. As explained in Chapter 4, unlike in traditional virtualization where the VMIs (virtual machine images) run on separate guest operating systems, all Docker images in a server run within a single OS kernel.

Docker is playing a significant role in the recent ascendancy and popularity of DevOps. Administration

When Docker Isn't Right for You

Sometimes Docker isn't right for you! Although you can run virtually anything you want within a container, you really don't want to run databases and any kind of stateful applications within Docker containers. Docker containers are ideal for applications such as microservices which don't maintain state. Since resizing the CPU and memory of a container involves restarting the container, applications that need dynamic resizing for CPU and RAM aren't ideal for Dockerization.



Docker is more suitable for applications such as microservices that don't maintain state.

Docker uses iptables to provide NAT between the host IP and the container's IPs. As a result, if your applications need high network throughput, they aren't ideal for Docker.

What Docker Consists of

In order to get going with Docker containers, you need to learn about the various components that play a role in Docker containerization. The following sections explain:

- The Docker Project
- The Docker Hub Registry
- Docker Images and Docker Containers
- The docker Command

The Docker Project

The Docker Project, of course is what developed the container format that we know today as Docker. The project has the twin goals of simplifying application development and distribution. The Docker Project provides a format for software containers, and also provides various tools to help you manage, provision and orchestrate containers. The Docker Project also manages the central repository called the Docker Hub Registry.

The Docker Hub Registry

The Docker Hub Registry (<https://registry.hub.docker.com>) acts as a store for you to save and develop your Docker container images. When you request a Docker container image that's not on your system, Docker by default checks the Docker Hub Registry for that image.



You store containers in registries and you can make the images available to download to systems running Docker.

The Docker Hub Registry hosts the official repositories for all Linux distributions and several application projects. You can create a Docker user account and maintain your own Docker repository where you can push your Docker images to. In order to store container images privately, you can create your own Docker Registry.

The Docker Service

The Docker service is the same as the Docker engine. It's the Docker service that grabs the images you pull from the Docker Hub Registry.

Docker Images and Docker Containers

Before we dive deep into Docker containerization, it's highly useful to clarify certain commonly used terms such as Docker images, containers, and so on.

Docker Images

A Docker image consists of all the components of an application, such as libraries, executables and configuration files – together, these components enable an application to run.

You can build a Docker image through a Dockerfile, or by downloading pre-built Docker images from the Docker Hub (<https://hub.docker.com/>). The Dockerfile helps you automate the creation of a Docker container image. A Dockerfile consists of instructions regarding the software to be downloaded, the commands to be run, networks, environment variables, and the required files and directories that must be added to the container's filesystem.

Docker Containers

A Docker container is the running instance of a Docker image, and thus the instance is what performs the actual work. A Docker container is an isolated environment where the Docker image will run. Each Docker container will include its filesystem and you can copy files from the host server to the Docker containers.

Applications can require other software, and this means that Docker containers can be linked together, with software from one container made available to another Docker container. A good way to manage a container with other containers is to use a product called Kubernetes to orchestrate the containers into pods.



A Docker image is a stored version of a container. A container is an instance of an image.

You start a container with the docker run command. You can run a container in the background (detached) mode, so the instance can keep running after the docker run command exits.

The Docker Command

The docker command is what you use to manage containers. You can run the command from the command line or run it as a service daemon that handles requests to manage containers.

There are separate commands for managing images and for managing instances. A command such as `docker images` will show you all the Docker images on your system. To view the containers that are actually running, you use commands such as `docker ps`. The `docker start` and `docker stop` commands let you start and stop a Docker container instance.

The docker command is quite versatile, and lets you do all the following:

- Create and remove containers and images
- Manage running containers
- Manage images and work with Docker registries
- Watch Docker events and log messages, and the CPU and memory usage statistics for containers

What Linux Administrators should know in order to support Docker

Docker offers a different paradigm from what most of us are used to, in terms of how various features of the Linux operating system are made available to Docker containers. In order to work with Docker, you should know the following:

- Host privileges
- Networking
- Storage

Host Privileges

Containers use host privileges to directly access a limited set of OS features such as the process table, the IPC namespace, specific CPUs, and devices. Super privileged containers are allowed to not only access, but also change the system.



You use regular containers for running applications, and super privileged containers to add tools to help you access the host system.

Docker Networking

You have to follow specific rules for managing the host network interfaces from inside the Docker containers. You can use ambassador containers (proxies) and service discovery solutions (explained in detail later in this chapter) to connect services running on different hosts. Since these solutions require you to expose ports through the hosts and also don't scale well. You may need to provide full cross-host networking solutions by providing IP connectivity between containers.

Following are the three commonly used Docker networking modes:

- **Bridge:** The Docker Bridge, called `docker0` is used to connect containers. Docker instantiates a veth pair connecting `eth0` in the container to the Docker bridge (`docker0`). It uses IP forwarding and iptables rules used for IP masquerading to provide the external connectivity. The bridge networking mode is good for development, but it's not very efficient, making it unsuitable for production purposes.
- **Host:** Under the Host networking mode, the container shares the host's networking namespace, thus exposing it to the public network. This is much more efficient than the bridge mode.
- **Container:** In this mode, a container uses the networking namespace from a different container. This mode allows for reuse of efficient network stacks but suffers from the drawback that all containers sharing a network stack need to use an identical IP stack. Kubernetes uses this mode.

Earlier on, Docker used links to network containers, but the current way is for networks to be created and managed separately from containers. Two key objects – network and services play a key role in networking. When you launch a container, it can be assigned to a specific network. Containers can publish services that let them be contacted through their name, without needing the links.

For cross-host networking of clusters of containers, there are several networking solutions such as the following:

- **Overlay:** This is the “batteries included” Docker solution for cross-host networking, and is probably the best solution during development and for small cloud-based environments.. Overlay uses VXLAN tunnels to connect hosts with their IP namespace and for external connectivity, relies on NAT.
- **Weave:** This is a more complete solution for networking, and includes WeaveDNS for service discovery and load balancing. Weave is easy to use and hence good for development purposes.
- **Flannel:** Flannel is a networking solution meant mainly for CoreOS Docker containers. Flannel assigns subnets to each of the hosts and the subnets are used to assign IPs to individual containers. In Kubernetes, Flannel is used to assign

unique IPs to each pod. Flannel offers a good simple solution in many production scenarios.

Docker Storage

You can use bind mounts to connect the host storage to a Docker container. Docker storage volumes are normal Linux directories on the host server that are bind mounted to the container. The following example shows one way to initialize a volume. You bind mount a volume at runtime with the `-v` flag as shown here:

```
$ docker run -it --name test-container -h CONTAINER -v /data/ Debian /bin/bash
```

This docker command will make the directory `/data` (inside the container named `test-container`) into a volume.

Alternatively, you can create a volume inside a container by specifying the `VOLUME` instruction in the container's Dockerfile:

```
FROM debian:wheezy
VOLUME /data
```

This does the same thing as the specifying of the `-v` option in the docker command I showed earlier.

Setting up the Docker Container Run-Time Environment

Now that you've learned the basic terminology used in Docker environments, let's get started with Docker containers by learning how to get them to run on a Fedora Linux distribution that's sponsored by Red Hat. In order to use Docker, you need to set up the Docker engine on a real or virtual Linux server. In order to get started with using Docker, all you need is a software package named `docker`, which contains the all-important `docker` command.

The best way to get started is by installing a standard Linux distribution with a desktop interface so you can develop and debug the containers you create. Once you're ready to go to production, you can deploy them through a container oriented OS such as CoreOS or Project Atomic.

Let's learn how to set up the Docker environment on a Linux CentOS server. Before you can install Docker, of course, you must download and install CentOS on either a real host or on a virtual machine (I'm using a virtual server that I created with Vagrant and VirtualBox). Once you do this you're ready to perform the Docker installation. On a RHEL based system, the Docker container's package is named `docker`. Here's how you perform the installation:

1. Update all the installed packages to their latest versions.


```
# yum update
```

1. Restart the server.

```
# reboot
```

1. Install the Docker package (this contains the docker command that you'll use to manage Docker containers).

```
# yum install docker
```

1. Start the Docker service. Before starting the service, you must enable the service.

```
# systemctl enable docker.service  
# systemctl start docker.service
```

You're done at this point and you can check the status of Docker in the following way:

```
$ [vagrant@localhost ~]$ systemctl status docker.service  
docker.service - Docker Application Container Engine  
Loaded: loaded (/usr/lib/systemd/system/docker.service; enabled; vendor preset: disabled)  
Active: active (running) since Sun 2016-03-13 13:44:25 EDT; 19h ago  
Docs: http://docs.docker.com  
Main PID: 881 (docker)  
Memory: 2.9M  
CGroup: /system.slice/docker.service  
└─881 /usr/bin/docker daemon --selinux-enabled  
[vagrant@localhost ~]$
```

Note the following:

- The docker service is enabled and is active
- The service name is the docker command (/usr/bin/docker)
- The `-d` option means that the Docker service runs as a daemon process
- The Docker service is SELinux enabled.

Now that you've installed Docker, you're ready to play with it, using the docker command.

Getting Information about the Containers

You can execute the `docker info` command to verify that Docker is installed, and get a bunch of useful information, as shown here:

```
[root@localhost ~]# docker info  
Containers: 19  
Images: 73  
Storage Driver: devicemapper  
Pool Name: docker-253:0-525784-pool
```

```

Pool Blocksize: 65.54 kB
Backing Filesystem: extfs
Data file: /dev/loop0
Metadata file: /dev/loop1
Data Space Used: 5.878 GB
Data Space Total: 107.4 GB
Data Space Available: 31.66 GB
/metadata
Library Version: 1.02.107-RHEL7 (2015-12-01)
Kernel Version: 3.10.0-327.10.1.el7.x86_64
Operating System: CentOS Linux 7 (Core)
CPUs: 1
Total Memory: 993.3 MiB
Name: localhost.localdomain
...
[root@localhost ~]#

```

As with any Linux service, you can configure Docker to start on boot. On an Ubuntu server, for example, you configure the docker daemon to start on boot, by running the following command:

```
$ sudo systemctl enable docker
```

Running Container Images

Once you get going with Docker by starting the Docker service as shown earlier, you can run a container without having any images on your system. This seems like sheer magic when you actually do this on your system! When you execute the docker run command, it goes and gets the image you are looking for from the Docker Hub (by default). The following example shows how to get and run the fedora image on your server.

```

# docker run fedora cat /etc/os-release
Unable to find image 'fedora:latest' locally
00a0c78eeb6d: Pull complete
Status: Downloaded newer image...

NAME=Fedora
VERSION="22 (Twenty Two)"
ID=fedora
...

```

When you run the docker command, this is what happens:

- The command locates the container image (fedora in our example), and downloads it to your laptop or server. The command looks in your local server first for the fedora: latest image, and if it fails to find it there, will download the image from the docker.io registry (same as the Docker Hub Registry).
- The command starts the container using the image.

- The docker command will finally run the cat command, which you passed as an attribute to the docker command.

Once you run this command, Docker stores the image on your local system. A container, as I explain shortly, is ephemeral in nature and disappears off the face of the earth when you stop it. When you restart the same container, Docker uses the stored image on the local server. Thus, the very first run will take much more time since Docker needs to download the image from the Docker Hub – subsequent runs are quite fast, since all Docker needs to do is to spin up the container using the image that you already have.

Managing Containers

In a previous section where I showed the first example of the docker run command, the command started the Fedora container and executed the Linux cat command – and then it exited. Often, a container such as one that runs a web server service, for example, needs to run continuously. So, it's good to learn how to run a container in the background.

Running Interactive Containers

Use the `-d` (for detached) option with the docker run command, to run a container in the background. In order to run it in the foreground in an interactive fashion, specify the `-i` option and to open a terminal session, add the `-t` option.

A common use case for an interactive container is when you want open a shell to do things inside the container, just as you'd log into a Linux server to view, modify and execute various things. The following example shows how to open a shell in the Fedora image that Docker has already downloaded on to my server.

```
# docker run -it fedora /bin/bash
bash-4.3#
You can view the processes running inside this Fedora container by issuing the ps command:
bash-4.3# ps -e
  PID TTY          TIME CMD
    1 ?            00:00:00 bash
    7 ?            00:00:00 ps
#
```

Wow! Instead of the usual hundreds of processes that run on a full-fledged Linux server, this Fedora container has but two processes! The process with the ID 1 isn't the usual init (or systemd) process on your Linux server – instead it's the bash command. The only other process is the ps command you just used – that's it! This is the amazing thing with containerized systems – they're lean to start with and you can make them as lean or as fat as you want it, to suit your needs. This is the big difference between a container and a whole virtual machine.

If you're curious to see the list of installed packages on this new server, issue the familiar `rpm -qa` command:

Making your base Image Heftier

In the previous section, I showed how to create a very light Fedora container using a fedora base image. There are two ways to create a heftier container with more capabilities (more software) – you can simply download an alternative fedora base image with more stuff already baked in. Or, you can add software to a simple base image such as what I've got here. Docker lets you add software to a running container, as shown here:

```
bash-4.3# yum install iproute net-tools bsd-games words \
          vsftpd httpd httpd-manual -y
Resolving Dependencies
--> Running transaction check
...
Complete!
bash-4.3# exit
```

This example shows how you can add any software you want to a container. Use the `yum` or `dnf` (RHEL and Fedora) or the `apt-get` (Debian and Ubuntu) commands from within a container, in order to add software.

Committing a Container

If I want to use my fedora container again, with all the software I've just added, of course the container must have these software already added to it. I just exited the running container but didn't stop that container, so it's still running, as you can see here:

```
# docker ps -a
CONTAINER ID   IMAGE          COMMAND                  CREATED
STATUS        PORTS         NAMES
88f6c09523b5   fedora:latest  "/bin/bash"             3 hours ago
Exited (0) 7 seconds ago          trusting_heisenberg
#
```

You know that this running container has been fortified with additional software. Why not save the container as a new image that you can use whenever you wish? You can do this with the `docker commit` command, as shown here:

```
# docker commit -a "Sam Alapati" 88f6c09523b5 Myfedora
```

The new image `Myfedora` is now stored as a separate image on your server, and is part of the Docker image set:

```
# docker images
REPOSITORY TAG          IMAGE ID          CREATED          VIRTUAL SIZE
```

| | | | | |
|---------|--------|--------------|---------------|----------|
| fedora | latest | 834629358fe2 | 1 month ago | 422.2 MB |
| testrun | latest | 226f7543f12a | 3 minutes ago | 431.6 MB |

Running Commands within a Container

Now that I've added the `ip-route` and the `net-tools` packages to the bare bones base `fedora` image, I know I can run the `ip` and the `route` commands inside the container, as shown here:

```
# ip addr show docker0
5: docker0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue
state UP group default
    link/ether 56:84:7a:fe:97:99 brd ff:ff:ff:ff:ff:ff
    inet 172.17.42.1/16 scope global docker0
        valid_lft forever preferred_lft forever
#
```

The `ip` command reveals the IP address of the host's `docker0` interface – `172.17.42.1/16`. The host automatically assigns IP addresses to each container as it's spawned, using DHCP to do so.



Containers use a router through the host server's `docker0` interface to access networks outside the local server.

Linking Containers

You use Docker links to enable containers on the same host to communicate with each other. By default, communications among the containers in a Docker network aren't exposed to the host network, since they use an internal Docker network.

Running Services inside a Container

Earlier, you learned how to issue various commands from inside a container, depending on the software that's available in the container. However, simply being able to run those types of commands (`ps`, `route`, `ip`, etc) isn't going to do you a whole lot of good, if that's all you can do! The great power of containers comes through the ability to run a service such as a web server from the container.

Running containerized services lets you preconfigure everything you need to run that service and thus make the service truly portable. In addition, containerization of a service helps you to easily scale the service – if you need more web servers to manage the workload, simply start up more containers, that's it - all you need to do is to make sure that each of the container you expose as a service uses a different IP address and/or port.

Creating a Containerized Web Server

Let's learn how to run an Apache web server inside a Docker container. An Apache web server serves web content from the `/var/www` directory by default. Also by default, it listens on the TCP ports 80 and 443 (secure). In the following example, I have my web server serve the basic content (`index.html`) from the `/var/www/html` directory. Here are the steps to get this done:

1. Create a directory for holding the web content.

```
# mkdir -p /var/www/html
```

1. Set the appropriate SELinux context (since I'm using Fedora, which is a RHEL compatible system).

```
# chcon -R -t httpd_sys_content_t /var/www/
```

3. 4. Add some content to the `index.html` page.

```
# echo " Apache is running, yeah!!" > /var/www/html/index.html
```

Now that you've prepared the ground for the Apache container, run the `httpd` service using a Docker image. I use the `mycont` image I had created in the previous section.

```
# docker run -d -p 80:80 -p 443:443 --name=MyApacheServer \
  -v /var/www:/var/www mycont \
  /usr/sbin/httpd -DFOREGROUND
```

Note the following:

- In this example, the `-d` option tells Docker to keep the newly created container running in the background until I tell it to stop running it – remember that by default the container runs a command and quits! The `-DFOREGROUND` option runs the `httpd` daemon in the foreground.
- The `-p` option publishes the container ports to a host port. The port to the left of the colon belongs to the host and the one after the colon is the container port. So, in this example, I'm exposing both the TCP HTTP port (80) and the HTTPS (443) port to the same port numbers on the host server.
- The important option you should know about here is the `-v` option. The `-v` option shows the bind mounted volumes. Using this option, you mount directories on the host system to directories on the container. In this case, I mount the default Apache directory for web content (`/var/www`) on the host to an identically named directory on the container.
- The image name is `mycont` – the name of the container created from this image is specified by the `--name` parameter. The name parameter is optional, but quite

useful in practice. So, if you want to open a shell in the new MyApacheServer container, you can do so by issuing the normal docker command:

```
# docker exec -it MyApacheServer /bin/bash
```

Now, the web server is running in the Docker container. You can test that the ports are configured correctly by issuing a simple curl command such as this:

```
# curl http://localhost/index.html
```

I see the contents of the index.html file stored in the /var/www/html directory in the host server, since port 80 from the container is exposed to port 80 on the host, and by using bind volumes, the container also uses the same Apache directories as the host.

Controlling the Container Resource Usage

When you run multiple containers on a host, resource allocation becomes a significant issue. You can control the usage of the host resources such as RAM and CPU with options such as `--memory`, `--memory-swap`, `--cpu-shares` and `--cpuset-cpus`.

Running Privileged Containers

By default, a container has limited access to a host's capabilities. For example, it can't access namespaces such as the process table and the IPC namespace on the host. Although, in general, you want to limit a container's access to the host and to other containers, the concept of a privileged container (also called a super privileged container) lets you grant a container greater access than what's allowed by default.

Building Docker Images

You can build Docker images in two ways:

- A developer can build an image and push it to a repository
- You can automatically build images with a CI/CD system, following each code push

In a production system, the best strategy would be to use a CI/CD system such as Jenkins to automate the building of images following code pushes. As each container is built, it's uploaded to a repository from where the CI/CD system can download and run the image

Building Images with a Dockerfile

You can automate the creation of images through a Dockerfile, which is analogous to the Vagrant configuration file named Vagrantfile, which lets you configure VMs easily. The Dockerfile is a text file that consists of instructions, and the Docker

builder reads this file and executes the instructions one by one to create the image.

2.1.1 Let me build a simple image to show how to create and use a Dockerfile. Here are the steps:

First, create an empty directory

```
$ mkdir sample_image
```

Next, create a file named Dockerfile and add the following to the file.

```
# Pick up the base image
FROM fedora
# Add author name
MAINTAINER Sam R. Alapati
# Add the command to run at the start of container
CMD date
```

Finally, Build the image with the docker build command.

```
$ cd sample_image
$ docker build -t /fedora/test .
```

The `-t` option specifies a tag name.

As you can see, the docker build command saves the intermediate images so it can use them in later builds to speed up the build process. After running each of the instructions in the Dockerfile, Docker commits the intermediate state image and runs a container with that image for executing the next instruction. It then removes the intermediate containers from the previous step. Once the last instruction is executed, it creates the final image.

Understanding the Dockerfile

A Dockerfile has the following format:

- **INSTRUCTION** arguments: It's customary to specify the instructions in upper case. A line with the `#` at the beginning is a comment. A Dockerfile usually has the following types of instructions:
- **FROM**: This must be your first instruction in the file, and it sets the base image. The default behavior is the latest tag:

```
FROM <image>
```

You can have multiple FROM instructions if you want to create multiple images. If you just specify the image name, such as FROM fedora, the build program will download the image from the Docker Hub. If you want to use your own private images or a third-party image, you need to specify it as shown in this example:

```
FROM registry-host:5000/
```


- **MAINTAINER:** sets the name of the author for the image:

```
MAINTAINER <name>
```

- **RUN:** you can execute the RUN instructions in the shell, or directly as an executable:

```
RUN <command> <param1>... <paramN>
Or
RUN ["executable", "param1", "param2", ..., "paramN"]
```

- **LABEL:** Use the LABEL instruction to tag a distribution, as in:

```
LABEL distro=fedora21
```

- **CMD:** when starting a container, the CMD instruction offers a default executable:

```
CMD ["executable", "param1", ..., "paramN" ]
```

- **ENTRYPOINT:** This instruction lets you configure the container as an executable:

```
ENTRYPOINT ["executable", "param1", ..., "paramN" ]
ENTRYPOINT <command> <param1> ... <paramN>
```

- **EXPOSE:** exposes the network ports on the container on which to listen at runtime:

```
EXPOSE <port> [<port> ... ]
```

Alternatively, you can expose ports at runtime when starting the containers.

In addition, there are additional instructions such as those that let you set the environment variables (ENV), ADD (copy files into the image), and COPY (lets you copy files from source to destination). The `docker build --help` command shows all the possible instructions you can use.

Building an Apache Image using a Dockerfile

Let me use a simple example to show how to build a Docker image through a Dockerfile. You can get the Dockerfile by installing the `fedora-dockerfiles` package in a Fedora system (in the `/usr/share/fedora-dockerfiles` directory). Or, you can get the Dockerfile for this from the `Fedora-Dockerfiles` GitHub repo, as shown here:

```
$ git clone https://github.com/nkhare/Fedora-Dockerfiles.git
```

Once you clone the repo, go to the `apache` subdirectory and view the Dockerfile:

```

$ cd Fedora-Dockerfiles/apache/
$ cat Dockerfile
FROM fedora:20
MAINTAINER "Scott Collier" <scollier@redhat.com>

RUN yum -y update && yum clean all
RUN yum -y install httpd && yum clean all
RUN echo "Apache" >> /var/www/html/index.html

EXPOSE 80

# Simple startup script to avoid some issues observed with container restart
ADD run-apache.sh /run-apache.sh
RUN chmod -v +x /run-apache.sh

CMD ["/run-apache.sh"]

```

The run-apache.sh you see in the last three instructions in the Dockerfile refer to the script that runs HTTPD in the foreground.

Now that you've the Dockerfile, it's easy to build a new image:

```

$ docker build -t fedora/apache .
Sending build context to Docker daemon 23.55 kB
Sending build context to Docker daemon
Step 1 : MAINTAINER "Scott Collier" <scollier@redhat.com>
Removing intermediate container 2048200e6338
Step 2 : RUN yum -y update && yum clean all
.... Installing/Update packages ...
Cleaning up everything
Step 3 : RUN yum -y install httpd && yum clean all
.... Installing HTTPD ...
Step 4 : RUN echo "Apache" >> /var/www/html/index.html

Step 5 : EXPOSE 80

Step 6 : ADD run-apache.sh /run-apache.sh
Step 7 : RUN chmod -v +x /run-apache.sh

mode of '/run-apache.sh' changed from 0644 (rw-r--r--) to 0755 (rwxr-xr-x)
Step 8 : CMD /run-apache.sh

Successfully built 5f8041b6002c

```

The following is what our little docker build command has done for us:

- It installed the HTTPD package in our base image
- It created a HTML page
- It exposed port 80 to server the web age

- It set the instructions to start up Apache when you start a container based on this image

You can test your new image by running a container from it:

```
$ ID = docker run -d -p 80 fedora/apache
```

Get the IP address of the container with the inspect command:

```
$ docker inspect -format='{{.NetworkSettings.IPAddress}}' $ID  
172.17.0.22
```

Use the curl command to access the web page:

```
$ curl 172.17.0.22  
Apache
```

Image Layers

Every instruction in a Dockerfile results in a new image layer. Each of these image layers can be used to start a container. The way a new layer is created is unique – it's created by starting the containers based on the image of the previous layer, executing the Dockerfile instruction, and saving the new image.



Multiple pre-built containers are layered together to build an application image.

Docker Image Repositories

You can store Docker images in a Docker image repository, or maintain your own image repository in your data center. Docker's hosted image repo hub isn't known for its reliability and hence, you're better off running your own image repository. It's much more reliable and faster as well to go this route, since there's very little network latency involved in the do it yourself approach.

Using a Private Docker Registry

By default, if a container image isn't already available on your own system, Docker fetches it from the Docker Hub Registry (<https://hub.docker.com>).

You can also set up your own private Docker registry. Maintaining a private registry isn't by any means mandatory – however, it offers the option of storing your images privately without having to push them to the public Docker Hub Registry. Once you set it up, you can push and pull images from the private registry without having to use the public Docker Hub Registry.

In this section, I show how to setup a private Docker registry and how to use it to manage your Docker images.

Setting up a Private Registry

Some Linux distributions such as Fedora have a `docker-registry` package that lets you start up the service that runs the private registry. You can install this package and start the `docker-registry` service with the `systemctl` command. On other Linux distributions, once you've a Docker service running, as explained earlier in this chapter, you must run the Docker provided container image named `registry` to set up the service. Note that Atomic Project Fedora doesn't include the `docker registry` package – therefore, you must use the `registry` container for this distribution as well.

The Docker Image Namespace

In a “pure” Docker system (something that only uses the Docker Project code, with no modifications) a command such as `docker run <image_name>` will always pull the same image from the Docker Hub to your local system. Similarly, a Docker system that you haven't modified with any patches will search only the Docker Hub when you issue a `docker search` command, assuming the image isn't already in a private registry on the server.

At the present time, you can't change the default registry to something other than the Docker Hub Registry. However, changes are afoot to modify this default behavior, as well as several other aspects relating to the image namespace. You can certainly change the default behavior by taking advantage of features that are already present in various Linux distributions.

Note the following key facts about setting up a private registry:

- You can use the same registry for multiple Docker client systems
- By default, the Docker registry stores images in the `/var/lib/docker-registry` directory. Since a larger number of images can consume a lot of storage space, ensure that you've plenty of storage available in the host where you set up the registry.

Managing the Private Registry

Once you pull an image from the Docker Hub, you can push that image to the local Docker registry on a server. Before pushing it, name the image using the `docker tag` command:

```
# docker tag hello-world localhost:5000/hello-me:latest
```

Once you tag the image, push the image to the local Docker registry:

```
# docker push localhost:5000/hello-me:latest
The push refers to a repository [localhost:5000/hello-me] (1 tags)
...
Pushing tag for rev [91c95931e552] on
{http://localhost:5000/v1/repositories/hello-me/tags/latest}
```

You can remove the image from the local private registry with the `docker rmi` command. You can retrieve an image from the private registry with the `docker pull` command. The following example shows how to do this:

```
# docker rm myhello
# docker rmi hello-world localhost:5000/hello-me:latest
# docker pull localhost:5000/hello-me:latest
Pulling repository localhost:5000/hello-me
91c95931e552: Download complete
a8219747be10: Download complete
```

The `docker images` command shows all the images stored in the local private registry on a server.

```
[root@localhost ~]# docker images
```

| REPOSITORY | TAG | IMAGE ID | CREATED | VIRTUAL SIZE |
|------------------------------------|--------|--------------|---------------|--------------|
| docker.io/fedora | latest | 7427c9af1454 | 9 days ago | 204.7 MB |
| docker.io/ubuntu | latest | 56063ad57855 | 10 days ago | 187.9 MB |
| docker.io/hello-world | latest | 975b84d108f1 | 5 months ago | 960 B |
| docker.io/sequenceiq/hadoop-docker | 2.7.0 | ea842b97d1b8 | 10 months ago | 1.76 GB |

```
[root@localhost ~]#
```

In an Ubuntu system, you don't install the Docker registry from a package – instead, you download the registry container from the Docker Hub Registry. You can pull the registry image from the Docker Hub with the `docker pull` command as shown here:

```
# sudo docker pull registry:latest
```

The rest of the commands to test the creation of the registry, to pull and push images are identical to those I described above for the Fedora system.

New Operating Systems Optimized for Docker

Earlier on, I showed how you run Docker in a traditional Linux distribution such as Fedora or Ubuntu. The current trend however, is to use a new generation of operating systems that have been expressly designed for Docker. Two things separate these operating systems from traditional systems:

- Both of these operating systems are preconfigured to run Docker
- Only containers are expected to run on the server (pull an image and run a container)
- They use an atomic upgrade mechanism to manage the servers

The three features mentioned here mean that these operating systems are lightweight in nature – they contain just the absolute minimum capabilities to run a container. The systems don't have tools such as yum or apt-get to download packages separately from the containers. In this section I explain the following lightweight container oriented operating systems:

- **CoreOS** – You can install CoreOS, available on most cloud providers, on bare-metal and test it locally through Vagrant (to be explained in Chapter 6). **Atomic Hosts:** These include servers such as CentOS Atomic, RHEL Atomic and Fedora Atomic.
- **Atomic Host:** Red Hat Enterprise Linux Atomic host is a variation of Red Hat Enterprise Linux 7 that's optimized for running Linux containers in the Docker format.

CoreOS and Atomic Host aren't everything that's available to you – you can also use other operating systems and other techniques such as RancherOS, and VMWare Photon.

Using CoreOS

CoreOS is a popular new Linux distribution that's geared to running applications within containers. CoreOS is highly scalable and is an easy to manage operating system, which provides a distinct marking off of operational and application related functions.



Both of these lightweight operating systems let you deploy containers on a bare metal system, or in a cloud environment such as Amazon EC2 or the Google Cloud Platform.

CoreOS offers an ISO that lets you copy its image to a partition and boot it up with no sweat. If you're using CoreOS to deploy to a cloud environment, you can use a specialized tool such as cloud-config.

You can create a cluster of CoreOS servers and launch containers in that cluster using schedulers. Flannel is a network overlay technique that's a part of CoreOs, and serves as the way containers can communicate through a provider IP space across multiple servers.

Working with Atomic Host

Atomic Host is the product of the Atomic Project and is an RPM-based Linux distribution builder expressly designed to work with Docker containers. Fedora, RHEL

and CentOS versions can run as Atomic hosts. Depending on your environment, you can set up an Atomic Host in any one of the following ways:

- Download an Atomic qcow2 image and add configuration information to it with the cloud-init tool.
- Use a Vagrant file to spin up a CentOS Atomic VM (for CentOS Atomic)
- RHEL Atomic offers the Atomic Installation ISO and Fedora offers the Fedora 22 installer. Once you download and install the Fedora Atomic ISO installation image (assuming you want to setup a Fedora Atomic Host), you can start up the ISO image and start installing it. Once you complete the installation, run the following atomic command to ensure you have the latest Docker version.

```
# atomic host upgrade
```

Once you do this you are ready to set up a Docker registry, or start running docker commands.

The Docker Stack in Production

As with a virtualized environment, setting up Docker in a production system involves several architectural components, each of which performs a specialized task, such as building images and running the images as containers. Following are the typical tasks involved in managing a Docker system in production:

- Building the image
- Sending the image to a repository
- Downloading the image to a new host
- Running the image as a Docker container
- Connecting the container to other containers (clusters of containers)
- Sending traffic to the containers
- Monitoring the running containers in production

In order to manage these tasks, you must set up architectural components such as the following:

- A build system
- An image repository
- A configuration management system
- A deployment mechanism
- An orchestration system to tie multiple containers together

- A monitoring system

The following principles help you move containers to a production use from development:

- Keep production and development environments as similar as possible
- Make sure to employ a continuous integration/continuous delivery (CI/CD) system
- Keep the Docker setup simple
- Automate as much as possible

It's normal to use the term deployment when you want to get something into production. Deployment involves the building and testing of Docker images and deploying the official tested images to the production server(s). Once the image is on the server, of course, you must run the containers on the server. However, in order to carry this off on multiple servers, you must use a repeatable process that can handle the container configuration during each deployment.

Provisioning Resources with Docker Machine

The best way to providing new resources to run containers is to use Docker Machine. Docker Machine can do all the following:

- Create the servers
- Install Docker on the new servers
- Configure local Docker clients to access the new containers

In order to deploy at scale, the simple Docker client won't cut it – you need to have an orchestration tool to manage the complexities of talking to different servers and coordinating the container configuration in different environments. While you can create scripts to talk to the Docker daemon's Docker Remote API, it's somewhat like reinventing the wheel. A lot of work has already been done for you as regards orchestration, and you can take advantage of various approaches to handle your deployment needs. Two basic approaches here would be to use orchestration tools or to use distributed schedulers.

Docker Orchestration

Deploying large number of containers involves a lot of decisions and a lot of work. You need to do some or most of the following as part of deployment:

- Organizing containers into clusters

- Scheduling the server resources
- Running the containers
- Determining the hosts where the containers should run
- Routing traffic to the containers
- Allowing the containers to expose and discover services

Container orchestration tools such Kubernetes provide most or all of the services listed here. Besides Kubernetes, you can also use Docker Swarm, Mesos and Flannel for orchestration. One of the big questions you must answer is which of the orchestration tools you must use – I shouldn't really use the word “must” here, so, let me modify the preceding sentence this way: which of the tools you may want to use – since orchestration tools add to your work as well, and small teams may not really need them.

Docker Orchestration and Clustering Tools

It's ridiculously simple to create Docker images and run Dockerized containers for various services. However, the real benefit and the real fun, so far as administrators are concerned, is when you run a Docker cluster, consisting of a large number of containers together. An easy way to jump into production with Docker containers is to use an orchestration tool. An orchestration tool coordinates the configuration and deployment of applications to multiple Docker containers simultaneously. Orchestration tools require minimal modifications to your system and you can get to your goal of zero-downtime deployments quite easily with these tools.

The most popular orchestration tools for Docker are Centurion, (New Relic), Helios (Spotify) and a Docker tooling set from Ansible, the popular configuration management system.

Distributed Schedulers for Docker Containers

Distributed schedulers also allow for zero downtime deployments, by letting you run the old and new application versions side by side until you're ready to migrate fully to the new versions. You simply specify policies regarding how you want to run your applications, and the scheduler figures out where to run it, and restart failed services on default containers.

There are several distributed schedulers you can consider, including the native Docker clustering tool named Swarm, which lets you create and manage coordinators. Google's Kubernetes is quite popular in this area, and there are other alternatives such as Apache Mesos and Fleet from CoreOS. Let's take a quick look at the various distributed schedulers in the following sections.

Using Docker's Swarm as a clustering tool

Docker Swarm is much simpler than a powerful scheduling tool such as Apache Mesos or Kubernetes, but allows you to create, deploy and manage containers across fairly large Docker clusters. Swarm doesn't really focus on configuration of the application or deployments, since its main purpose is to cluster the computing resources so Docker can use them.

When you use Swarm, the Docker client will still see a single interface, but that interface will be representing a cluster of Docker containers rather than just a lone Docker daemon as you've seen earlier in this chapter. You deploy Swarm as a single Docker container. Through this container, Swarm manages the Docker cluster. You deploy the Swarm container as an agent to each Docker host. This will allow you to merge all your hosts that run Docker containers into a single cluster.

You can install Docker Swarm manually or through the Docker Machine tool. When first starting out, the recommended approach is not to install Swarm via the manual method. It's far easier to install Docker Swarm using Docker Machine. Docker Machine also automatically generates the required TLS certificates to secure Swarm.

Setting up Swarm for Production Use

For production usage, administrators can manually create Swarm on a network by first pulling the Docker Swarm image. You can then configure the Swarm manager and all nodes to run Docker Swarm, all through using Docker. In order to do this you must install Docker on all nodes. In a production setting, you'd need to create a Swarm cluster by running multiple (at least two) Swarm managers in a high-availability configuration.

You can build a Swarm cluster within your data center or on the cloud. Let's say you want to deploy the Swarm cluster on the Amazon Web Services (AWS) platform. Here's an outline of how you can build a high-availability Swarm cluster for production.

1. Add network security rules to the AWS security group so required network traffic is allowed on the VPC network.
2. Create multiple Linux hosts that are part of the default security group. Let's say you create five hosts – of the five hosts, the Swarm primary and secondary managers would be nodes 1 and 2. The Swarm node would be on nodes 3 and 4. The discover backend would be on node 5, and you'll run a consul container on this host.
3. Install the Docker Engine on all the nodes, which enables the Swarm manager to address the nodes.

4. Set up a discovery backend, which helps the Swarm manager and nodes to authenticate themselves as cluster members. The discovery backend also helps the Swarm manager learn which nodes are available to run containers.
5. Create the Swarm cluster – this means you create the two Swarm managers in a high availability configuration. The first manager you start will be the primary manager and the second manager, a replica. When the primary manager fails, the replica becomes the primary manager.

Interacting with the Docker Cluster

Now that you've set up Swarm to manage the Docker container, you can interact with the entire Swarm-managed Docker cluster instead of a single host. Just set the `DOCKER_HOST` environment variable to point to the Swarm manager so you can run Docker commands against the Docker cluster.

Swarm has the advantage that it uses the standard Docker interface and this is easy to use and integrate into your current workflows. It's not designed however, to support complex scheduling. Fleet uses etcd to enable communications between machines and to store the cluster status.

How Fleet can help you

Fleet comes from CoreOS, the lightweight container operating system, and is designed to form a foundation or base layer or more advanced clustering solutions such as Kubernetes.

Fleet is built upon systemd and extends the system and service initialization capabilities of systemd to a cluster of machines.

Apache Mesos and Marathon as a Cluster Manager

Apache Mesos is a cluster manager that can scale to cluster of hundreds or thousands of nodes. Mesos is unique in that it can support diverse workloads – both Docker containers and Hadoop jobs can coexist for example. High availability and resilience are the main benefits offered by Mesos. Mesos is a low level scheduler and supports various frameworks for orchestrating container, such as Marathon, Kubernetes and Swarm. Mesos can support very large systems with thousands of nodes, but may be an overkill for small clusters of just a handful of nodes.

The Big Heavy – Kubernetes

Kubernetes is a Google-built container orchestration tool based on Google's production experience with containers. Kubernetes uses the concept of pods, which are groups of containers that you schedule and deploy a unit. Each pod, consisting of up to five containers, provides a service. Kubernetes uses other containers besides the

service providing pods, to provide logging and monitoring services. Kubernetes isn't ideal for all applications, but for microservices with little state to maintain, it offers a scalable service with little configuration work involved.

Docker Containers, Service Discovery and Service Registration

System architectures are changing rapidly over time, with computing environments becoming ever more dynamic in nature to support newer architectures such as service-oriented architectures and microservices, which we discussed earlier in Chapter 3. If you think a really amazing infrastructure such as Amazons' EC2 (Chapter 9) is dynamic, wait until you learn about Mesos (Chapter 16), which is an even more dynamic computing framework,

Docker containerization has contributed to a much more dynamic computing environment. Modern computing environments involve a large number of hosts and there are many services running on these hosts. Services are taken off and brought back online continuously, making service discovery a critical part of architecting applications.

Ideally, regardless of the size of your environment, you should be able to automate the service discovery process and minimize the effort involved in configuring and connecting this multitude of services. Configuration management tools are designed to solve a totally different problem altogether and they don't scale very well when you tack on the responsibility of service discovery to their main functions.

The fundamental idea that underlies service discovery is that new instances of applications (clients of services) should programmatically be able to easily and automatically identify the current environment details, so they can connect of a service without manual intervention. All service discovery tools are implemented as a global registry that stores information about currently operating services. The registry is usually distributed among multiple hosts to make the configuration fault tolerant and scalable.



Service discovery lets client discover service instances and networking takes care of making the connections work – that is, it connects containers together.

Service discovery platforms are designed to service connection details, but they are based on key/value stores, which means they can store other types of data as well. Many deployments piggyback on the service discovery tools to write their configuration data to the service discovery tool.

In the fast proliferating container and microservice world, services are distributed in a PaaS architecture and an infrastructure that's supported by containers (or VM images) is immutable. Predefined metrics determine whether services will be scaled up or down based on the workloads and metrics. In this environment, the address of a service is dynamically assigned of course, and isn't available until the service is deployed and ready for use. It's this dynamic setting of the service endpoint address that is at the heart of modern service registration and discovery.



Service discovery is the automatic provisioning of the connection information of a service to the clients of that service.

Service registration and discovery works essentially in the following manner: each of the dynamic services registers with a broker, to whom it provides various details such as its endpoint address. Services that want to consume this service will then query the broker to find out the service location and invoke the service. The broker's function is to allow registration by the services and the querying of those services. You can use various brokers such as ZooKeeper, etcd, consul, and Kubernetes (Netflix Eureka as well).

Connecting Containers through Ambassadors

In a production network, you can connect containers across hosts through the use of ambassadors, which are proxy containers. Ambassadors just forward traffic to the actual services. Ambassadors are simple containers that setup connections between an application and a service (such as a database service). Ambassadors offer the benefit that you don't need to change your development code to make it work in a production environment. System administrators configure the application to use different clustered services without required code modifications.

In a development setup, developers use Docker links to link the application to a database container. In a production environment, operation teams can use an ambassador to link the app to the production database service. Operations configure the ambassador so the traffic flows through it and use Docker links to connect the app to the ambassador.

Ambassadors do require extra work in configuring them and are potential points of failure. They also tend to be complex and when you require multiple connections, can overwhelm operations teams. In order to scale, it's better to use networking and service discovery solutions to discover and connect to remote services.

Let's trace the evolution of service discovery and service registration in the following sections, starting with the concept of zero-configuration networking, which was the initial step in the effort to automate the discovery of services.

How Service Discovery Works

Service discovery tools (to be discussed in the following sections) provide APIs that application components use to set or retrieve data. The discovery service is implemented as a reliable, distributed key-value store accessible through HTTP methods. Thus, the service discovery address for each component must either be hardcoded in the application, or supplied as a runtime option.

When a service comes online, it registers itself with a service discovery tool such as ZooKeeper, etcd or consul (all three to be introduced later). The service also provides the information that other services and application components would require in order to consume the services provided by the service that is registering itself. For example, an Oracle database will register the IP address and port for the Oracle Listener service, and optionally the user credentials to log into the database.

When another service that consumes the services provided by the first service comes online, it queries the service discovery registry for information and goes on to establish connections to those services. For example, a load balancer can query the service discovery framework and modify its configuration so it knows exactly across which set of backend servers it should distribute the workload.

Most service discovery tools provide automatic failure detection. When a component fails, the discovery service is automatically updated to indicate that the service is unavailable. Failure detection is usually done through regular health checks and periodic heartbeats from the components. The service discovery tools also use configuration timeouts to determine if components should be yanked out of the data store.

Zero-Configuration Networking

Zero-configuration networking uses several network technologies to create networks consisting of connected devices based on TCP/IP without requiring administrator intervention, or even specialized configuration servers. DNS allows resolution of host names and the Dynamic Host Configuration Protocol (DHCP) automatically assigns hostnames. Zeroconf takes things further: it automatically assigns network addresses, distributes and resolves the hostnames, and locates various network devices, all without DHCP and DNS.

Zeroconf has eventually led to the idea of service discovery, which is our next topic.

Service Discovery

Service discovery tools enable processes and services in a cluster to locate and talk to one another. Service discovery has the following components:

- A consistent and highly available service directory that lists all services
- A mechanism to register services in the directory and also monitor the service health
- A mechanism to locate and connect to services in a directory

Service discovery depends on identifying when processes are listening on a TCP (or UDP) port and then identifying and connecting to that port by name.

But, Doesn't DNS Already Do This?

One's first thought might be, "but DNS already looks up hostnames, no?" That's most certainly true, DNS does perform name resolutions but DNS was never meant for systems with real-time name resolution changes. DNS was primarily designed for an environment where services are assigned standard ports such as the well-known port 80 for HTTP, port 22 for SSH, and so on.

DNS can give you the IP for a service host, and that would suffice in these environments. In today's modern environments, services often use non-standard ports, with multiple services running together on a server. How do you discover these services automatically? This is the question that service discovery tools address.

DNS partially addresses the modern service discovery concerns through its service records (SRV), which provide both the port and the IP address in response to queries. Unfortunately, APIs and libraries can't do SRV record lookups, so DNS becomes the old simple DNS that can only resolve hostnames to IP addresses.

Distributed Lock Services

Around 2006, Google engineers created a distributed lock service named Chubby and started using it for internal name resolution instead of DNS. Chubby implemented a distributed consensus based on Paxos, an algorithm for ascertaining the consensus opinion among the members of a cluster), and is a key-value store that was used for locking resources, and coordinating leader elections. Paxos, however, isn't easy to implement, and eventually a more sophisticated tool named ZooKeeper (Out of the Apache Hadoop project) took over as the standard distributed lock service.

Zookeeper

ZooKeeper is a highly available and reliable distributed lock service that coordinates distributed systems. ZooKeeper has become the industry standard for coordinating

distributed systems and many projects such as Apache Hadoop, Apache Storm, Apache Mesos and Apache Kafka depend on ZooKeeper for distributed coordination, which is critical to their functioning. The fact that ZooKeeper is a highly available key/value data store as well made it a candidate for storing the cluster configuration and for serving as a directory of services.

ZooKeeper is a centralized service that can maintain configuration information, naming, provide group services and synchronization of distributed applications. Services can register with ZooKeeper using a logical name and the configuration information can include the URI endpoint and other information such as QoS (quality of service, discussed in Chapter 2).

ZooKeeper is among the most popular service discovery mechanisms employed by microservice architectures, so let's get a quick overview of this service. Here are essential concepts of ZooKeeper:

- **Znodes:** ZooKeeper stores data in a hierarchical namespace that consists of data registers called znodes, which are similar to files and directories.
- **Ensemble:** Multiple ZooKeeper instances running on separate servers are together known as an ensemble. The instances are aware of each other and must be odd numbered, meaning a set such as 3, 5, or 7 servers. The odd number requirement for the instances is to ensure there's a quorum when decisions are made for selecting a leader.
- **Node Name:** each node in a ZooKeeper namespace is identified by a path and the node name is a sequence of path elements.
- **Configuration Data:** Each node in a ZooKeeper namespace stores coordination data.
- **Client and Server:** Distributed clients connect to a single ZooKeeper Server instance, of which there are several in every distributed environment managed by ZooKeeper. The client maintains a TCP connection through which it sends heart beats and requests and receives responses from the server instance. If the connection to a ZooKeeper instance breaks, the client automatically connects to a different instance.

When you use a service such as ZooKeeper for service discovery and registration, each service in an application (for example the Catalog and Order services in a sales related service) registers and unregisters the service as part of its lifecycle initialization.

ZooKeeper is the most mature of the service discovery tools, but lacks several sophisticated features offered by newer tools such as etcd and consul.

Etcd and Consul

While ZooKeeper is pretty easy to install, configure, and manage from an administrator's perspective (it's the only coordinator used by Apache Hadoop clusters, which use it to support high availability), it's quite heavyweight, and requires highly sophisticated developmental efforts to implement it so it can manage service discovery.

Recently, a much simpler consensus algorithm, Raft, has come into being as an alternative to the Paxos algorithm. CoreOS engineers have used the Raft algorithm to come up with etcd, a distributed key-value store for distributed systems similar to ZooKeeper, written in the Go language. Etcd is a highly reliable key-value store for storing the most critical data of a distributed system.

Consul is another new distributed service discovery and configuration tool that makes it simple for services to register themselves and discover other services. In addition to service discovery and registration, Consul also focuses on failure detection through regular health checking that prevents routing requests being made to unhealthy hosts. Consul also provides several advanced features such as ACL functionality and HAProxy configuration. Although Consul contains a DNS server, you can use the SkyDNS utility to provide DNS-based service discovery when you use etcd.

Besides ZooKeeper, etcd and Consul, there are also a large number of projects that build on basic service discovery, such as Crypt, Confd, Eureka, Marathon (mainly a scheduler), Synapse and Nerve. One can also “roll their own”, if they require more features than what's offered by the available service discovery tools out there.

Service Registration

When you use SkyDNS and Consul, you need to perform registration, which is the final step of service discovery, by explicitly writing code to register a services (such as a Redis database service) with SkyDNS and Consul, for example. Or, the Redis containers can have logic to automatically register upon start. You can however, use a service that automatically registers containers upon their start by monitoring Docker events. The Registrator product from GliderLabs works with Consul, etcd an, or SkyDNS to automatically register containers. It performs registration by monitoring Docker event streams for container creation.

Automating Server Deployment and Managing Development Environments

This chapter focus on automating server deployment and managing development environments.

I start off by discussing Linux package management tools and then move on to Fully Automatic Installation (FAI), which is the way to go when installing large number of Linux servers.

Setting up consistent virtual environments is a big concern in many organizations. Vagrant is an amazing tool that's quite easy to use yet very powerful, and one that helps you effortlessly spin up consistent development environments.

Managing a few servers at a time through shell scripts is fine, but when handling large and complex environments with various services running on them, you need different strategies. I discuss various tools that'll help you perform parallel command execution such as PDSH, as well as more sophisticated parallel execution frameworks such as Fabric and Mcollective.

Automated server provisioning tools are very helpful in installing large number of servers and managing them with configuration tools such as Chef and Puppet (I discuss these two well-known configuration management tools along with other popular CM tools such as Ansible/Salt) in Chapter 7.

The chapter concludes with a brief discussion of two popular server deployment automation tools: Razor and Cobbler.

Linux Package Management

Linux systems can contain thousands of software packages and adding, updating and removing those packages is a common task for system administrators. Linux software packet managers are essentially commands that you can run to install software binaries on a Linux server by fetching the binaries from a binary repository. Red Hat compatible systems use a package format called RPM (Red Hat Package Manager), with package installers that end in .rpm, and Debian based systems use the DEB (.deb) format.

Using the rpm and dpkg commands

You can install software packages with the rpm command in a Red Hat based system, as shown here:

```
# rpm -ihv nmap-6.40.4.e17.x86_64.rpm
You can remove installed RPMs with the -e option:
# rpm -e nmap
```

On a Debian system you can add packages with the dpkg -i command and remove them with the dpkg -r command.

Although the rpm and dpkg commands work fine and do what they're designed for, they're quite problematic when installing packages that have dependencies. In such cases, when you try to install package abc, you're prompted to first install a required package such as xyz, which in turn may ask you to install another required package and so on. Package managers are utilities that handle all the dependencies for you and make installing and removing software a breeze.

Red Hat systems use the YUM (Yellow Dog Linux Updated, Modified) utility, invoked with the yum command. Debian systems on the other hand use the utility named APT (Advanced Package Tool), which you invoke with the apt command. Let's learn a bit about these tools next.

Why use a packet management system? The big benefit is that the package manager takes care of all software dependencies, installing any dependent packages automatically – this makes installing and upgrading packages a breeze.

In the absence of more sophisticated deployment tools, you can log into a server and use yum or aptget commands to install or upgrade software. The yum upgrade command for example will fetch all the latest packages from the binary repository and install them for you.

Package Management with YUM and APT

The YUM utility is a package manager for RedHat Linux systems. It helps you check for and automatically download and install the latest RPM packages. The YUM pack-

age manager also obtains and downloads the dependencies, automatically prompting you as needed.

Here are some examples that show how to use YUM by invoking the yum and the apt commands.

```
# yum install puppet
# yum remove puppet
Yum update      # upgrades all software on a server
# yum list available | grep puppet      # query list of packages available through YUM
# apt-get install puppet                # installs Puppet
# apt-get remove puppet                 # remove the Puppet software
# apt-get update                        # sync the server's software with APT repositories
# apt-get upgrade                       # upgrades all installed software that's different # apt
# /var/cache/apt/archives directory
```

Fully Automatic Installation (FAI)

Installing Linux on dozens or hundreds of servers at a time is a task that you must automate. Manual installs don't scale, and are error prone to boot (no pun intended). Fully Automatic Installation (FAI), which is probably one of the oldest such systems, lets you install a Debian Linux OS unattended on one or more servers. FAI is non-interactive and lets you install and customize Linux systems on physical machines as well as virtual servers. Internally FAI uses a set of shell and Perl scripts to get the job done.

The key to using FAI, which is a way to perform a network installation, is the PXE – Pre-Execution Environment. The PXE server acts as a network boot server. Network installation with an installation server lets you install Linux on a set of systems using the network boot server. By doing this, all the systems you configured to boot using an image provided by this server will do so, and will automatically start the installation program.

You don't need physical boot media to plug into the client in order to start the Linux installation. You can install Linux on multiple systems over the network using the network boot server. Let's therefore learn how to set up a PXE server first.

How FAI Works

FAI can install Linux OS not one or a very large number of machines. Here are the key components of FAI:

- Install server (also called the faiserver): runs the necessary DHCP, TFTP and NFS servers and provides configuration data for the clients on which you want to install a Linux OS.

- Configuration space: a subdirectory where you store the installation configuration files. The files include information about the hard disk layout (similar to the `fstab` file), software packages, time zones, user accounts, printers, etc.
- NFS-Root: a file system on the `faiserver` that serves as the complete file system for the target servers.

What you need for a Network Installation

You'll need a server that runs the following:

- A DHCP server to assign IP addresses
- A TFTP server to server the boot files
- An HTTP, FTP or NFS server to host the installation image

How it Works

As mentioned earlier, under a network installation, you don't need any physical boot media to be plugged into the client servers on which you are installing Linux. When you begin the installation, this is what the client servers do:

- Query the DHCP server
- Get the boot files from the TFTP server
- Download the Linux installation image from a HTTP, FTP or NFS server, depending on which server you're using.

In order to configure a network installation, you must configure the network server that holds the package repositories that are needed for the installation. Next, you need to configure the PXE server.

Setting up the Network Server

As mentioned earlier, you can choose NFS, HTTP (or HTTPS), or FTP to export the installation ISO image or the installation tree from the network server to the clients.

Once you set up the network server, you'll have made the ISO image accessible over NFS for clients to use as a network based installation source. The next step would be to configure the PXE server.

Setting up the PXE Server for a Network Installation

The PXE server contains the necessary files to boot the RHEL and start the network installation. In addition to the PXE server, you must also configure a DHCP server

and enable and start all necessary services. You must configure the PXE server by performing the following tasks (on a RHEL system):

1. Install the tftp package
2. Modify the firewall so it accepts incoming connections to the tftp service:
3. Configure the DHCP server to use the boot images packaged with SYSLINUX.
4. Extract the prelinux.o file from the SYSLINUX package in the ISO image file of the installation DVD.
5. Under the tftpboot directory, create a pxelinux/ directory and copy the prelinux.o file into that directory. Under the pxelinux/ directory, create the pxelinux.cfg/ directory and add the configuration file named default to that directory.
6. Copy the boot images under the tftp/ root directory. `# cp /path/to/x86_64/os/images/pxeboot/(vmlinuz,initrd.img) /var/lib/tftpboot/pxelinux/`
7. Start (or reload if they're running) the xinetd and dhcp services. `# systemctl start xinetd.service dhcpd.service`

The PXE server is now ready to start the network install. Start the client on which you want to install Linux and select PXE Boot as the boot source and start the network installation. This will start the installation program from the PXE server.

Using Kickstart

For RPM-based distributions such as Red Hat Linux and Fedora, using Kickstart rather than FAI is probably the better approach to perform automatic installations for RHEL and CentOS7 systems without the need for user intervention. You still use the PXE-boot to get the servers started, but the configuration information is provided by Kickstart files read from a local FTP server rather than FAI.

In order to perform a Kickstart installation, you configure the PXE server in the same way as shown in the previous section – all you need to do is to add the Kickstart file to the mix. You thus use the PXE server together with Kickstart in this case.

.Creating a Custom Kickstart File

You can deploy RHEL and CentOS simultaneously on a large number of servers by using a Kickstart installation. The Kickstart files contain all the answers to the interactive actions that are posed by the installation program, such as the disk partitioning schemes and the packages the installer must install. Kickstart files thus help you automate the installations when dealing with a large number of server deployments by letting you use a single Kickstart file to install RHEL on all the machines.

In the following sections, I briefly explain the steps involved in performing an automatic RHEL Kickstart installation.

Creating a Kickstart File

It's best to manually install the RHEL Linux software (in in this case RHEL 7) on one system first. The Kickstart file is a simple text file and you can name it anything you want. The choices you make during this manual selection are stored in the news server in the `/root/anaconda-ks.cfg` file – this will serve as your Kickstart file for automatic installations.

You can modify the Kickstart file as you please. ON RHEL, there's also a Kickstart Configuration Tool available, that lets you walk through server configuration and create and download a Kickstart file – the only drawback is this file won't support advanced disk partitioning.

Verifying the Kickstart File

Before you can use the Kickstart file to automatically install the Linux binaries on a whole bunch of machines, it's a good idea to verify its validity using the `ksvalidator` command line utility, which is part of the `pyKickstart` package.

Here's how you do it:

```
# yum install pyKickstart
# ksvalidator /pth/tp/Kickstart.ks
```

Making the Kickstart File Available

In order to make the Kickstart file available to the client machines, you must place it on a removable media (DVD or USB drive), or on a hard drive connected to the client server, or even better yet, a network share that the client machine can access. Since normally the new systems boot using a PXE server, it's probably a good idea to let the clients also download the Kickstart file from a network share.

Making the Installation Source Available

The Kickstart file contains a list of software packages required for the install. The installation process must then access an installation source such as a RHEL installation DVD ISO image for an installation tree to install those packages. Assuming that you're performing a network based (NFS) installation, you must make the installation tree available over the network, using the steps I described in the previous section.



An installation tree is a copy of the binary RHEL DVD with an identical directory structure.

Automating the Kickstart Installation

You can start a Kickstart installation manually with some user interaction at the system prompts. However, that's no fun, since I want you to see how to automate the whole darn thing! During the installation, you must specify a special boot option (inst.ks=) when booting the system.

Following is a typical Kickstart file.

```
#version=RHEL7
# System authorization information
auth --enablshadow --passalgo=sha512

# Use network installation
url --url="ftp://192.168.1.25/pub/"
# Run the Setup Agent on first boot
firstboot --enable
ignoredisk --only-use=sda
# Keyboard layouts
keyboard --vckeymap=us --xlayouts='us'
# System language
lang en_US.UTF-8

# Network information
network --bootproto=dhcp --device=eno16777736 --ipv6=auto --activate
network --hostname=localhost.localdomain
# Root password
rootpw --iscrypted $6$RMPTNRo5P7zulbAR$ueRnuz70DX2Z8Pb2oCgfXv4qX0jkdZLaMnC.CoLheFrUF4BEjRIX8rF.
```

Configuring the Clients to Automatically Install the Linux Software through Kickstart

You must instruct the clients to boot from the network from the BIOS, by selecting the Kickstart option from the PXE menu. Once the kernel and ram disk load, the client detects the Kickstart file and automatically installs the Linux software without any user intervention. If you want, you can connect to the installation process with a VNC client from another host in order to monitor the installation.

Automatically Spinning up Virtual Environments with Vagrant

Vagrant is a powerful open source tool that simplifies the task of spinning up virtual machines (VMs) and running them, with the help of simple command-line utilities. Vagrant makes it easy to distribute and share a virtual environment and it supports all major virtual platforms such as VirtualBox, VMWare and Hyper-V. It also supports all the well-known software configuration tools such as Chef, Puppet, Ansible and Salt.

Development environments can be notoriously hard to configure. Vagrant helps you, the Linux sysadmin, to take over the responsibility for setting up the development environment from the developers. New developers can be easily on boarded as a result, and it also makes it easy to update the environment used by the developers.

Vagrant is a configuration tool for VMs and helps developers quickly spin up new VMs easily. However, you can use it for other purposes. Generally, Vagrant is used to create VMs that help develop and deploy software. When you use Vagrant to set up development environments, it's a piece of cake to mimic a production environment, and also employ the same provisioning tools and strategies that you use in your production environment.

Developing consistent environments is the key to effective deployment of software. If your deployment pipeline is already in place and is working well for you, you can use Vagrant to recreate the same processes in the development environment. If your development pipelines need to be improved, Vagrant is ideal as an environment within which to develop the processes that make development environments consistent.

Vagrant is a big help in managing the configuration of VMs that you run out of Oracle Virtual Box, and is ideal for testing various things. Vagrant wasn't designed for heavy duty configuration across the data center. You can use a serious enterprise configuration system such as Ansible along with Vagrant, with Vagrant helping you test the configuration code that you're deploying with Ansible.

A few years ago, web applications were mostly all PHP and MySQL. Today, you have several web application frameworks such as Ruby on Rails, various databases, web servers, application servers and backend services. Installing all these components and configuring them correctly is a nightmare at times when you do it locally by hand.



Vagrant works with Oracle Virtual Box, and commercial alternatives such as VMWare, and also remote environments such as Amazon's Elastic Compute Cloud (EC2).

Beyond the nightmare of installing locally, problems like misconfiguration, differences between dev and prod environments, the difficulties of managing multiple projects and syncing development environments among all team members, are all problems that Vagrant elegantly solves. Working with multiple projects? Just create a separate VM for each machine. New team member to be on boarded quickly? Hand him/her a laptop and ask them to run just one command – `vagrant up`, to be up and running within minutes.

If you're an operations engineer, you can work on system automation scripts and can test them on a full-fledged sandbox that mimics production, so you can test real-

world scenarios. If something gets messed up, simply destroy the VM and recreate a fully functioning environment at the snap of your fingers.

Although Vagrant is mostly used in a web application environment, you can use it for anything you want, as long as you want to work with VMs. Vagrant is simply amazing – no two ways about it! Regardless of the complexity of your virtual environment and regardless of the type of virtualization you’re using, you can easily bring up the complete virtual environment with the following simple vagrant command.

```
$ vagrant up
```

With just this one command, you get all the following done:

- Create a virtual machine based on any Linux distribution you like
- Change the properties of this VM such as RAM, storage, etc
- Set up the network interfaces for this VM so you can access it from a different server, local or remote
- Set up the shared folders
- Set the hostname for the server
- Boot up the VM – you can see in VirtualBox that there’s a new VM in the running STATE NOW
- Provision software on the new VM via shell scripts, or CM tools such as Chef and Puppet

All this in a New York minute – wow!

The Vagrant commands let you ssh into the new VM, start, stop and resume it, and when you want to get rid of it, “destroy” the machine by deleting it from your hard drive – all virtual hard drives (file folders) relating to the VM are removed. Any data that you don’t save in a shared folder is lost for good, so be careful when you “destroy” a Vagrant VM. You can also package the machine state and distribute it to other developers.

he creator of Vagrant, Mitchell Hashimoto, has described Vagrant as the “Swiss army knife for development environments” since it does everything you need to create and manage development environments and also automates things and helps set up dev environments that parallel production environments.

Vagrant helps you quickly spin up full-fledged but disposable working environments.

Some Background

In Chapter 4 you learned about the different types of virtualization. As you know by now, a virtual machine runs within a software process (the hypervisor) which runs on a host computer, and mimics a physical computer. It’s the hypervisor that provides

the computing infrastructure such as CPUs and RAM to the virtual machines. You also learned about the two types of hypervisors – Type 1 and Type 2.

A Type 1 hypervisor runs on the bare metal host machine hardware and controls access to the computing resources, and their allocation to the VMs. VMWare ESX/ESXi and Oracle VM Server are examples of Type 1 hypervisors. Type 2 hypervisors on the other hand, sit on top of the OS and use the OS to control the computational resources. Vagrant environments typically use Type 2 hypervisors as the hosts for VMs that you create with it. The two most popular Type 2 hypervisors are Oracle VirtualBox and VMWare Workstation/Fusion family.

I use the freely available Oracle Virtual Box to show how to use Vagrant to spin up guest machines that run within the hypervisor. You can create guest machines that run an entirely different OS from that being run by the host computer. Since Vagrant employs the same API to run VMs on different hypervisors, sharing virtual environments is very easy between teams that are working with different OS platforms.

If you’ve ever set up a virtual development environment, you know you’d have to do most or all of the following to set up a new VM:

- Download the VM image
- Boot up the VM
- Configure shared folders, network connections, storage and memory
- Use a configuration tool such as Puppet/Chef, Salt or Ansible to install any required software.

The simple `vagrant up` command does all these tasks for you! The command sets up the entire development environment and a developer can also, with equal ease, destroy and recreate the virtual environment, all within a few short minutes as well.

Vagrant and its Alternatives

You can use alternatives such as plan desktop virtualization with VMWare and VirtualBox to do some of the things you can do with Vagrant – however these virtualization solutions don’t have the unique workflow of Vagrant. While you can do everything you can do with Vagrant with a regular virtualization solution, it isn’t an automated process as with Vagrant.

Containers are somewhat of an alternative but they really don’t provide full virtualization – they are instead isolated environments running the same kernel. As you saw in Chapter 5, containers do offer numerous benefits, but a big downside is that they can run only the same OS that runs on the host. If you use containers, all members of a team need to use the same host OS.

Containers are great for production settings since they securely isolate resources without a performance overhead. In a development environment the limitations posed by containers outweigh their benefits.

The cloud is another alternative but you normally incur a higher financial cost going that route, compared to using a local Vagrant environments to manage your development needs.

Getting Started with Vagrant

In order to run Vagrant on a server (or your laptop), you need to first install Oracle VirtualBox, which is open source, and thus free. You can download Virtual Box from <http://virtualbox.org>.

A virtualization system such as VirtualBox is called a provider and although I use VirtualBox since it's free and easy to get going, there are alternative providers such as a VMWare provider.

Virtual Box has minimal system requirements, but since a single workstation will be running multiple VMs at any given time, you need to ensure that the laptop or server where you're playing with Vagrant has sufficient RAM. There's no hard and fast rule here as to how many VMs you can run on a server – it depends on the software running in each of the VMs hosted by the server.



You can run VirtualBox in a Windows or a Linux environment.

Once you've VirtualBox installed, you are ready to install Vagrant, which you can download from <http://vagrantup.com>.

Vagrant doesn't create and host VMs – it's the hypervisor (VirtualBox in my case) that does those things. Vagrant simply manages the VMs.

In a production environment, you can use hypervisors other than Oracle VirtualBox – you can use VMWare Desktop Applications (Fusion and Workstation) if you need to deal with a large number of VMs. If you want to simplify things, you can go with an external hypervisor such as Amazon EC2 or DigitalOcean.

Although Vagrant supports other virtualization providers, Oracle VirtualBox is the most popular provider for developers.

On a RedHat system such as Fedora, install Vagrant binaries as shown here:

```
$ yum install 'vagrant*' 
```

You can check the installation by typing the following at the command line:

```
$ vagrant -version
Vagrant version 1.1.4
```

You're off to the races at this point.

Spinning up a New VM

You spin up new VMs with the `vagrant up` command, but first you must initialize the image. The `vagrant init` command creates a Vagrant configuration file from a template. In the following example I use an Ubuntu base image, also referred to as a Vagrant box.

The `vagrant init` command initializes a new Vagrant environment by creating a Vagrantfile.

```
$ vagrant init hashicorp/precise64
A `Vagrantfile` has been placed in this directory. You are now
ready to `vagrant up` your first virtual environment! Please read
the comments in the Vagrantfile as well as documentation on
`vagrantup.com` for more information on using Vagrant.
$
```

Start the new virtual machine with the `vagrant up` command:

```
$ vagrant up

Bringing machine 'default' up with 'virtualbox' provider...
==> default: Importing base box 'hashicorp/precise64'...
==> default: Matching MAC address for NAT networking...
==> default: Checking if box 'hashicorp/precise64' is up to date...
==> default: Setting the name of the VM: SG0221771_default_1461505264672_17333
==> default: Forwarding ports...
    default: 22 (guest) => 2222 (host) (adapter 1)
==> default: Booting VM...
==> default: Waiting for machine to boot. This may take a few minutes...
    ==> default: Machine booted and ready!
==> default: Checking for guest additions in VM...
    default: Guest Additions Version: 4.2.0
    default: VirtualBox Version: 5.0
==> default: Mounting shared folders...
$
```

Here's what happens when you run the `vagrant up` command:

- Vagrant creates a new VirtualBox machine based in the base image within the box specified in the Vagrantfile. Vagrant does this by copying the virtual hard disk files.
- VirtualBox randomly generates a MAC address when it creates a new machine. Vagrant matches the MAC address for NAT networking.

- Vagrant sets the name of the virtual machine
- Vagrant forwards port definitions (more on this later in this chapter)
- Vagrant boots the new VM
- Vagrant mounts the shared folders that you can use to share data between the VM and your laptop or server. As you know, I'm using the Oracle VirtualBox in this example, but the procedure for spinning up new VMs is exactly the same when you bring up VMs using VMWare or AWS as a provider.

The `vagrant up` command starts and provisions the Vagrant environment. You now have a full featured 64-bit 12.04 LTS virtual machine running in the background. The Vagrant machine instance is up and running, although you won't see it, since it's headless (that is, it has no GUI). You can connect to this machine by using the following command:

```
$ vagrant ssh
Welcome to Ubuntu 12.04 LTS (GNU/Linux 3.2.0-23-generic x86_64)

 * Documentation:  https://help.ubuntu.com/
New release '14.04.4 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Welcome to your Vagrant-built virtual machine.
Last login: Fri Sep 14 06:23:18 2012 from 10.0.2.2
vagrant@precise64:~$
```

When you log into the new VM, you do so as the default user `vagrant`:

```
$ vagrant@precise64:~$ whoami
vagrant
$ vagrant@precise64:~$
```

You really don't need the root password but if you want to, you can do so by issuing the command `sudo passwd root` as the user `vagrant`. You can then set the root password to anything you like and login as the user.

The `vagrant ssh` command connects to the machine via `ssh`. It drops you unto the SSH console within the new VM. You can do everything on this machine as you'd in a normal server, such as installing software, creating Docker containers, etc. You log out of the VM by typing `exit`:

```
vagrant@precise64:~$ exit
logout
Connection to 127.0.0.1 closed.
$
```

The `exit` command brings you out of the VM and puts you back into your terminal on the host server. Run the `vagrant destroy` command to delete your new VM – you can create it again, no problems, with the `vagrant up` command if you need it again.

```
$ vagrant destroy
    default: Are you sure you want to destroy the 'default' VM? [y/N] y
==> default: Forcing shutdown of VM...
==> default: Destroying VM and associated drives...

$ exit
```

The Vagrantfile

When working with Vagrant environments, you configure Vagrant per project, with each project having its own specific work environment. Each of the vagrant projects has a single Vagrantfile that contains its configuration. The Vagrantfile is a text file that Vagrant reads to find out the configuration of your working environments, such as:

- The operating system,
- The CPUs and RAM per VM
- How the VM can be accessed
- Provisioning of various software on the VM

It's a good idea to put the Vagrantfile under version control so all members of a team get the same software, settings and configuration environment for their work environment.

When I ran the `vagrant init` command in the previous section, it created an initial Vagrantfile. The file is in the directory where I ran the `vagrant init` command. Read the Vagrantfile since it gives a very good idea of how to configure the essentials of a vagrant project. In my case, everything is commented out except the following:

```
Vagrant.configure(2) do |config|
  config.vm.box = "hashicorp/precise64"
end
```

In my case the file essentially contains a block of Vagrant configuration that contains a configuration value for `config.vm.box` – this is all it took to create the new VM! Since Vagrantfiles are portable, you can take this file to any platform that Vagrant supports, such as Linux, Windows and Mac OS X,

Vagrant Box

As I explained in the previous section, the simple text file named Vagrantfile contains the configuration for spinning up your new VM. But you also need something from which to create this VM from – and that something is called a Vagrant box. A box is simply a base image for an OS that Vagrant will clone to quickly create a functioning VM. You can look at the box file as a template for a Vagrant-created and managed virtual machine. Using the box is the key strategy that makes it possible to spin up

VMs in seconds – imagine all the work involved in creating even the simplest of VMs from scratch.



Vagrant uses base boxes to clone new VMs

In the Vagrantfile, the value `config.vm.box` specifies the name of the box that Vagrant will use to create your new VM. In our case, it was:

```
config.vm.box = "precise64"
```

Multiple environments can share the same underlying box, with each environment using that box as the template for creating new VMs. However, you use a separate Vagrantfile for each project.

You can view all the Vagrant boxes you have on a server with the following command:

```
$ vagrant box list
centos/7          (virtualbox, 1602.02)
hashicorp/precise64 (virtualbox, 1.1.0)
rails-v1.0.0      (virtualbox, 0)

~
$
```

Vagrant VMs use the shared folders feature, which means that your development teams can continue to use the development tools (IDEs etc) they're most familiar, and hence the most productive with.

Vagrant Networking Vagrant automatically configures networking with the VMs you create by letting teams communicate with the VM. There are several networking options as I explain shortly, but here's an example that illustrates the basics of how Vagrant handles networking.

In the example, I use a forwarded port. A forwarded port exposes a port on the VM as a port on the host server. As you know, port 80 is the default port for web services, so let me expose port 80 so I can access any web services. In my Vagrant file, then, I add the following line:

```
config.vm.network "forwarded_port", guest: 80, host: 8080
```

I then issue the command `vagrant reload` to restart the VM with the new network port settings:

```
$ vagrant reload
==> default: Attempting graceful shutdown of VM...
```

```
==> default: Clearing any previously set forwarded ports...
==> default: Clearing any previously set network interfaces...
==> default: Preparing network interfaces based on configuration...
    default: Adapter 1: nat
==> default: Forwarding ports...
    default: 80 (guest) => 8080 (host) (adapter 1)
    default: 22 (guest) => 2222 (host) (adapter 1)
==> default: Booting VM...
...
```

You can see how Vagrant has forwarded port 80. You can test the forwarded port by starting a simple web server from within the VM and connecting to it from a browser on the host server:

```
$ vagrant ssh
vagrant@precise64:~$ cd /vagrant
vagrant@precise64:/vagrant$ sudo python -m SimpleHTTPServer 80
Servicing HTTP on 0.0.0.9 port 80...
```

I started a basic web server on port 80 with the command shown here. If you now open a browser and point it to localhost:8080 on your laptop (or wherever the VM is running) you'll see the directory listing for the /vagrant directory, served from your new VM.

Vagrant offers three different ways to set up networking:

- Forwarded ports: I've shown an example of this method, which is quite simple to setup you need to enumerate each and every port number, which becomes onerous in a complex environment
- Host-only networking: Creates a network that's private to your host and the VM is running on the host. While this is a secure way to network, it limits access from other machines outside of the host server.
- Bridged networking: Bridges the VM onto a device on the physical host, making the VM appear as another physical machine on the network. It offers you the benefit of having an OIP to access the VM, but you can't specify a static IP for a bridged network as the IP addresses are serviced via DHCP/

Provisioning Vagrant Boxes

In earlier sections, you learned how to spin up a VM with Vagrant's help, but it was a bare base box. You can use Vagrant boxes that already have software installed, such as Hadoop and Ruby on Rails, for example. You can also automatically install software as part of the creation of your development environments – this process is called provisioning.

While you can install any software you need after booting up a new VM. It's inefficient to do when you've a bunch of machines to configure. Vagrant lets you automate provisioning with shell scripts, or a CM tool such as Chef or Puppet.

How Automated Provisioning Helps You

Automated provisioning removes one of the biggest issues in setting up consistent development environments – configuration drift. Automated provisioning of software is easy, repeatable and also helps keep development and production systems in sync, avoiding unpleasant surprises when you move apps from development to production settings.

As a system administrator or an operations engineer, you can use Vagrant for quickly testing infrastructure changes before moving the changes to production

You can use Vagrant as a good training ground for learning CM tools such as Puppet and Chef, but if you're a novice, with CM tools, just start with simple shell scripts to see how to provision software with Vagrant. Once you become adept at provisioning, it's a simple jump to a full-fledged CM tool.

You can perform Vagrant provisioning with Chef through Chef Solo or the Chef client. Chef Solo is great for testing and for small environments. If you already have a Chef Server, Chef Client is going to be what you want to use. Similarly, you can do provisioning both with and without a master when using Puppet, and the master less approach is simpler and better when starting out with Puppet.

Besides Puppet/Chef as build-in provisioners you can extend vagrant to use additional provisioners through easily available plug-ins.

An Automated Provisioning Example

Let me show you how to configure a provisioner to set up an Apache web server. Before I start automating the installation of Apache, I must first set it up manually on a Vagrant box, to capture the correct procedure to create the web server.

In this example, I use the precise64 Ubuntu 64 bit Linux distribution, and I want to export port 80, as explained in an example earlier on. My Vagrantfile for this box will look as follows:

```
Vagrant::Config.run do [config]
  config.vm.box = "precise64"
  config.vm.forward_port 80, 8000
end
```

Use SSH to log into the VM, so you can install and configure Apache on the Ubuntu Linux server.

```
$ vagrant ssh
```

Install Apache on the Ubuntu server

```
$ sudo apt-get update
$ sudo apt-get install apache2
```



Vagrant runs all commands in a provisioning shell script as the user root, so there's no need to use sudo. If for any reason you think you need to be root in a Vagrant environment, do the following to set the root password, and then login as root:

```
$ sudo passwd root
```

By default Apache servers web content from /var/www directory. To avoid having to modify the Apache configuration files I'm going to modify /var/www directory to be a symbolic link to /vagrant which is the default shared folder directory.

```
vagrant@precise64:/vagrant$ sudo rm -rf /var/www
vagrant@precise64:/vagrant$ sudo ln -fs /vagrant /var/www
```

Now, Apache will by default serve any files you placed in the shared folder. At this point, if you go to <http://localhost:8080>, you'll see the directory listing of the shared folder:

Next, create an index.html (default file that Apache serves) file in the shared directory (/vagrant), as shown here:

```
vagrant@precise64:/vagrant$ logout
Connection to 127.0.0.1 closed.
$ echo "<strong>Hello<strong>" > index.html
```

If you refresh your browser now, you'll see the following:

```
Hello
```

Each time you bring up the VM with the vagrant up command, you'll need to perform all these steps, which is a waste of time. With automated provisioning, you can work much smarter.

Automated Provisioning

You saw how to set up and configure Apache in the previous section. Let's see how you can automate the same through a shell script, a Chef Recipe and a Puppet manifest.

Automating Provisioning with Shell Scripts

Using a shell script simply means that you put all the commands you had to issue manually into a single file. Here, I call the file provision.sh, but you can name it any-

thing you want. Add the following lines to the `provision.sh` file and save it in the Project directory.

You’ve your provisioning shell script ready – all you need to do now is to let Vagrant know where to find it. In order to do this, stick the following line in the Vagrantfile you created earlier:

```
config.vm.provision "shell", path: "provision.sh"
```

I didn’t have to specify a path to the `provision.sh` file since relative paths are relative to the Project root directory.

In order to test whether you’ve correctly automated the setup and configuration of Apache as part of a new VM creation, firsts destroy the current VM and run the command `vagrant up`, to start with a clean state:

```
$ vagrant destroy
$ vagrant up
Bringing machine 'default' up with 'virtualbox' provider...
==> default: Importing base box 'hashicorp/precise64'...
...
==> default: Running provisioner: shell...
    default: Running: C:/cygwin64/tmp/vagrant-shell20160424-6872-1xaawtl.sh
Installing Apache and setting it up...
$
```

As you can see, I’ve created a basic web development environment with no human interaction at all. You can use more complex scripts to create and configure databases, cron jobs, and whatever else you need to set up a full-fledged development environment.

Automated Provisioning with Chef

You’ll learn a lot about Chef (and Puppet) in Chapter 7. For now, just let me say that you can use either `chef-solo` or `chef-client` to provision Vagrant environments. `chef-solo` is easier, so I use that to show how to provision software in Vagrant VMs.

In order to provision our Apache web server using Chef, you need to create a cookbook and a Chef Recipe. By default Vagrant will look in the `cookbooks` directory relative to the Vagrant project directory for its cookbooks. So, first create the `cookbooks` directory under the Project directory and then create the following Chef recipe:

```
execute "apt-get update"
package "apache2"
execute "rm -rf /var/www"
link "/var/www" do
  To "/vagrant"
end
```

Save the recipe file to `cookbooks/mydir/recipes/default.rb`.

Once you do this, stick the following line in your Vagrantfile:

```
config.vm.provision "chef-solo", run-list: ["mydir"]
```

Now that you're all set up, run the `vagrant destroy` command to remove all traces of the previous installation and then execute the `vagrant up` command to create a new VM with automatic provisioning of the Apache server by Chef.

Automated Provisioning with Puppet

Automated provisioning with Puppet is quite similar to provisioning with Chef. I use a simple Puppet set up without a master. Just as Chef uses cookbooks and recipes, Puppet uses manifests. In order to let Puppet handle the provisioning of our Apache server, you must create a manifest folder under the project root directory. Next, create the Puppet manifest that'll do the job for us:

```
exec { "apt-get update":  
  Command => "/usr/bin/apt-get update",  
}  
  
package { "apache2":  
  Require => Exec["apt-get update"],  
}  
  
file { ["/var/www":  
  Ensure => link,  
  Target => "/vagrant",  
  Force => true,  
]
```

Remove the VM with the `vagrant destroy` command and run `vagrant up` again.

Finally, you don't necessarily have to choose among the available provisioners, in the sense that you can use more than one provisioner by simply specifying multiple `config.vm.provisioner` directives in your Vagrantfile, as shown here:

```
Vagrant::Config.run do [config]  
  Config.vm.box = "precise64"  
  ...  
  Config.vm.provision "shell", inline, "apt-get update"  
  Config.vm.provision "puppet"  
  ...  
End
```

Vagrant will provision your new VM by installing and configuring the software provided by each provisioner, in the order you specify the provisioners in the Vagrantfile.

Creating Vagrant Base Boxes

A Vagrant box that contains just the minimum software that allows Vagrant to function, but nothing more, is called a base box. A base box, such as the Ubuntu box pro-

vided by the Vagrant project *”precise64”) that I used earlier on, isn’t repackaged from any other Vagrant environment, and that’s why it’s called the “base box”.

The Vagrant project and others offer numerous base boxes. If you’re starting out, get familiar with Vagrant by using the available base boxes. Later on, once you get really comfortable, you can create your own base boxes to serve as a starting point for fresh development environments.

Components of a Base Box

Vagrant base boxes include a bare set of binaries. Here’s a list of the minimum components in a base box:

- Vagrant package manager
- SSH
- An SSH user to allow Vagrant to connect

Note that the provider you’re going to use also determines what goes into a base box. For example, if you’re using VirtualBox as the provider, you’ll need the highly useful Guest Additions so that you can use the shared folders capability of the VM.

It’s quite easy to build a custom box from a preexisting Vagrant box such as ubuntu/trust64. Just boot the predefined box and modify its configuration per your requirements. Export the box to a new file (with the .box extension) by executing the vagrant package command.

Packaging a base box into a box file differs, based on the provider. In the following example, I’m using Virtual Box, and I need do the following to package the new base box:

```
$ vagrant package --base Ubuntu-14.04-64-Desktop
```

The package command creates a file named package.box.



There are no rules as to how you should distribute the base boxes you create. However, Vagrant recommends that you add the new box to HashiCorp’s Atlas to supporting versioning, push updates and lots of other reasons.

You test the new base box by spinning up the new VM:

```
$ vagrant up
```

If the VM is spun up correctly, your new base box is good!

Using Packer and Atlas for creating base boxes

In the previous section, I showed how to create a Vagrant base box from scratch using a manual method. Vagrant strongly recommends using HashiCorp's Packer (and chef/bento or boxcutter templates) to create reproducible builds for a Vagrant base box. It also recommends that you use HashiCorp's Atlas to automate the builds.

Packer is a command line tool that lets you automate the creation of VMS by generating Vagrant boxes. Packer is the recommended way to create Vagrant boxes. In a nutshell, here's how Packer works:

- It downloads the ISO CD or DVD image of the OS you want to use
- It executes the installation program and sets up the server with the default configuration
- It runs provisioners such as shell, Chef, Ansible and Puppet to customize the new system
- It exports the custom system and packages it into a .box file.



Packer not only lets you create VMs of different operating systems, such as Ubuntu and CentOS and Windows 7, it also lets you create boxes for various providers such as VirtualBox and VMware.

Using Packer, you can specify the configuration of your Vagrant boxes (memory, disk size etc) in a JSON file. Once you specify the needed automation parameters (such as Ubuntu preseed file), Packer will do the initial OS deployment for you.



You can find details about Packer at <http://www.packer.io>. You can find examples of Packer definition files by going to the repository (called bento) maintained by Chef at Chef's Github account. <https://github.com/chef/bento>

You can download Packer distributions from <https://www.packer.io/downloads.html>. The chef/bento project contains numerous Packer definitions. You can clone the chef/bento project from GitHub (<https://github.com/chef/bento>). Once you clone the project, in the bento/packer subdirectory, you'll find a number of JSON files that contains definitions for building various operating systems. You can use Packer and these JSON files to build base boxes for Ubuntu, Fedora, etc. For example, you can build an Ubuntu 14.04 box for the VirtualBox provider by running the following packer command:

```
$ packer build -only=virtualbox-iso ubuntu-14.04-i386.json
```


Note that this packer command creates a new Ubuntu box without using a preexisting Vagrant box. It does this by downloading an ISO image from the Ubuntu distribution site. When the packer command finishes executing, you'll see the following file:

```
bento/builds/virtualbox/opscode_ubuntu-14.04-i386_chef-provisionerless.box
```

What you have here is a bare bones system with just enough software for the system to function. You can customize this box to your heart's content.

You can test this new Vagrant box created by Packer by using the same procedures as in the previous section where I showed the steps for manually creating a base box. Here's how to do it for our example:

```
$ vagrant box add bento-ubuntu opscode_ubuntu-14.04-i386_chef-provisionerless.box # install the box
$ vagrant init -m bento-ubuntu          # initialize the new VM
$ vagrant up                            # boot the new VM
$ vagrant destroy                        # destroy the box
```

Parallel Job-Execution and Server Orchestration Systems

In a large environment especially, it's critical that you use techniques other than old-fashioned shell scripts to simultaneously execute tasks across a number of servers. Remote command execution tools help significantly in performing near real-time parallel execution of commands.

Automated server provisioning tools such as Razor and Cobbler let you automate the installation of servers and the management of those servers with configuration tools. Using one of these tools, it's easy to go from a bare metal or VM server with nothing on it and come up with a fully configured system in a very short time.

Working with Remote Command Execution Tools

When dealing with a set of servers, administrators typically write programs that can parallelly execute their commands. It's easier to manage these types of operations by using tools and frameworks explicitly designed for this type of work. In this section, I explain some of the more popular tools for server orchestration and parallel job execution. Specifically, the following are the tools I discuss:

- Parallel Execution Processors (PDSH, DSH, CSSH)
- Fabric
- Mcollective

Parallel (SSH) Execution Processors

If you're handling just a handful of servers and don't think that you need the overhead of a configuration management (CM) tool such as Puppet/Chef/Ansible/Salt, you can look at other simpler tools that help you manage multiple Linux servers. Most of these tools also work very well with Puppet and other CM tools, to help make your life easier.

The tools I've in mind are all based on SSH, so you must first setup SSH key based logins on the servers (please see Chapter 2). Let's quickly review these remote execution tools in the following sections.

PSSH (Parallel SSH)

Pssh is a program for executing ssh in parallel in a set of nodes. It can send input to all processes, passwords to ssh and save output to files. Following is an example that shows how to copy a file to my home folder on two servers.

```
pscp.pss -v -H "alapati@host1 alapati@host3" test.gz /home/alapati
```

DSH (DISTRIBUTED SHELL)

You can install DSH (distributed shell) or pdsh (parallel distributed shell) with apt-get or yum. In order to execute a command on multiple nodes, you add the node list to a file such as the following example:

```
vi /tmp/test.list  
host01  
host02  
host03
```

Once you have your server list ready in a file, pass the node list file to pdsh when executing a command:

```
# pdsh -aM -c uptime
```

This command will show you the uptime on each server in your node list file (test.list in this example).

CSSH

CSSH (ClusterSSH) is somewhat different from the pssh and pdsh (or dsh) tools – it opens an xterm terminal to all the hosts you specify, plus an admin console. When you type a command in the admin console, it's automatically replicated to all windows. You can see the command actually execute on the various windows for separate servers. Here's an example showing how to use the CSSH tool.

Create a cluster configuration file named clusters in /tmp and define your servers in that file.

```
vi /tmp/clusters
clusters = 1204servers
1204servers = host1 host2
```

Run the following command.

```
# cssh alapati@host1 alapati@host2
```

You'll now see two terminal windows and the as you type your commands in the admin console, you'll see the command executing in both terminals.

Fabric

Fabric is a command-line tool that uses SSH to help you orchestrate various configuration operations. Fabric automates communications with remote servers by working on clusters of machines. It can help start and stop services on the clusters. You can use it for both application development and systems administration. Fabric is handy for uploading and downloading files. It can also prompt users for input and cancel execution of the tasks.

MCollective

MCollective (Marionette Collective) is a sophisticated framework that helps you build server orchestration and parallel job-execution systems. The purpose of Mcollective and similar orchestration frameworks is to allow you to parallelly execute configuration changes across multiple systems in a controlled fashion. MCollective is a tool designed to orchestrate and synchronize changes across a large number of servers simultaneously. The name Marionette alludes to a classic marionette controlling puppets.

Later in this book, you'll learn a lot about configuration management (CM) tools such as Puppet, Chef and Salt. How does a server orchestration tool such as Mcollective relate to those tools? Here's the difference: CM tools are for what they say they do – achieving consistent configuration in your data center. A tool such as MCollective has a far narrower scope – it lets you orchestrate changes in a parallel fashion.

How Mcollective differs from traditional tools

System administrators are used to rigging up scripts to perform simultaneous updates of a cluster of servers. However the scripting approach suffers from the following drawbacks:

- They work through a list of servers, one at a time
- They can't handle unexpected outcomes or responses
- They can't handle fatal error messages output to the screen
- They don't integrate well with or complement other management tools

MCollective contains the following features that make it quite different from the other parallel execution tools:

- You can use custom authentication and authorization mechanisms
- You can parallelly execute changes on thousands of servers without a master
- Make it possible to diagnose result codes because full data sets are returned as result codes
- Actions can be taken on the responses by processors
- In addition, Mcollective integrates very well with CM tools such as Puppet and Chef.

Often, an administrator runs something on a bunch of servers at the same time by running code that looks like the following:

```
$ for host in bunch of hosts
do
    scp config-file $host:/some/path
    ssh $host "service apache restart"
done
```

When you run this code for a large number of hosts, several issues might crop up – for example, you can't easily keep up with the output of the commands, and errors in the middle of the sequence can be missed. Your goal is to make sure that you do this thing fast and know for sure that the commands worked on all the target servers.

While Puppet and Chef and similar CM tools can help you make changes in systems, they really aren't tools designed to perform massive simultaneous deployments. The tools ensure eventual compliance of systems over a period of time. Something like puppetmaster can process only a few systems at once. Mcollective is a good complementary tool for CM tools such as Puppet (Mcollective ships as part of Puppet 4), and helps you achieve true parallel execution with consistent results.

The key difference between Mcollective and a tool such as Puppet is that Mcollective is a tool with a narrow focus – it lets you perform small changes across a huge number of nodes at precisely the same time. Puppet can perform numerous changes to ensure consistent configuration among a bunch of nodes, but it takes its time to get this done.

How Mcollective Works

Mcollective avoids two key drawbacks of other parallel execution tools:

- First, there's no central master, which helps you eliminate potential bottlenecks with a master/server centralized architecture for single server managing deployments over a bunch of servers.

- Secondly, it avoids drift among the servers it's updating, since it doesn't process the clients in an ordered loop as a shell script might, for example.

MCollective uses a unique concept of what the terms server and clients mean. A node you want to control through Mcollective is deemed a Mcollective server and it runs `mcollectived`. A node with the `mco` command-line client installed is deemed a MCollective client which is capable of issuing commands. You install the client software just on those systems from where you want to send requests. You can use management hosts or a bastion host (or your own laptop) for this.

You can use the `mco` command-line client in scripts or interactively. You can also write your own custom clients in Ruby to use as back ends for GUI applications.

The key to the scalable and fast parallel execution capability of MCollective is the Publish-Subscribe infrastructure it uses to carry requests to the servers. The `mcollectived` process on each server registers with the middleware broker (ActiveMQ or RabbitMQ for example) and immediately grabs and evaluates requests sent to the middleware broker by the clients. The `mcollectived` process uses an agent with which it immediately executes the command to satisfy the client requests. It's the agents that you install which process the requests and perform actions. The Publish-Subscribe mode permits simultaneous execution on hundreds or thousands of servers at the exact same time.



Puppet Labs (it maintains both Puppet and Mcollective) recommends ActiveMQ as the best middleware for high performance and scalability

You can control which systems execute specific commands using filters on host-names, operating systems, installed packages, and similar criteria. The Mcollective agents report back with the status of the process initiated by them. MCollective agents are available for Puppet, Chef, CFEngine, Ansible and Salt.

Mcollective and Puppet

You can use MCollective to control Puppet agents on various nodes. You can use MCollective to:

- Start and stop the Puppet agents
- Run the Puppet agent with various command-line options
- Make changes to nodes using Puppet resources

- Select the node to modify based on the Puppet classes or facts (key-value pairs with information about the nodes – Puppet’s `facter` program is the most common way to get the facts) on the node

Installing and Configuring MCollective

You can manually install the various binaries that you need to make MCollective work (including the middleware broker), but the recommended way to install and configure MCollective is through a CM tool such as Chef, Puppet, Salt, CFengine or Ansible. The reason for this is that it’s tedious to install MCollective manually on numerous servers, and using a CM tool makes it easy to maintain MCollective over time, as well as customize its settings on some servers.

Now that you’ve seen how parallel command job execution frameworks can help you, let’s take a look at two highly useful automatic server provision tools – Razor and Cobbler.

Server Provisioning with Razor

Razor is a popular automatic provisioning tool that lets you install servers across your data center. Razor also integrates very well with CM tools, so you can easily go from a bare metal server or VM with nothing on it to a fully configured server with the help of Razor and a CM tool.

You can install Razor by itself, and Puppet Enterprise also bundles it with its installation media. Razor’s broker component also supports Chef, and you can write plugins for other CM tools such as Ansible and Salt as well.

How Razor Works

Razor uses TFTP, DHCP and DNS to support automated server deployment. Razor, written in Ruby, runs as a TorqueBox web application and uses PostgreSQL as its backend, and it has its own network. When you connect a new server that’s configured for network boot to the Razor network, Razor does the following:

- The new server does a network boot and connects to Razor and loads its microkernel
- It detects the new node and learns about its characteristics by booting up with Razor’s microkernel and collecting information about it
- The new node gets the microkernel and registers with Razor, and will be bootstrapped by Razor by loading the installation files
- If the new server’s “facts” match a tag, the tag is compared to policies stored in the Razor server and the first matching policy is applied to the new server

- The new server is configured with an operating system, based on what the policy specifies
- If the policy contains the relevant broker information, Razor performs a handoff to the Chef or Puppet CM tool

Razor's Architecture

You can look at Razor as consisting of the following components:

- The Razor server
- The Razor microkernel
- The Razor Client

The Razor Server

The Razor server is Razor's main component. TorqueBox, which is application platform based on the JBoss application server, hosts the Razor server. You can interact with the Razor Server through the Razor client or through RESTful APIs.

The Razor Microkernel

The microkernel is just what it sounds like – it's a tiny Linux image that is used to boot on the new nodes and inventory the nodes. During the discovery stage, if the microkernel can't find matching policies for the new server, the server will present the microkernel's login screen, but normally you don't long into the microkernel.

The Razor Client

The client lets you access the Razor server and it's probably a good idea to run it on the same machine where you run the server, since the client provides no authentication by itself.

Working with Razor

If you're using Puppet Enterprise (PE), you already have Razor – however you must still configure the required database (PostGres), DHCP server, DNS server and an image repository.

To install Razor when you've Puppet running, you need to do very little:

```
puppet module install puppetlabs/razor
puppet apply -e 'include razor'
```

There are a number of prebuilt Vagrant environments that you can download and use. Or, you can manually install Razor yourself.

In addition to the PostgreSQL database as a backed server, Razor needs both DHCP and DNS servers. Razor assigns IP addresses to the new nodes that it installs through DHCP.

If you're dealing with a small test environment, you can use dnsmasq for both DHCP and DNS.

Server Provisioning with Cobbler

Cobbler is a Linux installation server that helps you simplify server provisioning by centralizing and automating tasks involved in the configuration and administration of an installation server. Cobbler helps you quickly setup a network installation environment and also serves as a tool that helps you automate tasks in Linux environments. Originally it was bundled with Fedora, but since 2011, it's being also bundled with Ubuntu. Cobbler glues together several related Linux tasks so you don't have to worry about each of them when performing a new installation or modifying an existing installation.

Cobbler has a built-in configuration system and integrates well with other systems such as Puppet, for example. You mostly use the commands `cobbler check` and `cobbler import` to perform the initial setup from the command line but can use a web application later on.

Cobbler is ideal for network installs which you can configure for PXE, media based network installations and virtualized installations with Xen, KVM, configuration management orchestration, etc.

Cobbler is also helpful in managing DHCP, DNS and the yum package mirroring infrastructure. It has its own lightweight CM system and you can integrate it with Puppet and other systems as well.

Where Cobbler Can Help You

As I showed earlier in this chapter, when performing a network environment, for server installations you must perform all the following tasks:

- Configure DHCP, TFTP, DNS, HTTP, FRP, NFS and other services
- Customize the DHCP and TFTP configuration files
- Create the automatic deployment files such as the Kickstart file
- Extract the installation binaries to the HTTP/FTP/NFS repositories

This sequence of steps involves a lot of manual work: you must manually register each client machine, and any change in the provisioning of a server means a manual change in the configurations and the automatic deployment files. The TFTP directory and other files can get quite complex when dealing a large number of machines.

Cobbler was designed to tackle the system related issues head on by acting as the central management point for all machine provisioning tasks. It lets you install machines without manual intervention. You simply run a command such as add new repository or change client machine operating system, and Cobbler will:

- Reconfigure services
- Create the repositories
- Extract OS media to newly created directories
- Control power management
- Restart servers

What Cobbler Offers

Cobbler sets up a PXE boot environment and controls everything related to the installation. When you use Cobbler to create a new machine it:

- Uses a template to configure the DHCP server
- Mirrors a repository or extracts a media to register the new OS
- Creates an entry in the DHCP configuration file with the IP and MAC addresses you specify for the new machine
- Creates necessary PXE files under the TFTP service directory
- Restarts the machine to begin the installation, if you enable power management

Cobbler knows how to extract necessary files from the distro ISO files and adjust the network services to boot up the new machines.

Kickstart

Kickstart templates let Red Hat and Fedora based systems automate the installation process. Cobbler works with Kickstart. You can use the Kickstart file to install machines in different domains and with different machines, by using configuration parameters (variables) such as `$domain` and `$machine-name`. A Cobbler profile can then specify, for example, `domain-mydomain.com` and specify the names of the machines that use this profile with the `machine-name` variables. All new machines in this Cobbler configuration will be installed with the same Kickstart configuration, and configured for the domain `mydomain.com`.

Fence Scripts

Cobbler can connect to power management environments such as blade center and ipmitool through fence scripts. When you reboot a new system, Cobbler runs the appropriate fence script for you.

Cobbler Architecture

Cobbler uses a set of registered objects to perform its work. Each registered object points to another object, inheriting the data of the object it points to. Cobbler uses the following types of objects.

- **Distribution:** represents an OS and contains information related to the kernel and initrd
- **Profile:** points to the distribution, a Kickstart file and other repositories
- **System:** this object represents the machine you're provision. The system object points to a profile or an image and includes specialized information such as the IP and MAC addresses and power management.
- **Repository:** this object stores the mirroring information (mirror URL) for a repository such as yum
- **Image;** this image can replace a distribution object for files that don't fit in that category. An example is where you can't device the files into a kernel and initrd.

Deploying an Operating System using Cobbler and PXE Boot Cobbler comes with a great deal of functionality out of the box, although it can get complex due to the many technologies it can manage. You need to know PXE (see in this chapter) and the procedures for automatically installing the Linux distro you want to install. Installing Cobbler is quite easy through the yum utility (`yum install cobbler`).



Installing and configuring the Cobbler web interface is a good strategy if you need to perform regular activities with Cobbler.

Infrastructure as Code, Configuration Management and Orchestration Tools

This chapter starts off with an introduction to the concept of Infrastructure as Code. Infrastructure as Code, also referred to as programmable infrastructure, involves the writing of code to manage infrastructure configuration, deployments and automatic provisioning.

Infrastructure as Code is a philosophy that takes you further than traditional change management automation practices, which involve only the replication of install/configure steps on multiple servers without your intervention. Infrastructure as Code ensures that all of your servers and other infrastructure components are provisioned consistently and effortlessly.

Under an Infrastructure as Code methodology, you go beyond the writing of mere automation scripts to using tested and established software practices that application developers have been successfully using for ages. These practices include version control, testing, using design patterns, and so on, all of which have revolutionized software development in recent years. This means that you go way beyond the mere automation of your change management processes.

The bulk of this chapter is devoted to how you use various tools for configuration management and infrastructure automation coding. Using tools such as Chef, Puppet, Salt and Ansible, you can plan the setting up, packaging, and delivery of applications in your infrastructure as code. Using this code, you can build an entire infrastructure application stack in a few minutes instead of the traditional days and weeks.

I discuss four major server configuration management tools in this chapter:

- Chef

- Puppet
- Salt
- Ansible

Infrastructure as Code

In today's cloud driven environments where automation tools proliferate, you can't expect to continue to employ the same strategies and procedures you used to manage infrastructure and software. In earlier days (just 5-10 years ago), it was common to take days and weeks to provision infrastructure components – today you need to do that in minutes, or in seconds.

Legacy change management is out of the window, as nimbler and more agile competition forces every organization to handle cloud environments and modern automation tools using infrastructure as code.

Large web operations run by companies such as Facebook, Twitter, Google and the rest are using radically new approaches to the designing, building, and troubleshooting their IT infrastructures. The key to all these and other companies that are successfully tackling the challenges of modern day large scale web architectures is their strong underlying adherence to the Infrastructure as Code approach.

By running your infrastructure as code, you can give operational responsibilities to development teams. As system administrators, you can then focus on the design and software lifecycles at a higher level, so to speak. Those who work with the systems will also operate them, and the admin teams can focus on building robust and scaled infrastructure architectures. The Infrastructure as Code approach lies at the very heart of the modern DevOps philosophy.

Infrastructure as Code is a way of managing infrastructure automation using well established and successful software development practices. Infrastructure as Code principles enable you to manage infrastructures so that you can make changes frequently, routinely (without drama), safely, and in a consistent fashion. You use consistent and repeatable procedures to provision and configure systems. Any changes made to the definitions you roll out to the infrastructure components using automatic, validated procedures that don't require your involvement.

While scripting and automation to create and update servers has been around for a very long time, virtualization and the move to the cloud has made a new class of tools popular. These tools make it easy for you to create a large number of server instances with minimal effort which you then need to configure and update.

Infrastructure as code essentially means the following two things:

- Infrastructure can and should be treated as code.
- Those who work with infrastructure should follow similar principles as software developers

Infrastructure as code in essence means that your environment resembles an application in many ways, namely:

- Its versionable
- Its testable
- Its repeatable

Benefits of Treating Infrastructure as Code

When you treat infrastructure as code to manage the configuration of your environment, you reap the following benefits:

- Automated CM tools are run any time you make changes to a component's configuration definition, or to any process that's part of provisioning and configuring infrastructure components.
- You can provision servers and other components blazing fast, in mere minutes.
- You can automatically provision various infrastructure components in response to events (such as a spurt in demand due to holiday season)
- The provisioning and configuration management changes are repeatable,, self-documented and consistent
- Your server configuration changes are applied without your involvement.
- Any change is applied uniformly to all the infrastructure components that are relevant, regardless of when they were provisioned
- Configuration changes are versioned and applied to various environments such as staging and testing, to support controlled testing and staged release strategies

Dealing with infrastructure as code will result in your getting away from routine, repetitive tasks. Your teams can provision and manage their resources, without your intervention. The teams will also be able to recover quickly from failures. Best of all, are made continuously in an incremental fashion, rather than waiting for infrequent, mammoth high risk deployments.

When you start treating infrastructure as code, you can take advantage of tools that software developers have been using forever, such as Version Control Systems, automated testing libraries, and deployment orchestration. Not only that, you'll be laying

the foundations for highly useful development practices such as Test Driven Development (TDD), Continuous Integration (CI), and Continuous Delivery (CD).

Cloud Environments and Infrastructure as Code

Infrastructure as code really came onto vogue after the emergence of cloud infrastructures, where dynamically managing infrastructures is a must. Cloud environments such as Amazon Web Services must not only be automatically provisioned, of course, but you should have the ability to easily destroy and rebuild servers, as well as to add and delete servers based on demand. It's infrastructure as code that makes infrastructure management truly dynamic.

Virtualized environments also can use dynamic infrastructure management, since servers always grow much faster than your ability take care of them. Well known problems such as server sprawl, where you contend with too many servers, and configuration drift with multiple software versions and different configurations across servers wreak havoc on your infrastructure management scripts.

Tools for Treating Infrastructure as Code

There are two broad sets of tools you can use when dealing with infrastructure as code:

- Infrastructure definition tools
- Server configuration tools

The following sections explain these two sets of tools.

Infrastructure definition tools

Infrastructure definition tools help you specify the infrastructure resources you want to allocate, and how they ought to be configured. AWS's Cloud Formation is a good example of a tool such as this, and I'll discuss it in detail in Chapter 10 ("Managing an Amazon Web Service Cloud").

Once you specify your infrastructure requirements in configuration definition files, the tools interacts with the underlying dynamic infrastructure platform and implements your requirements so the infrastructure matches your specifications. The dynamic infrastructure APIs will then create the server using a server template (a base image, or a proprietary format VM image, or an OS installation disk image from a vendor. may c

HashiCorp's Terraform and OpenStack Heat also fall under this category of tools. Although you mostly use Chef, Puppet and related tools to configure servers rather than provision them. Provisioning also falls under this category.

Since the tools use a declarative (what you want) rather than a procedural (how you want it done), the results are idempotent (repeated execution of the code results in identical results).

While you can execute the infrastructure definition tool manually, most of the time these tools pass the configuration information to a server configuration tool, when creating servers.

The definition tools network configuration (DNS addresses) and may also pass along information such as the server roles (database versus web server, for example), and the CM tools will then install the required software and configure it.

For a simple infrastructure, the configuration definition file that a configuration definition tools uses will suffice, but for larger environments, you may use configuration registries. A registry is simply an information directory for all your infrastructure components.

In Chapter 6, you learned about Zookeeper, Consul and etcd –all of these tools are well-known configuration registry products. In addition to these registries, all CM tools, such as Chef (Chef Server), Ansible (Ansible Tower), and Puppet (PuppetDB) come with their own built in configuration registries.

Server Configuration Tools

Server (and other infrastructure component) configuration tools such as Chef, Puppet, Salt, and Ansible are designed explicitly to configure infrastructures with the Infrastructure as Code approach. Using configuration definition files, the tools apply required configurations to the infrastructure components.

Agents on the nodes in an infrastructure may either pull the latest configuration definitions from a repository and apply them, or a central server can trigger updates to the nodes. The first approach is called pull-based, and the latter, push-based. Under the latter model, nodes may not need agents running on them. I explain the two models in the context of four popular CM tools later in this chapter.

Test-Driven Infrastructures

Well, if you base infrastructure management on the same principles as traditional application code development processes, then you ought to also embrace probably the most crucial part of an application development cycle – testing. Just as developers test their code before deploying, you can also put infrastructure code through testing, using similar methodologies.

A test driven infrastructure approach guards against the risks of implementing the Infrastructure as Code philosophy. In order to ensure that your code does actually produce the infrastructure changes you intended, and also that your code hasn't

caused any undesirable side effects, you put automated testing in place. The best approaches to follow here is to simply adopt the best practices in software development.

Test and Behavior Driven Development

The modern agile software movement that came about to overcome the pitfalls of traditional software development produced a testing practice called test-driven development (TDD). Although it's hard to implement TDD based testing in infrastructure development, it's worth the effort.

The phrase Red, Green, Refactor underlies the TDD testing approach. Under TDD, you follow these steps in an iterative fashion:

- Write a test
- Run the test and watch it fail.
- Write just enough code to make the test pass
- Run the test and watch it pass.
- Enhance code to make it perform well, readable, and reusable, without changing its behavior
- Rinse and repeat

Basically what you doing is writing cod to make failing tests pass – test, code, refactor is the strategy here. While repeated testing takes effort, it reveals design problems and builds trust among the teams.

Behavior Driven development focuses not on the reliability, correctness and maintainability of code, but rather addresses the question as to whether the code helps solve problems of your customers. Regardless of whether your clients may be the teams such as application developers and testers, or the final users of your company's; products, your goal would be deliver business value by providing infrastructure that meets the needs of your customers.

Testing Tools

There are several testing tools that help test infrastructure code. Unit testing only shows that the code works as expected, but it doesn't know how the code should "behave".

Tools such RSpec and Cucumber are behavior-driven, and focus on whether the code behaves as specified. These tools require that the specification of how code ought to behave should itself be code that you can execute.

With this introduction to the philosophy of treating infrastructure as code behind us, let's turn to the topics of configuration management automation, which are designed to configure infrastructure components with the infrastructure as code approach in mind, and a tour of the most popular CM tools being used today.

Configuration Management Automation

Configuration management is the set of practices involved in managing all components involved in delivering software application services to users. Thus, hardware, software, infrastructure (network etc) processes are all part of configuration management.

When you are a system administrator who's managing just a handful of servers, you can manage change pretty easily by doing things manually or with the help of some shell scripts. However, when you work with a team of admins/developers and need to configure hundreds (or thousands) of servers and applications, you need a configuration tool.

All the configuration tools I discuss in this chapter follow the essential principles of automation: they ensure that a system is configured in a consistent and reliable fashion. The tools can also automatically detect and fix systems without you having to tell it to do so.

Need for Configuration Management Tools

Following are the main reasons for the widespread use of configuration management tools.

- **Consistency:** automating routine tasks through CM tool ensures that you're setting up your servers in a consistent manner.
- **Robustness:** building systems manually leads to fragile systems that can break down easily during change management.
- **Efficient Change Management:** since you can reproduce servers in a consistent manner, you can perform frequent changes – change management isn't the nightmare it otherwise is, with all its potential for errors and downtime.
- **Understanding of the Change Process:** CM tools help teams understand the change process clearly, through their auditing and reporting features. Configuration changes made by all the system administrators are logged, for everyone to see.

Why Use Configuration Management Getting a web application to run doesn't end with just starting up the web application server and starting the web application. If you ever need to migrate the app to a different server or environment, you'll need to

muck with a bunch of configuration files, install software, configure the software, check for required packages, and so on.

As applications get increasingly complex, you find yourself having to deal with services that run on numerous servers, and include ancillary services such as web servers, databases, application servers, in-memory caching systems, message queues, load balancers and the list keeps growing every day!

On top of the actual services that make your applications run, you also need to setup logging and monitoring and as well as interact with third-party services.

Obviously, you don't want to go through this painful exercise every time you want to run your app on a different server. The correct way to do this would be to use configuration management. CM captures all the knowledge that's required for your application to run, and CM always stays up-to-date.

If you're thinking of setting up an OpenStack cloud for example, using a CM is really a must (although not technically) from a practical point of view.

Declarative Versus Imperative Programming

When performing system changes through scripts, system administrators typically use a programming style best known as imperative programming. Imperative programming means you provide the list of the specific operations and the commands (and the order in which they ought to be run) that change a target's state. Even when you use procedural programming with procedures that perform the state changes, you still define each operation and procedure that needs to be executed, as well as the order in which to execute them.

Configuration tools such as Puppet and Chef force you to think quite differently from what you're used to. These and other similar CM tools use a programming style called declarative programming. Declarative programming is easier to write, read and maintain. Under this programming style, you simply state what you want the end state of a system to be.

Let's take a simple and common change in systems, which is the creation of a user. You use a command such as `useradd` to add a user to a Linux system. Let's say the user already exists – when you run the command, of course it'll fail. To avoid this you can put in an if/then piece of logic in your shell script and it'll take care of this particular case where the user exists already.

However, if you need to check other things such as the uniqueness of the new user's UID/GID, and whether the password expiration is set, you're in for a lengthy script to take care of all those things. Imperative programming simply can't handle the various possibilities, since it needs to describe the actions required to reach the desired state.

Declarative programming on the other hand, simply describes the desired state – it leaves the actions to the interpreter. The interpreter here is something such as a Puppet or Chef agent – it's up to these agents to evaluate the current state and apply the necessary changes to bring the configuration to the desired state. Since you aren't writing out each of the steps in detail to handle every situation that may arise, declarative programming means simple, compact code.

A simple example here using the Puppet configuration language shows how to express the desired state of configuration:

```
user ( 'sam':  
  ensure    => present,  
  uid       => '1001',  
  gid       => '1000',  
  comment   => 'Sam Alapati',  
  managehome => true  
)
```

Note the difference that imperative programming makes: you only need to specify the username, uid, gid, etc. This is how you want this resource (user) to end up as (final state). How to create the user and ensure that the other desired configuration is done is up to the CM tool (Puppet, Chef, Salt, etc.). All the popular CM tools in use today use declarative programming as their foundation.

In the following sections, I introduce and briefly explain four popular configuration management tools: Chef, Puppet, Salt and Ansible.

Configuration Management with Chef

Chef is probably the most popular CM tool right now (late 2016). Chef turns your infrastructure, both on-premises and in the cloud, into code. In the following sections, I explain the basics of Chef's architecture, as well as its key features.



Chef isn't just for managing and configuring servers – you can use Chef to manage several components of your infrastructure, such as switches, routers and storage.

Chef's Architecture and Key Components

It's easy to understand Chef's architecture by looking at it as a combination of server and client components. The Chef Development Kit is a requirement to develop Chef code and the Chef Server is needed when you want to manage large environments.

The Chef Development Kit

The Chef Development Kit helps you write Chef code. It consists of the two main components – the Chef Client and a Test Driven Infrastructure.

The Chef Client consists of the following components:

- Cookbooks: A Cookbook is a collection of automation scripts
- Knife : Knife helps you manage your servers.
- Solo: Chef Solo allows you to run Chef code locally without needing a Chef Server.
- Zero: Chef Zero is an in-memory, fast-start Chef server.

The Chef Test Driven Infrastructure consists of the following components:

- Test Kitchen: Helps create sandbox environments that simulate production set-ups
- Bookshelf: Is used to store cookbook content such as files and templates.
- Chefspec and Serverspec: Chef's built-in testing tools

The Chef Development Kit is what you use to start writing Chef code. Since the Kit is written in Ruby, you must install the Ruby native scripting engine to run the development tools. The Kit's installer comes with the right version of the Ruby scripting engine. In a production environment, you install the Chef Client on every system you want to manage with Chef.



A Ruby gem is a supporting app or library written in Ruby. The gem is very similar to an installer. For example, to install Test Kitchen on Linux, you run the gem install command shown here, which installs the test-kitchen gem:

```
$ gem install test-kitchen -no-ri -no-rdoc
```

Note that you use the chef knife bootstrap command to install Chef Client on production nodes, as Test Kitchen isn't meant for production use.

The Chef Server

The Chef Development Kit encompasses the Chef Client. Chef Development Kit helps you write Chef code. Chef Server provides the additional support to scale your configuration management capabilities to a large number of servers.

Chef Server helps you scale configuration management to a large number of servers.

Chef Server relies on these components:

- Server Core: On premises or hosted
- Web Interface: management console and LDAP/AD integration
- Actions: reporting jobs, pushing jobs

Configuration management means you use recipes to manage the configuration of nodes. System configuration is in n most cases stored in files stored on disk.

Chef uses templates to dynamically create configuration files. It may take the values from data bags or attributes or calculate them on the fly and pass them to a template.

Writing Chef Recipes

Chef uses text files called recipes to configure your infrastructure. You declare how you want the state of an infrastructure component in a Chef recipe, using Chef code. Since Ruby is the language you program Chef with, I start off with a very brief intro to Ruby and follow it up with an outline of Chef's syntax for its recipes. Once you review these preliminaries, you'll be in a good shape to appreciate the small Chef recipe I create later in this section.



A recipe refers to the code that consistently reproduces an infrastructure component. A recipe is a file that contains Chef code.

Ruby and Chef

The Chef Development Kit (and the Chef Client) are implemented in Ruby. Although you don't need prior background in programming with Ruby, you do need to learn its basics and learn how it's related to Chef syntax.

Ruby is an object-oriented language that's quite similar to Python and is designed for usability and programmer productivity (not to speak of fun!). You can quickly learn what you need of Ruby, as its syntax for variables, arrays, regular expressions and conditionals (if/then) are quite similar to other languages. You define Ruby methods with the `def` keyword and classes with the `class` keyword and modules with the `module` keyword.

Chef Syntax

Chef accesses the power of the Ruby language by using a Domain Specific Language (DSL), which is a subset of Ruby. So, everything you can do in Ruby, you can do with Chef code. The DSL contains Ruby-like statements that express Chef system administration directives. The use of customized expressions for system administration in the

DSL makes it easy for you to get started with using Chef to perform system administration tasks.

Chef uses resources as its building blocks to define an infrastructure's components, and to configure things on a system. Resources help you define actions that Chef will perform. Here's the syntax of invoking a DSL method in Chef:

```
resource 'NAME' do
  parameter1  value1
  parameter2  value2
end
```

In this method, here's what the key parts stand for:

- resource is the name of the resource, such as package, service, or template.
- NAME stands for the name of the resource
- do and end: the keywords do and end are the beginning and end of a block. You declare resource parameters and their values inside a block.
- parameters: these help assign attribute values to the resources

If you wish to look at this from an object-oriented point of view, you can say that that the DSL creates new resource objects and configures them with various attributes for the resource's parameters, and finally, executes the resources upon evaluation of the code by Chef. That is, the DSL method can be viewed as:

```
resource = Resource.new('NAME')
resource.parameter1 = value1
resource.parameter2 = value2
resource.run!
```

The following chunk of Chef code uses the package, service, and template resources to perform infrastructure configuration.

```
template '/etc/resolv.conf' do
  source 'my_resolv.conf.erb'
  owner  'root'
  group  'root'
  mode   '0644'
end

package 'ntp' do
  action :upgrade
end

service 'apache2' do
  restart_command '/etc/init.d/apache2 restart'
end
```

And here's what this Chef code snippet does for you:

- It declares and compiles (from `my_resolv.conf.rb`) a template resource and places it on the target machine, where Chef will evaluate this code.
- Declares a package resource named `ntp` that will upgrade the `ntp` package
- Declares a service resource named `apache2` that will restart `apache2`. Chef will manage this service.

In addition to the three resources you’ve seen (package, service, template), there are many others, such as `cron`, `directory`, `file`, `mount`, and `execute`, for example. The following three examples show how you manage a file, create a symlink, and create a group.

```
# Delete the /tmp/ops file
file '/tmp/ops' do
  action :delete
end

# Link /tmp/ops to /tmp/LinuxAdmins
link '/tmp/ops' do
  to '/tmp/LinuxAdmins'
end

# Create the ops group
group 'ops'
```

A Simple Chef Recipe

A Chef recipe is a file that contains Chef code. Let’s create a recipe that prints the like “Welcome to Chef!”. You create a Chef recipe that does this for you by creating a file named `hello.rb` as shown here (the `.rb` extension means it’s a Ruby file).

```
file 'hello.txt' do
  content 'Welcome to Chef'
end
```

The resource here is of course, a file (refers to the `hello.txt` file) resource, and the recipe tells Chef to do the following:

- Create a file named `hello.txt`
- Write the statement “Welcome to Chef!” to the file `hello.txt`



A chef recipe indicates just the desired configuration to Chef, and not the actual steps to perform the configuration on a machine.

Once you have your recipe ready, how do you get Chef to perform the actions inside the recipe? I haven't shown you how to use a Chef Server yet, so let's use Chef Solo to execute the Chef code in the recipe. Chef Solo lets you run Chef recipes without a Chef Server.

The `chef-apply` tool is a wrapper on top of Chef Solo, to make it easy for you to run Chef code. Using `chef-apply`, you execute the code as follows:

```
$ sudo chef-apply hello.rb
Recipe: (chef-apply cookbook)::(chef-apply recipe)
  * file[hello.txt] action create
    - create new file hello.txt
    - update content in file hello.txt from none to 40a30c
    --- hello.txt      2016-04-30 20:18:28.000000000 -0700
    +++ /tmp/.hello.txt
    -0700 @@ -1 +1,2 @@
    +Welcome to Chef
```

What I've done with this command is to automate the creation of the `hello.txt` file through Chef.

Specifying the Desired Configuration Rather than How to achieve it

Chef's DSL focuses squarely on your describing the configuration of a machine rather than telling Chef how to do the actual configuration. It's up to Chef to figure out how to do it.

In our example recipe (`hello.rb`), I asked Chef to create a new file named `hello.txt`. If the `hello.txt` file exists, `chef-apply` won't do anything. It'll create the file with the content I specified, only if the file doesn't exist.

What if the `hello.txt` file exists but its contents were changed? Chef is smart enough to check the file's contents and if the contents are modified, it reverts the file's contents to the content you specified in your recipe. There's no way to change the file's content unless you specify different content in the `hello.rb` recipe. This, in a nutshell, is how Chef helps you combat configuration drift in your environment – written in cement it indeed is!

Using Test Kitchen to Manage Test Environments

Although you can test Chef code on your private server before launching it in a production environment, note that Chef will make changes to the configuration of the server when you run the code.

A smart way to go about testing Chef code during development is to set up a sandbox environment that closely resembles a production environment. This will give you a safe place to test your Chef recipes. Chef comes with Test Kitchen, which helps you

create a sandbox environment for testing. Test Kitchen utilizes Vagrant and Virtual Box to get its job done.

Test Kitchen runs on Vagrant, and you create your sandbox environment on top of Test Kitchen. Test Kitchen is installed as part of the Chef Development Kit, and you need to install it separately if you're using Chef Client.

In order to create a virtual environment using Test Kitchen, you use the `kitchen create` command:

```
$ kitchen create default-centos65
```

This example shows how to create a virtual environment running CentOS. This command downloads the Vagrant base box and configure and boot the VM instance. Test Kitchen will pull base boxes that Chef Software makes available on the Internet via VagrantCloud. The CentOS instance created by the last command will set up a bare-bones CentOS installation with just enough stuff to let Chef run.

You can login to the CentOS VM by doing this:

```
$ kitchen login default-centos65
Last login: Fri May 28 10:41:48 2016 from 10.0.1.1
Welcome to your Packer-built virtual machine.
```

Run all your test Chef code in this Test Kitchen supported sandbox environment.

Test Kitchen uses YAML file format for its configuration files. YAML files work with two types of data – key-value pairs and lists.

Using Chef Client to Manage Nodes

Any machine that Chef manages in your environment is called a node. A node need not be a physical server- it can be a VM, a cloud instance, or a Docker container. You just need to install Chef Client on the node for Chef to manage the node. You manage the nodes through Chef by running Chef recipes.

The chef-client evaluates Chef cookbooks that contains recipes with Chef code, on each target node.

Earlier, I explained how Test Kitchen is the way to go to test Chef code in a sandbox environment. In order for Chef to manage the test environment, you make the VM running Test Kitchen a node for Chef.



chef-client is an agent (or service) on a server managed by Chef. Chef server is a centralized repository for storing information you need to manage the entire infrastructure with Chef.

Assuming that you're still working with Test Kitchen in a test environment, you can use the following command to set up the chef-client:

```
$ kitchen setup default-centos65
```

The kitchen setup command installs chef-client by running Chef Solo as the provisioner.

I showed you how to use the chef-apply tools to execute Chef code in a Chef recipe. However, in a production setting, you use the chef-client tool rather than chef-apply to run the code. One of the big reasons is that real life Chef recipes can be long and you may use multiple recipes so as to make configuration maintenance easier. Chef-client is best suited for running Chef code across multiple recipe files. You can still use chef-apply for one time adhoc management tasks.

When you execute a coding recipe through chef-client, it's called a Chef run. During a Chef run, chef-client reads recipes that indicate the desired configuration through resources, and Chef performs the tasks necessary to configure the node into the state you desire. Here's an example:

```
$ chef-client --local-mode hello.rb
```

Chef Client Modes

You can run the Chef Client in three different modes: local, client, and solo. Here are the key differences among the three modes:

- In the local mode, the chef-client tool simulates a Chef Server instance in memory by using Chef Zero, the in-memory Chef server. Instead of writing data to a server, it writes to a local directory. This process of writing data to the local file is called writeback.
- In the client mode, it's expected that you've a Chef Server running on one of your systems in your network. In a production setting, this is the mode in which you use Chef most of the time.
- The solo mode is really a legacy mode, and lets you run Chef locally – it's better to use the local mode than the solo mode.

You use the knife bootstrap command to install Chef Client on your production nodes- Test Kitchen really isn't for production usage.

Ohai and Collecting System Information

Ohai is a command line tool used by Chef Client during its runs. Ohai collects system information for Chef such as network configuration, memory and CPU details, OS system type, etc. Here's an example that shows a small part of the information collected by the ohai tool:

```

{
  ...
  "ipaddress": "10.0.2.15",
  "macaddress": "08:00:27:1C:AD:B6",
  ...
  "os": "linux",
  "os_version": "2.6.32-431.el6.x86_64",
  "platform": "centos",
  "platform_version": "6.5",
  "platform_family": "rhel",
  ...
  "virtualization": {
    "system": "vbox",
    "role": "guest"
  }
  ...
  "hostname": "default-centos65"
  ...
}

```

The chef-client reads the ohai information and converts it into a node object, which the Chef code can access. Things such as IP addresses in the output, for example, are attributes, which become variables that Chef maintains. It's the detailed information gathered by ohai that helps Chef to perform node configuration correctly.

Chef Recipes and Chef Cookbooks

Earlier, you learned how you implement Chef code through recipes. In order to configure components of your infrastructure such as a web server or application, you use Chef cookbooks. Each cookbook contains a set of instructions that'll configure/deploy an infrastructure component. A cookbook is a collection of recipes and includes things such as configuration information, archives, images, and libraries, in addition to the all-important recipes.

If you're using the Chef Development Kit, you can use the `chef generate` command to create cookbooks:

```

$ chef generate cookbook abc
Compiling Cookbooks...

```

If you're using Chef Client instead for development, use Chef's knife tool to generate the cookbook directory structure:

```

$ knife cookbook create test --cookbook-path .

WARNING: No knife configuration file found
** Creating cookbook test
** Creating README for cookbook: estd
** Creating CHANGELOG for cookbook: test
** Creating metadata for cookbook: test
$

```

How Chef Performs its “Runs” When you run Chef in a production setting, chef-client runs as a service on each node, and at frequent intervals (say hourly), polls the Chef Server for any changes in the cookbooks or recipes that the client must implement on that node.

Following are the stages/steps of a Chef run.

1. The chef-client process starts on the remote nodes, through a cron job or a service.
2. The chef-client process runs ohai to gather the node’s attribute information and constructs the node object (in memory).
3. The Chef Server sends a run list containing the recipes that should be executed on the target nodes.
4. Cookbooks and Ruby components are loaded. Things such as recipes (files), attributes, resource providers, and libraries are loaded at this point.
5. During the converge phase, the Chef recipes are executed on the nodes. Here’s where the actual work is performed, such as the installation of software packages, copying of files, etc.
6. Notification and exception handlers send emails and messages to pagers regarding the success/failure of the Chef run.



A chef run can involve dozens of cookbooks and hundreds of recipes. The run list points to all the files in the cookbook.

MANAGING MULTIPLE NODES WITH THE CHEF SERVER

You can make do without a Chef server if you’re dealing with a single node, but in any type of production environment, you need the Chef Server to manage multiple nodes. The Chef Server acts as a repository for all configuration data in your infrastructure.

Chef Server is implemented in Erlang and uses a backend database to store its data. The nginx web server acts as the interface to the Chef Server and performs load balancing as well for it. The Web UI is also how you interact with the Chef Server.

There are three types of Chef Servers:

- Hosted Enterprise Chef: This is a cloud-based “Chef as a Service” offering. A free tier lets you use 5 nodes.

- **Enterprise Chef On-Premises:** This is the Chef Server you deploy within your own infrastructure and runs within your organizations’ firewall. You can manage up to 5 nodes for free with this setup.
- **Open Source Chef Server:** Lightweight free open source version of the commercial Chef Server. Let’s you manage an unlimited number of nodes without paying a penny.

It’s interesting to note that the Chef Development Kit also comes with a bare-bones version of the Chef Server named `chef-zero` which uses a very light footprint (20 MB instead of 2 GB RAM usage) to simulate a production Chef Server environment. Chef-zero runs in-memory and helps you easily check out the full-fledged Chef Server features locally without installing the real thing.

Using Puppet for Configuration

Puppet handles complex distributed components to ensure consistency and availability. Puppet is po

Installing Puppet

While I don’t show the Puppet installation details, it’s useful to know that you install Puppet by installing and enabling the latest version of the Puppet Labs Collection repository in your system. Puppet 4 and other core dependencies are tested and shipped together in new Puppet Collections, so as to ensure that they work together.

As with Chef, you install the Puppet agents on all the target nodes you want to manage through Puppet.

In addition to the main Puppet 4 package, there are several other key programs, as summarized in the following sections.

- **Facter:** Evaluates systems and provides you facts such as the architecture, IP name, hostname etc about a system. These system facts are used by Puppet as variables for customization.
- **Hiera:** A simple lookup tool Puppet uses to you use to load configuration data used by Puppet manifests and modules.
- **Marionette Collective (MCollective):** You’ve learned about MCollective in Chapter 6. This orchestration framework is integrated with Puppet and helps you manage the Puppet agent.

Puppet Manifests

Puppet uses declarative programming and a manifest, which is a text file, is how you describe the desired state of configuration of the infrastructure resources. Manifests

are what you use to define the configuration policies you want to enforce in your infrastructure. Writing manifests is the most important part of creating Puppet configuration policies.

How you Declare Resources in a Manifest

The smallest building block of the Puppet configuration language is a resource. Puppet uses several built-in resource types that help you manipulate system components such as users, files, packages and services on your servers. In addition, you can create custom resources.

Resources create, alter, or revoke the objects belonging to a specific resource type, such as users, group, and file.

Puppet uses the following standard format for declaring resources:

```
resource_type { 'resource_title':
  ensure      => present,          # usually 'present' or 'absent'
  attribute1  => 1234,             # number
  attribute2  => 'value',          # string
  attribute3  => ['red','blue'],   # array
  noop        => false,            # boolean
}
```

And here's a simple example that shows how to declare a resource such as a file.

```
file { '/tmp/testfile.txt':
  ensure => present,
  mode   => '0644',
  replace => true
  content => 'Puppet is Amazing!',
}
```

In this example, I declare the file resource. Notice carefully the declarative tone of this policy. I just declare the name of the file that I want to exist and its contents. I'm not bothered about how to make the changes to get this done.

Creating a Manifest

You use the Puppet Configuration language to write Puppet manifests. This language has the usual data types, operations, conditionals, etc as other programming languages. Let me use the simple notify resource to write a simple Puppet manifest, which outputs the greeting message Hello World!

```
notify {'greeting' :
  message => 'Hello World!'
}
```

This manifest will of course output the greeting message “Hello World!”. You put this code in a file named helloworld.pp. The .pp file extension represents a Puppet file, in this case a manifest.

Applying a Manifest

In order to make things happen, you must tell Puppet to execute the contents of the manifests you create. You use the puppet apply command to tell Puppet to apply a Puppet manifest. Here's an example that shows the puppet apply command in action – I use the manifest file named testfile.pp and it has the declaration of the file resource I showed earlier:

```
$ puppet apply /vagrant/manifests/helloworld.pp
Notice: Compiled catalog for client.example.com in environment production
Notice: Hello, world!
Notice: /Stage[main]/Main/Notify[greeting]/message:
        defined 'message' as 'Hello, world!'
Notice: Finished catalog run in 0.01 seconds
```

Here's what our simple Puppet manifest does for us:

- It determines the evaluation order based on the dependency and ordering information
- Utilizes system facts as customization variables
- Evaluates if the changes should be applied on the target node
- Executes immediately and creates the resource (in our case a notification message) on the local system
- Lets you know that it finished the catalog run – the Puppet catalog is built from the manifest(s) – if there are multiple manifests, they belong to a single catalog that applies all the manifests together.

Dependencies and Ordering

A catalog is a set of actions that are part of various manifests. During the convergence step of the apply process, Puppet evaluates the dependencies of each resource and will apply no actions if a dependency required to be applied first fails. Puppet provides for complex dependency management using its dependency metaparameters.

If a service can't be started until a package that contains the application is installed, you can order the resources using metaparameters such as `before` and `require`.

```
package { 'puppet':
  ensure => present,
  before => Service['puppet'],
}

service { 'puppet'
  ensure => running,
  enable => true,
  require => Package['puppet'],
}
```

You could use either of the two sets of code here to ensure the dependencies are satisfied. In addition to `before` and `require`, you can also use the `notify` and `subscribe` metaparameters.

When it's building a catalog, Puppet creates a Directed Acyclical Graph (DAG) based dependency graph. It's considered a best practice not to rely on Puppet's implicit dependencies, but to explicitly define all the resource dependencies.

Creating Puppet Modules

While you can create and use Puppet manifests by themselves, it's a best practice to include individual manifests inside a Puppet module. A module is a bundle of Puppet code, related files, and data. Puppet modules offer many of the same benefits as modular programming. They help you organize code and data, share reusable code among teams, and provides files and functions for using in other modules.

Puppet Environments

When you install Puppet, it creates a production environment, as shown here:

```
$ ls -l /etc/puppetlabs/code/environments/production
total 4
-rw-r--r-- 1 root root 688 May 28 11:24:36 environment.conf
drwxr-xr-x 2 root root  6 May 28 11:24:48 hieradata
drwxr-xr-x 2 root root  6 May 28 11:24:48 manifests
drwxr-xr-x 2 root root  6 May 28 11:24:48 modules
```

It's a good idea to create a separate test environment to test new or updated modules before moving them to production.

Separating Data from Code

It's considered a good practice to separate out the code from data. That is, the manifest should contain only code, which makes them reusable code blocks that can use different sets of data for different environment or purposes.

Using the look up tool `Hiera`, you can provide node specific data to a module so it can create custom

- Company wide data common to all layers of the hierarchy
- OS specific changes
- Site-specific data
- Application specific details

Using the multilevel hierarchy which includes global data parameters common to all environments, common data, and node and environment specific data, you can use the common code blocks throughout your enterprise.

You don't really need to create all your Modules!

One of the greatest things about using Puppet is that you don't need to create custom modules for everything. It's quite likely that the module you're thinking of writing has already been written by someone else and is yours for the taking!

Both Puppet Labs (the company behind Puppet), and the global community of Puppet module developers offer numerous Puppet modules for every purpose you can possibly think of.

Puppet Forge is the largest repository for Puppet modules. Puppet Forge has a web interface you can use to find Puppet modules contributed by others. By default, the puppet module command uses the Puppet Forge repository. Here's an example:

```
$ puppet module search apache
Notice: Searching https://forgeapi.puppetlabs.com ...
NAME      DESCRIPTION      AUTHOR
puppetlabs-apache  Installs, configures, and manages ... @puppetlabs
example42-apache   Puppet module for apache               @example42
evenup-apache      Manages apache including ajp proxy ... @evenup
theforeman-apache  Apache HTTP server configuration       @theforeman
snip many other results
```

Many of the same modules hosted by Puppet Forge as well as other modules are available on GitHub (to be explained in Chapter 8).

You can also maintain your own internal repository of Puppet modules and the puppet module search command can be used with that repository as well.



On Puppet Forge or GitHub, the quality of the Puppet modules can be uneven. A module that's listed as being Puppet Supported is written and supported by Puppet Labs.

Puppet Supported modules are tested, supported, and maintained by Puppet and Puppet Approved modules are modules that are deemed to meet Puppet's quality standards and have been reviewed by Puppet Labs. Puppet Approved modules tend to be very high quality modules in general.

You can judge the module quality by checking out the Quality Score, which can reach a maximum value of 5.0. The quality score is arrived at based on the quality of the code, its compatibility with Puppet and the quality of the metadata. To further help in your selection of a module, you can look at the community rating of the modules, which incorporate user feedback.

Installing Puppet Modules

You can install Puppet modules from Puppet Forge, GitHub or from a developer's code repository. Let's look at a couple of examples that show how to install modules.

The following example shows how to install a module from Puppet Forge (or any repository, for that matter). This example installs the useful Puppet supported stdlib module:

```
$ puppet module install puppetlabs-stdlib --environment test
Notice: Preparing to install into
/etc/puppetlabs/code/environments/test/modules ...
Notice: Downloading from https://forgeapi.puppetlabs.com ...
Notice: Installing -- do not interrupt ...
/etc/puppetlabs/code/environments/test/modules
└─ puppetlabs-stdlib (v4.8.0)
You can install a module from GitHub in the following way:
$ git clone https://github.com/jorhett/puppet-mcollective mcollective
Initialized empty Git repository in
/etc/puppetlabs/code/environments/test/modules/mcollective/.git/
remote: Counting objects: 183, done.
Receiving objects: 100% (183/183), 51.13 KiB, done.
remote: Total 183 (delta 0), reused 0 (delta 0), pack-reused 183
Resolving deltas: 100% (98/98), done.
```

Writing Custom Puppet Modules

When creating custom modules, you create what are called class manifests. A class is a manifest with some special properties. You can also add plugins to Puppet modules. Plugins help provide new facts, functions and data that can be made a part of the Puppet catalog.

A useful plugin for a module is one that provides custom facts to extend the built-in facts provided by Facter. The Puppet agent gets the node special facts (custom facts) from the module and uses them when the catalog is being built.

Once you build a custom module, you set it up for testing. Testing makes your modules robust by testing all the input parameters, files, packages, choices based on various operating systems and for both valid and invalid inputs.

The final step in building a custom module is to package it. At this point you can also upload a module to Puppet Forge and/or publish the module on GitHub, if you so desire, so as to share your modules with others.



In order to get the Approved Status from Puppet abs, you need to satisfy stringent approval requirements that include extensive unit and acceptance testing,

Using a Puppet Server

Puppet Server provides a central infrastructure for the Puppet agents. It stores the configuration data and distributes the code to the nodes. It also stores backups of all the Puppet-managed files. Finally it evaluates and builds catalogs for the nodes.

There are two types of Puppet servers:

- Puppet master: This is the original Puppet server, and is available with a basic Puppet installation.
- Puppet Server: This is the newer Puppet Server, and is a separate product. This server is a high performance, scalable alternative to the build-in Puppet master.



Puppet server could mean either the Puppet Server (separate product) or the Puppet master (built-in, invoked with the puppet master command).

HOW THE PUPPET SERVER CHANGES THE BUILD PROCESS

In order to fully appreciate the role of a Puppet Server, you need to learn how Puppet builds the resource catalog. A node is any device or system that runs a Puppet agent or a system that's managed by an agent. The Puppet agent evaluates and applies Puppet resources on a node.

When you issue the puppet apply command, Puppet gathers the node data (runs Facter among other things), builds a catalog of resources for the node and evaluates the catalog. Evaluation determines which resources in the catalog are to be applied on a node. Puppet will create a dependency graph based on the automatic dependencies and ordering attributes. Finally it compares the resource state and makes appropriate changes in the states to bring the resource in compliance with the configuration policy.

A Puppet server offers several benefits over what you can do with the simple puppet apply command which doesn't involve a Puppet server. These benefits include the following:

- The server can validate the lists of facts transmitted by a node to the server, using the concept of “trusted facts”.
- Using what are called trusted server facts, you can prevent the clients from overriding certain facts.



When you use a server, client –provided facts are referred to as just facts and server validated facts, as server facts.

How the Puppet master and Puppet Agent Work Together

When you use a Puppet server (Puppet master or Puppet Server), the server is responsible for evaluating the data and compiling a catalog for the node.

When you use a server, it performs the core processing and policy evaluation instead of the agent performing these tasks. Here's how the agent and the server interact:

- The Puppet agent submits the environment and facts to the Puppet server
- The Puppet server builds the Puppet catalog
- The Puppet agent applies the catalog's resources in the node where the agent runs
- By using a server, you don't need to distribute code and data to each of the nodes in your environment. This of course means a more secure deployment environment and also reduce network access requirements for data sources.

Puppet master or Puppet Server?

Puppet master comes with the built in WEBrick application server that accepts only two concurrent connections. To scale, you need to use a Rack application server such as Phusion Passenger.

Puppet Server is newer and extends the services provided by the older Puppet master. Puppet Server is faster and offers more visibility into the server processes.

Puppet master is going to go away in Puppet 5, the next release of Puppet. So, if you're new to Puppet, use the Puppet Server, and if you're currently using Puppet master, plan your migration to Puppet Server.

Using MCollective to Manage Puppet

You can use MCollective, the orchestration tool I discussed in Chapter 6, to control Puppet agents running on various nodes. Using Mcollective, you can:

- Start and stop the Puppet agents
- Make changes to the nodes using Puppet resources
- Control the concurrency of Puppet runs when you run Puppet across a set of nodes

Server Provisioning and Orchestration with Salt

Salt is a remote execution framework and CM system, and also serves as a server provisioning and orchestration tool that is increasingly becoming popular. Salt helps you install and configure operating systems, as well as deploy and configure custom applications. SaltStack is the company behind Salt.



The name Salt was chosen by Salt's creator to denote the fact that salt makes everything (including system management) better!

Salt offers several benefits:

- It's easy to set up and manage even in large data centers
- Offers flexibility by providing multiple ways for master and client nodes to communicate and synchronize
- Salt uses the message broker ZeroMQ and faster network traffic through msgpack for fast communications (msgpack is a small and efficient binary serialization format similar to JSON)
- Can perform high speed parallel execution for rapid deployments

At its core, Salt is a sophisticated and fast remote execution framework that's highly extensible. Salt is based on Python and lets you take advantage of a large number of Python utilities. Extensibility is a strong feature of Salt – you can easily customize various aspects of Salt. You use Jinja, a powerful templating engine, to extend Salt.

Architecture of Salt – the Basics

In Salt, all your hosts are called minions, and the central server that performs actions on the minions is called a master. When the master daemon starts up, it provides a socket to which minions can bind and wait for commands (events). Minions are identified by unique IDs such as hostnames.

Salt Topologies

You can setup Salt with a single master and multiple minions. In a large environment, you can group the minions and have each separate group talk to an intermediate host called a syndication master. The syndication master will act as a proxy for the master in this architecture. Geographically dispersed set of hosts benefit from this architecture.

It's also possible to cut out the master altogether, and use Salt directly with a master-less minion setup.

Salt offers two basic ways for you to use it: execution modules and states. Let's review these in the following sections.

The Remote Execution Framework

Salt uses a remote execution framework that relies on the publish-subscribe messaging model. This remote execution framework is the foundation upon which all the functionality of Salt is built upon.

Once both the minions and the master daemons are up, they communicate via the ZeroMQ data bus, which

When the master publishes commands and places them on the ZeroMQ bus, the minions identify the command via the target.

The minions send back data to the master, including information indicating whether the command(s) succeeded.

Salt States

Salt states are ways to configure your infrastructure. Using salt states, you declare the desired configuration for a host. You use the YAML format to describe the configuration, such as the list of packages you want installed on a server. You can say that a salt state is a way to define a sort of a template to set up hosts.

Communications via Encrypted Channels

When the minions first connect to the server, they are validated, and the minions and the master exchange encryption keys, with the master storing the public keys of all the minions. Further communications between the master and the minions will be thus, secure.

Examples of Salt usage

You can use Salt as a remote execution framework or for configuration management and orchestration. When using Salt as a remote execution framework, one of the most common tasks is to install or upgrade packages across an entire infrastructure. The following simple salt command will install the Apache web server on all hosts in an environment.

```
salt '*' pkg.install apache
```

In this example I specify a target (* = all hosts) and install the package directly through the execution module.

You can use Salt to perform CM tasks by describing the desired configuration in the form of states, which are stored in YAML based files called SLS (salt states). For

example, you can specify the state to manage the main index file of the Apache web server as follows:

```
webserver_main_index_file:
  file.managed:
    - name: /var/www/index.html
    - source: salt://webserver/main.html
```

The state here shows that the file is managed by Salt and that file is then stored on the Salt master.

In this example, I used a file state module, Salt comes with several built in salt modules to help you easily describe the configuration of a host.

Unlike in the first example where I specified the target hosts, in the latter example I didn't specify any target minions. The state system uses an abstraction called a top file that specifies the hosts. The state specifies the recipe for how to configure the host and the top file determines the hosts where that recipe must be used.

Grains and Pillars

Salt uses two types of data:

- Grains are data coming from the minions (for example, the OS) to the master. Grains are stored on the minions. Here's an example that shows a list of all grains on a minion, with the help of the grains.ls command:

```
$ sudo salt master.example grains.ls
master.example:
  - SSDs
  - cpu_flags
  - cpu_model
  - cpuarch
  - domain
  - fqdn
  - fqdn_ip4
  - fqdn_ip6
  - gpus
  - host
```

- Pillar data is data about the minions stored on the master, and is available individually to the minions.

Grains mostly deal with the host metadata such as the number of CPUs on a host. You can set your own metadata for a host by using grains. Pillar data on the other hand is data required by a host to do something, such as the password for a database.



Minions inform the master about its grains and ask the master for its pillar data

Basic Salt Commands

You use command-line tools to interact with Salt. There are several Salt commands, of course, but the following four commands are the most commonly employed:

- salt
- salt-key
- salt-run
- salt-call

The first three commands are run on the master, and the salt-call command is run on the minions.

The salt command

The command salt is heavily used on a day to day basis. This command will run a command on multiple hosts. The following salt command runs the test.ping command on the minion named minion1.example. This command runs the argument, test.ping, which is called an execution module, on a single server. Here, I'm using Salt as a remote execution tool.

```
$ sudo salt minion1.example test.ping
minion1.example:
    True
```

Salt comes with a bunch of useful execution modules such as those for package management, user management, and iptable editing. In addition to these pre-built modules, you can write your own execution modules in Python. Your custom modules have the advantage that they use Salt-specific data structures, thus making them far more powerful than standalone system administration Python scripts.

The salt-key command

During the initial handshake between the master and the minions, they need to establish a trust relationship. When a minion first connects, the trust relationship is in the unaccepted state. Using the salt-key command, you can change the state of the trust to accepted, as shown here:

```
$ sudo salt-key -accept—all
```


The salt-call command

The salt command runs the remote execution modules on the minions. You can use the salt-call command to run an execution module directly on a specific server. When run in the local-only mode, this command works only on the local hosts, and is a big help during debugging.



Its simplicity has led a lot of folks to call Ansible a for-loop over SSH scripts, as a way of knocking it for its simplicity– the reality is far from it – Ansible is a quite powerful tool.

The salt-run command

The salt command executes the execution module asynchronously on every minion. If you wish to run a command sequentially across minions, you can use a class of modules called runners. The code runs only on the master.

The salt-run command is useful for coordination commands across a set of hosts. You don't need to specify the target set of minions. Here's an example:

```
$ sudo salt-run manage.up
master.example
minion1.example
minion2.example
minion3.example
minion4.example
```

The manage.up module uses the test.ping execution module you've seen earlier, to check the health of a minion.

Execution Modules

Salt comes with a bunch of execution modules (and functions) that help you perform routine administration tasks. All these modules are written in Python.

You can use the sys.list_modules function to check out the available modules on a minion:

```
$ sudo salt-call sys.list_modules
local:
  - acl
  - aliases
  <snip>
```

Here's a brief description of the main Salt modules:

Sys: provide information and documentation about modules

- **cmd:** A powerful module, the `cmd` module contains functions that run arbitrary commands. Here's an example:

```
$ sudo salt \* cmd.run 'grep root /etc/passwd'
minion3.example:
  root:x:0:0:root:/root:/bin/bash
minion2.example:
  root:x:0:0:root:/root:/bin/bash
  operator:x:11:0:operator:/root:/sbin/nologin
<snip>
```

- **pkg:** This module manages packages. It abstracts several package managers (`yum`, `apt`, etc.) into what Salt calls a virtual module. Here's an example that shows what version of the `nginx` web server would be installed on each of 5 minions (master is also a minion!):

```
$ sudo salt \* pkg.list_pkgs
minion2.example:
  -----
  MAKEDEV:
    3.24-6.el6
<snip>
```

- **user:** The `user` module helps manage users. Here's an example where we add a user with the `user.add` function:

```
$ sudo salt minion2.example user.add wilma
minion2.example:
  True
```

- **saltutil:** This module contains functions that help you manage Salt itself, rather than the infrastructure. It helps you manage jobs, keys, etc.

Salt States and Configuration Management

You describe the desired configuration of a server, database, application, or web server with a state (SLS) file. States are in essence a recipe for a desired state, and you combine states in order to build complicated states.

State modules and Execution Modules

A state is a way for you to describe how you want a host to look.

Following is a simple SLS file that adds a new user named `sam`:

```
user_sam:
  user.present:
    - name: sam
```

```
- fullname: Sam Alapati
- uid: 2001
- home: /home/sam
```

In this example, I do two things:

- Use a state module named user
- User the function named present

A state module looks similar to the execution module I showed earlier. Earlier, when using the `user.add` execution function, I added a user. Now, I use the `user.present` state function to ensure the user exists.

A state module goes beyond the execution functions – in this case, it checks whether the user exists and calls the `user.add` function only if I need to add the users. If the desired state already exists, there's no need to call the execution module (and functions) – thus, state module calls are always idempotent. You can run the same state call multiple calls and the state remains the same if the real and desired states are identical.



You can extend Salt with custom states, modules and grains.

Highstate and the Top File

In order to achieve a desired configuration, you must apply several states together, in a particular combination. When you combine multiple states into complex “highstates”, you use a state file called the top file (usually named `top.sls`). You start by defining individual states such as adding users, or installing packages, and put it all together with the help of the top file.

Often you want one action to precede another – you can order the way in which Salt implements your described configuration by using the `require` declaration, which forces the state you specify to execute first. The following example shows how to ensure that the Nginx web server package exists before the Nginx instance is started.

```
$ sudo salt minion1\* state.highstate
minion1.example:
-----
      ID: packages_vim
      Function: pkg.installed
<snip>
-----
      ID: user_www
      Function: user.present
<snip>
-----
      ID: roles_webserver_packages
```

```

    Function: pkg.installed
<snip>
-----
        ID: roles_webserver_start
    Function: service.running
        Name: nginx
        Result: True
        Comment: Started Service nginx
        Started: 02:14:13.267952
        Duration: 371.647 ms
        Changes:
            -----
            nginx:
                True

Summary
-----
Succeeded: 4 (changed=1)
Failed:    0
-----
Total states run:    4

```

Using Ansible for CM, Orchestration, and Provisioning

Ansible is a configuration management tool, like the other tools you’ve learned about in this chapter – Chef, Puppet and Salt. In addition to CM tasks, you can use Ansible for a bunch of other purposes:

- You can use Ansible for deploying home-grown applications as services. In this regard, Ansible works similar to Capistrano and Fabric (explained in Chapter 6), two popular open source deployment tools.



The name Ansible refers to a fictional sci-fi communication device that can transfer information at speeds faster than the speed of light.

- The orchestration of deployments means you have to perform operations in a specific order, such as bringing up a database before you start up a web server. Ansible is quite good as an orchestration tool.
- Ansible also serves as an excellent provisioning tool. Provisioning is the art of spinning up new VM instances. Ansible can provision servers in EC2, Google Compute Engine, Rackspace, and any other clouds based on the OpenStack APIs.

How Ansible is Different from the other CM Tools

Ansible contains several good features that set it apart from the other CM tools:

- The syntax is easy to read (YAML based playbooks).
- Ansible playbooks are simple, whether you write them for one remote server or a thousand.
- Ansible provides a thin layer of abstraction which enables you to use the same scripts on hosts running various operating systems
- You don't need to install agents or other software on the hosts
- Ansible ships with many powerful declarative modules that help you perform various server administration tasks.
- Whereas Chef and Puppet are pull based (agents periodically pull down config information from the master), Ansible is push-based. When you create or modify a playbook, you run the playbook (with the `ansible-playbook` command) and Ansible connects to the servers and executes the modules.

As with Salt, Ansible uses the YAML file format and the Jinja2 templating language, so it helps to pick up the rudiments of these two technologies.

The Guts of Ansible

To make Ansible do anything, you script the actions in a YAML based file called a playbook. Playbooks describe the hosts you want to configure and a set of task to perform on those hosts.

Let's say you want to install an Nginx web server on some Linux hosts. You need to perform the following set of tasks to do this:

- Install the Nginx binaries
- Generate the nginx config file
- Copy a security certificate over
- Start the new Nginx service

In order to do this with Ansible, I execute a playbook I've named `webservers.yml`, using the `ansible-playbook` command:

```
$ ansible-playbook webservers.yml
```



In Ansible, a playbook is a configuration management script.

My playbook (webservers.yml) for creating and starting the Nginx web server looks like the following:

```
name: Configure webserver with nginx
hosts: webservers
sudo: True
tasks:
  - name: install nginx
    apt: name=nginx update_cache=yes

  - name: copy nginx config file
    copy: src=files/nginx.conf dest=/etc/nginx/sites-available/default

  - name: enable configuration
    file: >
      dest=/etc/nginx/sites-enabled/default
      src=/etc/nginx/sites-available/default
      state=link

  - name: restart nginx
    service: name=nginx state=restarted
```

In this Ansible playbook, the first task is to install the nginx binaries. So, Ansible will log into all the servers you specify and installs the nginx package on those servers. It's important to understand that Ansible does this step parallelly on each server that you specify.

Once it completes the installation of the nginx binaries, Ansible will step through the rest of the steps in the playbook sequentially, waiting to start the next step until the preceding step has completed, again in a parallel fashion.

Here's the output when I run the webservers.yml playbook:

```
$ ansible-playbook webservers.yml
PLAY [Configure webserver with nginx] *****

GATHERING FACTS *****
ok: [testserver]
TASK: [install nginx] *****
changed: [testserver]
TASK: [copy nginx config file] *****
changed: [testserver]
TASK: [enable configuration] *****
ok: [testserver]
TASK: [copy index.html] *****
changed: [testserver]
```

```
TASK: [restart nginx] *****
changed: [testserver]
PLAY RECAP *****
testserver                : ok=6    changed=4    unreachable=0    failed=0
```

Note the GATHERING FACTS step at the very beginning of Ansible's play execution – that's Ansible collecting information about the server(s) it connects to, such as the operating system, IP addresses, etc. You can turn off fact gathering if you want to.

How Playbooks Work

Playbooks are the very essence of Ansible, since it's through playbooks that you configure your infrastructure. Let's explore the content of a playbook in the following sections.

Plays

A playbook is a list of plays. Plays are the way you perform tasks on the hosts. Each play must contain:

- A set of hosts that you want to configure
- A set of tasks that you want to execute on the hosts

Tasks

Each play in a playbook can have one or more tasks that the play performs. The following play installs nginx:

```
install nginx
apt: name=nginx update_cache=yes
```

====Roles Help Devise Complex Playbooks Ansible uses roles as a way to simplify the creation of complex playbooks. A role such as database role is something you assign to a host that server as the host for a database server.

Modules

Ansible ships with several modules, which are scripts designed to perform specific actions on hosts. For example, the copy module copies files from the local server to the hosts, the file module can set the attributes of a file or directory, and the service module starts/stops services.

You can associate each task with only one module.

Inventories Describe your Infrastructure

An inventory is the set of hosts that Ansible is aware of. By default, you use simple text files called inventory files to list the hostnames in your environments.

Grouping of Hosts

Ansible automatically defines a group that includes all the hosts (all or *). You can also use inventory files in the .ini format to define your own groups. An .ini file groups together related entries together in a configuration file.

Let's say you want to deploy a Django web application that needs to support an nginx web server, a RabbitMQ message queue, and a Postgres database. You may have multiple web servers and database servers in production. You may also have a staging environment. You may use VMs for production and staging, and a Vagrant environment for local testing.

Here's how you'd define your Ansible inventory file, grouping servers by environment and by function:

```
[production]
server1.example.com
server2.example.com
server3.example.com
server4.example.com
server5.example.com
server6.example.com
server7.example.com
server8.example.com
server9.example.com
server10.example.com

[staging]
server11.example.com
server12.example.com

[vagrant]
vagrant1 ansible_ssh_host=127.0.0.1 ansible_ssh_port=2222
vagrant2 ansible_ssh_host=127.0.0.1 ansible_ssh_port=2200
vagrant3 ansible_ssh_host=127.0.0.1 ansible_ssh_port=2201

[lb]
delaware.example.com

[web]finished catalog
server2.example.com
server4.example.com
server5.example.com
server11.example.com
vagrant1
```



```
[rabbitmq]
server8.example.com
server12.example.com
vagrant3
```

```
[db]
server9.example.com
server10.example.com
server12.example.com
```

Using Numerical or Alphabetical Patterns to list the Hosts

Instead of using a complex inventory file with a large number of hosts, you can use a numeric pattern to make things more scalable. Assuming you have 3 web servers named web1.example.com, web2.example.com and so on, you can use the numbering strategy in your inventory file:

```
[web]
web[1:20].example.com
Alternatively, you can list the servers using an alphabetical range:
[web]
web-[a-t].example.com
```

Version Control and Source Code Management

Version control is how you can ensure that you can recall older versions of a file when you need them. A version control system records changes to one or more files and lets you easily recall any older version you want.

In the previous chapter, you learned about the concept of Infrastructure as Code. A Version Control system can act as the foundation for Infrastructure as Code. Everything you do in your infrastructure, you can capture it in scripts, configuration files and definition files and check them into a VC such as Git.

In this chapter, I explain the essentials of Git, which is the most popular version control system today.

In addition to Git itself, I also explain how to use GitHub, a popular Git “repository”. GitHub is a website to which you can upload your Git repository, the goal being easier collaboration with other people working on the same project.

Version Control - An Introduction

Version control (VC) is mostly used for development proposes. However, there’s no reason why you can’t use it to maintain infrastructure related code as well. VC allows you to revert files back to older versions compare changes between versions, and find out who modified a piece of code. In a nutshell, you can recover quickly when you make mistakes, while paying very little in terms of overhead.

Types of Version Control Systems

Version Control systems come in three broad flavors:

- Local version control systems: These are purely local in the sense that they aren't shared with others.
- Centralized version control systems: These systems maintain a single “central” copy of the project, and developers commit their changes to this centrally stored copy.
- Distributed version control systems: Every user “clones” a copy of a repository and has the full history of the project on their own local server.



Pulling is the act of getting changes from a code repository. Pushing is the act of making your own changes to a repository. Git is a distributed version control system. This means that all the folks working with a Git based project have a copy of the full history of the project – and not just the latest state of the files.

Revision Control and Infrastructure Management

One might wonder how revision control, which has traditionally been used for source code management, helps with infrastructure management. System administrators can use revision control in exactly the same way as developers do!

For example, you can use Git for storing the BIND configuration you use in your environment. Similarly, you can store the configuration details for software such as Postfix, iptables, Hadoop and Mysql, etc in revision control.



Configuration management tools such as Puppet and Chef integrate well with Git

You can change all your infrastructure related scripts and configuration files in a VCS. When a system administrator wants to make changes, he/she checks the necessary files out of the VCS, makes their changes, and commits them back into the VCS.

These committed changes are then available for application to the infrastructure. If you're using a change management pipeline (see Chapter 11), the changes are automatically applied to and tested in a test environment. If different members of the admin team are making incompatible configuration changes, the VCS will force you to reconcile the differences by showing up the changes as being in conflict with each other.

In this context, it's interesting to know that a lot of open source software projects such as the Linux kernel, OpenStack, and Yum, store their source code in Git.

Why Use Source Control?

Source control systems protect code from being messed up – either due to human errors or catastrophic events. When working within a team, source control isn't an option – you must use it.

Mercurial and Git are the two most popular open source systems available for source control management. Subversion is an older system that ran on a single server. Mercurial and Git, being distributed VC systems, use multiple copies of code repositories.



Git is a creation of the same fellow that wrote the Linux kernel - Linux Torvalds.

Our First Foray into the World of Git

In this section, I get your feet wet, by providing a whirlwind tour through basic Git operations. This way, you get a good flavor of Git commands and Git operations and what they do. I'll follow this section up with a brief explanation of Git's fundamental features: branching, committing, and cloning.

Installing Git is quite easy, and you can look it up. Let's say you are in an empty directory such as /home/testdir. From this directory, create a local git repository with the git init command:

```
$ git init
Initialized empty Git repository in /home/salapati/testdir/.git/
$
Let's create a simple Python file named test.py with the following content:
print('Hello Git!')
Let's add this new Python file to the git repository I initialized earlier.
$ git add test.py
Check the status of your new repository with the git status command:
$ git status
On branch master
Initial commit
Changes to be committed:
  (use "git rm --cached <file>..." to unstage)
    new file:   test.py
```

The output of the git status command reveals that there are “changes to be committed”. Our new file test.py is sure part of the local repository but not yet committed. Why don't we commit the changes and see what happens?

```
$ git commit -m "simple print program"
[master (root-commit) 52d60d7] my first commit
```

```
1 file changed, 1 insertion(+)
create mode 100644 test.py
```

Check the current status of the git local repository.

```
$ git status
On branch master
nothing to commit, working directory clean
```

OK! My test.py file is safely tucked away now in my local repository. Just for the heck of it, I make a change in the test.py file and see what git says.

```
$ git status

On branch master
Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git checkout -- <file>..." to discard changes in working directory)

    modified:   test.py
```

```
no changes added to commit (use "git add" and/or "git commit -a")
```

I can execute the diff command to see the differences between the original test.py file and the modified version of that file. The git diff command shows the changes since the last commit of the test.py file:

```
$ git diff
diff --git a/test.py b/test.py
index 76b8c39..62782b2 100644
--- a/test.py
+++ b/test.py
@@ -1,1 @@
-print('Oops')
+print('Ops!')
```

In order to commit the changes I've made to the test.py file, I need to first add the modified test.py file:

```
$ git add test.py
```

Once I add the new test.py file, I commit the change:

```
$ git commit -m "my first change"
[master e1e11ec] my first change
1 file changed, 1 insertion(+), 1 deletion(-)
```

The git log command will let you see all the changes you've made to the test.py file.

```
git log test.py
commit e1e11ecf802ae1a78debe6193c552dcd15ca160a
Author: William Lubanovic <bill@madscheme.com>
Date:   Tue May 13 23:34:59 2016 -0500
    change the print string
commit 52d60d76594a62299f6fd561b2446c8b1227cfe1
```

Author: William Lubanovic <bill@madscheme.com>

Date: Tue May 13 23:26:16 2016 -0500

simple print program

```
git clone https://github.com/craig5/salt-essentials-utils
Cloning into 'salt-essentials-utils'...
remote: Counting objects: 523, done.
remote: Total 523 (delta 0), reused 0 (delta 0), pack-reused 523
Receiving objects: 100% (523/523), 83.00 KiB | 0 bytes/s, done.
Resolving deltas: 100% (167/167), done.
Checking connectivity... done.
```

```
SG0221771@D4T11J72 ~/salt
```

```
$
```

Introduction to Git

Git is a popular version control (revision control) software. Modern apps are created by large teams of developers working together in hundreds of files simultaneously. A source control tool such as Git makes it easy for groups to work on text files used in creating software.

What Sets Git Apart from other CVS Systems?

Following are the main features that set git apart from other similar systems:

- **Branching:** Branching is probably git's most powerful feature and this alone would be a good reason to choose git over other systems. The git branching model has revolutionize collaboration among developers, who can create, merge and share branches easily with others working with them.
- **Data integrity:** and introduces all changes as revisions which can be inspected provides strong integrity by tracking gall files and directories introduces all changes as revisions that can be easily inspected. Data integrity ensures that there's no chance for a change to go unnoticed by the system.
- **Distributed model:** Unlike centralized VC systems that use a client/server model, git is a distributed VC system. This means that all repositories are fully functional and can send and receive content.
- **Locality:** You can use git locally which means that you aren't reliant on the network to execute most git commands. Most git commands such as commit, branch and merge don't involve any data transfer and can be performed locally where you run git.

Installing Git is pretty trivial (use yum or apt to install git) and if you've Cygwin installed, you already have git. The more important thing to know is that git expects you to configure it before you can use it. Before you can commit anything within a git

repository, you must configure the user.name and user.mail settings. Here's how I did it:

```
$ git config --global user.name "Sam Alapati". Let's create a text file with some content and save it
$ git config --global user.email sam.alapati@gmail.com
```

Now that some preliminaries are out of the way, let's start exploring git and see how you perform various VC tasks with it.

Creating Local Repositories with Git

Git is a distributed version control system. In a traditional centralized version control system such as Subversion or CVS, the repository isn't local – it's stored on a remote server. When you commit your changes, they're sent over the network and committed to the VC system so others can view the changes.

You can choose to send your commits from another repository, store all your commits, branches and

Creating a Repository

You create a new Git repository with the git init command, as shown here:

```
$ mkdir my-git
$ cd my-git
$ git init May24
Initialized empty Git repository in /my-git/May24/.git/
$
```

Note that May24 is the working directory for this project.

Checking the Repository Status

When you first create a repository, of course it's empty! You can check the status of the repository with the git status command to see what git shows you:

```
$ cd May24
$ git status
On branch master

Initial commit

nothing to commit (create/copy files and use "git add" to track)
$
```

In the comments, Initial commit means that the repository is ready for its first commit. Let's create a text file and save it:


```
$ cat test.txt
This is my first Git file.
I'm going to commit it now!`
```

Once I save my file, it becomes a part of the Git repository. Now, when I issue the git status command, I see something – the file I created, test.txt, is now part of the new repository! All changes made to that file will be henceforth tracked and all the versions of the file will be saved as well.

```
$ git status -s
?? test.txt
$
```

The `-s` flag for the git status command prints a shortened output of the git status command.



A clean repository has no files that need to be committed. A dirty repository is one where there are one or more uncommitted files.

COMMITTING A FILE

Is my new file committed so I won't lose it? Well, the two question marks (??), are Git's way of telling you that the test.txt file isn't yet committed. The repository is dirty, meaning it contains files that were modified by not committed.

In order to commit the test.txt file, I use the following two commands:

```
$ git add -A

$ git commit -m "First Revision [test.txt]"
[master (root-commit) 4a6fdd7] First Revision [test.txt]
 1 file changed, 2 insertions(+)
 create mode 100644 test.txt

$
```

Here's what the two commands do:

- **add:** Adds file contents to the index (updates what will be committed)
- **commit:** Records changes to the repository

Committing the file means there are no more uncommitted files lying around in the working directory, which leads to the git repository being assigned the status of clean (as opposed to dirty). If you issue the git status command now, it confirms that the repository is clean:

```
$ git status
On branch master
nothing to commit, working directory clean

$
```

Checking the Logs and the Status of a Repository

When I check the log of the repository, it shows that I've one revision:

```
$ git log
commit 4a6fdd7bf247db95f7d8d67085d5dfffc02eba3ea
Author: Sam Alapati <sam.alapati@gmail.com>
Date: Tue May 24 05:53:50 2016 -0500

    First Revision [test.txt]
```

```
$
```

The git log command shows commit logs.

Now let me make a change to the test.txt file by adding a third line to it. I check the status afterwards:

```
$ git status
On branch master
Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git checkout -- <file>..." to discard changes in working directory)

        modified:   test.txt

no changes added to commit (use "git add" and/or "git commit -a")

$
```

I decide this time to issue just the git commit command (without preceding it with the git add command) to commit my git changes:

```
$ git commit -a
[master dfce066] Added a third line now
1 file changed, 1 insertion(+)

$
```

I can issue the git status or the git commit command to check if the directory is clean.

```
$ git commit -a
On branch master
nothing to commit, working directory clean

$
```

Git lets you restore a project to the very first revision, or any other revisions you've committed since then.

Staging and Committing Files

Git never automatically register any changes you make in the working directory. In the previous section, I showed how you need to issue the `git add -A` and `git commit -m` commands to create a new revision.

Although it seemed like this part of commands constituted a single operation, they actually take care of two distinct steps:

- The `git add` command selects the files for the revision. These files are called staging files, and a list of the staged files is stored in the index or staging area, in the file `.git/index`.
- The `git commit` command creates the revision with the files that have been selected.

All files you store in a Git repository fall into three groups:

- **Untracked:** When you first create a file, it's an untracked file, meaning git doesn't store the file in its repository and also doesn't track the file's contents. This is the reason you see the two question marks (`??`), with the first `?` indicating that the file is unknown in the staging area and the second `?` indicating that the file is unknown in the working directory as well.
- **Staged:** These are files that are in the staging area and whose list is stored in the `.git/index` file. When you issue the command `git commit -m`, all the changes in the staging area go into the revision.
- **Unmodified:** These are files that are stored in a repository following a commit. As you saw earlier, the `git status` command doesn't show any information about unmodified files.

Committing Multiple Files

In the examples I showed earlier, I staged and committed a single file at a time. Let's see how you can stage and commit multiple files at a time.

In this example, I have several files such as `new.txt`, `deleted.txt`, `modified.txt`, and `old.txt`. I delete the file `deleted.txt` and rename the file `old.txt` to `revised.txt`. When I run the `git status` command, I see the following:

```
$ git status
On branch master
Untracked files:
```

(use "git add <file>..." to include in what will be committed)

```
modified.txt
new.txt
revised.txt
```

nothing added to commit but untracked files present (use "git add" to track)

\$

Git recognizes that there are “untracked” files present in the working directory. It also suggests running git add to track the files. I run the git add -A command to stage all my changes in one fell swoop:

```
$ git add -A
```

```
$ git status
```

```
On branch master
```

```
Changes to be committed:
```

```
(use "git reset HEAD <file>..." to unstage)
```

```
new file:   modified.txt
new file:   new.txt
new file:   revised.txt
```

\$

Finally, I commit all of my changes:

```
$ git commit -m "Commit all my changes"
```

```
[master e2df127] Commit all my changes
```

```
3 files changed, 3 insertions(+)
```

```
create mode 100644 modified.txt
```

```
create mode 100644 new.txt
```

```
create mode 100644 revised.txt
```

\$

When I run the git status command, there's no output since the working directory is clean and there's nothing to commit: \$ git status -s

Cloning Repositories

Sometimes you may need to copy an entire repository, with the copy containing all the branches stored in the original repository. You can use the git clone command to clone repositories:

```
$ ls
```

```
May24
```

```
$ git clone May24 May25
```

```
Cloning into 'May25'...
```

done.

```
$ ls
```

```
May24 May25
```

In this example I used the git clone command to clone the May25 repository named May24 from the original repository named May24.

Remote Git

While it's possible to work with Git as a pure local version control system, and never have to share your changes, you'll often want to share changes with others, usually as part of collaborating with a team. You use a remote repository for sharing changes with others.

You can use various types of remote repositories, including GitHub. Since multiple repositories can be connected to your local repository, the remote repositories aren't named, but if you have a single remote repository, it's usually named origin.

You use the git push command to push changes from a local repository to a remote repository.

The Git Branching Model

One of the things that makes Git far superior to alternative VCS systems is its branching model. The power offered by just this one capability may justify a switch from an alternative VC system to Git.

New branch means creating a copy of a project. You can work with this copy without affecting the project. If your changes don't work out you can simply abandon the branch and return to the original, which is always called the master branch. If your changes work out, on the other hand. You can easily incorporate the changes into the master. You also can revert to the state of the project before the merge, at a later point, should you so desire.



By convention, users expect to find the latest code for a project in its master branch.

A branch is an independent line of development – It's a pointer to an arbitrary commit in the database. You can preserve any point in your project's history by creating a branch. Branches are independent of each other, and modifying one branch doesn't affect the other branches. Your commits in a branch don't affect other branches. Only you can modify the branches you create – Git won't ever modify them.

Creating and switching branches lets you fork your project into independent development lines. For example, one may create a branch to add new features to your application, or to implement bug fixes. The best practice when performing work on a new topic is to perform it in a dedicated branch.

Creating Branches

By default, when you initialize a new repository (as I showed earlier), the repository will just have a single branch named master:

```
$ git branch
* master
```

The asterisk denotes the current branch where you are. If you commit anything now, it'll be part of this current branch.



You can use the `git branch` command to list, create and delete branches.

You can create a new branch with the `git branch` command, as shown in the following example, where I create the new branch named `test`.

```
$ git branch test

* master
  test
```

As you can tell from the asterisk, your current branch is still `master`. You can switch to the new branch `test` by executing the `git checkout` command:

```
$ git checkout test
M      virtual-machines/Vagrantfile
Switched to branch 'test'
Now, when I issue the git branch command, I can see that you're now on the new branch test.
$ git branch
  master
* test

$
```

You can also create a new branch that points to an existing branch with the following command.

```
$ git checkout -b branch2 branch1
```

The new branch named `branch2` will contain the same revision as the existing branch named `branch1`.

Merging Branches

As I mentioned earlier, you create branches to spawn alternate development lines, such as adding functionality to your application. Once you complete the work on the new features(s), you'll want to incorporate your work into the main line of development. This process is called the merging of branches.

Here's an example showing how to merge branches:

```
$ git merge test12 test2 test3
warning: old-style 'git merge <msg> HEAD <commit>' is deprecated.
Already up-to-date.
```

Resolving Conflicts

In real life projects, there are often conflicts when you merge branches. A conflict during the merging of branches occurs when the branches include different versions of the same line in a file. In these cases, git won't automatically merge the multiple versions – it simply has no idea what's the “correct” or “best” version. You need to manually choose between the different versions.

You resolve conflicts when staging a file. The git add command is what you use to stage files, as I showed in the earlier part of this chapter.

Setting up Git on a Server – Hosting Git Repositories

As useful as Git is for a single user, Git's power really comes into play when you use it to manage a project with a large number of contributors. Colleagues who are dispersed all over the world can work independently on a project and collaborate through pushing their code to a Git repository.

You can host a Git Server yourself, or use a third-party service such as GitHub or Bitbucket. In this section, I explain the options for hosting your own Git server.

Multiple Protocols

The way in which you host your Git repository depends on which protocol it uses to transfer data. Git can use one of four protocols to transfer data: local, ssh, HTTP, and Git.

THE LOCAL PROTOCOL

When you use the local protocol, the remote repository is in another directory on a shared filesystem (NFS mount), or on the same server. Using a shared mounted file system, you can clone, push to and pull from a local repository. In order to clone a local repository, for example, just use the path to the repository as the URL:

```
$ git clone /opt/git/project.git
```

In order to add a local repository to an existing project, you do this:

```
$ git remote add local_proj /opt/git/project.git
```

Once you do this, you can pull from the remote repository as if you were doing so over the network.

In the local option, you don't need to set up a server explicitly since users access the repositories over the network. Existing file permissions and network access permissions control the access to the networked filesystem. To fetch (get a remote repository) and clone repositories, users need read permissions and to push (copy a local repository to a remote) they need write permissions.

The drawbacks of the local protocol are that it's not easy to share up shared access to a networked filesystem, and that it doesn't protect the repository against accidental damage.

The SSH Protocol

The SSH protocol is widely used for self-hosting Git repositories. SSH is an authenticated and encrypted network protocol that's pretty easy to set up. However, SSH doesn't [permit anonymous access to Git repositories.

You can grant access to the Git repositories over SSH in two different ways: create separate accounts on the server for all the clients who need to access the repository or create a single shell account to access the Git repositories and use public-key login to authenticate the users who want to access the repositories.

You can also have the SSH serve authenticated from a central authenticating server such as an LDAP server.

Anonymous Git Protocol

The Git protocol (git://) is commonly used for fast anonymous and unauthenticated read-only access to Git repositories. This protocol is served by a simple TCP daemon listening on a dedicated port.

To set up this the Git protocol you run the Git protocol server, git daemon.

The Git protocol has the drawback that it lacks authentication. By default this protocol is used only for read access (fetching and cloning repositories) and you combine this protocol with SSH or HTTPS for pushing repositories.

SMART HTTP(S) PROTOCOL

Just as the SSH protocol does, using HTTP(S) offers encryption and server authentication. You set up the "smart" HTTP(S) protocols with a CGI server, such as Apache

web server, which performs the authentication. You can set up the authentication so it allows unauthenticated anonymous read-only access and requires authentication for push. Smart HTTP works very similar to the SSH and Git protocols but runs over standard HTTP/S ports. This means you can use basic username/password authentication, instead of having to set up SSH keys.

Setting up a Git Server

You set up a Git service on your own servers using one of the protocols I explained in the previous section. At a simple level, it's straightforward to set up a Git Server. You take a Git repository, create a “bare” version of it and place the repository on a server to which you and others in the team have SSH access. Bingo! You've a Git Server in place that facilitates collaboration on projects with a team.

Once you enable SSH access by users to a server, all you need is to simply get your bare repository somewhere on that server where all the users have read/write access. That's all you have to do to your Git server going.

CREATING A BARE REPOSITORY

In order to set up a Git server, the first thing you need to do is to you need to export an existing Git repository into a new “bare” repository. A bare repository means that it doesn't contain a working directory.

Here's an example that shows how to create a new bare repository by cloning a repository named `my_project`.

```
$ git clone -bare my_project my_project.git
cloning into bare repository 'my_project.git'...
done.
```

A bare repository directory has the `.git` suffix, as in `my_project.git`.

PUTTING THE BARE REPOSITORY ON A SERVER

Once you've your bare repository ready, move it to the server you want to use it as the Git server, and set up the protocols you chose to use for user access to the server.

Let's say your server is named `git.example.com` and you want to store all your Git repositories under the `/opt/git` directory. You set up the new repository on the Git server by copying the bare repository to this server:

```
$ scp -r my_project.git alapati@git.example.com:/opt/git
```

Users with SSH access of the server can clone the new repository by running the following:

```
git clone user@git.example.com:/opt/git/myproject.git
```

If the users have write access to the `/opt/git/my_project_git` directory, they will also have push access to the repository.

Enterprise Deployment of Git

At a very simple level, you can run a single Git server that hosts all your Git repositories. Users will clone repositories from this server and push their own changes to it. If you're trying to serve a geographically distributed clientele, this mode isn't very efficient, because network connections may not be fast or reliable.

You can run multiple Git servers on a per site basis, with clients cloning and pushing to their local site's Git server. While this model is faster and more reliable than the first one, you're tasked with keeping the multiple Git master repositories in sync.

If money isn't a problem, you can go the hosted Git route, and use either GitHub or Bitbucket or Atlassian Stash to maintain all your repositories. Clients will pull directly from the service, and you don't have to maintain the Git server, as GitHub (or an alternative service) takes care of all of that for you!

Tools for Managing Git Repositories

It could be quite tedious to create your own solution to manage Git repositories. Fortunately, there are several excellent third-part solutions that help administrators, as I summarize here:

- **Git Repository management:** Tools such as Gitolite help manage access control, and make it easy to add repositories and control the permissions.
- **Web Interfaces:** You normally use Git from the command line. There are several excellent web interfaces that allow you to view Git repositories through a web browser. Two such tools are the `gitweb` Perl script distributed with `git` and `cgit` for the Linux kernel repositories.
- **Code review and collaboration tools:** Tools such as Gerrit Code Review help developers review each other's changes through a web interface
- **Git management solutions:** These are web based tools for taking care of the administrative side of repository management, such as adding users and creating repositories. A good tool here is GitLab.

Finally, you really don't need to host your own Git repositories. There are several well-known third—part Git hosting options, such as GitHub and Bitbucket.

Using GitHub

Instead of maintaining your own Git servers, you can use a third-party service such as the well-known GitHub to manage all your code repositories. GitHub is a website

where you can upload your Git repository to enable sharing and collaboration with team members and others.

GitHub facilitates collaboration by offering a central location to share the repository, and a web interface to view it. It also offers features such as pull requests, forking, issues and wikis, all of which enhance collaborative work on projects.

Although there are several other git hosts, GitHub is the largest, with over a million repositories. Git is very useful, but it's not always easy to understand and using GitHub eases the pain.

Setting up your Account

The first thing you need to do to use GitHub is to set up your (free) user account, which you can do by going to <https://github.com>. Once you set up your account, you can connect to Git repositories using the https:// protocol, using your credentials for you just set up to authenticate yourself.

The GitHub Workflow

The heart of the GitHub flow is the collaboration workflow based on Pull Requests. Here's the high level summary of the workflows:

- Create a topic branch from master
- Make your improvements (commits)
- Push the branch to your own GitHub project
- Open a GitHub Pull request
- Discuss with the project owner and optionally, commit more changes
- The project owner either merges or closes the Pull Request

Forking Projects

If you don't have a push access to a project and you wish to contribute to it, you can fork the project. This means that GitHub will make a private copy of the project for you and you can then push your changes to it. A fork is nothing but the original project living in your own namespace.

Forking lets projects add users without having to grant them push access. You can fork a project, push to it and submit changes back to the original repository by doing a Pull Request. ===== Creating Pull Requests Let's say you find a good program you want to use on GitHub. Upon viewing the program, you think you can improve it by making some changes to it. Your first move is to fork the original project. You can

then make the changes, test them, and submit it back to the project through a Pull Request, as a proposal for change.

This'll automatically open a discussion thread with the project owner, and both of you can then discuss the change until the owner is satisfied with your changes, at which point the owner can merge it with the original repository.

Continuous Integration, Continuous Delivery, and Continuous Deployment

In earlier chapters, you learned about several automatic infrastructure configuration tools. While these tools are indeed great, even more important from an organization's perspective are tools and techniques that let you deploy error-free applications to serve your customers in a fast manner.

This chapter deals with three important topics that are part of every administrator's work: Continuous Integration (CI), and Continuous Delivery (CD):

- Continuous Integration is the focus on building code, running tests, and making sure that code is in a “clean” state. The goal is the evolution of the latest version of the software into something that you can work with – older builds and changes become less important as you transition through the changes.
- Continuous delivery is the getting ready of tested code for deployment anytime you're ready.
- Continuous deployment is the deployment of a set of changes to one or more environments. You usually have a pipeline with multiple decision points, parallel flows, and go forward and rollback steps. In other words, continuous delivery is a complex workflow with specialized actions such as rolling deployments and contacting the servers.

DevOps is a difficult concept to define, in my opinion. If you apply an expansive definition, just about everything in this book comes under the purview of DevOps. However, DevOps originated in the area of continuous integration and development and most of its sphere of influence lies in this area.

Builds and Deployment

As the number of services and servers grow over time, manual deployments become unwieldy and software releases as well as the testing and integration cycles that they depend on, get longer. Continuous integration, continuous delivery, and continuous deployment lets you automate your entire build, test, and deployment pipelines.

The first step in automating deployments is continuous integration. Anytime developers commit changes to their VCS, automated tests are run on the changes, leading to the early detection of integration issues. CI ensures that the code can be built and packaged for delivery and that the new changes pass testing.

The second step in automating deployments is continuous delivery. While the CI testing is performed on an integration server, the CD pipeline deploys the software to test environments, variously called testing, or staging environments. The key here is that the entire process is automated. Once the code gets through to the end of the CD pipeline, it's ready for production deployment, usually either with the mere click of a button or by issuing a single command.

The final stage of the deployment pipeline is continuous deployment, which involves the testing, building, deployment to production of code without your manual intervention. Every commit to the VCS results in an automatic deployment to production. This is how deployments can be made throughout the day rather than once every month or every quarter, as is the practice in manual deployments. A release cycle under this setup could take as little as 10 minutes from the time code is committed to the time the new code is deployed in production.

Setting up effective CD pipelines requires skills from the development and operations areas. You can certainly build your own CD pipeline, by using the popular open source CI/CD software Jenkins, or the commercial tool Atlassian Bamboo easily control the pipeline.

Remember that for CD to work well, you need not only a tool such as Jenkins, but also a server configuration tool such as Chef or Puppet to build the servers per the configurations you store in your VCS. You may still need to write some custom scripts to specify the order of deployment of the various servers such as web servers, database servers, and cached servers. In addition, you'll need to include commands such as those that restart the servers, or place and take the servers out of a load balancing service.

Let's see how a CD pipeline works in a cloud environment, where you employ something like AWS as your cloud hosting provider that creates the server images for you. In this pipeline, I assume you want to use Amazon Machine Images (AMIs) to create new instances during your deployment. This also lets you use the auto-scale and automated server replacement features of AWS.

In this scenario, here's how your CD pipeline might look like:

1. The development teams commit new code to your CVS such as GitHub.
2. Your CI/CD tool such as Jenkins is automatically notified by a post-commit GitHub Git web hook.
3. The CD pipeline swings into action:
 - Unit tests are run
 - The build is assembled
4. The build artifacts and the test results are zipped and stored in storage (say, Amazon's S3 storage service)
5. AWS's APIs create a new server instance and then upgrade it using CM tools, so all packages are up-to-date.
6. A new AMI image is created from the new server instance, to act as the source of new server instances.
7. The new AMI server image is deployed on the testing (or staging) environment and tested end-to-end, using testing tools such as Selenium.
8. The tested AMI images are now deemed ready for production. Let's say this is a web server image. Amazon's Elastic Load Balancer service recreates all the server instances from the tested new AMI image from the previous step.

CD pipelines like what I've described let you deliver new software very fast, often letting you incorporate new features several times daily. Simply run A/B tests on the users to validate the new features.

But, can I really send all my commits to production straightaway?

Often people are uncomfortable with the idea that new code commits are going to be set to production right away. Use the following CD best practices to make sure that things don't break when you use a CD pipeline.

- Ensure that all code must undergo unit tests.
- Create end-to-end tests cases for all the critical paths of an application using a testing tool such as Selenium. You run these tests before deploying any code to production.
- Use feature toggles that let you enable/disable features on the fly, without having to redeploy the application.
- You can employ A/B testing (and the rolling out of new features in a hidden mode) and feature toggles together to test all new features on a subset of users,

thus minimizing your risk. You can see if the feature usage by the sample set of users is leading to an enhancement of your business metrics.

- Use sophisticated monitoring tools with useful metrics so you can get prompt alerts on critical business metrics.

How Build Tools Help – Maven, Ant and Gradle

As system administrators and devops team members, you're going to be closely involved with the build efforts of development teams. It's therefore a good idea to familiarize yourself with some common build tools. Three popular tools are Maven, Ant, and Gradle.

The Maven Build Tool

Maven is a build and project management tool designed to achieve the following goals:

- Simplify the application build process
- Provide a uniform build system
- Provide useful project information
- Provide guidelines for development best practices

Maven simplifies builds and makes it easy to add new components to existing projects.

In order to reduce the length of the build files, Maven uses a convention over configuration strategy. A Maven build consists mostly of declarations and configuration – there's very little scripting.

Maven introduced the capability to download dependencies via the network and this revolutionized the way software is developed.

Creating a Simple Project

You can check out how Maven creates projects, by using the Maven Archtype Plugin. Archtype is a skeleton template for a new project, and you can include the archetype plugin as follows:

```
mvn archetype:generate
```

Maven uses a dependency resolution mechanism to download dependencies. When you issue this command, Maven starts downloading jar files to your server.

The Structure of the Build File

Maven uses the “Project Object Model” or pom to define its projects. The definition is stored in a file named pom.xml. The pom.xml file contains the following information:

- The name and description of a project
- The packaging of the projects
- The version of the project
- The build configuration (directories, plugins)
- Dependencies (such as libraries) of the project
- The report configuration (test coverage, static analysis)
- The project developers and contributors
- The Maven repositories used for this project
- The project’s infrastructure (such as the source control repository, and the continuous integration server)

Any Maven project you create will be a descendent of the maven “master pom”, which defines the standard build layout and configuration for a maven project. So, you just modify and extend this build layout/configuration to fit your needs.

A POM File

I’ve talked long about a pom file and it’s now time to show an actual pom.xml file. A pom file can be pretty long and complex, but ours is pretty simple:

```
<project xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/XMLSchema-  
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/maven-v4_0_0.xsd"  
  <modelVersion>4.0.0</modelVersion>  
  <groupId>org.scala-lang.demo</groupId>  
  <artifactId>scala-test</artifactId>  
  <packaging>jar</packaging>  
  <version>1.0-SNAPSHOT</version>  
  <name>Demo of maven for Scala Lang website</name>  
  <url>http://scala-lang.org</url>  
  <repositories>  
    <repository>  
      <id>scala-tools.org</id>  
      <name>Scala-tools Maven2 Repository</name>  
      <url>http://scala-tools.org/repo-releases</url>  
    </repository>  
  </repositories>  
  <pluginRepositories>  
    <pluginRepository>  
      <id>scala-tools.org</id>  
      <name>Scala-tools Maven2 Repository</name>
```

```

        <url>http://scala-tools.org/repo-releases</url>
    </pluginRepository>
</pluginRepositories>
<dependencies>
    <dependency>
        <groupId>org.scala-lang</groupId>
        <artifactId>scala-library</artifactId>
        <version>2.7.2-rc2</version>
    </dependency>
    <dependency>
        <groupId>junit</groupId>
        <artifactId>junit</artifactId>
        <version>3.8.1</version>
        <scope>test</scope>
    </dependency>
</dependencies>
<build>
    <plugins>
        <plugin>
            <groupId>org.scala-tools</groupId>
            <artifactId>maven-scala-plugin</artifactId>
            <executions>
                <execution>
                    <goals>
                        <goal>compile</goal>
                        <goal>testCompile</goal>
                    </goals>
                </execution>
            </executions>
            <configuration>
                <sourceDir>src/main/java</sourceDir>
                <jvmArgs>
                    <jvmArg>-Xms64m</jvmArg>
                    <jvmArg>-Xmx1024m</jvmArg>
                </jvmArgs>
            </configuration>
        </plugin>
    </plugins>
</build>
</project>

```

Let's see what our little pom.xml file contains:

First of all it contains a project XML element, within which you specify all other elements. Let me summarize the key elements in this pom.xml file:

- **<name>**: Your project's name
- **<groupId>**: Usually the inverted domain-name of the project is what you specify here
- **<artifactId>**: The unique ID of this project

- `<version>`: The current version of the project
- `<repositories>`: Alternative locations where Maven looks for satisfying the dependencies
- `<packaging>`: This element specifies the artifacts you're building such as an ear, jar, and war files
- `<dependencies>`: The dependencies of the project, which are retrieved from a Maven repository
- `<build>`: The configuration specifics of how the project is to be built
- `<goal>`: A goal represents a specify type of action defined by a plugin, such as `scala:compile`, and `jboss:deploy`.

As you can see, POM files can be pretty detailed but they do ensure that the developers get their desired results.

The Default Directory Layout

Maven's default directory layout is used by convention, although you can customize the defaults. Following is a typical project layout.

```
project/
  pom.xml    - Defines the project
  src/
    main/
      java/ - Contains all java code that will go in your final artifact.
      scala/ - Contains all scala code that will go in your final artifact.
      resources/ - Contains all static files that should be available on the classpath
                  in the final artifact.
      webapp/ - Contains all content for a web application (jsp, css, images, etc.)
    site/ - Contains all apt or xdoc files used to create a project website.
  test/
    java/ - Contains all java code used for testing.
    scala/ - Contains all scala code used for testing.
    resources/ - Contains all static content that should be available on the
                  classpath during testing.
```

The Build Lifecycle

A build lifecycle is at the heart of how Maven functions. Maven believes in clearly defining the process of building and distributing the artifacts (project components). There are three ways to specify build lifecycles:

- The default lifecycle takes care of your project deployment
- The clean lifecycle takes care of cleaning up the project
- The site lifecycle handles the creation of the documentation for the project sit

Each of the three lifecycle describe here contains a set of build phases, which are basically the stages in that lifecycle. For example, the default lifecycle contains the following phases:

- **Validate:** Validates the correctness of the project
- **Compile:** Compiles the project's source code
- **Test:** Tests the compiled source code with a unit testing framework such as JUnit.
- **Package:** Packages the compiled code into a format that can be used to distributed the code, such as a JAR (Java Application Archive), or WAR (Web Application Archive) file
- **Verify:** Checks the results of the integration tests to assure compliance with quality criteria
- **Install:** Installs the package into a local repository
- **Deploy:** Copies the final package to the remote repository for sharing with other developers/packages.

A plugin goal is finer than a build phase – it represents specific tasks that contribute to the building (or managing) of a project.

The Build Plugins

There are build plugins for most activities required in Java and Scala projects. These plugins attach their goals to the appropriate build phases.

Handling Dependencies

Projects can just declare their dependencies and it's up to Maven to automatically ensure that the dependencies are in place. You deal with the dependencies and the plugins to use, and so on. With Maven, you never need to specify the tasks that need to be performed.



Maven relies on conventions and provides goals out-of-the-box.

Using Ant to Perform Builds

Apache Ant is the first modern build tool, and has similarities to the well-known Linux make tool that helps build binaries. Ant is simple with a small learning curve, and is especially useful for small projects. Maven was built to overcome several problems with Ant.

For beginners, Ant build files are easy to understand – you can figure out what the build script is doing by reading the configuration XML.

Ant uses the Ivy tool to download the dependencies for a project. Ant uses a procedural approach, and its use of the hierarchically structured XML for its scripting is problematic. Ant's XML can become unwieldy for larger projects.

The big advantage with Ant is the control you get over the build process. Following is a typical ant script, named build.xml. This script compiles a JAR file, and uses Ivy for downloading the dependencies.

```
<project xmlns:ivy="antlib:org.apache.ivy.ant" name="java-build-tools" default="jar">

    <property name="src.dir" value="src"/>
    <property name="build.dir" value="build"/>
    <property name="classes.dir" value="${build.dir}/classes"/>
    <property name="jar.dir" value="${build.dir}/jar"/>
    <property name="lib.dir" value="lib" />
    <path id="lib.path.id">
        <fileset dir="${lib.dir}" />
    </path>

    <target name="resolve">
        <ivy:retrieve />
    </target>

    <target name="clean">
        <delete dir="${build.dir}"/>
    </target>

    <target name="compile" depends="resolve">
        <mkdir dir="${classes.dir}"/>
        <javac srcdir="${src.dir}" destdir="${classes.dir}" classpathref="lib.path.id"/>
    </target>

    <target name="jar" depends="compile">
        <mkdir dir="${jar.dir}"/>
        <jar destfile="${jar.dir}/${ant.project.name}.jar" basedir="${classes.dir}"/>
    </target>

</project>
```

This build.xml file creates a JAR file, once you issue the command `ant jar`. Notice how, unlike in the case of Maven, I need to specify the tasks involved in compiling the JAR file.

Using Gradle as a Build Tool

Gradle seeks to combine the best features of Maven and Ant. It combines Ant's power and flexibility with Maven's build life cycle features. Google uses Gradle as the default build tool for its Android operating system.

Gradle uses its own DSL based on Groovy programming language, rather than XML files, as is the case with both Maven and Ant.

Gradle's build scripts are much more compact and easier to understand than the Ant and Maven build scripts. Gradle tries to combine Maven's use of conventions with Ant's flexibility. Following is a typical Gradle build script.

```
apply plugin: 'java'
apply plugin: 'checkstyle'
apply plugin: 'findbugs'
apply plugin: 'pmd'

version = '1.0'

repositories {
    mavenCentral()
}

dependencies {
    testCompile group: 'junit', name: 'junit', version: '4.11'
    testCompile group: 'org.hamcrest', name: 'hamcrest-all', version: '1.3'
}
```

Gradle has a large number of tasks available out-of-the-box, or through its many plugins. Gradle has a somewhat higher learning curve compared to Maven and Ant, but has a simpler, easy to understand syntax.

An Overview of the Software Development Process

The evolution of software engineering practices over time gave rise to continuous integration. In order to really understand CI you must learn the basics of how software development practices have evolved over time.

Documented methodologies such as the waterfall model were the basis for large software development strategies for many years. The evolution of the agile software development process underlies the move to CI. In order to understand how agile software development practices work, you need to understand something about a software development life cycle.

The Software Development Life Cycle

The software development cycle (SDLC) is the process software goes through in the planning, development and testing phases before deployment of the software. Following are the phases of a SDLC:

- Requirements analysis: Gathering the business/customer requirements
- Design: Designing of the software and creation of the project plan
- Implementation: Code development by the development team
- Testing: Testing of the software's features and raising defects (bugs/errors)
- Deployment: Deploying the software in the production environment
- Maintenance: Also called the evolution phase, this phase uses feedback from users and customers to patch and upgrade the software

Following the Waterfall Model of Development

The waterfall model is a classic software development methodology that follows a sequential development process. This is really a model patterned after well-established manufacturing processes. The waterfall model uses the same steps as those for a software development life cycle. That is, it starts with a requirements analysis and ends with the maintenance phase.

When organizations had all the time in the world, and the pace of technological development was slower, the waterfall model worked quite well. Typically the SDLC phases took a year or so in most cases, and the model was suitable for software with static requirements. This model can't deal with the modern requirements of websites that demand very frequent modifications and extremely short development cycles.

The Agile Software Development Process

Current software development follows the agile software development process, as opposed to the traditional waterfall model. The agile model is based on the following 12 agile principles:

- Satisfy the customer through early and continuous development of software
- Welcome changes in requirements at any time, to harness the change for the customer's competitive advantage
- Deliver software frequently – from a couple of weeks to a couple of months
- Have the business and development teams work together daily
- Center projects around motivated individuals and give them the environment and support them

- Face-to-face conversations as the most efficient and effective mode of exchanging information
- Use working software as the primary measure of progress
- Promote a sustainable pace that can be maintained indefinitely
- Continuous attention to technical excellence and design
- Simplicity is essential to the process
- Self-organizing teams lead to the best architectures, requirements, and designs
- Teams regularly analyze ways to become more effective and change their behavior accordingly

Under an agile software development process, you break the software into multiple features or modules. The goal is to deliver the modules in iterations. The iterations cover short time frames – each iteration may last for 1-3 weeks, during which all the functional teams – analysts who gather the requirements, designers, coders, and testers work simultaneously (or parallelly) in their areas of expertise. Thus, unlike in the traditional waterfall model, teams aren't waiting on one another and keeping idle for long stretches of time during the software development process.

The strategy is to develop and release a tested feature in a single iteration. If the customers like the feature during the demo phase, the feature goes live. Else, the feature goes back into the development cycle for further iterations and refinement. Teams can also develop and test multiple features simultaneously.

Agile development is suitable for modern day web applications whose requirements change frequently. The stress is on delivering working software quickly rather than spending time on requirements, planning, and documentation.

The Scrum Framework

Scrum is a framework designed for developing complex software, and is based on the agile software development process. Scrum involves roles, tasks, and teams.

Scrum relies on the development team to determine how a feature is developed and uses self-organizing and cross functional teams. Self-organizing means there isn't a team leader determining who will do what and how to solve a problem. The cross-functional nature means that everyone gets involved in moving a feature from a concept to its final delivery to the customer as a finished product.

Scrum uses the following terms:

- **Sprint:** A sprint denotes a software iteration – it's a period of time during which a releasable software increment is created by the development team. Sprints follow each other and may last for a week to a month.

- Product backlog: List of all the requirement features and keeps changing.
- Sprint backlog: Set of items from a product backlog that you select for a sprint.
- Product owner: Acts as the face of the Scrum team and mediates between the Scrum team and the rest of the organization.
- Scrum Master: Person responsible for ensuring that everyone understands Scrum and acts on ensuring that the teams follow the Scrum principles and rules

The Scrum master, the product owner, and the Scrum team work together to follow the Scrum principles to deliver crucial software features quickly to the clients. Following are the key components of a typical Scrum software development process.

Sprint Planning

Developers create the sprint plans and explain them to the Scrum master and the product owner. Sprint planning allows the Scrum team to plan the software features in the current sprint cycle.

Sprint Cycles

During the sprint cycle, developers work on the backlogs they agreed upon in sprint planning. The sprints may last for a week to a month.

Scrum Meetings

Scrum meets occur daily and allow the development team to discuss what was accomplished the day before, and what will be accomplished today. This scrum meeting also allows developers to discuss obstacles that are slowing them down. Daily scrum meetings also provide an opportunity to measure the project's progress.

Sprint Reviews

The sprint review is where the developers demo the features to the customers. The product owner updates the product backlog based on the product performance.

Sprint Retrospective

The sprint retrospective follows the sprint review and occurs before the planning of the next sprint. During this meeting, the team discusses things that went well, as well as things that they need to improve on during the next sprint.

The Agile Software Development Process and Continuous Integration

Continuous Integration helps agile software development processes achieve their goal of fast delivery of tested software. Developing new features involves numerous code changes that require various tasks such as the following:

- Checking the new code in
- Polling the VCS for changes
- Building the code
- Performing unit tests
- Integrating the code
- Packaging the code
- Deploying the new feature

CI helps by automating these processes, resulting in fast as well as mostly error-free deployments that are common in a manual process. Teams are also automatically notified of any build, integration, or deployment failures, thus making remediation virtually instantaneous. It's hard to think of agile software development processes without the use of CI techniques.

Continuous Integration (CI)

Continuous integration is the set of processes and practices that are involved in regularly integrating the work of multiple developers in a shared repository (or mainline). Integration is the act of submitting a developer's work to the common work area – this is usually done by merging a developer's work (personal branch) with the common work area (integration branch). Continuous Integration is running frequent integration builds so they become a non-event in your project.

The CI Process

CI involves the compilation, rebuilding, and execution of automated tests and inspection. CI isn't merely the putting together of software components – it's the heart of software development, and ensures quality software by running a build with every change.

CI is a practice where team members frequently integrate their work, with each person integrating at least daily – this means that there usually are multiple integrations every day. An automated build process verifies and tests the integrations, to catch errors. This leads significantly lower integration issues and helps your teams to develop effective software rapidly.

Integration in this context is the act of combining multiple source code artifacts to determine who it all functions together as a one entity. Build processes result in committing of the new changes to the CVS. Devops teams often use CI systems to automatically build and test source code. You can think of continuous integration as continuous staging.

When organizations develop a high degree of confidence in their build strategy, they can have new builds automatically placed into production – this called continuous deployment (CD).



You use the same tools for both CI and CD. For example, you can use Jenkins both CI and CD.

Traditional software development moves code through code commits, binary builds, QA and testing and staging before moving on to production. A lot of the processes are manual, thus making automation of code deployment a key factor. CI systems and build pipelines can start with a pipeline that begins with source code development and ends with the automatic staging of tested code for deployment in production.

Problems Detected by CI

CI detects the following problems early in the software development cycle:

- Build failures (before the code integration)
- Integration issues
- Build failures (after the code integration)

Integration issues are usually due to developers infrequently merging their own code with the code in the integration branch. Long gaps between code integrations make it difficult to find and fix problems. Integration issues caused by infrequent code integration can delay a project and may even lead to the project's failure.

Building Software at Every Change

A build is the heart of a CI system. A build is the compilation, testing, inspection and deployments of software. The build encompasses the processes that put the source code together and verify the software works as it's designed.

The idea behind continuous builds is to get quick feedback that helps you find and fix problems throughout the software development lifecycle.

The term build plays a significant role in CI. A build is the set of actions that generate, test, inspect and deploy your software. In a CI context, there can be three types of builds:

- Integration Builds; An integration build is the combining of software artifacts or components (such as programs and files) to create a working software application

- **Private builds:** These are builds run locally on a workstation by the developer before committing the changes to a VCS. These builds minimize the chance that your changes break the integration build.
- **Release builds:** This build prepares the software for a release to the users.

Benefits of a CI System

A well-managed CI systems offers several important benefits:

- **Risk reduction:** Software defects are detected and fixed early.
- **Deployable software at any point in time:** Any changes you make to the application's code are integrated with the code base regularly. Any problems are applied to the software immediately. Thus, you end up with software that's ready to deploy at any point in time
- **Reduction in manual repetitive processes:** Automated tools enhance the efficiency and frequency of the CI related work.
- **Better project visibility:** Members of all teams are aware of the build pipelines and know the current status of a project
- **Establish greater confidence in the software projects:** Team members develop greater confidence in the projects, since they are aware of all the steps of the build pipeline and also are aware of the current status of the project.

Best Practices for CI

CI works best when you follow established best practices, as explained in the following sections.

Use Private Workspaces

Developers should work in their own private workspaces, so their private copy of the code is isolated from changes being made on the mainline branch.

Developers can work in private workspaces through branching, or by cloning the code repositories (as explained in Chapter 7, where I discuss Git and GitHub)

Rebase Frequently from the Mainline

When multiple developers are working on their private repositories or private branches, it could lead to potential merge issues. Rebasing is when developers update their private code branch with the latest version of the integration branch.

CI involves the frequently merging and rebasing of code – integration is continuous and frequent.

Perform Frequent Check-Ins

Developers should check in their working branches at least once a day. If they do infrequent check-ins, say once a week, there's the potential of merge issues. Daily commits or merges help detect conflicts early and resolve them quickly.

Perform Frequent Builds

A CI tool should ensure that all commits or merges are built immediately, to test the changes. CI tools can do this by frequently polling the integration branch in the CVS for changes. The tool should automatically build and test any changes it finds. You can alternatively schedule nightly builds.

Write Automated Tests

A continuous build can inform you about any build failures. Continuous testing can let you determine if the build is ready for production. Many of the tests that involve real world scenarios require a heavy amount of scripting and maintaining.



All CI tools such as Jenkins can instantly notify you about build and deployment failures, as well as communicate the latest results. Instant notifications often save you time since you can fix minor issues immediately.

Don't Commit Broken Code

Performing private builds before coming to a VCS is a good strategy. This way, if the developer's private build fails, they don't check in their changes. Alternatively, you can set up the VCS to trigger a build through the CI tool, and check in the code only if the CI tool returns positive results. Some VCS tools have a gated check-in feature to do this. Fixing broken builds immediately is always a good strategy.



Static code analysis is the analysis of software that doesn't involve the actual execution of the program. Static code analysis is usually part of a Code Review (aka white-box testing), and is a distinct part of the implementation phase of a software development lifecycle.

You can also add steps such as performing a static code analysis to the gated check-in strategy. You do this by integrating the VCS system with the CI tool, which in turn is integrated with a tool such as SonarQube that performs static code analysis.

Automate the Deployment to the Testing Environment

Automating code deployments to the testing environment saves considerable time compared to a manual deployment strategy. When developers check in their code, it should automatically undergo a compilation check and code analysis before being checked into the integration branch. All the code on the integration branch is frequently built, packaged, and deployed to the testing environment.

Requirements for a CI System

At its core, a CI system requires the following four things:

- A connection to a version control system (VCS)
- A build script, where build refers to a set of activities that generate, test, inspect, and even deploy software.
- A feedback mechanism from the CI sever to communicate the status of the builds (email will do)
- A way to integrate the source code changes – manual, custom scripts, or a CI server such as Jenkins or TeamCity

Features of a CI System

Following are some of the essential features of a CI system.

Automated

CI is mostly a hands off process. It's an automated process. Many problems with application development and deployment are due to compilation issues or bugs of various types (feature bugs, and logic bugs for example). Geographically spread out dev teams make it harder to keep the application errors out of the picture.

Automating lets you frequently test the code, thus reducing the risk in the software development process.

Source Code Compilation

A common feature of a CI system is the continuous source code compilation. A CI system lets developers run different types of tests such as unit tests, component tests, load tests and performance tests.

Testing

Automated and continuous testing is at the heart of CI. Testing ensures the following:

- New code for new features works as expected

- Current features work fine without new problems
- Bugs that were taken care of earlier doesn't reappear (regression testing)

Development teams test frequently, ensuring that the code passes all tests before they commit to a CVS. All the testing seems a bit overwhelming, but CI tools such as Jenkins, Go Continuous Delivery (GoCD), and Travis Continuous Integration (Travis CI) let you create pipelines that consist of the build steps for all the unit tests that you need to perform.

A CI system uses automated dynamic code inspections to enhance software quality through the enforcement of rules. For example you can impose a rule that limits a class to no more than 400 lines of code (not including comments). The CI system will automatically run your rules against the code base. Or, you can use a static code analysis tool such as Checkstyle, which inspects Java code to ensure that it adheres to your coding standards.

A Typical CI Scenario

A team of developers, a version control system (VCS), and a CI server (such as TeamCity, Go, or Jenkins) are the essential components of a CI system.



An integration build can't happen without committing changes to the VCS.

Following is how the developers, the VCS and the CI server interact in a typical CI scenario.

Developers Commit their Changes to a VCS

When a developer is ready to move on with a set of changes to code, he or she runs a build that's private to himself or herself, which integrates changes from the other team members. They then commit these changes to the version control repository.

CI requires you to automate your builds, so you can run them with a single command. Automatic builds help developers confirm the software components work together, and that the tests are successful after the latest changes.

Following is an example of a private build with Ant, a popular build tool.

```
> ant integrate
Buildfile: build.xml
clean:
svn-update:
all:
```

```
compile-src:
compile-tests:
integrate-database:
run-tests:
run-inspections:
package:
deploy:
BUILD SUCCESSFUL
Total time: 3 minutes 13 seconds
```

The VCS and the Role of the CI Server

A VCS is a must for a CI system. Once developers commit their code, the CI server detects the changes in the VCS, and retrieves the latest copy of the code from the VCS and executes a build script. This build script integrates the software.

While you can use your own scripts, or even run manual integration builds when you apply changes to the VCS, it's highly recommended that you use a CI server. The CI server reduces the number of custom scripts you'll need to write otherwise. There are several good open source CI servers, and I discuss the most popular one (Jenkins) later in this chapter.

You can configure the CI systems to poll the SCM system regularly for changes, or what is more commonly done, let the SCM itself tell the CI systems to generate new builds following each commit of new source code by the developers.

The CI server runs integration builds every time a developer commits a change to the VCS. You usually configure the CI server to check for changes in the VCS every few minutes. This is what's responsible for ensuring the “continuous” part of CI. The integration process is a continual process that's constantly looking for changes in the VCS, so it can execute a build script wherever it detects a changes.

NOTE: Compilation is the creation of executable code from a source.

Feedback from the CI Server

The CI server sends feedback regarding the build results to appropriate project team members, in the form of reports. If there are no errors in the builds, the compiled code is classified as build artifacts and the CI/CD system will store them. The CI server will await the next changes to the VCS.

The CI server can reveal various things about your software, such as quality metrics, code standard adherence, code duplication.

It's the testing tools in the CO process that lead to the success of a DevOps system, by catching errors early and minimizing the chance of deploying bad code.

Tools such as Jenkins and TeamCity save you significant amounts of time both for the development and the operations teams. Developers can catch bugs early, before they

turn the code over to QA and operations teams will know that they're dealing with applications that have been validated before being staged.

Continuous Delivery

Continuous Delivery (CD) is reliable software releases through automating the test, build, and deployment processes. Once your organization comes up with a great idea, how do you deliver it to your users quickly? CD addresses this question: what's an effective pattern for getting your software from development to release? The goal of CD is to implement this effective pattern.

There are several well-known software development methodologies that approach software design, development, and testing in different ways. Efficient software building starts with identifying the requirements, designing solutions, developing, and testing them. CD aims to make the entire software release process efficient and reliable. Its goal is to enable developers, testers, the build team, and the operations team to cooperate effectively.

Continuous delivery (CD) is a strategy in which you fully automate all software testing and trigger the tests to run for each build. After each build, the automated tests are run and a new release is “delivered” or created, poised for release in other environments. You may or may not deploy software automatically – however, the changes are proven to be deployable whenever you wish.

You reap similar benefits from CD as you do from CI, such as small and frequent changes that reduce your risk exposure, and problems that are found early and hence are easier to fix.

CI encompasses CI, plus a bunch of testing, such as user acceptance tests and automated tests. CD delivers releases to beta environments for manual testing, if some of the tests aren't automated.

CD and CI

CD is the making available of applications to be deployed into production at any time. CD flows from CI, and is an extension of the latter, with CI being more concerned with the activities of the development teams.

CD involves bringing together the development, operations, and business teams together to ensure applications are released into production in a timely manner.

Once again, CD here means continuous delivery, not continuous deployment. CD ensures that every build is readily available for deployment into production. CD in general ensures that builds are deployed to the user acceptance testing (UAT) environments so teams can test the app and determine if and when it ought to be put in pro-

duction. The sum of the business and operations factors determines if a build is deployed into production or not.

The CI process generally ends with the creation of the artifacts necessary for deploying the application.



CD involves much more than a set of tools and processes – it requires close collaboration among developers, testers, operations teams, and business groups.

The Build (or Deployment) Pipeline

A build pipeline, also called a deployment pipeline, is a strategy of effectively getting your software from development to its release. It's an automated implementation of the workflow involved in building, deploying, testing, and release your applications.

A build pipeline starts with the development of source code and ends with tested and verified code that is automatically staged for deployment in the production environment. The CI system performs all prior steps, including testing, automatically.

Here's are the stages of the deployment pipeline: **Commit strategy: compile, unit test, build installers** Acceptance testing: Automated testing processes **Capacity and Stress testing: Automated testing processes** Manual testing: exploratory testing ** Release: Let the software out in the wild

Deployment pipelines have their roots in continuous integration and are in effect the culmination of the CI process. A build pipeline lets you separate a build process into multiple stages. This allows multiple builds to run simultaneously, with the builds in various stages of the pipeline.

Stages of a Build Pipeline

Build pipelines have multiple stages. The first stage of the build pipeline is usually fast, to ensure the team gets quick feedback. This stage compiles the code to produce the binaries and runs unit tests on them. It passes the resulting artifacts to the subsequent stages.

The latter stages perform other types of testing (other than unit testing that is), such as acceptance and functional testing, as well as performance testing. These stages deploy the software to environments such as test environments to perform the tests described here. It might also deploy the software to a User Acceptance Testing (UAT) environments or manual testing.

Culmination of the Build Pipeline

The build pipeline culminates in an automatic deployment to an environment that's quite similar to production, such as a staging environment. Subsequent deployments to production environments may be manual.

Only those builds that successfully passed through all the testing stages are fully promoted and deployed to production.

Continuous Delivery and DevOps

Continuous deployment and DevOps go hand in hand and have a symbiotic relationship. Whereas CD helps with quickly shipping quality software, DevOps helps disparate teams that deliver and support the software come together. Together, CD and DevOps, by streamlining and improving work processes, help you ship quality software quickly.

Continuous Integration with Jenkins

Jenkins is a popular CI system that you can use to automatically build and test source code. You can integrate Jenkins with SCMs such as Git, letting code be submitted automatically to Jenkins for compiling and testing when developers commit code to the SCM system.

Jenkins isn't limited to continuous integration – you can use it for continuous delivery and continuous deployment as well. Jenkins provides job histories and its console lets dev team members track the stages of the automated test and build cycles.

Architecture of Jenkins

Jenkins is very powerful, owing to its extensible platform which accesses all types of plugins to enhance its capabilities. For example, you use a plugin to enable Git support.

Jenkins uses a master-worker architecture. While you can configure the master to perform builds and create reports, you want to let the master run on a separate node and let the users interact with the master. You can configure multiple build processes (executors) on each of the worker nodes.

When users contact the Jenkins master, it sends instructions to the workers running on the build servers for performing builds. The workers perform the builds and send back the build artifacts and build reports to the master. Users access the artifacts and reports from the master server.

Jenkins works as an orchestrator of various Devops tools to achieve Continuous Integration. Jenkins uses plugins to communicate with the VCS server, the repository tool (for example, Artifactory), and the static code analysis tool.

Components of the Jenkins Framework

Jenkins consists of the following components:

- **Job:** A Jenkins job consists of a description of the job, parameters, build steps and post-build actions.
- **Parameters:** Parameters can include environment variables and pre-defined values and triggers, whose purpose is to trigger the before and after build activities and to help perform the builds
- **Builds:** These are different from the software builds and could range from simple commands to complex Perl, Python, Shell, and Ruby scripts. They can also include Maven and Ant builds. You can have multiple build steps inside a Jenkins job, as in the case when you perform a Maven build and follow it up with a script to merge the code in a Git repository.
- **Post-Build Actions:** These are actions to be performed after a build, and can trigger new Jenkins jobs, as in the following example:

```
Projects to build: Upload_Package_To_Artifactory
```

```
Trigger when build is: Stable
```

- **Pipelines:** A Jenkins pipeline is a set of Jenkins jobs that run together, sequentially or in parallel. Pipelines can be used to build, test, and deploy code.

```
Static Code Analysis, Integration-Testing
```

```
a minute ago    18 sec
```

```
Publish to Artifactory
```

```
a minute ago    1 sec
```

```
Deploy to Testing Server
```

```
a few seconds ago  2 sec
```

```
User Acceptance Test
```

```
a few seconds ago  24 sec
```

```
Performance Test
```

```
...
```

Jenkins and other Continuous Integration Tools

Jenkins acts as the CI server in a continuous integration setup, but it needs to work with a vast array of other tools to get its job done. The tools vary according to the type of application you're integrating. Here are the typical tools used for CI with Jenkins, assuming you're dealing with a Java-based web application besides Jenkins itself, which of course is your main CI tool.

- **Maven:** Build tool
- **JUnit:** java application testing tool to perform both unit and integration tests

- Eclipse: IDE for developing Java apps
- Git: Version control system
- SonarQube: Static code analysis tool
- Artifactory: Package repository

And here's how all these tools interact to get CI going:

- Developers configure their Eclipse IDEs with the Git server so they can clone the feature branches from the VCS to their workstation.
- The Jenkins master server is connected to the Git server through its Git plugin, so it can poll the Git server for code changes
- Jenkins uses the Maven build tool to build the code checked into the Git server
- Using the SonarQube and Artifactory plugins, Jenkins is connected to SonarQube for performing static code analysis and to Artifactory, which is the package repository

Let's say you use GitHub as your SCM. You can setup a Git web hook on GitHub so GitHub can immediately trigger a new build on the Jenkins server when new code is committed to GitHub. If you don't setup this Git web hook, the Jenkins server will be periodically checking and will soon notice the change in your GitHub repository and trigger a build on its own.

The Need for a Binary Repository Tool

You'll need some place to store all the binary code that is produced by the frequent builds during continuous integration. A version control system holds the code that's under development and testing. A binary repository tool is for storing the binary files – the end product of your continuous integration pipelines. Binaries include code packages, executables, and similar artifacts. The binary repository tool can also help in the version control of the binaries.

A binary tool helps you track the builds such as what triggered the build, and which version of the code in your VCS was the source for a build. It also records all the dependencies and stores the deployment history.

Popular binary tools such Artifactory and Nexus are integrated with a CI tool such as Jenkins to store all build artifacts. Jenkins publishes tested code to Artifactory through a plugin and Artifactory stores the builds.

A Typical Continuous Integration Strategy using Jenkins

Although you may use a tool such as Jenkins to help with your CI work, CI strategies include many software configuration management decisions such as the strategies for version controlling your code, your branching strategies. There's no single set of standards for achieving CI, since everything depends on the nature of the project and its requirements. Here, then, are some of the typical strategies followed by CI practitioners.



All development work happens on the feature branches and all integration work occurs on the integration branch.

Branching strategies

Branching helps organize your code and makes it easy to separate the working code from the code that's being currently developed. In designing your CI strategy, you can use the following types of branches:

The Master Branch

The master branch, also called the production branch, has the working copy of the code that has been developed, tested through various stages, and delivered. You don't perform any development on the Master branch.

The Integration Branch

The integration branch, is also referred to as the mainline branch, and it spans out of the master branch. As with the master branch, no development activity occurs on the integration branch. As its name indicates, this is the branch where you integrate, build, and test all features. Developers can create feature branches from the integration branch to work on various features.

A merging of code from the feature branch (see the next section) creates a commit on the integration branch. Following this commit, the code undergoes the build, static code analysis, and the integration testing stages. Once it passes all these tests the code is uploaded to a binary repository such as Artifactory.

The Feature Branch

The feature branch is where the developers work on their stuff. You can have multiple feature branches spanning out of the integration branch.

Developers perform code check-ins (same as a code commit) on the feature branches they're working on. This code will then pass through the build and unit test phases. Once it successfully moves through these testing stages, the code is merged with the integration branch.

Creating the Continuous Integration Pipelines

A Jenkins pipeline, as explained earlier, consists of a set of jobs that can run either sequentially or parallelly. CI pipelines automate the process of the continuous building, testing and integration of all new code, with appropriate communications to the concerned parties during each stage. In my example here, I create two pipelines, to poll the feature and integration branches for new code committed to those branches.

A Pipeline that Polls the Feature Branch

The first of our two Jenkins pipelines works closely with the feature branch. When developers commit code to the feature branches this pipeline gets triggered. The pipeline consists: of two Jenkins jobs:

- Job 1: The first job in this pipeline regularly polls the feature branch for changes, and performs the following tasks: Perform builds on the committed code Run the unit tests
- Job2: The second job in the pipeline merges all successfully built and tested code to the integration branch. Once the tested code is merged with the integration branch, the second pipeline is triggered.

A Pipeline that Polls the Integration Branch

The second pipeline in my example works with the integration branch. This pipeline swings into action whenever there's a commit to the integration ranch. In my example, I again have two Jenkins jobs in this pipeline:

- Job 1: The first job regularly polls the integration branch for changes and perform the following tasks: Performs static code analysis Builds and executes the integration tests
- Job 2: Uploads the built package to the binary code repository (such as Artifactory)

Continuous Delivery with Jenkins

You can use the same Jenkins pipelines that you've created earlier for your CI exercise for performing continuous delivery. A CD pipeline produces production ready features that passed not only all the CI tests (unit tests, integration tests), but also user

acceptance testing (UAT), end-to-end testing and performance testing (stress testing). During CD, you'll be performing automated testing using testing tools such as Selenium and TestNG for performing acceptance testing, as well as tools like JMeter for performing performance testing.

When you use Jenkins to perform CD, you'll need to add additional pipelines on top of those you've created for CI. The jobs in these new pipelines will perform the user acceptance tests and performance tests, generate the test results reports, and send put e-mail notifications about the test results.

Earlier, you learned how Jenkins uploads the code packages after successfully integration testing, to Artifactory (or a similar binary repository tool). When you use Jenkins to continue with the CD process, it deploys the package on a testing server. Jenkins will then use testing stools such as JMeter (open source tool for automated performance testing), TestNG and Selenium to perform the user acceptance and performance tests on the binaries.

At this point, the artifacts are considered production ready, and are marked as such in Artifactory.

Continuous Deployment with Jenkins

Continuous delivery ends with the labeling of artifacts as production ready. The deployment of the artifacts to production isn't automatic – you choose when to deploy the builds, based on your business requirements. So, CD doesn't involved continuous deployment.

A Continuous Deployment cycle is one where the deployments are automatic – any production ready features are immediately put into production without pausing for your input. Often, business groups need to weigh in as to when the right time is to put a new feature into production – so, continuous delivery and not continuous deployment is the goal in most organizations.

Continuous deployment, unless continuous delivery, has several challenges since you're performing releases to the production servers. Deployments to production web servers or application servers means downtime. In order to avoid frequent downtimes and unavailability of services, a common strategy is to use clustered environments, with banks of web servers and web application servers. You could then deploy in production with a zero downtime strategy, by applying changes to the services running on each node, one by one.

Pros and Cons of Jenkins Jenkins is quite popular as a CI tool, and is an open source tool that has a large number of plugins. Plugins are a source of strength and at the same time, a cause of issues in using Jenkins. The problem is that the plugins are written by various authors, and your build pipeline often involves using plugins such as the Build Pipeline plugin and the Copy Artifact plugin.

There are sometimes multiple plugins that do the same thing, making you choose one of them, and you also need to configure each of the plugins. Since the plugins were created independently, and aren't tested to work together, the plugins turn into a pin point when setup up Jenkins build pipelines.

Jenkins isn't ideal for all users and therefore one might investigate other CI tools such as ThoughtWorks's Go CD/CI tool, and TeamCity, another popular CI/CD tool.

Continuous Integration with TeamCity/GO

TeamCity is a popular CI server and you can configure the tool to perform build practices such as source code compilation and the running of integrated tests. The tool lets you follow the “build once and deploy everywhere” strategy and deploy the executables into testing environments for functional testing and by exposing the tested artifacts for downloading.

TeamCity comes with a lot of functionality out of the box and makes it easy to get started on the road to CI. There is a large number of plugins in for TeamCity, most of which are bundled with the tool. TeamCity provides first-class support for various technologies such as Java, .NET, Android, and iOS. TeamCity also has REST API, which help you perform remote actions such as triggering builds, and checking the status of running builds. It also comes with a great dashboard UI for viewing the projects and the status of the build configurations

Go is a formerly commercial tool that has been open sourced. Go originated in CruiseControl, an early CI server. Build pipelines are the string point of GO, which has excellent visualization and configuration features for pipelines. A pipeline's stages fit together automatically and you don't to cobble them together yourself. However, GO comes with a limited set of features and there are only a limited number of plugins that are available, unlike in the case of Jenkins and TeamCity. The limited amount of features owes partly to the developers of GO wanting to limit users to following certain CI practices.

Centralized Log Management and Analysis, and Handling Streaming Data

System administrators deal with logs on a daily basis. However, it's only recently that administrators have been realizing the true value of the enormous amount of logs that automatically flow into their systems. Traditionally administrators have looked at logs as something they reviewed during troubleshooting exercises. Now, log files are being increasingly seen as continuous streams of events or information that can yield valuable insights based on underlying data patterns.

Let's take the case of the ubiquitous HTTP server access logs. These logs contain both error messages (404, page not found) as well as successful transactions. In addition you have data such as client IPs and response time. You can use some of this data to dig deeper into the business processes, for example by getting the geographical locations of your users based on their originating IPs. Logs can be quite useful not only for troubleshooting, but also for understanding user behavior. Organizations use the log data for various purposes such as automating security scanning, and scaling their web services.

System administrators can use log data for debugging, root case analysis and understand customer behavior in terms of their usage or buying patterns. Machine data allows an organization to ask questions that they couldn't envision earlier to find out answers to fundamental questions about their IT infrastructure.

In this chapter, which is mostly dedicated to the capture and analysis of log data, I start off with a review of centralized log management, and segue into the analysis of the logs you capture, to derive useful information about the operating systems, middleware, and other parts of your infrastructure.

Many data driven organizations need to receive and process continuous streams of data to keep and gain competitive advantage. Thus, effective handling of streaming data has become a critical component of many information technology initiatives. Streaming data needs specialized data ingestion strategies to handle data flowing into an organization from various sources, and streaming analytics to mine the streams of data. Apache Kafka is a popular open source product that's being used by many organizations to handle streaming data, and I review the essential features of this product and explain how it helps you work with streaming data.

Introduction to Logging and Its Issues

Logs are a useful source of information for managing various computer resources, users and application management, and security. All kinds of applications and servers – web servers, middleware platforms, end-user applications – generate logs. In addition, operating systems and firmware generate logs a swell.

Managing log data can be overwhelming, especially in web based architectures, but you must devise effective strategies for handling the logs and derive useful information from them. Logs are raw data and the trick is in collecting, storing, analyzing, and mining this raw data in the logs to glean useful information from them.

What is a Log?

A log is a collection of event records. The term logs refers to a set of log messages that are generated by various entities such as a computer system or software to indicate an event of some kind. Logs can be classified into various categories, such as warning, error, and alert messages.

Regardless of the protocol used to transmit it, a log message contains the following basic entries:

- **Timestamp:** Denotes the time when the log message was generated
- **Source:** System that generated the log message (represented by the IP address or hostname)
- **Data:** The essence of the log message – includes entities such as source/destination IP addresses/ports, usernames, program names, objects such as filenames/directory names, number of bytes that were sent/received.

Varieties of Log Data

Traditionally, machine or log data was supposed to be for the system administrator, who needed to analyze the logs to find root causes of systems failures, or performance slowdowns. The advent of the internet and web applications and the use of big data

has changed the type of data enterprises want to process and analyze. Today, log data is used as a key source of real-time, or near real-time business intelligence.

Analyzing machine or log data has become part of mainstream data analysis, and enterprises seek to use this data to improve business decision making includes all the following types of data:

- Web log files: Capture activities occurring on the web sites and related web applications. Web logs (also called access logs) constitute the most widely analyzed form of machine or log data.
- Clickstream data: Tells you what your visitors did when they were on your website.
- OS logs: Used for performance and system monitoring
- Application logs: Show details about application execution and help you optimize your operations
- Firewall logs: Analyze security related incidents and issues.
- Social Media Data: This is the data from social media sources such as Twitter which you can use for marketing and sales purposes.

Issues in Performing Log Analysis

Every organization deals with logs from a wide variety of components and systems. The disparate sources of logs pose several problems. The problem is that everything in a log – the timestamp formats, the way the source is represented, as well as the log message content itself, can vary across systems and applications. This is a central problem in handling log data. Following is a brief review of the commonly encountered issues in handling log data.

Under-appreciation of Logs

Most system administrators don't appreciate logs anywhere near what the logs demand. Logs are pesky things that only cause problems by filling up the file systems without notice. Effective mining of log data requires more than writing some scripts to look for pieces of information in the log files.

Decentralized Logs

Your logs are produced by various applications running on different servers. You may be able to run a cat or tail command and pipe the results to the grep command when dealing with one or two servers, but when dealing with 50 servers?

Multiple Log Formats

You know that the log formats vary across applications and devices, with each log file representing different times and formats, and different ways to show the error messages. Here are two snippets from two different log files, the first from a Tomcat web server and the second from an Apache access log. Note the stark differences in the date formats as well as the way information is presented to the viewer.

```
#Tomcat Web server log entry
Jun 24, 2016 4:58:32 PM org.apache.catalina.startup.HostConfig deployWAR
INFO: Deployment of web application archive \soft\apache-tomcat-7.0.62\webapps\sample.war has finished

#Apache log entry
127.0.0.1 - - [24/Jun/2016:16:54:58 +0530] "GET /favicon.ico HTTP/1.1" 200 21830
```

Multiple Time Formats

Each component may use a different timestamp format, so it's going to be difficult to correlate events occurring at the same time. Here are some of the time formats you'll see in your logs.

- 188443726
- Oct 12 23:21:45
- [5/Jul/2016:08:09:10 +0000]
- Tue 01-01-2016 8:00
- 2016-06-30 T 05:45 UTC
- Sat Jul 02 08:24:48 2016
- 08:48, 04 July 2016 (UTC)

You may try to use some tools such as `grep`, `awk`, `sed`, `head`, and `tail` to trap specific information in the log files, but you need to know the specific text you're looking for, and the analysis will run long when looking at terabytes of log data. You also can't correlate the information within the numerous log files that are spread over multiple devices.

Centralized log management is a solution for all the issues I described here. For example, tools such as Logstash help you make sense across a variety of log formats. A central log collector provides the following advantages:

- One place to store all types of log messages
- A place to store backups of logs
- A place where you can analyze the log data

Centralizing Logs with Syslog

It's common for applications to write logs to various locations on disk storage. Often logs contain various information. Over time, logs may run out of space, so you need to make sure you have enough space to hold the logs and put in log rotation policies. Some logs may contain sensitive information such as credit card payments so you need to make sure they remain secure as well.

A good strategy to gain an upper hand over voluminous logging is to store them all in a central location. You can then backup, secure and work with the log data. Syslog is the standard *nix based system's logging mechanism and there are several implementations the standard system.

In order to use Syslog, you set up a central Syslog server and configure client logs to use the server. The syslog clients run on the system that's generating the logs and you can configure Syslog to send its logs to a local file, or to a centralized syslog server through rsyslog. The syslog daemon (syslogd) is a server-side process that receives logs from the syslog clients.

The syslog system has various drawbacks, such as lack of guaranteed delivery of messages (uses the UDP protocol and not TCP), insecure transport of log messages, and sometimes missing the source of the log information.

While traditionally the syslog (syslogd) daemon was used to provide the logging service, in recent times, syslogd has been slowly supplanted by other logging software. The open source rsyslog package and the logging component of systemd, called systemd-journald (journald) have become popular. For Ubuntu and many other Linux implementations, rsyslog is the default syslog package.

Using Rsyslog for Logging

Rsyslog, in the words of its creators, is the “rocket-fast system for log processing”. Rsyslog is a high performance log processor that's available on most Linux distributions, and is the default logger for the RedHat and Fedora systems. Rsyslog started out as a regular syslogd, but has evolved into a versatile logging tool that can accept inputs from a variety of sources, transform them, and output the results to various destinations.

You configure and manage rsyslog in a manner similar to how you do the traditional syslog daemon. The rsyslog daemon is compatible with the traditional syslog daemon but has enhanced capabilities.

Rsyslog can deliver over a million messages per second when processing messages to local destinations. When you send to remote destinations, performance is still lightning fast.

Since several Linux systems use compatible daemons, the rsyslog daemon offers a standardized method of logging. This makes rsyslog quite useful in large heterogeneous environments.

Rsyslog can send output to various destinations such as text files (/var/log directory usually), SQL databases, other hosts, etc.

An rsyslog log entry consists of a single line that contains the data, time, process name, PIP and the message from the process. Optionally you can use the logger command to generate entries in the logs.

Rsyslog Templates

Templates play a critical role in rsyslog based logging. Using the templates lets rsyslog generate logs that are formatted just like the traditional syslogd messages. By default, rsyslog uses templates that support the syslog message format.

Here's a sample template that supports the syslogd message format:

```
$ template TraditionFormat, "%timegenerated%, %HOSTNAME% %syslogtag% %msg%\n", <options>
```

Rsyslog Rules

You can use various filter criteria to create rules in the rsyslog.conf file. You can use filter criteria based on various criteria such as facility and priority (these are traditional filters), property (for example, HOSTNAME, msg) or use rsyslog's built-in scripting language, RainerScript, to construct custom filters.

Systemd-Journald (Journald)

Systemd-journald is a part of systemd that collects and stores log data. As with rsyslog, systemd-journald can trap logs from various sources such as user processes, the kernel, and standard output (STDOUT) and standard error (STDERR) of various system services.

Instead of using traditional flat text files, the journald daemon stores log data in structured and indexed journals, enabling you to use the logs in various ways. Many Linux distributions have made systemd-journald an integral part of their logging infrastructure, along with traditional log daemons such as rsyslogd.



Although you can run both systemd-journald and rsyslogd side-by-side, it's a good practice to use a single standardized logging framework.

Setting up a Remote Syslog Server

You can set up an easy to use log management system with rsyslog, ElasticSearch, and Kibana. In this system: **Rsyslog will receive the syslog data and format it, before sending it to ElasticSearch** ElasticSearch indexes the data and stores it to facilitate your searching ** Kibana provides a nice web UI where you can search the log data
NOTE: Rsyslog is easy to configure, and is installed on most Linux distributions.

SYSLOG-NG

SSYSLOG-NG grew out of syslog to provide users with more options and features. Syslog-ng seeks to eliminate the problems with classic syslog, such the lack of guaranteed delivery of messages and insecure transmission of log data.

The use of TCP instead of UDP means that syslog data can be reliably transmitted. Syslog-ng also natively supports the TLS protocol, which means that syslog messages are securely transmitted.

A big advantage offered by syslog-ng is that it lets you send syslog messages to a relational database, making reporting and log analysis a breeze. Syslog-ng is also geared towards handling very high message rates.

Syslog processes messages with a fast and flexible pattern matching engine to extract useful information from the raw message body. For example, syslog-ng takes the first message (login message from an OpenSSH Server and describes it with the second syslog-ng pattern:

```
Accepted password for sam from 10.48.0.248 port 37184 ssh2.
```

```
Accepted @QSTRING:authentication_method: @ @QSTRING:protocol_version: @. for @QSTRING:username:
```

There are several advantages to the pattern matching method used by syslog-ng. As the example shows, syslog-ng parses the log message using special variables. The names it assigns to the various parts of the log message become syslog-ng macros which you can use in filters, file names, and database column names.

Syslog-ng patterns are fast, and are easier to read and main. You can also trigger actions such as sending alerts when log messages match a pattern.

Rotating Logs with Logrotate

Linux's Logrotate utility lets you rotate, compress and remove logs based on configured Schedules. It's quite simple to configure Logrotate. The following example shows how to by creating a file in /etc/logrotate.d. This configuration rotates logs daily, keeping 31 days of logs, and compress the logs as well with gzip:

```
/usr/local/testapp/logs/*.log {  
    daily  
    rotate 31
```

```
        copytruncate
        compress
        notifempty
    }
```

So far, I've been discussing policies to help manage log acquisition and retention, and archiving of logs. However, you keep logs for a reason – to extract the data in those logs! Let's turn to a review of log analysis in the following sections.

Log Aggregation

Manual searching of logs through tools such as `tail`, or `grep` are out of the question when you are dealing with a large number of servers. For example, a request to a front-end server may result in numerous web service calls and you simply can't sift through all these logs on all the servers.

Possible Strategies

You need to collect all the logs in a central location, and you can do it different ways:

- Log to a data store directly instead of logging to a file. While this sure makes it easy to search through logs, the unavailability of the logs, or a slowdown of the data store will wreak havoc with the applications that send log data to the data store.
- You can have applications write to log files, and have the logs shipped to a central log server, with the help of log-forwarding agents running on the servers. Centralized log servers such as `Fluentd` are products that let you do this. Log forwarding agents `tail` the log files, perform filtering and transformations, and forward the resulting logs to the central log server. You can also standardize time formats to a single zone to avoid the hassles of dealing with logs in different time zones and formats.
- While shipping logs to a central server is pretty useful, searching through the logs is going to be time consuming. The next step then is to employ a log-indexing platform to speed up searches and make the logs easily available. Following then is ideally how you'd set up this log processing strategy.

Log Aggregation Workflow

Following are the steps in a log aggregation workflow.

- Install log-forwarding agents on all the servers from where you want to collect logs

- Configure the log-forwarding agents to let them know which logs to forward, and how to filter and transform the logs, and where to send them
- Logs are sent to the search engine servers where they are stored and indexed to enable easy searching
- A log processing platform provides a web-based interface to make searching and visualization easier

Log Aggregation Alternatives

Configuring and managing log aggregation involves significant effort, so if you can use a hosted solution at all, that should be the best way for you to go. Splunk is a tremendous log aggregation tool with great features, but its cost makes it out of question for many startups. Also look into Loggly, which provides useful logging functionality.

Using a hosted logging solution means you'll also shared your application logs with others and if you can't abide by it, you must look into self-hosting the logging solute, through an open source product. Here, Logstash, a powerful, scalable log-indexing platform that uses agent to ship logs to the Elasticsearch search engine, is a logical choice. Logstash comes with a web interface called Kibana that lets you perform free text search and visualize logs.

Effective Log Management and Log Analysis with Elasticsearch, Logstash, and Kibana

ElasticSearch, Logstash and Kibana (ELK) is a comprehensive log analysis solution, built with the three components that make up this product's name.

The ELK stack isn't something that just lets you store logs – it's designed especially for log analytics. It consists of the following three basic components:

- ElasticSearch: A powerful search and analytics engine
- Logstash: The central logging management provider that processes incoming logs and ships/forwards the logs from multiple servers to a source you choose, such as ElasticSearch
- Kibana: Versatile graphical web interface for querying log events, with great visualization capabilities

Logstash ships log data to Elasticsearch, which indexes the data in a searchable data store, and Kibana enables you to visualize the indexed data in a graphical format for analyzing the logs.

Companies such as LinkedIn have chosen ELK as their logging solution for the following reasons:

- It's a horizontally scalable solution – just add nodes when you need to do more work
- It is very fast and quite close to real time
- It's inexpensive since it's open source
- It has a sizable user community to go to for help



Although I discuss Elasticsearch as part of the ELK stack, Elasticsearch (and Logstash) is an independent tool that you can use by itself to explore data.

The triumvirate of Elasticsearch, Logstash, and Kibana form a powerful tool for log analysis, and are often referred together as the ELK stack. ELK helps you efficiently ship, analyze, and present data from your logs. Here's what the three components of an ELK stack do:

- Elasticsearch is a search server – it stores and indexes the log data
- Logstash is for cleaning and shipping data
- Kibana is a tool that you use for viewing and analyzing the logs

Each of these three tools is a powerful standalone tool in its own right – together, they make one check of a log managing framework. Logstash, Elasticsearch, and Kibana work together: You can use Logstash to impart order to your log data. Elasticsearch queries help you examine the data, and Kibana helps you display the data through its dashboards.

Searching Engines and Indexing - Elasticsearch

ElasticSearch is a real-time distributed search and analytics engine built on Apache Lucene, an open source full-text search engine. Elasticsearch lets you explore data at astonishing speeds and helps you extract actionable knowledge from your data. It has full-text search, structured search, and analytics capabilities.

ElasticSearch is a NoSQL distributed data store supported by a company named Elastic. As a database, it lets you store and search data. ElasticSearch makes it easy to search (unlike Lucene), with its powerful RESTful API. It also extends Lucene's capabilities by offering real-time analytics for both structured and unstructured data distributed across multiple servers.



Marvel is a free (for development use) management and monitoring tool that comes with Sense, an interactive console called Sense which makes it easy to work with Elasticsearch from a browser.

Two key things are baked into Elasticsearch's architecture: its distributed nature and its ability to easily scale. The distributed architecture protects against failures since the data is replicated.

ElasticSearch is a schema less data store, with data stored in the JSON format and partitioned into shards. An ElasticSearch index consists of multiple shards, which are the tiniest elements in Elasticsearch.

The basic operational unit of Elasticsearch is not a server, but a cluster. Although you can have a single node "cluster" in development, for production you want to set up a minimum of two nodes.

Talking to Elasticsearch

If you're using Java, you can use one of two built-in clients – the node client and the transport client - to talk to Elasticsearch. If you're using other languages, you can talk to Elasticsearch via a RESTful API/, through any web client. A request to Elasticsearch is similar to how you make any other HTTP request, with the following structure: `curl -X<VERB> <PROTOCOL>://<HOST>/<PATH>?<QUERY_STRING> -d <BODY>`

Basic Concepts of ElasticSearch

Let's review the key concepts and terminology of ElasticSearch, as a prelude to understanding how it stores index data.

Documents and Fields

Elasticsearch is document oriented – it stores entire objects or documents, and indexes the document's contents so it's easily searchable. By default, all data in every field is indexed, to facilitate fast retrieval. Unlike in a relational database where you store rows of columnar data, in Elasticsearch, you store, index, search, and sort documents.

A document is a JSON document which has a type and an ID to uniquely identify it. The term document in Elasticsearch refers to a top-level, or root object that's stored in the JSON format in Elasticsearch, under an unique ID. Most application objects can be serialized into JSON objects, with keys and values. A key is the name of a field and value can be anything – a string, a number or another object.

Following is an example of an Elasticsearch document:

```
{
  "_index" : "oreilly",
  "_type" : "linux",
  "_id" : "1",
  "_version" : 1,
  "found" : true,
  "_source":{
    book_name : "modern linux administration"
  }
}
```

A field is the basic unit inside a document, - it's a key value pair, such as:

```
book_name : "modern linux administration"
```

Indexing

Strong data in Elasticsearch is called indexing. Indexing a document stores the document and makes it searchable. An Elasticsearch index (analogous to the traditional term “database”) is a set of related JSON documents, with the indices broken up into shards. An index is a collection of documents that share similar characteristics. An index contains several logical partitions, called types, with each partition containing multiple documents. For example, Facebook indices can use post and comments as some of the index types.

Each document contains multiple fields and all documents belong to a specific type, and the various types are part of an index. An Elasticsearch cluster consists multiple indices. If you were to compare a relational database to Elasticsearch, here's how they stack up together:

```
Relational DB: Databases => Tables => Rows      => Columns
Elasticsearch: Indices  => Types  => Documents => Fields
```

Following is the URL to search and query indices:

```
http://localhost:9200/[index]/[type]/[operation]
```

A document is a JSON document stored in an Elasticsearch index.

Logstash for example, by default sends JSON documents to Elasticsearch with the default index pattern `logstash-%{+YYYY.MM.dd}`. This format partitions the index by date so you can search by day and delete older data easily.

Shards (Primary and Replica) You partition your documents into multiple containers, called shards. A shard is where the index data is actually stored. An index is really merely a logical namespace and it points to one or more shards, which are physical entities. Shards are the low-level worker units that each store a portion of an index's data. It's through allocating shards to various nodes in the Elasticsearch cluster that Elasticsearch distributes data around the cluster.

Elasticsearch automatically balances your data across a cluster's nodes by migrating shards among the nodes.

There are two types of shards: primary and replica shards. Each document in an index belongs to a single primary shard. The number of primary shards thus determines the total amount of data that an index can hold. You can start with just a single primary shard, but as read and index requests increase over time, you'd need multiple primary shards, on different nodes. In order to avoid having to constantly re-index data as you add new primary shards, it's a good idea to start with multiple primary shards.



You can move shards from one node to another in case of a node failure, or when you add nodes to the cluster. There are primary shards for each index. Each document in an index is stored on one primary index and multiple replica shards.

By default there are five primary shards for each index and each document in an index is stored on one primary shard and multiple replica shards.

The replica shards are for failover and provide redundancy for the data, and for load balancing purposes, by servicing read requests such as searching or the retrieval of documents. When a primary shard conks out the replica shard is promoted to a primary status to ensure the cluster continuity. When you index a primary shard, you also index the replica, and multiple replicas dramatically increase the search performance.

Nodes and Clusters

ElasticSearch groups indices into nodes – a node is a single running instance of ElasticSearch. Each node belongs to a cluster and by default all nodes join the “elastic-search” cluster. There are three types of nodes:

- **Master nodes:** these are lightweight nodes meant for managing the cluster. These nodes don't hold any data, index or perform any search requests. They perform cluster-wide operations such as the creation and deletion of indexes, or adding/removing nodes from the cluster. Multiple master nodes are recommended, to ensure redundancy. If you've a large cluster, use three dedicated master nodes, with one acting as the primary master and the other two serving as backups. **NOTE:** Any one of a cluster's nodes can become the master node.
- **Data Nodes;** These nodes contain the data (index documents) and serve all index and search requests by performing the searches on the indexed documents. You can increase the number of data nodes to scale your search capacity.

- Client nodes (also called routing nodes or load balancer nodes): these nodes perform load balancing during some processing steps, or route requests for index documents to the appropriate nodes. They share the burden of a data node's workload such as spreading research requests across nodes and gathering the response output. High volume searches or index operations benefit from having these type of nodes.

Understanding the Elasticsearch APIs

Although in an ELK setup Logstash and Kibana will take care of the Elasticsearch indices, it's important to understand how Logstash and Kibana use Elasticsearch RESTful APIs to perform the document storing, index management, and query searches.

The following curl command shows how to query the Elasticsearch cluster from the command line:

```
curl -XGET 'http://localhost:9200/logstash-2016.06.14/_search?pretty'
```

When you specify the Elasticsearch output plugin in the Logstash configuration file, it automatically creates the Elasticsearch indexes. However, it's good to know how to use the Elasticsearch API to create an index. Here's an example:

```
curl -XPUT 'localhost:9200/oreilly/?pretty'
```

Here, the index name is oreilly.

An Example Showing how to Create and Retrieve a Document

Here's a more concrete example that shows how to index documents and then query them. The following PUT command will do the following:

- Index a document with details for a single employee
- The document will be of the type employee
- The type employee is part of the index named mycorp

```
PUT /megacorp/employee/1
{
  "first_name" : "John",
  "last_name" : "Smith",
  "age" : 25,
  "about" : "I love to go rock climbing",
  "interests": [ "sports", "music" ]
}
```

Notice how you could directly index a document while creating the document, without having to specify any index creation specifics – Elasticsearch handles all that

for you behind the screen. In order to retrieve the data you've stored, simply execute an HTTP GET request with the address of the document (index, type, ID), as shown here:

```
GET /mycorp/employee/1
```

And here's the response that contains the original JSON document you've created, as well as some metadata about the document:

```
{
  "_index" : "megacorp",
  "_type" : "employee",
  "_id" : "1",
  "_version" : 1,
  "found" : true,
  "_source" : {
    "first_name" : "John",
    "last_name" : "Smith",
    "age" : 25,
    "about" : "I love to go rock climbing",
    "interests": [ "sports", "music" ]
  }
}
```

I've shown how to use the HTTP verbs PUT and GET to create a document and to retrieve the data. Similarly, you can use the DELETE verb to delete a document, and the HEAD verb to check if a document exists.

Searching Data

While GET retrieves a specific document, it's not as much fun as a search. You can employ the `-search` endpoint and a query-string search to search for documents, as shown in the following request, which returns the top 10 results:

```
GET /mycorp/employee/_search?q=last_name:Sam
```

The results of this query will include all the "sam" employee records.

A query-string search works fine for adhoc searches, but isn't powerful. ElasticSearch uses indexes to retrieve data, but its Query Domain Specific language based on JSON, called Query DSL, is what makes ElasticSearch querying so powerful. Query DSL is a rich, flexible query language that lets you build complex queries. ElasticSearch uses DSL to get querying results in the desired format. For example, using our Google stock ticker data, we can write a query using DSL to sort the data on the OPEN field, as shown here:

```
curl -XPOST 'localhost:9200/Logstash*/_search' -d '{
  "query" : {
    "match_all" {}
  },
  "sort" : { "open" : { "order": "desc" } },
}
```

```
    "size" : 3
  },
```

Instead of using query-string parameters, this DSL uses a request body built with JSON and uses a match query. It's useful to know that Elasticsearch has a great number of highly useful plugins that help you analyze the cluster, and easily execute complex queries. You can also add filters to execute more structured searches.

Elasticsearch is way more powerful than the search capability of any relational databases, which mostly are limited to finding matches for search criteria. Elasticsearch can perform powerful searches within full-text fields, and show you the most relevant results first.

Elasticsearch Marvel and Shield

You can integrate two additional components with ELK:

- **Marvel:** Let's you monitor the Elasticsearch deployment through Kibana
- **Shield:** Provides security features for ELK such as authentication, RBAC, encryption of data and audit logging.

Logstash

Production environments generate voluminous logs every day. You can't expect to analyze and parse the logs using traditional tools such as `tail` and `grep`. Centralization of all logs becomes a key imperative in larger environments. You can start off by using a dedicated rsyslog server and send all logs to that server.

However, as the log volume rises, you'll find yourself reaching out for more storage capacity. You'll also realize that archiving the data for extracting information isn't easy either. Correlating log data with multiple formats to generated events might turn out to be very difficult to do. At this point, your needs are the following:

- Parse logs efficiently
- Search the logs efficiently
- Index processed logs
- Expose the logs so folks can easily get at them

Logstash is an open source logging agent or tool that solves all these problems for you. Logstash acts as a data pipeline that collects, parses, and analyzes a huge variety of structured and unstructured data and events across a variety of system data – not only log data, but other types of data as well. Logstash sources log data from various types of nodes, and delivers it securely to a centralized log server for storing, index-

ing, and analyzing. Once you consolidate your logs with Logstash, you can gain easy access to the logs.

Logstash Plugins

Logstash comes with numerous useful plugins for input, filter, and output. The Logstash configuration file is nothing but a series of input, filter, and output plugins and the properties you specify for each of the plugins. The plugins are what determine how Logstash parses, processes, outputs, and ships the data in the format you want.

You can view all the available plugins for Logstash by issuing the following list command:

```
$ bin/plugin list
```

The Following sections describe the most important plugins.

Input Plugins

The input plugin helps you configure the set of events that will feed the logs events to Logstash. Here are some popular Logstash input plugins:

- **file:** Streams events and log lines from a file
- **redis:** Streams events and logs from a Redis instances
- **stdin:** Streams syslog messages over the network via standard input
- **lumberjack:** Receives events through the lumberjack protocol (lumberjack is now known as the Logstash forwarder)
- **twitter:** Twitter plugins are useful in reading events from the Twitter streaming API, and are useful in analyzing a Twitter stream for uses such as sentiment analysis and trending topics analysis.
- **elasticsearch:** Reads from an ElasticSearch cluster using the results of a search query

Filter Plugins

Following are some of the commonly used Logstash filter plugins:

- **date:** Used to parse date fields in events (used as Logstash timestamp fields)
- **drop:** Drops all data from events that matches filter conditions
- **grok:** Filter and parse unstructured data from logs and events into a structured format
- **mutate:** Renames, removes, and replaces fields in events

Output Plugins

Logstash's output plugins help integrate incoming events with a variety of destinations, of which Elasticsearch is one. Output plugins include the following:

- **file:** Writes events to a file
- **csv:** This plugin is used to write the output to a CSV file
- **e-mail:** Sends emails based on some conditions in the output data, and handles failure conditions
- **redis:** Writes events to a redis queue and acts as a broker for several ELK implementations
- **kafka:** Writes to a Kafka topic
- **ganglia:** Ganglia is a monitoring tool (I discuss this in Chapter 11) that uses the output to send metrics to Ganglia's gmond service based on log events
- **mongodb:** Writes to the MongoDB database
- **elasticsearch:** The most important plugin, and is used in the ELK stack. Kibana analyzes the data that Logstash outputs to Elasticsearch.



In addition to the input, output, and filter plugins, Logstash also comes with several Codec plugins that are used to encode/decode events sent or received by Logstash. In addition, you can write custom Logstash plugins in Ruby.

Logstash helps you collect, parse and modify any data format into a common format which you can then use to build analytical systems across applications. Log formats, data and time formats in logs vary across applications since the logs are often customized for an applications. Logstash comes with a large number of plugins that help overcome all the typical problems with disparate log formats and lets you build visualizations on top of any type of log data.

Logstash under the Hood

Logstash can receive various types of logs, such as TCP, UDP, Syslog, and files, and processes and filter them according to filters that you define. Pushing logs to Logstash is called shipping the logs.

Components of the Logstash Architecture

Logstash consists of the following basic components:

- The Logstash shipper: The shippers run on separate nodes from the Logstash server, and transfer the log and event files to the Logstash server.
- The Broker: Receives the collected event logs from the shippers and buffers them. Redis is commonly employed as a broker to hold data captured by agents running on the Logstash shipper. However, many High availability Logstash environments use AMQP (RabbitMQ) for message brokering.
- The Logstash Indexer Indexes events within the Logstash server. Later, ElasticSearch will provide a full-text index to make the logs searchable.

The Logstash Workflow

Logstash configuration involves the following three main topics:

- Input: defines how events are generated and sent to Logstash
- Filters: define the way you manipulate and customize events
- Output: defines how Logstash sends events to external systems

The Logstash agent consists of a processing pipeline with three distinct stages:

- Inputs generate events
- Filters modify the events
- Outputs ship the events to places such as ElasticSearch

Here's a typical Logstash configuration file that shows the input and output block configuration:

```
input {
  redis {
    host => "47.147.50.240"
    type => "redis-input"
    data_type => "list"
    key => "logstash"
  }
}

output {
  elasticsearch { host => "logstash.myserver" }
}
```

The input block specifies that the Redis broker plugin will listen to events on the specified host interface. The broker labels the received events with the redis-input type.

The output block can send the data to any destination, but in this example, sends the output to Elasticsearch. It uses a separate plugin named `elasticsearch` that's in charge of sending log events from Logstash to Elasticsearch.

Configuring Logstash – Some Examples

You can configure Logstash to filter events, process web server logs and syslog messages. You can specify conditionals to control the events that are processed by a filter or output block.

Configuring Logstash

You create a configuration file to specify the plugins and their settings, when you configure Logstash. You specify the event files and use conditionals to process the events only they meet certain criteria. Here's a simple Logstash configuration file:

```
input { stdin { } }
output {
  elasticsearch { hosts => ["localhost:9200"] }
  stdout { codec => rubydebug }
}
```

Typically, a Logstash config file has the following structure, with a separate section for each type of plugin you want to add to the processing pipeline.

```
input {
  ...
}

filter {
  ...
}

output {
  ...
}
```

Configuring the Plugins

The Plugin configuration consists of the plugin name following by a block that contains the settings for the plugin. In the following example, the `file` plugin is configured twice.

```
input {
  file {
    path => "/var/log/messages"
    type => "syslog"
  }

  file {
    path => "/var/log/apache/access.log"
```

```

    type => "apache"
  }
}

```

Configuring the Filters

Filters let you pick and choose from your log data to suit your needs. The following Logstash configuration sets up the grok and date filters.

```

input { stdin { } }

filter {
  grok {
    match => { "message" => "%{COMBINEDAPACHELOG}" }
  }
  date {
    match => [ "timestamp" , "dd/MMM/yyyy:HH:mm:ss Z" ]
  }
}

output {
  elasticsearch { hosts => ["localhost:9200"] }
  stdout { codec => rubydebug }
}

```

Logstash uses the grok filter to grab the correct messages and break them into chunks of information that help you query and analyze the log data. Logstash comes with a bunch of useful grok patterns.

The date filter parses out a timestamp and uses the timestamp as the event's timestamp, regardless of when you collect the log data.

Suppose you have a configuration file with the following input configuration:

```

input {
  file {
    path => "/tmp/access_log"
    start_position => "beginning"
  }
}

```

Logstash opens and reads the access_log file, processing each event it sees. Logstash captures the lines in this file and processes them as events and typically, stores them in Elasticsearch.

How Logstash gets its Logs

You can ship logs to the central monitoring server in the following ways:

- Use the Logstash agent running on the hosts to generate logs and ship them to the Logstash server

- Use syslog by activating the syslog plugin for Logstash. Logstash supports the syslog daemons Rsyslog, syslog-ng, and syslogd to send log messages to the Logstash server.
- Use the Logstash forwarder (previously called Lumberjack) to send log messages to Logstash. You need to install and configure a Logstash forwarder on each node in your environment. The forwarders can reduce the resource consumption and bandwidth usage. There's also an enhanced log forwarding tool to ship logs, called Log-courier

Kibana

Kibana is an open source analytical product that you can use to query, display, and report on data. It offers a GUI that lets you explore data, both for ad-hoc data exploration purposes as well as create powerful and detailed dashboards.

The true power of Kibana comes from its ability to use the Elasticsearch query language. Kibana was built and developed by Elastic. It's a visualization platform that's built to work on top of Elasticsearch and take advantage of Elasticsearch's capabilities. Logstash streams data into Elasticsearch, from where Kibana access it and creates visualizations out of it. Kibana uses the powerful search and index capabilities of Elasticsearch, to display graphics to you. It exposes the data through histograms, charts, graphs, tables, and similar formats.

In addition to letting you visualize data in various formats, Kibana lets you create dashboards that help you query data in real time. The dashboards, which are interfaces to the base JSON documents, help you understand the data without any need for coding.

There are four main tabs in the Kibana interface:

- Discover: Lets you perform free text and field based searches
- Visualize: Enables you to build visualizations such as pie/bar charts.
- Dashboard: The dashboard is a collection of saved searches and visualizations, and lets you filter the data and perform data aggregations
- Settings: Allows you to configure index patterns, data types, and scripted fields

It's important to understand that Elasticsearch drives many of the key functionalities of Kibana, which serves as the front end to Elasticsearch. For example, all of Kibana's visualization capabilities are built on Elasticsearch's aggregation feature. Similarly, Kibana's metrics and buckets aggregation feature uses Elasticsearch's aggregation functionality.

Buckets distribute subsets of indexed documents into groups based on various criteria, and offer similar functionality as SQL's GROUP BY function. Metrics can then be applied to the bucketed documents. Some common buckets are histograms and ranges. Metrics are computations that are performed on the values of fields inside each bucket, such as the Count, Average, Min, and Max metrics.

A Simple Example Showing How ELK Works

Let's use a simple example to understand how ELK works in practice. I use Google's stock price data covering the period from January 1, 2015 to December 31, 2015. The work begins with logs flowing into Logstash, from where they're sent to Elasticsearch for indexing, and finally to Kibana for visualization and dashboards.

The Input Data

The input data contains the typical stock related columns such as Open Price, Close Price, and Volume. Each row in the data set represents the stock price data for a specific day, separated by commas, as shown here:

```
$ head G00G.csv
2014-12-31,531.25244,532.60236,525.80237,526.4024,1368200,526.4024
2014-12-30,528.09241,531.1524,527.13239,530.42242,876300,530.42242
2014-12-29,532.19244,535.48242,530.01337,530.3324,2278500,530.3324
```

Configuring the Logstash Input

I'm using the file input plugin to read the stock data. The file plugin streams each line as an event from the source data file. The Logstash configuration for the input dataset looks as follows:

```
input{
  file{
    path =>"/test/Logstash/input/G00G.csv"
    start_position =>"beginning"
  }
}
```

Limiting the Input Data through Filtering

Often log data is too detailed, and you need to perform some type of filtering on that data to get the fields you really need.

A Logstash filter plugin will help you control the size and format of the input data by performing intermediate processing on the input events. You can make use of Logstash's conditional filters to select fields based on specific conditions.

Since our basic input is in the CSV format, I specify a simple csv filter, as shown here:

```

filter {
  csv
  column => ["date_of_record", "open", "ihg", "low", "close", "volume"]
  separator => ","
}

```

I need to use Logstash to perform some more intermediate processing. First I need to specify a date filter to handle the date format:

```

filter {
  date {
    match => ["date_of_record", "yyyy-MM-dd"]
    target => "@timestamp"
  }
}

```

I also need to specify the data type of the fields in my CSV, which by default, are all of the type string. I need to convert the data type to integers so I can perform computations and comparisons of the data. I use the mutate filter for converting the fields to a specific data type. Here's how my mutate filter looks like:

```

mutate {
  convert => ["open", "float"]
  convert => ["high", "float"]
  convert => ["low", "float"]
  convert => ["close", "float"]
  convert => ["volume", "float"]
}

```

The mutate filter's convert functionality lets me convert the price and volume fields to integers.

Outputting the Data to Elasticsearch

Now that I've my input data nicely filtered and with the correct data types for the fields, I now need to send the data to Elasticsearch. Once the data gets to Elasticsearch, it'll index the data based on the fields and you can then use Kibana to query and visualize the data.

The following simple output block is sufficient for now, to configure the sending of the output from Logstash to Elasticsearch.

```

output{
  elasticsearch {
    host => "localhost"
  }
}

```

Using Elasticsearch in a Production Setting

Typically, in an ELK data pipeline, logs are shipped to a centralized Logstash indexer through the Logstash shipper that runs on the source generating the logs. The Log-

stash indexer will then output the filtered logs to an Elasticsearch cluster, from where Kibana will query the data to display visualizations and build dashboards using the log data.

Several large companies such as LinkedIn and Netflix have been using ELK for building centralized logging stores and extracting actionable information from their voluminous logs. Following are some considerations for a production implementation of the ELK stack.

Scaling the ELK Solution

You can scale each of ELK's three major components to meet increasing demand. In a production setting, it's a good idea to use multiple Elastic search nodes – say three primary nodes, with one of them serving as the primary and the other two as back-ups. For high volume environments you can add load balancing nodes (also called routing nodes).

In addition to multiple Elasticsearch instances, you can also add multiple Logstash and Redis instances, and even multiple Kibana instances. So, your scaled architecture will look like the following:

```
Logstash (Shippers) => Redis (Cluster) => Logstash (Cluster) => Elasticsearch (Multiple Nodes) =>
```

Avoiding a Data Loss

A message broker is a requirement in a production setting, in order to avoid losing important log data. Logstash could be slow in indexing data to Elasticsearch. A message broker such as Redis can help you deal with large amounts of data when Logstash falls behind. The message broker buffers the data to help Logstash tackle large loads.

When you place a Redis server in front of the Logstash machine, that'll serve as the entry point for all log events into the ELK system. The Redis server will buffer the log events until TLK can index it. Logstash slows down when ElasticSearch is busy indexing stuff. Redis can help by holding documents that it can push to Elasticsearch later on, thus avoiding losing the log data due to a system overload.

Should Logstash's indexing to Elasticsearch fail for any reason, the broker will hold the data in its queue, thus providing resiliency for the data. You can also use ZeroMQ or RabbitMQ as alternative brokers instead of Redis.

When using a broker such as Redis, the ELK architecture looks like the following:

```
Logstash (Shipper) => Redis (Broker) => Logstash (Indexer) => Elasticsearch => Kibana
```

Securing Data

You must protect Elasticsearch indices from unauthorized access. In addition, you must protect the Kibana Dashboard, which you can do by setting up a reverse proxy such as Nginx to access Kibana. Kibana allows for SSL encryption for client requests.

To protect Elasticsearch indexes, you can create roles for Kibana in Shield (paid service offered by Elastic), and limit the access you want to grant to Kibana's users. You can also use the SearchGuard tool to secure Elasticsearch.

Scaling ELK

You can deploy ELK in your data center, or in the Amazon or Google Cloud. Regardless of where you deploy ELK, it's a good idea to make it scalable – scaling ELK means scaling all its components:

- Redis: Run multiple Redis servers, and use Redis labs as a service
- Logstash: In order to push to ElasticSearch, ideally you should run a Logstash instance on each Redis server.
- Elasticsearch: Run ElasticSearch instances on at least three master nodes (to get a consensus and avoid a “split-brain” situation, when there's a dispute between two nodes as to which is the actual master). You should have at least two data nodes so you can replicate data at least once.

Log Analysis with Splunk

Splunk is a great commercial tool for log analysis that's mature and enterprise focused. Splunk handles many types of data very well so it's not limited to analyzing logs. In recent years, the open source ELK stack has supplanted Splunk as the leading log analytics product. Splunk was started in 2003, so it's been around for a while. Splunk can consume logs from hundreds of third party applications. It employs sophisticated formulae, deep-dive searches, and keyword analyses to summarize the data into useful formats for you to act on.

Splunk is mostly used as an on premise tool, and there's also a hosted version of Splunk Enterprise, called Splunk Cloud. Splunk's pricing levels are designed for enterprise customers.

Splunk indexes log files (and other text data) that includes timestamps. This lets you perform real-time searches to find exceptional events, as well as perform root cause analyses. You can use Splunk as an operations dashboard, and configure it to send notifications.

Splunk is technically a time series indexer. It was originally designed to process machine data and it has three key functionalities:

- **Collecting the Data:** Splunk can collect data by monitoring changes to files or directories on a real time basis, as well as from network ports and programs, or relational databases.
- **Indexing the Data:** Splunk breaks the data it collects into events (lines of data) and updates its high performance index that point to the stored data.
- **Searching and analyzing the data:** You can use the Splunk Processing Language (SPL) to search for data and get the results in the form of alerts or reports.

Scaling and High Availability

In a test project, a single Splunk instance can serve as both the indexer and the searcher, but in real life you set up multiple servers to handle the indexing and search functionality. In addition, you can set up clustering to increase the resilience by automatically failing over the indexers.

Creating multiple Splunk instances to handle the data sent by hundreds of forwarders lets the forwarders distribute data across the instances, providing scaling benefits, as well as a failover capability, since forwarders automatically switch to other instances when a Splunk instance becomes unavailable. Splunk's scaling capability is highly flexible, since you can scale the indexing and search functions separately to suit your needs.

If a searcher or indexer instance becomes unavailable, users can't access the slice of data that's in that Splunk instance. You can create a cluster of Splunk instances so you can replicate the data over multiple Splunk indexer instances. When an indexer crashes the instances with the replicated data provide the data for the search. A cluster master coordinates the replication, failure, and other cluster related work. The nodes where the search and indexing occurs are called the peer nodes. The Splunk master only performs the coordination work and not any indexing or searching work. For example, you may set up 1 master, 6 indexers, a set of forwarders and one search head for a Splunk cluster.

Working with Data

The main Splunk user interface is the Search app, but Splunk is designed as a platform for others to build applications that extend Splunk's functionality. Third-party developers can build extensions through the Apps (using the Splunk Apps package), or through Technology Add-ons. Add-ons are simple components and apps contain additional functionality as well as standard features such as saved searches and dashboards, as well as their own user interfaces.

Sources of data

Splunk can collect and analyze data from various sources such as files and directories, and networks. You use the Splunk for Unix Technology Add-On for getting data from UNIX and Linux systems into Splunk. Splunk can capture the data inputs from various sources such as the Linux file system, as in /var/log and /etc. and also from scripted inputs such as your own shell scripts.

You can use the Splunk web which is Splunk's standard UI, the Splunk CLI, Apps or Add-ons, and configuration files to work with the various types of data sources.

The Need for Forwarders

You can't expect to run Splunk on all of its data sources, such as operating systems, databases, middleware, and application programs. Forwarders let you capture remote data. A Splunk forwarder is the same as any Splunk instance, but with just those components that are required forward data. The forwarders gather data and forward it to the Splunk instances or indexers, which will load the remote data.

Forwarders provide a bunch of benefits to Splunk: they can buffer data at the remote locations, which is helpful when the main Splunk instance goes down. They also support the secure transmission of remote data, with data acknowledgements. They let you scale and improve performance with their support for load balancing.

Splunk and Data Wrangling

Splunk is a great tool for working with large amounts, of data. Data wrangling is an important component of data science projects, and refers to the process of converting or mapping data from one form into another more useful format. Splunk is a great tool that lets you become a data wrangler and extract useful knowledge from your data.

Using the Splunk Processing Language

SPL lets you easily analyze the data you load into Splunk, through various types of commands that I summarize in the following sections.

Reporting Commands

Reporting commands help you answer questions such as which your top browsers are, and which are your top 5 IP addresses. For example, you can use the top command with one or more fields available in a Splunk event to performing the reporting queries. To get the list of your top five IP addresses that made requests to your website, you issue the following search command:

```
sourcetype=access_combined_wcookie  
| top limit=5 clientip
```

The pipe (|) sign in my search passes the result from the first Splunk command as input to the next search command. Since the clientip field holds the IP addresses from where the individual client requests were sent, I get the information that I'm after.

Sorting Commands

The sort command is the main SPL sorting command. You can issue a query such as the following and pipe the results to a sort command to get a list of your most popular products in the descending order:

```
sourcetype=access_combined_wcookie action=purchase
| stats count by productId
| sort -count
```

Filtering Commands

Filtering commands reduce the original set of events into a smaller set and thus help speed out an analysis. Splunk's head command lets you process a subset of events. The example shown here gets you the most popular products in the last 100 events:

```
sourcetype=access_combined_wcookie action=purchase
| head 100
| stats count by productId
| sort -count
```

Grouping Commands

Grouping commands help identify patterns in data. The SPL command Transaction lets you group events into transactions. This is in some ways akin to the way relational databases use the concept of a transaction. You can enter the following SPL code in the search bar to let Splunk group events into transactions:

```
sourcetype=access_combined_wcookie
| transaction JSESSIONID clientip startswith="addtocart" endswith="purchase"
```

Advanced Analytics with Log Files

Traditional data analytics requires you to know ahead of time the questions you want to ask of the data. Following this, you design the data schema, create your SQL code to query the data and feed the results to an analytical tool such as Cognos.

Splunk refers to the information it collects and processes as operational intelligence. Splunk seeks to provide operational intelligence through:

- Searching, which lets you perform root cause analysis during troubleshooting incidents
- Real-time visibility that you let monitor the system as a whole

- Historical analysis that lets you find trends and patterns.

Unlike traditional analytics that uses Early Structure Binding, Splunk uses a different strategy called Late Structure Binding, which involves writing the events to log files, collecting the log files, and creating searches, and reports using Splunk.

Note that there's no database involved in all of this, with all the headaches that part of managing it – just log file management. Unlike traditional analysis, it takes minutes and hours instead of days and weeks to get actionable information out of your logs.

Splunk versus ELK

There's really no simple answer as to the question of which is better – Splunk, or ELK. Both are great tools, but one of them is open source, and the other a commercial product. While a commercial tool costs licensing money, of course, there are hidden costs involved in running with an open source solution.

With ELK, you do need to spend time and effort to build the features that you need, features that Splunk already has. Splunk has many powerful built-in features that will take time and effort to configure when you're working with ELK.

Open source solutions also usually don't have the training and support that commercial vendors can provide.

SaaS solutions for Logging –Sumologic and Loggly

Sumologic is a popular log analysis software, and it offers a SaaS analytics product that's completely cloud-based. You can also choose the hybrid cloud solution under which an on premise Sumo server gathers and sends all logs to Sumologic's cloud servers.

Sumologic is referred to by some as “Splunk in the cloud”. It provides many of Splunk's capabilities, such as search refinement, and charting of the log data. Sumologic also has a great feature where administrators are notified when key metrics deviate from baselines after events such as software upgrades, or a network breach attempt.

Loggly offers a cloud-based services that mines large amounts of log data in real time to provide insights. Interestingly, Loggly doesn't require you to use proprietary collection agents on the machines to collect log data, instead using open standards such as syslog and HTTP to collect the logs. You use your existing syslog daemons to send your Linux logs to Loggly for analysis. Loggly catalogs your log data for you and continuously refreshes its indexes as applications generate logs. Loggly also provides monitoring, real-time alerts, and anomaly detection.

Handling Streaming Data

Most of us are familiar with data processing at two different levels: online and batch. And we're also familiar with the processing strategies for the two levels – we use online data to perform business transactions, such as airline ticket sales, and book sales at Amazon. Most organizations send the daily transactional data to massive data warehouses using mostly nightly batch processes. Business intelligence tools and related tools analyze the data stored in the data warehouses for the benefit of the business groups.

Streaming data architectures are an entirely different breed from traditional data processing formats – these architectures handle continuous streams of events, - machine metrics, GPS signals, and collect these events as streams of data. They can also use a not so continuous set of events such as website traffic.



Stream computing analyzes data without having to persist the data in files or a database.

The ability to make sense out of streaming data means that you can perform real-time analytics, the holy grail of modern data driven organizations. This is how companies such as Amazon are able to use your clicks on their web sites to drive you towards purchases based on recommendations and other data science based strategies.

Newer stream technologies can handle very high volumes of messages such as millions of messages per second, and even store the messages coming in at that fast clip. There are many areas besides business where the time value of information is crucial, such as evaluating streaming traffic data for delivering timely reports to drivers.



Stream computing lets you respond to events as they're occurring.

When you analyze the unique clicks on a website, all the clicks together constitute a batch, which is nothing but a set of data points. Batch oriented tools process a collection (or batch) of data points together, and its results are available only after the entire batch completes its processing. If you're dealing with large batch data you'll need to wait a considerable time to get actionable information from your data.

If you need immediate answers to your questions, stream processing, and not batch processing, is the correct way to analyze the incoming flow of streaming data, such as

unique web clicks on a website. Stream processing is ideal for cases where data flows into your system through user or system generated events and you need to analyze the data with sub-second latencies. Stream processors don't analyze the incoming data after collecting a certain amount of it. Rather, they continually ingest the stream of new data, which is continually fed into the stream processor.

Since the streaming is continuous, unlike in batch processing, there's no formal start and end points to the data inflow. There's usually a message bus directing the never ending stream of data in processing are possible in this type of data flow, because results can be available immediately upon the arrival of a new data point.

Stream processing latencies are usually measured in milliseconds or seconds, rarely going beyond a few minutes. Twitter's continuous streaming of its voluminous amount of tweets is a classic example of how someone may need to handle high volumes of data but manage to process the data a very low latency. Obviously, batch processing is not the right solution for analyzing Twitter's trending topics!

Messaging systems (also called publish/subscribe messaging) that can ingest streaming data from multiple sources and make them available to multiple consumers are at the heart of modern streaming architectures. Besides serving as a queue to hold messages so as to avoid losing them when you're hit with a burst of messages, messaging systems serve many more critical purposes.

A messaging system has three basic components:

- Producers (aka publishers): These are the sources of the event data that is continuously transmitted to the messaging system.
- Topics: A topic is a high level abstraction for a group of messages.
- Consumers (aka subscribers): The consumer (or groups of consumers) is the user of the messages – they get the messages by requesting or subscribing to a topic. Subscribing to a topic means consumers don't have to read a bunch of useless messages to find what they need. Producers can send messages to multiple customers simultaneously.

The key takeaway with the publish/subscribe messaging systems is that the publisher of the message doesn't send the message directly to a consumer. The publisher merely classifies the message and the consumer subscribes to one or more classes of messages. Often, the message system employs a broker which sits between the publishers and the subscribers, to facilitate the transmission of the messages.



In Chapter 3, I explained the trend towards microservices where larger systems are decomposed into simple single-purpose services that perform specific functions. Microservices communicate with other services using lightweight protocols such as REST, or a messaging systems such as Apache Kafka.

Apache Kafka

Apache Kafka is a publish-subscribe messaging system used instead of traditional message brokers such as JMS due to the higher throughput and reliability it offers. Kafka is highly fault tolerant and can handle huge amounts of messages for low-latency analysis in your Hadoop system, by allowing parallel data loads into Hadoop.

Kafka acts a central pipeline where multiple types of data such as application logs, operations metrics, business activities, and database and web server events are collected and delivered to customers for both offline reporting and real-time monitoring.

Apache Kafka's architecture resembles that of well-known messaging systems such as ActiveMQ and Rabbit MQ. However, Kafka is especially designed for handling streaming data. While it's often called a messaging system, and does provide many capabilities of a typical messaging system, it's fundamentally different in the sense it provides a different abstraction of a structured commit log of updates.

Apache Kafka started out of a project at LinkedIn, where it was used to facilitate the movement of data between services. The communications needed to be asynchronous and message based, and the messages had to be saved as well, instead of using them once and discarding them.

As with a traditional message queue such as MQ, in the Kafka messaging system, produces send messages to a message queue (topic). Consumers are notified when any topic they subscribed to has new messages so they can read new messages off the queue.

Following are the key things that Kafka brings to a messaging environment, to make it feasible as a message –passing layer for large service architectures.

- It requires all messages to be acknowledged in order. This removes the need to track acknowledgments on a per-message basis.
- Messages can be persisted for days or weeks.
- Consumers are required to manage the offset of the next message they will process.

The sum result of these principles is that Kafka writes messages sequentially to a file system, and the consumers read those messages sequentially as well, allowing Kafka to handle very high speed messaging.

Benefits offered by Kafka

Kafka offers the following features that make it highly desirable in situations where you need a high throughput with reliable message delivery. You can summarize Kafka's benefits thus:

- **Performance:** High throughput for publishing and subscribing message even when handling large amounts of stored messages
- **Data consistency:** You don't need to implement functionality for checking data consistency since Kafka takes care of it for you.
- **Scalability:** Kafka is a distributed system, without any downtime
- **Reliability:** Provides fault tolerant data replication and balances the message consumers when failure occurs, in addition to offering delivery guarantees
- **Durability:** Stores messages on disk
- **Real-time:** Due to the speed with which Kafka handles messages, it supports real-time use cases

Due to the benefits Kafka offers, it's widely used for work involving the tracking of website activity, the collection and monitoring of metrics and aggregating logs, in addition to stream processing. Some key use cases for Kafka are the following:

- Meter and sensor data from grid points
- Solar and wind energy production data (transient sources)
- Forecasting data for weather, energy production and energy markets
- Activity data such as log data
- Event messages
- Application performance tracing

Kafka Architecture

Instead of using a bunch of publish/subscribe systems to transmit messages from various components of a system, such as the frontend servers, database systems, and shopping carts, you can architect a centralized system for publishing all types of messages, thus helping you scale your messaging infrastructure. This is what Apache Kafka can do for you.

Apache Kafka uses a simple architecture: Multiple message producers send messages to a single Kafka broker, which persists the messages on disk for reliability. Multiple consumers read the messages using the broker. Since messages are stored on disk,

consumers can read the messages long after the messages were received by the broker, which stores and forward messages for various topics.

Message Producers and Message Consumers

A Kafka system contains two basic types of clients or users:

- Producers create new messages to specific topics.
- Consumers read the messages by subscribing to topics, and read messages in the order the messages were produced. Consumers keep track of messages they've read by tracking the message offset, which is a piece of metadata (an integer value) that Kafka adds to each message. The offset is stored reliably, thus allowing a consumer to restart its message consumption without ever losing track of where they were prior to the restart.

Topics and Partitions

A consumer reads the messages sent to a topic in the order that the broker receives them. In order to achieve a higher throughput, you can divide a topic into partitions so the consumer can use multiple threads to read the messages from that topic. The messages within a topic partition are ordered, and consumers can't acknowledge individual messages out of order.

Partitions can be hosted on different servers, thus helping scale the topic for better performance.

How Kafka Persists the Messages

Kafka, often described as a 'distributed commit log', stores data in order, and distributes the data to provide resiliency as well as to scale the messaging infrastructure.

Writing Messages in Batches

For efficiency, messages are written in batches to Kafka. A batch is a set of messages produced for a topic and partition.

In order to make the message content more structured, message schemas are imposed on the messages. It's common to use Apache Avro data format for this purpose. Avro is a serialization framework that's compact and allows Kafka's messages to be understood without coordination.

Brokers and Clusters

A single Kafka server is called a broker. It's the broker's job is to receive message from the producers, assign offsets to each message, and store all messages on disk. The

broker serves the consumers by responding to requests for messages and sending them the committed messages.

Kafka brokers are part of a larger entity – the Kafka Cluster. Within a Kafka cluster, there are multiple brokers, one of which acts as the cluster controller, in charge of administrative operations such as watching for broker failures. A topic's partitions may be assigned to multiple brokers for redundancy of the messages so other brokers can serve the messages if the original broker fails.

How Kafka Works

The key to understanding how Kafka works is realizing that its commit log abstraction makes possible a reliable and highly efficient way of distributing changes to consumers. Data producers send streaming records which are appended to the commit log and consumers can stream the updates to the log with very low latency. Since consumers advance through the commit log independent of each other, a reliable and ordered stream of updates is sent to each consumer.

Database updates need to be delivered in the order they occur. Kafka's commit log can be spread over a cluster, with each part of the log replicated for fault-tolerance. The cluster helps parallel and ordered transmitting of data to multiple customers. The Kafka cluster can grow in size or shrink, without applications being aware of the changes. Even when a Hadoop cluster working with data from Kafka goes down for a while, there's nothing to worry, as the Kafka cluster safely persists all changes. All these capabilities set Kafka apart from typical enterprise messaging systems.

As with typical messaging systems, Kafka's design involves four key components - topics, producers, consumers and brokers. Here's how the four key components interact:

- A topic is a user defined category to which messages are published.
- Producers publish messages to the topics
- Consumers subscribe to those topics for accessing the published messages.
- Brokers are servers that manage all the work related the messaging, which involves the persistence and replication of messages in various topics.

Consumers are responsible for keeping track of the messages they consume. Consumers keep track of the messages they consume by tracking what's called an offset, which is a sequential number identifying messages. Consumers can easily get to any message stored on disk by supplying the message's offset value. Unlike in traditional messaging systems such as JMS, the broker is relieved of the duty of tracking message consumption. This design feature is at the heart of Kafka's ability to scale with the number of consumers.

A Real-Time Messaging Solution

In addition to its high reliability and throughput, Kafka stands out for its ability to provide a real-time publish-subscribe solution for very large data volumes. You can see multiple producers, such as the web applications generating logs, producers generating web analytics logs, and so on. Producers send data to various types of consumers, including offline consumers who may store them in HDFS or in an RDBMS. Near real-time consumers may store data in a NoSQL database such as Cassandra and real-time consumers filter messages in the in-memory database and trigger alert events.

Kafka offers a way to combine both online and offline processing and not only load data parallelly into a big data system such as Hadoop, but also partition real-time consumption over a cluster of machines. LinkedIn uses Kafka for streaming activity data and operational metrics, as well as to stream data to Hadoop for offline analysis. Twitter uses Kafka as part of Storm which is an infrastructure for stream processing.

Creating a functioning Apache Kafka cluster requires you to set up the ZooKeeper server for storing the cluster coordination information such as status, configuration, location information, etc.

Creating a Kafka Topic

You create Kafka topics using a command line utility provided by Apache Kafka. The following example creates a simple topic with a single partition and replica.

```
# bin/kafka-create-topic.sh -zookeeper localhost:2181 -replica 1 -partition 1 -topic testtopic
```

I named our topic testtopic in this example.

Once the topic is successfully created, you start a producer to send messages to the Kafka cluster.

Starting the Producer

In a real Kafka cluster, you'll usually have multiple brokers, producers and consumers. In our case, it's enough to start a single producer, as shown here.

```
# bin/kafka-console-producer.sh -broker-list localhost:9092 -topic testtopic
```

You need the following two parameters to start the command line producer client:

- broker-list: The server and port information for the brokers(localhost and 9092 in our example)
- topic: Name of the topic (testtopic in our example)

What I created here is a command line producer client that accepts your input from the command line and publishes them to the cluster as messages. The consumer will

then consume the messages. But first we need to create the consumer, which is our next and last step in the cluster configuration.

The producer client is now running, so you can start sending messages. Type something like the following message on the terminal where the producer is running.

```
Hola Kafka, Como Esta?
```

Starting a Consumer

In order to consume the messages sent by the producer, you need a consumer. Start the consumer as shown here.

```
# bin/kafka-console-consumer.sh -zookeeper localhost:2181 -topic testtopic -from-beginning
```

The consumer runs with the default configuration properties from the `consumer.properties` file, as shown here:

- `groupid=test-consumer-group`
- `zookeeper.connect=localhost:2181`

As soon as the consumer starts running, you'll see the message you typed in the producer server appear on the screen.

System Monitoring, Event Processing, and Tracking Application Performance

The sea change in infrastructures that support applications, such as the increasing use of cloud containers, and the use of micro-services, has raised the bar for operational visibility. Monitoring strategies and monitoring tools have kept up with the change in the operational environments.

In the old days, monitoring was simple: it mostly meant the use of shell scripts that alerted you to exceptional events in your infrastructure, and applications that ran on that infrastructure. Today, you've a bewildering array of monitoring tools you can use, often employing several tools.

This chapter introduces the most popular mentoring tools used today, classified into the following groups of tools:

- Time-series databases
- System monitoring
- Log Management
- Web and User monitoring
- Event processing
- On call management
- Error tracking
- APM (Application Performance Management)
- Enterprise Suites

Monitoring solutions, as most of us know, collect performance statistics from your systems and applications and alert you when certain events occur on the monitoring systems and services. Monitoring can be as simple as running a few Linux shell scripts periodically and having them send you pages or emails.

However, the monitoring landscape has undergone a tremendous change, especially with the introduction of several monitoring solutions that use the service model to provide monitoring solutions without having to set up one. It's more than likely that you'll use more than one monitoring tool in your enterprise, as the tools have gotten pretty specialized, and you often need to employ the specialized tools to get the best information for a specific need.

Monitoring systems aren't really limited to merely "monitoring" systems and services any longer. The definition of the field and the range and quality of monitoring solutions has grown tremendously in the past few years. Today's monitoring solutions help you do a bunch of other things, such as comparing current performance with historical performance, process "events", mine the logs, manage application performance, user behavior, and web performance. In this chapter, I discuss monitoring solutions that are available today, by classifying them into the following groups:

- System monitoring
- Log management
- Event Processing
- Error Tracking
- Time-Series Databases
- APM
- Web and User Monitoring
- Enterprise Suites
- Special purpose monitoring tools

The proliferation of monitoring tools and the different types of tools doesn't mean that you use these tools in isolation – far from it. For example, you complement system monitoring and application performance monitoring tools with web and user monitoring tools. The reason is simple – while the web and user monitoring tools can alert you about a current user issue that you need to fix right away, the root cause for that issue may be somewhere in your databases or servers, or in the network. Therefore, once you find out about a new user issue, you turn to your system monitoring tools to track down the root causes and fix the issue.

Monitoring – an Overview

This chapter discusses various types of monitoring, such as systems monitoring time-series databases, event processing, real user and synthetic monitoring, and application performance management. However, it's useful to review some monitoring basics before we look at the methodology and tools used for the various types of monitoring.

How Monitoring is generally done

At their core, monitoring systems are very simple: they are systems whose task is to watch other systems. However, over time, monitoring systems have grown in complexity. Well implemented monitoring systems can not only send alerts, but also help you proactively fix problems and provide valuable capacity planning information. They also help you honor your SLAs and often fix problems without even bothering you. Poorly built monitoring systems on the other hand, send you too much alert data, can be backdoors into your infrastructure, consume time for their care and feeding, and use excessive network bandwidth for transmitting huge amounts of data as part of their health checks.

Traditional monitoring tools such as Nagios use a pull model which means you poll target systems regularly and pull monitoring data from them wherein polling agents connect to remote servers and applications to test their reachability. They also run remote commands on the hosts and use their SNMP output to find details about resource usage such as CPU and RAM, for example. Pull models rely on the monitoring of uptime and reliability.

In the push model, systems send metric data to a unified storage repository. These metrics can go way beyond the traditional metrics of availability to the quality of service. The goal of modern monitoring services isn't merely to respond to uptime and reliability issues, but to help make more informed business decisions as well. Unlike traditional monitoring systems, which were content with sending alerts and notifications, these monitoring systems do more. They can aggregate the metrics, store them in a database, and provide powerful visualization capabilities for analyzing the data over time.



Today, monitoring tools such as Nagios and Zabbix are considered legacy tools.

You can divide monitoring solutions into two basic groups: those that use agents and those that are agentless. The Simple Network Management Protocol (SNMP) is the

most common way for systems and devices to send metrics. SNMP provides a Management Information Base (MIB) that stores references to various bits of information about a server or device. In Linux systems, most commonly, SNMP is implemented via `net-snmp`.

SSH is the other common monitoring solution besides SNMP. Monitoring systems log into the target services via SSH to perform their health checks and gather metrics.

If systems have limited access and need to be infrequently monitored, SNMP works fine. If you have remote access and many monitoring tools, SSH may work better.

Typically you install monitoring agents on all services to implement monitoring and alerting. You configure the agents to collect metrics about the server as well as services that run on that server. Agents have plugin that can collect metrics from a variety of sources such as databases and message queues. The agents periodically publish the metrics they collect to a central monitoring service (server).

The monitoring service stores the metrics that it receives from the agents and uses them to support its dashboards and graphing features. In addition, when thresholds are crossed for a metric, it can alert you. When viewing the metrics on the central monitoring service, sometimes you're interested in a specific server and sometimes a server group, such as when you want to get a sum of all database connections across a set of replicated database service running on different servers.

Need for New Monitoring Tools

Trends such as cloud based infrastructures, containerization, and micro services demand new monitoring techniques. Organizations such as Netflix build internal monitoring tools to handle their large number of instances in their infrastructure, with large numbers of immutable instances created and destroyed by features such as auto-scaling. Google needs to handle millions of containers, and Amazon's AWS Lambda has containers that exist for a fraction of a second to handle a single request.

Do-it-yourself Versus using External Monitoring Services

There are numerous ways to monitor systems. You can use polling systems such as Nagios, or walk your networks with SNMP and Cacti, or use custom Perl scripts and cron jobs. You may even outsource the whole monitoring function.

Regardless of how you perform monitoring, the following are your broad goals in most cases:

1. Detecting faults: Using thresholds for system behavior, monitoring should detect when system components aren't performing at the appropriate level.
2. Alerting: if a situation requires quick responses you configure email or pager alerts to your team.

3. Capacity planning: Times series data (stored by tools such as Graphite) helps you study trends in the data and make decisions about adding capacity to your systems.

In addition to their ability to monitor everything that you can do on your own, external monitoring services can detect things that you can't. Since they connect to your services from external networks, they're in the unique situation (actually the same situation as your users, who are all external) where they can easily monitor and track network outages, DNS related problems, load balancer, and SSL certification expiration issues. These are somewhat tricky for you to track from inside your systems. Some external monitoring services also let you measure your service performance on a geographical basis, often using various devices such as simulating mobile networks and low-bandwidth scenarios, in order to maintain optimal user experience.

External monitoring services can act as complements to your current monitoring infrastructure and provide a great way to monitor performance and availability from the perspective of your user population.

Setting up a comprehensive monitoring system is no trivial issue – it's not a case where you install agents on the target servers and are done with it. Besides installing the agents of course, you need to configure the plugins for each service you want monitored. You must then configure the agents to publish the data they capture to the central monitoring service. On the central monitoring server, you must configure alert thresholds, graphs, and dashboards. It's no piece of cake to get all the various moving parts working together nicely. Third party cloud based solutions can save you a lot of bother in setting up and maintaining these monitoring frameworks in a changing landscape of ever increasing numbers of servers and services.

External or hosted services obviate the need for you to develop custom code for plugins to monitor things such as message queues and web servers. The services come with simple APIs that let you easily publish custom metrics from inside your applications. They also enable you to easily configure graphs, dashboards, and alerts.

Services such as Server Density provide external SLA monitoring as part of their services. Services such as Pingdom exclusively provide external SLA monitoring.

Monitoring and the Mean Time to Recovery

The goal of monitoring and alerting solutions is to reduce your mean time to recovery (MTTR). The MTTR for a system is the combination of the following four things:

- Problem discovery time
- Response time
- Investigation time

- Time for fixing the problem

Automated monitoring reduces time spent in the first three phases that contribute to the length of the MTTR.

How Metrics and Logs can help Proactive Behavior

Reducing MTTR is all about the speed with which you react to bad things that befall your system. However, by appropriate use of metrics gathered by monitoring systems you can identify several problems before they occur. If an application stops since a database or the Linux server storage is full, you can see it coming way ahead of time in most cases. The ability of the monitoring systems to correlate different metrics enables you to trace the impact of various systems that have contributed to the failure of a specific system.



Aggregated logs help reduce time for problem investigation.

Different Types of Metrics

Monitoring and metrics can do a lot more for you besides reducing your MTTR. In order to make the most of your monitoring infrastructure, you must gather the following types of metrics.

- Operating system metrics: These are traditional server and network related metrics that help sysadmins (and DevOps people) to debug performance issues.
- Service metrics: These are metrics from services such as databases, web servers, message queues, and caching servers. Here, you need to focus on metrics such as time lost due to locking, and the number of database transactions per second (TPS). Analyzing these metrics lets you find out which component in your infrastructure is the bottleneck right now.
- Application metrics: Application metrics are published by your applications to provide insight into their performance. Application metrics include calls made by applications to third-party services, databases, and object caches. They also include the time taken for each web service end point to generate its responses. Application metrics tell you a lot about the features that are being actually used, and what the code is exactly doing in production.
- Business metrics: Business metrics tell you how well a business is functioning. Metrics such as the number of users logins per minute, the number of users

added, the number of abandoned shopping carts, are all business metrics which tells you how well business functions are being performed.

Metric Tracking across Multiple Services

Suppose your website is seeing a big increase in the number of HTTP error codes per second. Or, the CPU usage of a catalog service jumped by 20% since 1 P.M. The only way for you to know if these spikes are bad is to gather metrics about your system behavior over a long period of time. Capturing metrics over a long period helps discover patterns in resource usage. You should be able to not only get aggregated metrics for an entire system such as average CPU usages, but also be able to get ahold of metrics for a single service or a single instance of a service. In order to do this, you should be able to associate the metadata with the metric you're interested in. A tool such as Graphite helps you analyze the response time for an entire system or a single service, or even a single instance of a service.

Graphite allows you to send metrics in real time and query those metrics to produce charts and other displays to easily understand things. In order to cut down on metric volume, you can configure Graphite so it uses fewer older metrics compared to more recent ones. This allows you to store metrics about your system behavior over a long time period without worrying too much about storage.

Devops and Monitoring

Whereas system administrators are concerned with health checks and status of the infrastructure components, DevOps people often have a different set of expectations regarding monitoring. Traditional monitoring tools are inadequate to meet these requirements. DevOps is hugely concerned about reliability, speed and the ability to quickly scale to meet the demand for services. DevOps also needs to correlate IT performance to business metrics.

With this review of system monitoring basics, let's turn to a quick review of the main monitoring systems, such as time series databases, system monitoring, event process, and application performance management.

Time Series Databases

Time-series databases let you store and visualize performance metrics such as the following:

- System and network performance (for example, Disk I/O)
- Application performance (for example, application latencies)

- Business KPIs (for example, ad impressions)

The hallmark of time-series database is scalability and performance, with an ability to consume millions of samples per second. Some successful products in this monitoring category are Circonus, Cacti, Graphite, and Librato. I introduce Cacti and Graphite in the following sections. As with many of the newer alerting/monitoring tools I review in this chapters, increasingly, the goal of the tools is to bridge the gap between raw performance and alert data to derive actionable insights.

Using Cacti

Cacti is a round robin database-based (RRDtool) network graphing tool that uses SNMP polling to gather information. Cacti stores its time series data in the RRDTool database.

Cacti uses templates called Host Templates for different types of devices. The Host Template defines the Graph templates associated with a device. For example, if you assign the Host Template named Cisco Router to a devices, Cacti automatically knows about the associated graph templates.

Using Librato

Librato is a cloud-monitoring platform metrics and monitoring dashboard for any metric that you choose, including user actions. You can create custom dashboards for different teams, each of them geared towards metrics that are of interest to the groups.

Librato correlates the metrics that it collects across time, creates visualizations, and also alerts developers and operations teams.

Librato is a SaaS services, and works as a hosted telemetry system for time series data. In order for you to use Librato's capabilities, you must send your metrics to the Librato API, in the form of key/value pairs along with a timestamp. Librato will store your metrics, and provides you visualizations and alerts.

The easier way to send metrics to Librato is to use a Librato Integrator, which offers an easy mechanism to send service metrics, or metrics from Docker, Redis, AWS Cloudwatch and other services. In order to send generic metrics, you can choose a collection agent, or a language binding. Note that these are libraries that allow you to instrument your code or use local agents such as statsd to collect custom metrics.

Librato's functional DSL enables you to apply mathematical transforms to your metrics time series. You can create composite metrics representing ratios, or combine a large number of metrics into a single line. Composite metrics allow you to define higher-level metrics from the raw metrics. In a way this is similar to running complex queries against a set of raw metrics.

Librato's alerts system allows you to link the alerts with notification services such as email, or pagers.

Graphite

Graphite is a leading graphing system for operational and business metrics. Munin and Ganglia produce nice graphs, but they aren't updated continuously. Graphite produces graphs in almost real time. While Graphite is similar in many ways to Munin and Ganglia, it doesn't use the RRD database. Rather, it uses its own Whisper time series library.

Graphite is a collection of Python components that fetch and store time series data from various hosts. Graphite is easy to learn owing to the simple interfaces and formats it employs. You can easily push metrics to Graphite from scripts, applications, or from the command line. Since it uses URL based APIs, you can easily embed graphs in websites, and in dashboards. New users can easily get up to speed with the prototyping of graphs with no training.

Time series data is nothing special – it's a sequence of values you collect for any entity at regular intervals. And a time-series database such as the one Graphite uses, is a software system that's optimized for storing and retrieving time series data. The key to Graphite is its ability to store and retrieve time-series data quickly and at high volumes. A Graphite cluster can process millions of data points per second, offering you great visualization in a scenario that calls for a time series analysis.

Graphite uses an extremely simple metrics format for submitting data to the Carbon listener. All you need is a dot-delimited metric name (for example `foo.bar`), a numerical value and a timestamp, called the epoch timestamp. Here's how you can send a metric value from the command line:

```
$ echo "foo.bar" 12121212 `date +%s` | nc graphite.server.mycompany.com 2003
```

The echo command in this example sends the output consisting of the string "foo.bar", along with a number and the epoch timestamp (generated by the date command) through the netcat utility. The netcat utility hooks up with the Carbon service at the port 2003.



A data visualization services such as Graphite is a must for most organizations, regardless of what other monitoring systems and tools you have in place.

Components of Graphite

Graphite buffers writes in memory to optimize the writes. In memory caches allow you to store frequently queried data for fast retrieval. Following are the Graphite time series dataset components:

- **Carbon:** Carbon a network service (actually, a Python daemon) that listens for incoming metrics (time-series data from monitored hosts) and stores them temporarily in a memory buffer cache
- **Whisper:** Whisper specifies the database file layout, and the programming library that Carbon and the Graphite web application use to interact with the database files. Whisper uses a flt file format for storing time series data. A Whisper file is created for each metric that Graphite tracks. You can alternately use Influx DB for storing monitored data, for performance reasons.
- **The Graphite web application interface (the Composer):** You use Graphite through the Graphite web application, which is Django based frontend, to view the data stored in the Whisper database. Alternatively, you can use Grafana as a dashboard for Graphite metrics since it's way more refined than the Graphite Web interface. Grafana offers you more capabilities for querying, visualizing, and annotating data.

Let's learn about these three basic Graphite components in some detail in the following sections.

Carbon

The Carbon component runs three daemons: the carbon-cache listener, carbon-relay, and the carbon-aggregator. Of these, carbon-cache is the real work horse. This service receives metrics from clients, temporarily buffers them in the cache, and periodically writes them to disk in the Whisper database files. You can run multiple carbon-cache processes in a busy system.



In summary, carbon-cache collects and stores the metrics, carbon-relay provides metrics routing, load balancing, and replication, and the carbon-aggregator service normalizes metrics so you get them in a predictable fashion.

The carbon-relay process relays metrics from one Carbon process to another. Carbon-cache helps you scale up – it can handle a single Graphite server with 16 parallel carbon-cache processes. In order to replicate metrics from one datacenter to another, you can also set up an external load balancer such as HAProxy to distribute traffic among several carbon-relay processes. The carbon-aggregator, as its name indicates, can aggregate source metrics and buffer them before publishing them with

a new metric name. For example, if you want to track the average latency of a bunch of application servers, Graphite can render an average of the server metrics.

Whisper

By default Whisper uses the SQLite database to store the metrics data but you can use a more serious database such as PostgreSQL or MySQL, for a production environment. Memcached is even better from a performance perspective.

Graphite Web UI

The Graphite web UI is where you go to create your graphs. The Graphite-Web instances reads metrics information from the Whisper database files and turns this data into graphs. Graphite can also read from Round-Robin Database (RRD) files. Ceres is a more advanced time-series database format that's designed to eventually replace Whisper as the default storage provider for Graphite.

Graphite offers both real-time monitoring and a fast analytical Datastore that lets you collect, store, and analyze data quickly.

Users approach Graphite through its web interface, the Composer. Composer helps you add metrics and remove them from graphs, and visualize and correlate your data. You get instantaneous feedback while making changes and the easy workflow helps you identify cause and effect and anomalies easily, which you can't do with most traditional monitoring systems.

Statistics and Transformations

Graphite's rich array of statistical and transformative rendering functions sets it apart from all the other monitoring tools. These functions allow you to perform various statistical computations on time series data stored in the Graphite database.

Graphite chains functions. It treats each time series on a chart as a data stream and passes the original data into a rendering function, and passes the output of the function to another function. When it's finished with all its transformations, it passes the series to a rendering format such as a graph, or a CSV/JSON response.

Event Processing Tools

Manual processing of data is out the window in most places due to the sheer amount of data that you need to sift through to make key operations decisions. Event processing tools help automate the incident resolution process. These tools consume alerts from monitoring tools and enhance the information they capture by adding the following processing steps:

- Correlating matching related alerts

- Enriching data through adding context and additional insight to events
- Routing and the funneling of events to various groups
- Suppressing noise by removing unnecessary events

Event processing tools share the following features:

- An events dashboard
- Alert routing
- Event correlation and enrichment
- Alert analytics
- Integration with collaborative platforms such as JIRA, and ServiceNow.

Key event processing tools include Riemann, Big Panda, and Bosun. I discuss both Riemann and Big Panda in this chapter.

Riemann

Riemann is a monitoring tool that aggregates events from hosts and services and send them to a stream processing language, where they're summarized and used for actions. It also tracks the state of events and allows you to build checks that depend on a sequence or combination of events. You can also send the events to other services or store them.

You configure Riemann with a Clojure based domain specific language (DSL) configuration file. Therefore, in order to get started with Riemann you need to pick up just enough Clojure programming skills.

Push versus Pull Based Monitoring Systems

Classic monitoring systems such as Nagios are pull (polling) based systems, where the focus is on checking uptime and availability. Push based monitoring systems are mostly concerned about measuring things. While you're still concerned about availability, it's not the main purpose of this type of monitoring. Your main goal is to measure the application performance, and user experience. Thus, the focus is squarely on throughput, performance, and value, rather than the mere availability of the systems. Riemann is a good example of a modern monitoring system – it's an event-centric push model based system.

Event Stream Processing

Riemann is an event stream processor that aggregates events for hosts and applications and manipulates, summarizes, and sends actions based on the events. When something important happens in a system, an event is submitted to Riemann. For

example, a HTTP request that the system handles is an event that'll be sent to Riemann, along with the time it takes to handle the request. When an exception occurs, an event is sent along with the stack trace. Lightweight daemons can watch servers and send events about CPU, RAM and disk usage. Riemann filters, aggregates, and acts on the flow of events to understand the system. Riemann can send output to Email, PagerDuty, Graphite, Librato, and other systems.

Riemann isn't an alerting system, although it can send alerts and notifications. Rather, it's a stream processing tool that can create alerts. Riemann's configuration is in the form of Clojure code, which enables you to build abstractions to describe potential error types, and limits the rate at which alerts are sent to you. For example, you can specify that a certain error alert be sent only if more than 5% of those errors occurred within the last 5 minutes. Monitoring tools such as Nagios enable you to do these types of customizations as well, but you'll need to use various plugins for the customizations. With Riemann, it's all code. And you can change it to your heart's content.

Configuring Riemann

You configure Riemann with a Clojure configuration file. On an Ubuntu system, the configuration file is `/etc/riemann/riemann.config`. Here's the default configuration file:

```
; -*- mode: clojure; -*-
; vim: filetype=clojure

(logging/init {:file "/var/log/riemann/riemann.log"})

; Listen on the local interface over TCP (5555), UDP (5555), and websockets
; (5556)
(let [host "127.0.0.1"]
  (tcp-server {:host host})
  (udp-server {:host host})
  (ws-server  {:host host}))

; Expire old events from the index every 5 seconds.
(periodically-expire 5)

(let [index (index)]
  ; Inbound events will be passed to these streams:
  (streams
   (default :ttl 60
    ; Index all events immediately.
    index

    ; Log expired events.
    (expired
     (fn [event] (info "expired" event)))))))
```

As you can tell, Riemann's configuration file looks quite a bit different from the traditional configuration files you're used to! Here are the key aspects in the Riemann configuration file:

- The first section of the file sets up how Riemann logs its data to a file (riemann.log).
- The second section configures Riemann's interfaces, such as TCP, and UDP, which are bound to localhost by default.
- The next two section configure streams and indexing. Streams are the heart of Riemann monitoring, and are functions to which Riemann passes events for aggregating, modifying, and escalating.

How Riemann Gets its Data

You can send monitoring data to Riemann using various tools such as those that perform basic health checks, and others that work with web servers such as Nginx, Apache, and cloud based services such as Amazon Web Services.

Riemann doesn't store its event data, although its index does store selected events in memory until they expire. You can store events by sending the events to a system such as Graphite or Elasticsearch.

Big Panda

Big Panda belongs to the brand new class of alerting solutions that helps you make sense of all the alert data flowing from your systems. It centralizes all the layers of your monitoring stack such as infrastructure, and application monitoring, web and user monitoring, and logging as well, under a single umbrella. The goal of tools such as BigPanda is an enhanced capability to sift through noisy alert environments where thousands of alerts pour into an environment, leading to a significant reduction in the meantime to resolution (MTTR).

Big Panda uses a sophisticated correlation mechanism to categorize alerts from multiple monitoring tools into a logical grouping of applications, micro services, or the cloud, so you can isolate and analyze incidents that are relevant to you. Unlike other tools, Big Panda clusters numerous IT alerts into high-level incidents so you reduce the alert noise and focus on fixing critical issues. Big Panda claims it can reduce alert noise by 95% with its automated incident detection strategies, and reduce MTTR (mean time to resolution) by an average of 86%.

Big Panda uses a data science platform to manage machine data across all of your monitoring tools such as Nagios to deliver timely and relevant actionable incidents so they can be taken care of by your support teams. It does this through suppressing

most of the irrelevant alerts by clustering related alerts, and visualizing patterns in the event streams.

Handling Alert Flooding

Traditional monitoring solutions such as the venerated Nagios tool tend to be very noisy. During an IT outage, there are hundreds of alerts from a monitoring system such as Nagios. During such a time, false alerts distract your attention and keep you from seeing the big picture.

Automated Incident Detection

BigPanda automatically correlates metrics from various monitoring tools into logical groupings based on custom criteria such as applications, microservices, or the cloud. Grouping and correlating the alerts into unified incidents enables you to easily zoom in on critical issues. This helps you focus your attention only on relevant incidents. Instead of getting many single data points (alerts), you get the full context of an incident.

How BigPanda Works

You can easily integrate your monitoring tools such as Nagios, Zabbix, AppDynamics, and New Relic, with BigPanda. BigPanda analyses and consolidates alerts from all these monitoring tools into a unified data model. You don't need to define rules since BigPanda automatically works to provide you results through its built-in correlation rules. You can however, customize the correlation rules to suit your needs.

BigPanda normalizes the raw data it collects from all the monitoring tools in your environment so you can visualize the lifecycle of an issue over time.

If a router crashes, it can result in hundreds of alerts, but BigPanda can correlate all the related alerts into one incident. Similarly, BigPanda groups all alerts from the same application regardless of which host or application server, or even data center, the alerts originate from.

Logical Correlation of Alerts

The logical correlation of alerts across topology, time, and context is at the heart of BigPanda's ability to provide actionable information to you by analyzing its raw data collected from all the monitoring tools. Here's what these three logical parameters mean:

- **Topology:** Topology refers to the host, service, application, or the infrastructure element that produces the alerts. Alerts are more likely to be corrected when they originate from the same topology component, such as the alerts emanating from the same databases or application server.

- Time: The time parameter helps correlate alerts occurring together over a time period, such as alerts produced within a few minutes of each other.
- Context: The context parameter helps find relationships among alerts. For example, if many hosts in a data center are sending out network related alerts, it's likely that there's a problem with a network switch.

System Monitoring

Traditionally, monitoring meant just monitoring systems – system administrators relied on various tools (including, of course shell scripts) to gain a visibility into the status of their production systems to learn about hardware, software, network related issues.

Modern system monitoring tools do a lot more than the traditional monitor/alert mechanism administrations relied on for ages. These systems also provide the following:

- A status dashboard for the entire infrastructure that shows your entire system universe in green/yellow/red colors
- Agents for executing periodic health checks
- Hierarchy and dependency mapping checks
- Collectors for servers and networks
- Plugin architectures that help monitor various types of infrastructure such as databases and web servers. Collectd, DataDog, Ganglia, Munin, Nagios, Sensu, Server Density, Zabbix, and Zenoss are all well-known system monitoring tools. I discuss Ganglia, Nagios, and Sensu in this chapter.



Monitoring can involve multiple things besides alerting about exceptions – modern monitoring can also include graphing capabilities, and automated fixes.

Nagios

Nagios, an open source monitoring tool, is the granddaddy of the current generation of monitoring tools, and serves as a good comparison to the other tools in the monitoring landscape today. Nagios has been around for about 20 years, and while newer and fancier monitoring tools abound, it still retains a special place in the systems monitoring world, and is considered by many to be a standard for open-source monitoring and reporting.

Nagios is a highly modular monitoring system (provided by special monitoring applets called plug-ins that provide support for various services and devices) that's extremely flexible and extendable, and is easy to set up and get going. On an Ubuntu system, for example, simply by installing the nagios3 meta package and its dependencies, you get a Nagios system that's preconfigured and fully functional, ready to monitor your environment.

Nagios doesn't try to do everything by itself – it has great interoperability with many other open source tools, extending its monitoring capabilities.

Compared to most commercial monitoring apps, Nagios has few built-in checks – in fact it has no checks at all. Nagios depends on tiny special purpose programs called plug-ins for the checking logic. You can add support for new types of services simply by adding new plug-ins to the Nagios system. Plug-ins can execute locally on the monitoring server or remotely, on the hosts that Nagios monitors.

Nagios Components

Nagios consists of three components: the Nagios daemon, the Web interface, and the plug-ins. While the daemon and the Web interface are part of the Nagios installation, you must download and install the plug-ins separately.

The main purpose of the Web interface is to let you view the current state of the monitored hosts of course. Beyond that, the web interface lets you do the following;

- Access historical reporting and trending tools
- Schedule downtime
- Enable and disable service checks and notifications
- Draw graphical maps of your environment



Nagios is a centralized polling system with no built-in monitoring logic. It has little built-in functionality beyond scheduling and notification.

The most powerful interface to Nagios is the Event Broker, which is a process that watches the event queue for events and passes relevant details about the events to the Event Broker modules.

The Host and Service Paradigm

Nagios is easy to understand and to extend. It's a scheduling and notifications framework that calls on plug-ins, which are small programs, to do the actual monitoring.

Since it's easy to retrieve data from the Nagios framework, you can access real time data for graphing programs such as RRDTool and MRTG. Nagios doesn't care what tools you use to monitor your systems. So, you can take your home grown scripts and such and incorporate them in Nagios without having to reinvent the wheel.

Nagios lets you define everything. So, each element that Nagios works with is user defined. Hosts and services are the two fundamental types of objects for Nagios. Nagios lets you define a single host check and multiple service checks for the services that run on the host. Based on the host and service check definitions, Nagios calls various plug-ins to obtain the host or service status. For example, the plug-in named `check_ping` may perform the host check for a server.

Nagios Plug-Ins

Plug-ins, the tiny monitoring programs, are what Nagios relies on to provide monitoring services. The Nagios Plugins project is a vast set of user-contributed plug-ins that contains all the functionality you need in your monitoring system. Some plugins such as port scanners and ICMP and SNMP query tools run on the monitoring systems. Other plug-ins, such as utilization checks, are run remotely.

You can also write custom plug-in and modify your monitoring scripts into Nagios plug-ins. A Nagios plug-in is simply a program that provides a specific exit code, as shown here:

```
Code 0: Okay
Code 1: Warning
Code 2: Critical
Code 3: Unknown
```

You can use any programming or scripting language that'll emit an exit code to write Nagios plug-ins. A plug-in does the following:

- Gathers run time statistics from a host, such as the current CPU load
- Compares the statistics with thresholds
- Emits an appropriate exit code

Since all plug-ins are self-contained little programs, you can execute them from the command line to troubleshoot problems with the plug-ins.

Monitoring with the Nagios Plug-Ins

Nagios monitors hosts and services in three different ways:

- Connect to remote hosts using remote execution techniques and run plugins on those hosts.
- It can wait for remote hosts to notify it by defining passive checks.

- It can launch plug-ins locally to check the availability of various hosts and services.

Let's use a simple example, a service check that makes an ICMP echo request, also called the ping command. In order to set up a ping check, you define the ping command in the check-commands file, as shown in the following code snippet. `define command{ command_name check_ping command_line $USER1$/check_ping -H $HOSTADDRESS$ -w $ARG1$ -c $ARG2$ -p 5 }` `define command{ command_name check_ping command_line $USER1$/check_ping -H $HOSTADDRESS$ -w $ARG1$ -c $ARG2$ -p 5 }`

```
define command{
    command_name    check_ping
    command_line    $USER1$/check_ping -H $HOSTADDRESS$ -w $ARG1$ -c
$ARG2$ -p 5
}
```

The `-w` and `-c` flags specify the warning/critical thresholds. The dollar signs surround the macros. When Nagios calls the `check_ping` command it replaces the macro names with actual values. In this case, it replaces the numbered ARG macro with values from the arguments in the service definition. I define the service as follows, as part of setting up the ping check:

```
define service{
    host_name        Server
    service_description    check_ping
    check_command     check_ping!500.0,20%!1000.0,60%
    max_check_attempts    5
    normal_check_interval    5
    retry_check_interval    3
    check_period    24x7
    notification_interval    30
    notification_period    24x7
    notification_options    w,c,r
    contact_groups    admins
}
```

Our `check_ping` command runs every 5 minutes, 24X7. Nagios notifies the admin group every 30 minutes 24X7 until the service is back online.

Remote Monitoring via NRPE

Several Nagios plugins such as `check_memory` need to collect data from the systems, and therefore, they can't remotely monitor systems. Nagios uses an indirect mechanism to gather data from remote servers. The Nagios Remote Plugin Executor (NRPE) helps collect the information indirectly from remote systems. NRPE sits on the remote machines and executes the commands and plugins that Nagios needs. Nagios would then collect the data from NRPE.

NRPE uses a plug-in called `check_nrpe` that's executed locally and a daemon that runs on the hosts you want to monitor. The `check_nrpe` plug in asks the daemon to execute a predefined command. The daemon runs the command and provides the output and the exit code back to Nagios. NRPE can securely execute any program on the remote server as long as it defines the program in the daemon's config file on the remote server.

Nagios can also collect data through SNMP, which is quite handy for monitoring network equipment such as routers that have built-in SNMP agents.

How Nagios Schedules Checks

Nagios places all of its processes and its host and service checks in a global event queue. Although you define the scheduling of the check events, you don't do so by specifying absolute time the way you do inside a crontab. Instead, you specify how long Nagios must wait after a plug-in exits before executing that plug-in again.

In order to define the time interval, you make use of two options:

- The interval length, which defines a block of time in seconds (60 seconds is the default value)
- A normal check interval, which is the number of interval lengths Nagios must wait.

If a plug-in execution time exceeds the normal check interval, Nagios will reschedule that plugin.

Distributing the Load

When Nagios starts up, it needs to contact a whole bunch of servers. In order to avoid overburdening the remote servers, Nagios uses an interleave factor to space the communications out. For example, with an interleave factor of 3, Nagios checks every third service. To avoid overburdening the Nagios server itself, Nagios inserts an inter-check delay between checks instead of executing all of them in parallel.

Ganglia

Ganglia is a distributed system monitoring tool for computing systems such as clusters and grids. It allows administrators to monitor the entire cluster through a single location. Ganglia employs a listen-announce protocol to monitor a cluster's state and use a tree of connections among the cluster's nodes to aggregate their state. Ganglia is engineered to add very low overhead to the node, as well as high concurrency, with an ability to scale to clusters with thousands of nodes.

Ganglia is a graphing and monitoring solution that can aggregate information to provide great overview display. As with Munin, Ganglia uses RDDs underneath to sup-

port database storage as well as graphing and thus, Ganglia graphs look similar to Munin graphs.

Ganglia architecture consists of the following components:

- **gmond:** Gmond stands for the Ganglia monitoring daemon. Gmond collects node information and runs on each of the servers that Ganglia monitors.
- **Gmetad:** This is the Ganglia meta daemon that runs on the master node and collects information gathered by all the Gmond nodes.
- **RRD:** RRD is a round robin database where Ganglia stores data in time series that are suitable for graphing.
- **Web frontend:** The PHP base web frontend displays the data RRD graphs for you.

Ganglia is something you want to consider if you've a large number of computers with general purpose operating systems from which you want to gather near real-time performance data.

Using Munin

Munin is a good product to start out with in this area. Munin provides graphs of server usage that help you notice resource allocation issues before they develop into full blown crises.

Using Collectd

Collectd is a daemon that collects periodic system performance statistics and allows you to store the values in many ways, including in RRD files. You can then use the stored statistics for finding performance bottlenecks, and for estimating the future system load. Collectd is written in C for performance reasons. It also contains optimizations that lets it handle hundreds of thousands of datasets.

Compared to most of the other tools I discuss in this chapter, Collectd is limited in its functionality – you can't use it to generate graphs and it's also quite limited in terms of its monitoring functionality. You can, however, send Collectd gathered metrics to Nagios through the collectd-nagios check.

Using Sensu

Sensu is an infrastructure and application monitoring and telemetry system that uses RabbitMQ as a message broker and Redis to store its data, and is considered ideal for monitoring cloud environments.

Sensu is an open source monitoring framework whose goal is to be 'simple, malleable, and scalable'. Sensu uses checks to monitor services and measure the resource status.

Checks are scripts that Sensu executes on the servers that host the Sensu clients, and report whether a certain condition is met, such as a running MySQL database. Sensu uses handlers for acting on the event data. Handlers are scripts that perform actions such as sending emails and generating alerts, as well as adding metrics to external stores such as Graphite.

Sensu is often called a “monitoring router”, since it connects the output from the check scripts running on the nodes with the handler scripts running on the Sensu servers. The output of the checks is routed through the handlers, which decide what to do the results of the checks. Handlers can send alerts to Email, HipChat, PagerDuty, IRC, and Twitter, among others, and can also send metrics to backends such as Graphite and Librato.



Sensu lets you reuse monitoring checks and plugins from legacy monitoring tools.

Sensu is a tool well suitable for monitoring cloud-based infrastructures. Unlike Nagios for example, which employs a client/server model, Sensu uses RabbitMQ, which is an AMQP-based messaging system that helps make it much more easily scalable than tools that rely on a single central master server.

Decoupling the submission of checks by the Sensu clients from the reading of the check results enables Sensu to offer a much greater throughput than architectures where clients directly submit check results to the master. When you check a service, the Sensu client writes the check results to a RabbitMQ queue, where the Sensu server process reads it.

Since you can run RabbitMQ as a HA cluster, resiliency is provided for as well. If the master server goes unavailable, messages will wait in the queues until the master comes back online.

Sensu was designed with the cloud in mind, and also to work with configuration tools such as Chef, Puppet, and Ansible. In order to work with Sensu, you must first install the Redis database, as well as the RabbitMQ message server. Sensu was designed from the outset as a replacement for an existing Nagios installation, so it lets you reuse all your monitoring checks and plugins from legacy monitoring tools such as Nagios, and Zabbix.

Sensu Components

Sensu consists of the sensu-server, sensu-client, and the sensu-api. You use the sensu-dashboard to get an overview of the Sensu infrastructure and to perform actions. Here's what each of the components does:

- **Sensu-server:** The server initiates checks on the clients, receives the output of the checks, and sends it to the handlers. The Sensu-server stores its data in a Redis instance, and depends on RabbitMQ for passing data to the sensu-client nodes.
- **Sensu-client:** the sensu-client runs check scripts and returns the results to the sensu-server through RabbitMQ.
- **Sensu-api:** The sensu-api is a REST API through which you access the data stored by the sensu-server in the Redis database.

Sensu uses a pub/sub model for communicating, which means that it maintains dynamic client registry, where ephemeral systems are automatically registered and deregistered. You therefore don't have to worry about false-positive alert storms when you scale an infrastructure up or down.

Sensu performs both major functions: monitoring and telemetry, that is, it monitors and measures. It monitors applications and system services through the execution of service checks, and also piggybacks on the service check mechanism to collect useful metrics from the system. The Sensu check format consists of an exit status code and an arbitrary payload such as a message string, JSON, and PerfData, making it easy to collect detailed metric data.

Checks and Handlers

You can make use of the many checks (in the form of Ruby scripts) available from the `sensu-community-plugins` repo on GitHub. Let's see how you can write a check to see if the cron daemon (`crond`) is running. You can use the `check-proc.rb` script for this purpose. The following example shows how to grab the `check-procs.rb` script from GitHub:

```
wget -O /etc/sensu/plugins/check-procs.rb https://raw.githubusercontent.com/sonian/sensu-community-plugins/master/plugins/check-procs.rb
chmod 755 /etc/sensu/plugins/check-procs.rb
```

You can store the check definition in a JSON file and place the file on both the Sensu server and the client nodes. Here's the JSON file:

```
{
  "checks": {
    "cron_check": {
      "handler": "default", "email",
      "command": "/etc/sensu/plugins/check-procs.rb -p crond -C 1 ",
      "interval": 60,
      "subscribers": ["webserver"]
    }
  }
}
```

The most commonly used handler type is “pipe” which tells Sensu to shell out and run the command you specify. You can define an email handler through a JSON file as well (handler_email.json, for example):

```
{
  "handlers": {
    "email": {
      "type": "pipe",
      "command": "mail -s 'sensu alert' your@address"
    }
  }
}
```

You need to restart Sensu on both the server and client nodes at this point. Shortly afterwards, you’ll see messages in the sensu server and client logs stating that the check request was published, that the server received the request, and that the check result was received by the server. Should the crond (cron daemon) go down for any reason, you’ll see the alerts in the Sensu Dashboard, as well as receive an email notification.

Web and User Monitoring

Modern web application based user monitoring tools are an entirely different kettle of fish from the typical “systems monitoring” tools you’re used to dealing with.

There are two key types of web and user monitoring tools:

- Synthetic monitoring tools simulate user traffic to an application
- Real user monitoring is when a tool tracks real failures arising in the user’s clients

Web and user monitoring involves the front end of your web applications. They try to find out exactly how long pages are loading for real users, and offer details that you need in order to diagnose and improve your web performance.

It’s likely that you’ve heard of the term “user experience”. Web and user monitoring tools are all about user experience. While regular monitoring tools track the technical metrics relating to performance, web and user monitoring measure something more directly useful – the level of ease with which users are able to access and use your web applications. Most commonly, tools in this category do the following:

- SLA (service level agreements) tests uptime
- Monitor the HTTP/HTTPS/SSH and generic TCP endpoints
- Geographical segmentation

Pingdom is a good example of a web end user monitoring (and SLA monitoring) tool, and I discuss this tool in this chapter, following a review of synthetic and real user monitoring.

Synthetic Monitoring

Synthetic monitoring is usually performed in datacenters using throttled connections to mimic the actual experience of users. The testing involves the loading of pages in web browsers and the collection of the performance metrics.

There are many advantages to synthetic monitoring exercises. Synthetic monitoring exercises:

- don't need to install anything
- can measure and compare the performance of any website
- can perform detailed analyses
- can create performance waterfall charts
- can make a video recording and film strip of the user experience

The drawbacks to synthetic monitoring are obvious: only average network conditions can be mimicked in a datacenter. The sample sizes are small, and the testing exercises are geographically limited.

Real User Monitoring

Real User Monitoring (RUM) captures performance metrics when the real thing is happening - that is, when real users are on your website. RUM can't measure detail to the depth that synthetic monitoring can. Its big advantage should be obvious: it monitors real users working with real connections from all over the world.

RUM tools work by instrumenting web site pages with a JavaScript beacon that writes back collected data. You can deploy the JavaScript by either manually instrumenting the code through inserting JavaScript in the page headers, or by dynamic injection of pages from an application server.



Limited performance metrics are a drawback of RUM.

Using Synthetic Monitoring and RUM Together

Ideally, you should employ both synthetic monitoring and RUM to gain insights into the performance of your websites and really understand how the front-end code and your front-end architectural components are impacting your web performance and user experience.

Pingdom

Pingdom is a popular web performance monitoring solution. Pingdom provides both website monitoring, as well as a real user monitoring service. In the following sections, I briefly review the key services offered by Pingdom. Pingdom offers RESTful and HTTP-based APIs that allow you to create your own scripts or applications with functionality that is provided by the Pingdom control panel.



The consulting firm Gartner estimates that by the year 2020, 60% of APM buyers will reside outside of IT operations.

Website Uptime Monitoring

Pingdom employs a network of over 60 probe servers from various locations in the world, to test your websites from across the world, often at very short intervals, such as every minute. It alerts you when something breaks, and performs root cause analysis so you can fix the issue, and prevent it from reoccurring.

Real User Experience Monitoring

Pingdom's real user monitoring capabilities let you get information about the experience of actual visitors to your web sites. You can check user experience over a period of time, or in real time.

Transaction Monitoring

While uptime monitoring helps in ascertaining that your web sites are actually up, it does no good to you if the sites aren't performing the way they've been designed to perform. Pingdom offers multi-step synthetic transaction monitoring that helps test the way the web sites are working. Your websites and web applications involve several key transactions, such as users registering for new accounts, users logging into your web sites, or customers placing items in their shopping carts. Several of these transactions involve multiple complex steps that need to work together without a hitch.

Pingdom's synthetic transaction monitoring automatically tests transaction steps by reproducing the steps at predetermined intervals. You'll be alerted of a transaction

step failure. Just as you can do for the uptime checks, you can get an overview of the current synthetic transaction monitoring checks and their outcomes.



Search engines such as Google penalize search result rankings for sites that are slow, or that produce a poor user experience.

Page Speed Monitoring

Page speed monitoring allows you to see how fast your webpages are loading from various countries, or from various devices. All you need to do is install a tiny JavaScript snippet into your applications, to enable Pingdom to start collecting performance data. You can immediately learn about spikes in response times and correlate them to recent changes you've made in your web applications, and fix the problem right away.

In order to monitor page speed, you enter the web page URLs and Pingdom will perform automatic testing with a real web browser (Google Chrome). You can view various details about your web pages, such as the size of the CSS files and images, and their load times. The full page tests and analysis let you understand the specific content types that are slowing down your websites.

Application Performance Monitoring (APM) Suites

Web and user monitoring is often the domain of what are known as Application Performance Monitoring (APM) Suites. The word “suites” connotes that these tools satisfy multiple feature requirements, and are sometimes offered as independent products under the umbrella of the APM product suite.

APM suites typically provide most of the following types of monitoring related capabilities:

- End-user experience monitoring (EUEM): This involves the capturing of data about how the end-to-end latency, and quality appear to the real users of your applications. In addition, synthetic transactions emulating the end user can also focus on the availability of your applications.
- Application topology discovery: Discovering the software and hardware components that are part of executing an application.
- User-defined transaction profiling: Tracing of user events comprising a transaction as they occur within applications during user interactions with your application infrastructure.

- Application component monitoring: Monitoring the resource consumption and other events occurring in the software and hardware components.

Application Performance Monitoring tools look similar on the surface to Web and User Monitoring Tools since you use them to detect and debug user experience related issues. However, there's a huge difference between these two types of tools. APM tools track key application metrics such as latency and throughput by tracking the complete transaction flow of an application, starting from the client and running through the entire stack, where the backend databases are usually situated.

APMs provide analytics around your application performance, such as the time it takes to execute various areas in the code, and to complete a transaction. The APM tools do this by instrumenting code, monitoring application logs, and including network and hardware metrics.

APM tools usually include the following features:

- Latency and throughput metrics
- Geographical segmentation
- Database query performance tracking
- Tracking code deployment related changes in performance metrics

APM provides end-to-end business transaction oriented management of distributed applications. The goal is to rapidly identify issues and resolve them so as to maintain the ideal user experience.

APM tools perform the following functions:

- End-to-end transaction tracing: Includes application mapping transaction profiles, real-time analytics, and alerting.
- Code-level visibility: Code level insights across servers, databases, caches, and queues.
- Dynamic baselining and alerting: Automatically determine normal performance levels and flexible alerting with the help of services such as PagerDuty and ServiceNow.
- Troubleshooting: Identify and fix code bottlenecks, inefficient database queries, and infrastructure blockages.

A good example of how APM tools are fundamentally different from traditional monitoring tools is the runtime application architecture discovery modeling features incorporated in the APM tools. These runtime modeling capabilities reveal the live topology of how end user requests traverse through the various runtime environments in your environment. In a typical microservices architecture, an end user can

communicate with hundreds of components and servers. Each of these points of contact is a contributor to latency of course, and the topology map shows you at which points the latency is high.

Most APM tools seek to provide a 360 degree view of your environment to ensure optimal performance of desktop, web, and mobile applications in both on premise and cloud environments.

New Relic APM

New Relic is a developer oriented APM tool that monitors production applications and provides insight into their performance and reliability. It saves you time in identifying and diagnosing performance issues. New Relic APM has subscription based APM service that consists of the following components:

- **Application monitoring:** Lets you view application performance trends such as page load times, application response times, error rates and slow transactions. External service instrumentation captures the performance of external services such as web services. Transaction tracing lets you view transaction details, and allows the tracking of slow HTTP requests all the way down to the SQL.
- **Database Monitoring:** Shows a detailed overview of database performance, including the response times and throughput, and a SQL query analysis to identify slow SQL statements.
- **Alerting and error monitoring:** You can set your own warning and critical thresholds to proactively track and resolve problems, an external ping service that your customer can reach your site, in order to provide availability monitoring.

AppDynamics

AppDynamics offers a comprehensive package of Web application and mobile monitoring, analytic, reporting, and alerts. AppDynamics enables you to perform RUM, Synthetic Monitoring, as well as mobile, database, and server monitoring.

AppDynamics attempts to provide visibility into the entire Web stack, so you can manage the user experience (UX) across browsers and devices, and perform root cause analysis to trace the causes of performance issues.

Unlike many other APM tools and suites, Appdynamics uses custom health rules and policies to measure UX and transaction performance, rather than standard Apdex (Application Performance Index, an open standard for measuring user experience, also called UX) metrics. Health rules are customizable, and you can create rules to reflect application or transaction performance.

Setting up alerts through health rules is, of course trivial, and all monitoring tools can do that for you. What's different with tools such as AppDynamics is that you can perform a deeper analysis into the significance of an issue, in terms of how it impacts an organization's bottom line. AppDynamics's Business Impact Analytics, for example, identifies failed website transactions and connects them to the users who were impacted. This and other similar data helps the company's marketing groups run a customer win-back campaign to get back the customers they lost due to the issues.

Comparison of New Relic and AppDynamics

Both AppDynamics and New Relic support Ruby and Python, as well as all the standard application languages and platforms such as Java, Scala, .NET, PHP, Node.js. On the user monitoring side, both support iOS, Android, and JavaScript.

There are many similarities in the functionality of the two tools, but there's a big difference in the way they distill performance metrics into performance indicators. New Relic lets you manually set the threshold indicating end user stratification to calculate their Apdex score index, which measures the ratio of satisfied requests over the total number of requests. AppDynamics automatically creates a dynamic baseline for application performance, which varies depending on the load on the system.

AppDynamics focus more on visualizing the stack end to end, whereas NewRelic concentrates more on the bottom line response times. A big difference is that AppDynamics is available in a SaaS, on premise and a hybrid mode, whereas you can get New Relic services only in a SaaS model. If you must have an on premise solution you need to go with AppDynamics.

APM suites are assembled from multiple tools, which adds to their complexity, especially with regards to maintenance. While On-Premise offerings are strong, SaaS based offerings are increasingly becoming available. APM suites are viewed as expensive, and thus the pricing mechanism may need alterations to appeal to increasingly dynamic IP environments. IT operations are the primary users of the APM services. DevOps teams (APM architectures that support near real-time processing helps with this group, of course) and the application owner/line of business are also significant users of these products.

On-Call Management Tools

Increasingly, companies are moving away from the traditional monitoring models of sending alerts only to the operations teams. It's quite common now for developers and infrastructure engineers (aka site reliability engineers) to directly respond to application related alerts. The goal is simple: cut resolution time by letting the creators and maintainers of the applications directly respond to the alerts.

On-call management tools route alerts automatically to the person who is on call, usually via mobile notifications. If the person doesn't respond in a within predefined interval, the alerts go to the next person in the escalation chain.

PagerDuty is a popular On-call Management tool. Its goal is to help you resolve incidents in real time. With PagerDuty, you can minimize disruptive events and enhance the availability of your environment. As with BigPanda, PagerDuty automatically centralizes and groups alerts from an application or service, and correlates actionable events into groups, while reducing the event noise. By triaging application issues with contextual insights, PagerDuty centralizes alerts from your monitoring stack (such as Nagios, New Relic, AWS, ServiceNow) and services, and provides you visibility into your entire infrastructure.

Following are some of the key features of PagerDuty:

- Event Grouping: Aggregation, classification, and correlation of events.
- Reliable alerting: Guaranteed delivery of alerts to the appropriate persons.
- Scheduling and escalations: Configuring of on-call schedules, rotations, and escalations