

Real-time Stream Processing with Apache Flink



Timo Walther

Flink committer
twalthr@apache.org

Stream Processing



- **Data stream:** Infinite sequence of data arriving in a continuous fashion.
- **Stream processing:** Analyzing and acting on real-time streaming data, using continuous queries



Streaming landscape



Apache Storm

- True streaming, low latency - lower throughput
- Low level API (Bolts, Spouts) + Trident



Spark Streaming

- Stream processing on top of batch system, high throughput - higher latency
- Functional API (DStreams), restricted by batch runtime



Apache Samza

- True streaming built on top of Apache Kafka, state is first class citizen
- Slightly different stream notion, low level API



Apache Flink

- True streaming with adjustable latency-throughput trade-off
- Rich functional API exploiting streaming runtime; e.g. rich windowing semantics

What is Flink




A "use-case complete" framework to unify
batch and stream processing



Apache Flink



- True streaming with adjustable latency and throughput
- Rich functional API exploiting streaming runtime
- Flexible windowing semantics
- Exactly-once processing guarantees with (small) state



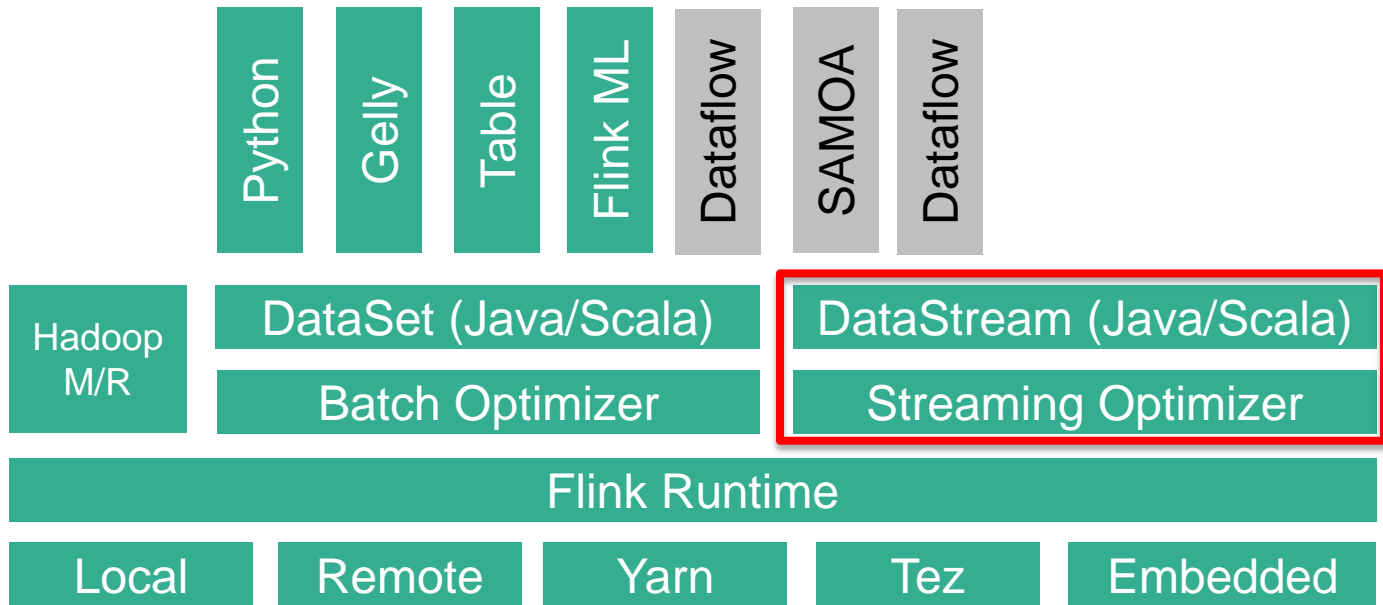
Issues

- Limited state size
- HA issue

Flink stack



**current Flink master + few PRs*



Overview of the API



- Data stream sources

- File system
- Message queue connectors
- Arbitrary source functionality


- Stream transformations

- Basic transformations: *Map, Reduce, Filter, Aggregations...*
- Binary stream transformations: *CoMap, CoReduce...*
- Windowing semantics: *Policy based flexible windowing (Time, Count, Delta...)*
- Temporal binary stream operators: *Joins, Crosses...*
- Native support for iterations

- Data stream outputs

- For the details please refer to the programming guide:


- http://flink.apache.org/docs/latest/streaming_guide.html




Use-case: Financial analytics



- Reading from multiple inputs
 - Merge stock data from various sources
- Window aggregations
 - Compute simple statistics over windows of data
- Data driven windows
 - Define arbitrary windowing semantics
- Combine with sentiment analysis
 - Enrich your analytics with social media feeds (Twitter)
- Streaming joins
 - Join multiple data streams
- Detailed explanation and source code on our blog
 - <http://flink.apache.org/news/2015/02/09/streaming-example.html>



Reading from multiple inputs




```
case class StockPrice(symbol : String, price : Double)
val env = StreamExecutionEnvironment.getExecutionEnvironment
```

```
(1) val socketStockStream = env.socketTextStream("localhost", 9999)
(2) { .map(x => { val split = x.split(",")
               StockPrice(split(0), split(1).toDouble) })
```

```
(3) { val SPX_Stream = env.addSource(generateStock("SPX")(10) _)
      val FTSE_Stream = env.addSource(generateStock("FTSE")(20) _)
```

```
(4) val stockStream = socketStockStream.merge(SPX_Stream, FTSE_Stream)
```

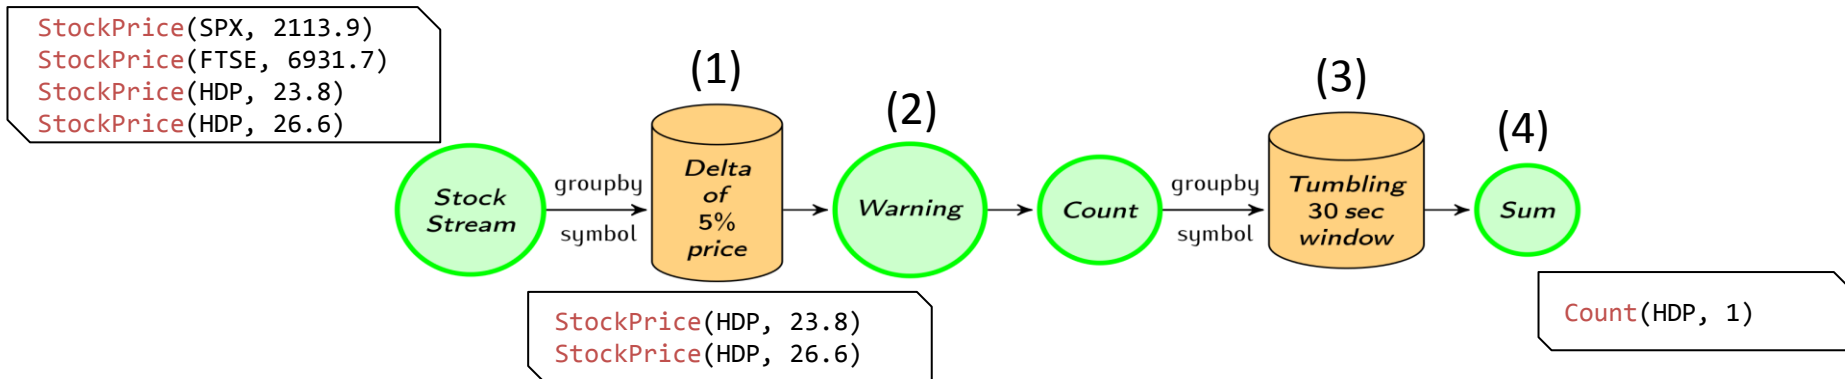
Window aggregations



```
val windowedStream = stockStream
(1) .window(Time.of(10, SECONDS)).every(Time.of(5, SECONDS))

(2) val lowest = windowedStream.minBy("price")
(3) val maxByStock = windowedStream.groupBy("symbol").maxBy("price")
(4) val rollingMean = windowedStream.groupBy("symbol").mapWindow(mean _)
```

Data-driven windows



```
case class Count(symbol : String, count : Int)
```

```
val priceWarnings = stockStream.groupBy("symbol")
```

```
(1) .window(Delta.of(0.05, priceChange, defaultPrice))
```

```
(2) .mapWindow(sendWarning _)
```

```
val warningsPerStock = priceWarnings.map(Count(_, 1)) .groupBy("symbol")
```


```
(3) .window(Time.of(30, SECONDS))
```

```
(4) .sum("count")
```

Combining with a Twitter stream



"hdp is on the rise!"
"I wish I bought more
YHOO and HDP stocks"

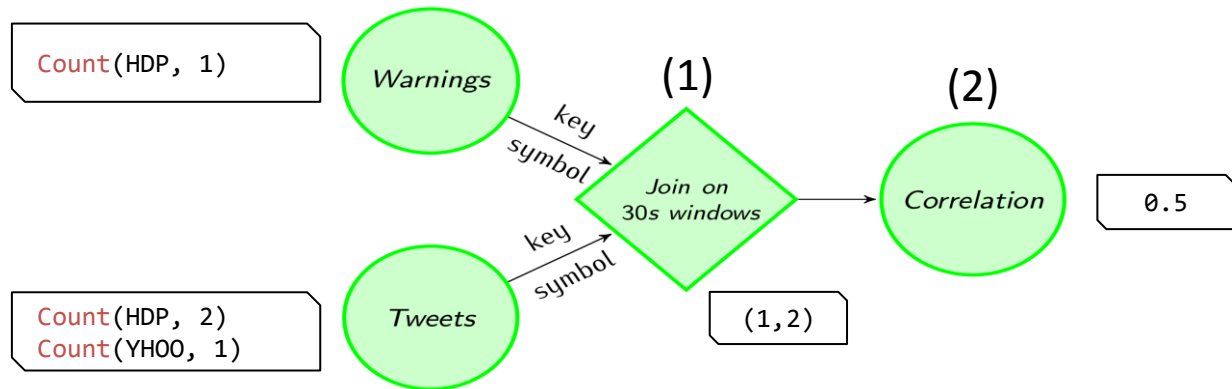


```
(1) val tweetStream = env.addSource(generateTweets _)
```

```
(2) {  
    (3) val mentionedSymbols = tweetStream.flatMap(tweet => tweet.split(" "))  
        .map(_.toUpperCase())  
        .filter(symbols.contains(_))  
}
```

```
val tweetsPerStock = mentionedSymbols.map(Count(_, 1)).groupBy("symbol")  
(4) .window(Time.of(30, SECONDS))  
    .sum("count")
```

Streaming joins



(1) {

```
val tweetsAndWarning = warningsPerStock.join(tweetsPerStock)
      .onWindow(30, SECONDS)
      .where("symbol")
      .equalTo("symbol"){ (c1, c2) => (c1.count, c2.count) }
```

(2) {

```
val rollingCorrelation = tweetsAndWarning
      .window(Time.of(30, SECONDS))
      .mapWindow(computeCorrelation _)
```



flink.apache.org
@ApacheFlink